What Prices Cannot Tell a Market Supervisor: The Cost of Identifying Trader Groups and the Value of Class Disclosure

Pratik Niroula* Prashant Niroula † Sujan Khatiwada‡

Working paper, 5 November 2026

Abstract

What can prices reveal about trader groups, and when is additional disclosure worth its cost? In a linear feedback market, prices identify a merged distribution of trading responses; equal-memory groups remain inseparable, while nearby distinct groups can require sharply increasing data. We translate identification difficulty into sample, execution and capital requirements under explicit observation contracts. A two-group supervisory example separates an aggregate prudential constraint from a targeted halt decision: nearly indistinguishable price responses can imply different feasible actions. General-population extensions require a separately verified dominant-group gap. We replace an unrestricted strategic-response claim with a bounded-response condition and connect the example to information fees, intervention losses and decision errors. A reproducible synthetic accounting study exhibits both positive and negative values of paid information and failures under bias or underestimated responses. Existing position-report studies are descriptive diagnostics; they do not identify learning memory, validate the intervention model, or establish welfare gains. Proof review, the observation-to-market bridge and the precise incremental contribution remain explicit research obligations.

*Department of Computer Information Science, Minnesota State University, Mankato, MN 56001, USA. pratik.niroula@mnsu.edu.

†Faculty of Management, Tribhuvan University, Kirtipur, Nepal. prashantniroula20@gmail.com.

‡Department of Electrical, Computer and Robotics Engineering, Minnesota State University, Mankato, MN 56001, USA. sujan.khatiwada@mnsu.edu.

Working-paper scope. The contribution is a conditional identification-and-decision result under explicit market and observation contracts. It does not establish universal market safety, real-market welfare gains or identification of proprietary algorithms from published position stocks. The new core proofs are built and audited; external mathematical ingredients, empirical mappings and independent review are distinguished from that verification.

Keywords: algorithmic trading, identification, trader heterogeneity, market microstructure, system identification, minimax sample complexity, value of disclosure, audit trails.
JEL: C14, C32, C58, D83, G14, G18.

Reproducibility and verification. The appendix maps historical tables and figures to generating code and result files. Existing empirical results are retained with their original specifications, failures and limitations; this revision does not independently rerun every historical study. "Pre-specified" denotes the authors' dated specification and repository record, not independent registration. The new accounting study reproduces six archived response rows and records both the normal-quantile and conservative gates, including adverse configurations. The pinned Lean build and strict axiom audit verify typed declarations; the exact supervisory witness, conservative Gaussian event bound and acquisition algebra have explicit coverage below. Historical name mappings or declaration counts do not establish full manuscript equivalence. The stationary IV central limit argument, external approximation results, general-K supervisory construction and economic observation assumptions retain their stated limits. Code and receipts accompany this working paper; independent proof review remains a separate task.

Contents

1 Introduction **4**

1.1 Related literature 14

1.2 Outline 21

2 Model and motivating examples **21**

2.1 Discussion of the assumptions 25

2.2 The observer, its options and its costs 27

3 The identification boundary **29**

4 Probing resources, memory resolution and money **31**

4.1 Quantity 32

4.2 Execution cost 34

4.3 Capital supplied by a counterparty 34

4.4 Capital paid by the probe trader 36

4.5 A worked example: two Kalman desks 37

4.6 What prices hide, in economic terms 37

4.7 Natural excitation and the cost of ambiguity 41

5 Account-level capital **41**

5.1 Observation design and the tri-region partition 41

5.2 Account-level capital and the gap to the lower bound 43

5.3 Robustness: heteroskedastic impact and probing protocols 48

5.4 From par units to Bitcoin dollars 48

6 Unknown mixtures and passive data **49**

6.1 Why naive pole fitting fails 49

6.2 Local information: the Cramér–Rao benchmark 50

6.3 A global two-point bound for any estimator 50
6.4 Many unknown cohorts 52
6.5 Do the rates survive real market noise? 58
6.6 Near-colliding rates: a uniform rate for the merged measure 59
6.7 General memory: the minimax rate is logarithmic 61
6.8 Passive data 65

7 Robustness to nonlinear and adaptive trading rules 67

8 What disclosure buys 70

8.1 How much disclosure is worth 72
8.2 Disclosure in dollars 75
8.3 When disclosure is worth acquiring 78
8.4 Mandated disclosure: the CFTC's trader-class reports 86
8.5 Voluntary disclosure: Binance's trader-class ratios 92

9 A spectral blind spot 97

10 Empirical feedback and its measurement 99

10.1 Estimand and filtering mechanics 99
10.2 Data 100
10.3 Clock alignment dominates reduced-form feedback 100
10.4 Feedback beyond filtering mechanics 100
10.5 Order splitting and the estimation of feedback 101
10.6 Feedback across 30 perpetual markets 102
10.7 Two fee events 102

11 Implications for supervision and market design 103

12 Conclusion and an open problem 106

A Proofs **115**

B Notation **128**

C Internet appendix **129**

C.1 A finite-window price-surveillance bound 129

C.2 Passive distinction of a known pair in finite samples 130

C.3 Blind inversion of the closed loop 131

C.4 Probing protocols 131

C.5 Coordinate-weighted output 131

C.6 Real-noise injection: requirements at each spacing 132

C.7 Numerical stress tests of results that are not machine-checked 132

C.8 Probing a running market without resets 132

D Reproducibility: source of every number **133**

Returned advanced experiment: evidence and limits **175**

Research work register: 30 October 2026 **179**

1 Introduction

Markets are increasingly run by algorithms. Trend-following funds, market-making programs, execution algorithms and, more recently, AI trading agents all read recent prices, decide what to trade, and by trading move the prices they read. When something goes wrong, a sharp sell-off, a sudden loss of liquidity, a crowded trade unwinding, the first questions supervisors and risk managers ask are who was trading and how. Which group of traders drove the move? How quickly do they react, and how long do they remember what prices did? Could the same group do it again?

This paper asks a question that supervisors now face when they decide which trading records to require: which questions about who drives prices can public data answer, which can they answer only at a prohibitive cost, and what is it worth to be told the rest?

The question matters now. Trading is passing to learning algorithms and AI agents (Colliard et al., 2022; Dou et al., 2025), and regulators have begun to require firms to say which algorithm sent each order. Under MiFID II, EU trading venues must keep order records that identify the person or algorithm responsible for each investment and execution decision (European Commission, 2017; ESMA, 2016). The US Consolidated Audit Trail (CAT) is designed to record each order in US-listed stocks and options, and the account behind it, across exchanges and broker-dealers; its revised 2026 budget, published in March 2026, is about $147 million (SEC, 2012; CAT LLC, 2026). Such systems serve many purposes, among them enforcement, rebuilding market events after the fact, and surveillance. One question they raise, and the one we study, is how much of the information about who drives price moves could instead be learned from public data, and at what cost.

The question. A supervisor wants to know which group of traders drives prices. What can public prices tell it, what does it cost to learn the rest, and what are trader-class reports worth?

The answer.

1. *Prices see the crowd, not the groups.* They reveal one combined summary of all the groups. Two groups with the same memory look exactly like one larger group, however much price data is collected (Theorem 3.1).
2. *So supervision splits in two.* A market-wide measure depends on an aggregate response, subject to the observation and tail conditions. A targeted decision can require more information: two markets can look almost identical in prices while halting the largest group meets the prudential ceiling in one and not in the other (Proposition 8.5, machine-checked for two groups).
3. *Learning the groups from prices is very expensive.* The data needed grow as $(\Delta\eta)^{-(4K-2)}$ when the memories of K groups are $\Delta\eta$ apart: halving the gap multiplies the data by 64 for two groups and by 1,024 for three (Theorems 6.4 and 6.5). If memory is a general mixture, accuracy improves only with the logarithm of the data (Theorem 6.9).
4. *Class records change the information requirement.* When each class's flow is reported separately, the problem stays well conditioned however similar the groups are (Proposition 8.1). In 20 futures markets, the CFTC's weekly class report pins down the persistence of speculative positions 9 to 71 times more precisely than prices; a hypothetical common-root-T conversion gives about 1,350 at the median, but does not identify an equivalent sample requirement (Figure 1).
5. *What the empirical failures reveal.* The tests do not establish that fitted persistence is learning memory, and coarse reporting buckets can mix opposing strategies. Ten of the twenty-two predictions we tested failed (Table 1): in equity indices, basis arbitrageurs trade against trend-followers, causing within-class sign cancellation and motivating further tests of more granular records without establishing their necessity.

We keep three statements apart throughout: what cannot be learned at all, what can be learned but only at high cost, and where estimates become unstable.

Precise comparison with existing identification research. Clustered exponential recovery already explains ill-conditioning and separation-dependent rates (Batenkov et al., 2021; Katz et al., 2023); those exponents are not claimed as a general new discovery. Data informativity already distinguishes information sufficient for a control task from unique system identification (van Waarde et al., 2020); the broad distinction is likewise not new.

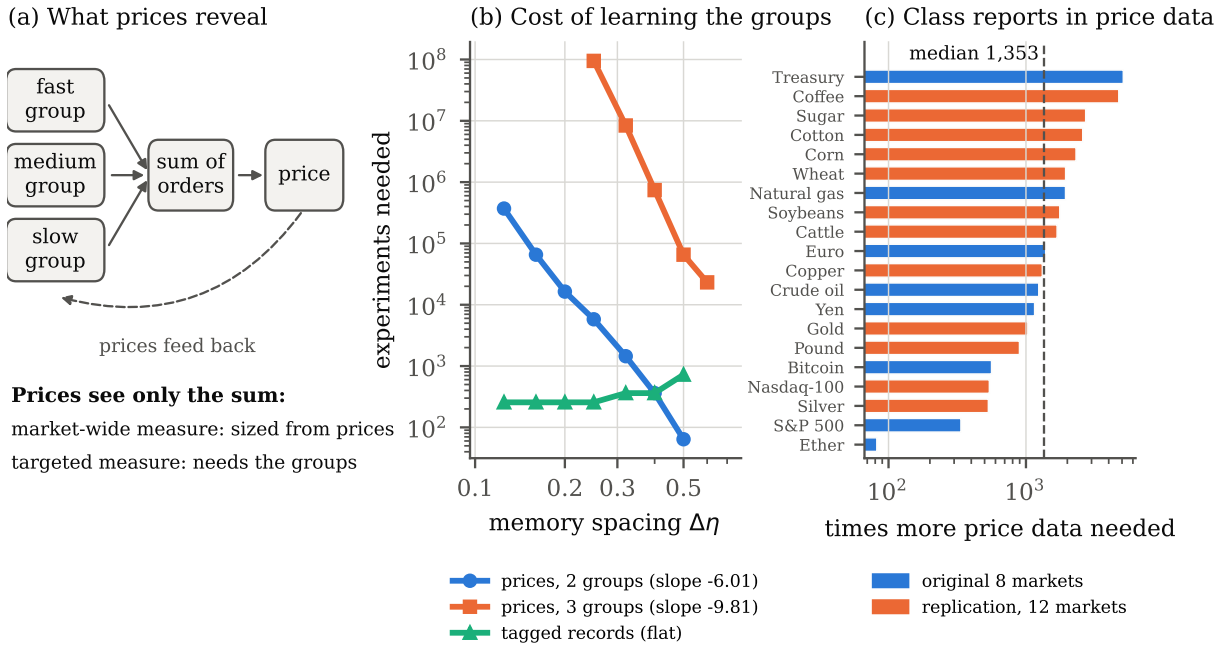


Figure 1: The paper in one figure. (a) Groups of traders reach the price only through the sum of their orders, so prices reveal the combined behaviour: enough to size a market-wide measure, not to decide whether halting one group is enough. (b) Pre-specified real-noise injection test (Section 6.5): experiments needed to recover the groups' weights against the spacing of their memories, with real one-minute Bitcoin price noise; price-only requirements grow with slopes -6.01 (two groups) and -9.81 (three groups; theory -6 and -10), tagged requirements stay flat. (c) Post hoc: how many times more data a price-only analysis would need to match the precision of the CFTC's weekly class report under standard root-T scaling, in the 20 markets of the registered tests (Table 24); the conversion assumes common root-T scaling and comparable estimands, which these data do not establish. Kind: registered results (b) and a calculation from registered results (c), `figs/make_fig_summary.py`.

Our candidate contribution is their explicit combination in this market feedback model: the public-response equivalence class, an execution/capital interpretation under a specified probing contract, and a two-population supervisory decision whose feasible targeted action differs despite close public responses. The all-K action extension needs an additional weight-gap construction. Classical system identification (Ljung, 1999) provides the observation/estimation framework; a price-only record here is neither exogenous order flow nor a revealed portfolio preference. Novelty must be judged on these model-specific results against the cited theorems, not on AI terminology or a claim to have invented supervision with-

out identification. Task-optimal exploration already supplies quantitative, task-dependent lower bounds and efficient experimental designs in linear systems (Wagenmaker et al., 2021). Accordingly, neither quantitative task-dependent learning nor choosing information for a decision is claimed as new here. Our comparison must turn on the latent population map, the financial execution/capital contract, and the specified class intervention; no general dominance over task-optimal exploration is asserted.

Why prices hide the groups. The reason is the feedback loop. Algorithms trade on recent prices, their orders move prices, and those prices feed back into the algorithms. We model each group as tracking the price with a fading memory: recent prices count most, older prices count less, and one number, the group's memory rate, sets how fast the past fades. A desk behaves exactly this way if it uses a Kalman filter or a regularized online learning rule (Lemma 2.1). Because every group reaches the price only through the sum of all their trades, the price records the total and loses the parts. Think of several people pushing the same door: from the door's movement one can learn the total push and how it builds and fades, but two people who push in the same rhythm look like one stronger person.

Three ways to learn more, and what each costs. An observer that wants to know more, whether a supervisor, an exchange's surveillance team or a large trading desk, has three options (Section 2.2).

(a) *Wait and watch prices.* This costs time. In a stationary Gaussian market observed without error, a passive observer needs at least as many observations, in order of magnitude, as an active one needs experiments (Theorem 6.11).

(b) *Trade against the groups* and watch how the price responds, much as one taps a wall to find a hidden beam. Every test trade moves the price against the trader, so learning this way costs money. To tell one group from two groups whose memory rates differ by 2δ, the quantity traded grows like $1/\delta^2$ and the capital at risk like $1/\delta^4$: halve the

difference and the capital needed rises sixteen-fold (Theorems 4.1 and 4.3). If memory is a general mixture of fading memories rather than a few exponentials (with the slowest bounded away from a random walk), accuracy improves only with the logarithm of the data, so halving the error squares the data needed (Theorem 6.9).

(c) *Acquire labelled records*, such as audit-trail data or class-level position reports. This costs whatever the reporting system costs, but it removes the explosion in data: labels save a factor of order $(\Delta\eta)^{-(4K-2)}$ at fixed accuracy, where K is the number of groups and $\Delta\eta$ the spacing of their memory rates (Theorem 8.2).

Proposition 8.4 says when each option is the cheapest.

How large are the costs? A hypothetical illustration. To give a sense of scale we price the trading option in a hypothetical benchmark. The model parameters are illustrative; only the size of the measurement noise is estimated, from Bitcoin data. An outside participant that re-learns two unknown algorithms each week with test trades would lose at least $15 million to $356 million a year per market; with three unknown algorithms it would lose more than $40 million a year per market even when their memories are far apart (Section 8.2). These are lower bounds for one stated way of learning, not estimates of what any regulator spends: regulators do not trade to learn. They are best read as the price of the information that labelled records deliver without any trading.

Why this matters for finance. The results point in a constructive direction for several groups of readers.

- *Supervisors and policymakers* get a way to decide what data to require. If the goal is only to watch overall price risk, public prices are enough, because the combined price response is learned at the ordinary rate (Remark 6.3). Labelled records earn their cost when supervision must be targeted, for example when one algorithm must be paused while the others keep supplying liquidity, and when the reporting system

covers enough markets to share its fixed cost (Proposition 8.4). In our stylised cost model this illustrates a possible economy of scale for market-wide reporting; it does not explain the legal design of regimes such as CAT. For the narrow task of attributing behaviour to classes, a persistent pseudonymous label is enough, which bears on how much personal data such a task needs but not on what enforcement needs (Section 11). Proposition 8.5 gives a stylised decision in which this matters: the size of a market-wide measure depends only on the combined response, while deciding whether halting one group is enough needs that group's own intensity and inherits the full cost of separating the groups.

- *Exchanges and trading venues*, which already see account-level flow, learn which questions about their participants public prices cannot answer, and why publishing class-level positioning, as some crypto venues already do, adds information that prices do not contain.
- *Risk managers and trading desks* get a realistic limit on what can be inferred about other traders or about crowding from market data, and a warning that common feedback measures computed from public data mostly reflect filtering and timing rather than true feedback (Section 10).
- *Researchers* who estimate models with several trader types get one reason why the parameters that separate the types are often imprecise, and a number that says how precise they can be before any estimation (Section 1.1).

What the data show. Two real reporting systems let us check the main prediction, that class records pin down what prices leave open. The U.S. Commodity Futures Trading Commission collects positions from large futures traders and publishes a weekly snapshot of them, aggregated by trader category. In a pre-specified test in eight futures markets (CME Bitcoin and Ether, the E-mini S&P 500, 10-year Treasury notes, the euro, the yen, crude oil and natural gas), these class positions pin down the persistence of the designated

speculative class's positions to a range 9 to 71 times narrower than weekly prices allow, while prices alone leave it almost anywhere in the admissible range; a replication registered later on twelve new markets (metals, grains, softs, cattle, the Nasdaq-100 and the pound) holds in all twelve (Section 8.4). Converted into data, matching the precision of the weekly class report from prices alone would take, at the median market, about 1,350 times as much price data, the equivalent of thousands of years of weekly prices; this is hypothetical arithmetic, not a lower bound, because estimands and effective horizons differ and intervals reach parameter bounds. On Binance, in a pre-specified test on 30 markets never used before that matches the memory horizon of the two routes, the class data give the narrower interval in 29 of 29 evaluable markets, by a median factor of 7.3 (Section 8.5).

Robustness, a new failure and real noise. Rerunning the CFTC test with CME futures prices, without contract-roll weeks or with longer bootstrap blocks leaves its primary result unchanged. Read as weekly flows rather than positions, the same reports failed a second pre-specified test: they pin down how each class's flow responds to prices, but not how persistent that flow is. A third test asked whether positions carry a second, faster root that could be learning memory, as the model with inventory costs predicts; it failed too (2 of 8 markets), so these data cannot separate learning memory from the slow turnover of positions. Another pre-specified test gives synthetic groups real Bitcoin price noise instead of Gaussian noise and finds the predicted rates, a slope of -6.01 against -6 for two groups and -9.81 against -10 for three, while class-tagged requirements stay flat; exact multinomial empirical resampling is maintained across all sample sizes without a normal approximation (Section 6.5). These comparisons are descriptive: a position series reveals its own persistence directly, while the price route must infer it. A replication registered on 31 October 2026 reproduces the primary CFTC result on twelve new markets (12 of 12). Two tests on fresh October 2026 Binance data failed: price changes did not forecast a class position beyond its own persistence, and aggregate taker flow did not pin down the combined response, although its uncertainty shrank

at the predicted root-T rate. Several predictions failed, and Table 1 lists them next to the ones that passed. The data therefore show what prices cannot pin down. They do not show that this persistence is the learning memory of the model.

Two supplementary results. Neither is needed for the main results. Markets with the same eigenvalues, the same one-step volatility and the same long-run price variance can still have very different worst-case short-run losses, so watching eigenvalues or volatility cannot replace knowing who trades (Theorem 9.1). And standard feedback measures computed from public data mostly reflect filtering and clock alignment: two fee changes that moved trading volume by factors of 2.6 and 4.9 left measured feedback unchanged (Section 10).

Contribution. The candidate contribution is the market-model connection between a public-response equivalence class, the sample and financial resources needed for finer identification, and a specified supervisory action. Equal-memory groups are exactly inseparable; distinct but nearby groups face a finite-noise difficulty, not permanent nonidentification. Proposition 8.5 establishes a two-group separation example and a conditional general-K extension. The statistical clustering exponent and the distinction between task-relevant information and full identification have precedents; the precise comparison below states what is inherited and what requires model-specific proof. Historical position-report evidence is descriptive and does not establish the action model or the necessity of granular reporting. Formal receipts cover particular typed statements; semantic correspondence and independent proof review remain separate obligations.

How to read the paper. Every main result is stated formally and then explained in a short paragraph headed "In words", which says what it means for a trader or a supervisor; proofs are in the online supplement. Section 4.5 works through two desks that use Kalman filters, in numbers, and Table 5 maps each question to its answer. Readers who want the finance first can read this introduction, Section 2.2, the worked example and Section 8, and

Table 1: What the empirical tests found: passes and failures

Test	Prediction	Result
CFTC, 8 markets, H1	Class positions pin down a persistence that weekly prices leave open	Pass, 8 of 8
CFTC, H2	The designated speculative class trend-follows	Fail, 4 of 8 (three markets have the opposite sign)
CFTC, H3	That class has the less persistent positions	Pass, 6 of 8 (two indistinguishable)
CFTC flows (weekly changes), H1	Class flows pin down their own persistence (95% interval narrower than 0.5)	Fail, 3 of 8 (6 required); flow response nonzero in 6 of 8
CFTC positions, two-root test, H1	Positions carry a second, faster root (memory) distinct from the slow turnover root	Fail, 2 of 8 (6 required); where positive, the second root has a half-life under half a week
Real-noise injection, $K = 2$	Price-only requirement grows as $(\Delta\eta)^{-6}$; tagged requirement stays bounded	Pass: slope -6.01 (exact multinomial resampling at all N); tagged max/min 2.8 (≤ 8 required)
Same test, $K = 3$ (secondary)	Price-only requirement grows as $(\Delta\eta)^{-10}$	Holds: slope -9.81
Binance, 30 fresh markets, (i)–(ii)	Class data restrict persistence; price-only intervals wide	Pass, 30 of 30 (one-hour price kernel)
Binance, same test, (iii)	Prices pin down the combined response	Fail, 5 of 30; the test failed its 80% criterion
Binance, 30 new markets	Horizon-matched: class interval narrower	Pass, 29 of 29 evaluable (median factor 7.3)
Binance, same test, secondary	Class persistence is a stable memory	Not supported: close to a random walk after bias correction
Binance, 105-day follow-up	Combined response far cheaper to learn than the separate rates	Fail, 10 of 30 (23 required)
Binance, 58-day follow-up	Tagging advantage grows with the sample; parametric shrinkage	Pass, 22 of 29; fail, 20 of 29 (22 required)
Binance, return forecasts	Class memory states improve 30-minute return forecasts	Fail ($\Delta R^2 = -1.09$ points, inconclusive)
Binance, state tracking	Class data track class positions better than prices	Pass, but explained by persistence of the target
CFTC, 12 new markets, H1 (replication)	Class positions pin down a persistence that weekly prices leave open	Pass, 12 of 12 (tagged widths 0.014–0.042, price-only 0.85–0.96)
Same, H2	The managed-money or leveraged class trend-follows	Pass, 11 of 12 (Nasdaq-100 has the opposite sign)
Same, H3	That class has the less persistent positions	Fail, 3 of 12 (6 required)
Binance, Oct. 2026, 30 markets, N2	Prices improve out-of-sample forecasts of the top-trader position beyond its own persistence	Fail, 1 of 30 (18 required); median R^2 gain -0.01 points
Same, N3	Taker flow pins down the combined response to past returns	Fail, 5 of 30 (24 required); interval excludes zero in 17
Same, N3 secondary	The sign of that response is stable within the month	Fail, 21 of 30 (24 required)
Same, N3 secondary	Its interval shrinks at the root-T rate	Pass, 30 of 30 (median ratio 1.69, theory 1.73)

take the statistical results as given. Table 2 explains the main symbols.

Table 2: The main symbols in plain words

Symbol	Meaning
K	Number of trader groups ("cohorts")
η_i	Group i's learning rate: the weight it puts on the newest price. Large η_i means a short memory and fast reactions; small η_i means a long memory
$r_i = 1 - \eta_i$	Group i's memory rate: the share of last period's view that survives
q_i	Group i's trading intensity: how strongly it trades on its view. Positive q_i means it buys when the price is above value, as trend followers do
$\Delta\eta$, 2δ	How far apart the groups' learning rates are ($\Delta\eta$ for unknown groups, 2δ for a known pair)
a, κ	How persistent pricing errors are, and how far one unit of net order flow moves the price
M (merged measure)	The combined summary of all groups that prices reveal: where the memories are and how much trading sits at each
N	Number of experiments (test trades from a fresh start) or observations
Par units	The model's own unit of quantity and money, before conversion to USDT (Table 10)

What we do not claim. The identification results hold for the stated linear model and for the stated observation contracts, that is, for what the observer is assumed to see. The model treats price impact and the persistence of pricing errors as known primitives rather than solving a dealer equilibrium, and it assumes the observer knows the fundamental value; real data meet neither assumption. They are not general limits on surveillance with participant-level records, which can be far more informative (Propositions 4.4 and 8.1). The matching upper bound holds on a class with separated rescaled rates and weights bounded away from zero. Capital figures depend on the observation-noise calibration, which we estimate rather than assume. The empirical results are reduced-form diagnostics; they do not identify structural learning parameters. Our pre-specified real-data test of the identification boundary failed its criterion on one of its three parts: on 42 days of data the merged response was precisely estimated in only 5 of 30 markets (Section 8.5). Binance's class ratios are positions, not tagged order flow, so they are the closest public analogue to the disclosure of Proposition 8.1, not an instance of it. We hold the algorithms' memories and intensities fixed across disclosure regimes; disclosure that reveals positions may lead informed desks to trade less or to split

orders across accounts, which we do not model.

1.1 Related literature

Algorithms, AI and supervision. Recent work in finance asks how learning algorithms and AI agents behave in markets: reinforcement-learning market makers and informed traders can sustain wide spreads or collusive outcomes (Colliard et al., 2022; Dou et al., 2025), prices formed against Q-learners have distinctive dynamics (Banerjee and Szydlowski, 2026), and algorithmic markets can be manipulated (Tremacoldi-Rossi, 2025). Traders who differ in horizon or speed are studied in equilibrium by Cespa and Vives (2015) and Biais et al. (2015), and Crouzet et al. (2020) model investors who trade at different frequencies. These papers model or simulate the algorithms. We take the supervisor's side and ask what can be learned about them from market data, and what disclosure is worth. A related literature measures the cost of regulation and of regulatory reporting (Trebbi et al., 2024); our dollar comparison with the Consolidated Audit Trail is a statement about the benefit side: what the audit trail buys that prices cannot.

Price formation and impact. Our price equation is a discrete linear impact model in the tradition of Kyle (1985), Glosten and Milgrom (1985) and Hasbrouck (1991), in which order flow moves prices and prices revert toward value. Transient impact and the cost of trading follow Almgren and Chriss (2001), Obizhaeva and Wang (2013), Gatheral (2010) and Bouchaud et al. (2004); Gatheral (2010) shows that absence of dynamic arbitrage requires a positive-definite decay kernel, Alfonsi et al. (2012) that excluding transaction-triggered manipulation requires convex, nonincreasing decay, and Alfonsi et al. (2016) extend positive definiteness to several assets. Positive definiteness is the property Lemma 4.2 delivers in our cohort economy. Bouchaud et al. (2018) document order splitting and long memory in order flow, which is the main confounder in Section 10. Relative to this literature we treat the impact kernel not as a primitive but as the closed-loop output of adaptive traders, and ask

what it reveals about them.

Algorithmic trading and feedback. Algorithms that react to recent prices can amplify or dampen moves. Brunnermeier and Pedersen (2009) and Shleifer and Vishny (1992) study liquidity spirals and fire sales, and Kirilenko et al. (2017) document the behaviour of automated intermediaries in the 2010 Flash Crash. Sornette et al. (2023) show that non-normal interactions can generate transient bubbles. Our contribution is on the measurement side: we show what a closed-loop observer can and cannot learn about the feedback, and that standard reduced-form indices are dominated by filtering and timing.

Learning algorithms and what can be inferred about them. Empirical game-theoretic analysis treats the choice among simulated trading strategies as a game (Wellman, 2006). Empirically, algorithmic trading has been found to improve liquidity and price discovery on average (Hendershott et al., 2011; Brogaard et al., 2014). These studies simulate the algorithms or observe them through regulatory data; we ask what an outside observer can infer about them from market data alone, and answer in samples and in capital. Our stance is that of partial identification (Manski, 2003): when the structure is not identified, characterize what the data do determine, here the merged measure and, through it, the intensity-weighted distribution of the desks' information quality (Proposition 4.6).

Control without identification. We do not claim that acting on a system without identifying it is new. Moallemi et al. (2008) study strategic execution against an observer who learns from prices, which bears directly on our assumption that cohorts do not react to being observed, and Park and Van Roy (2015) study execution that learns price impact while trading, with regret guarantees; our object is different, the latent trader cohorts behind the impact rather than the aggregate impact coefficients. Active identification of linear systems has matching bounds and optimal experiment design (Wagenmaker and Jamieson, 2020), and task-directed exploration asks what must be learned for a decision (Wagenmaker et al.,

2021); our information matrix differs because the unknowns are zeros of a closed-loop response and the input is paid for in trading losses. van Waarde et al. (2020) characterize when control is possible without unique system identification. Our results concern a specific market observation contract and its costs in executed quantity, capital and money.

Closest work. The closest literature is super-resolution and moment-based mixture estimation, which gives rates for recovering clustered point sources and mixing measures from noisy moments (Moitra, 2015; Heinrich and Kahn, 2018; Batenkov et al., 2021; Liu and Ammari, 2022; Wu and Yang, 2020; Kong and Valiant, 2017). Our K-cohort rate is, in that language, the square of the amplitude condition number $(\Delta\eta)^{-(2K-1)}$ for K clustered nodes (Heinrich and Kahn, 2018; Wu and Yang, 2020), and we do not claim the exponent as new. What changes when the mixture sits inside a closed-loop financial market is fourfold (Table 3): first, in mixture estimation and open-loop system identification the unknown rates are locations or poles of the observed law, whereas in a closed-loop market the cohorts feed back into the price, making the memory rates transmission *zeros* of the observable transfer function $H(z)$ that must first be extracted from the loop (Theorem 3.1); with a and κ known this extraction is a smooth invertible map, so it decides which method is valid (pole fitting on prices is not) rather than the rate; second, passive price observations face at least the same $(\Delta\eta)^{-(4K-2)}$ sample-complexity barrier as active probing, uniformly in the spacing, because $H(z)$ is minimum phase and the Whittle likelihood inherits the moment gap (Theorems 6.10 and 6.11); third, in a market statistical sample size translates into executed quantity $\Theta(\delta^{-2})$ and round-trip trading loss or margin capital $\Theta(\delta^{-4})$ via the self-impact Toeplitz quadratic form $\frac{1}{2}v^\top T_h v$ (Theorems 4.1–5.4); and fourth, class-tagged disclosure breaks the feedback loop cohort by cohort and cuts the sample requirement by $(\Delta\eta)^{-(4K-4)}$ (Proposition 8.1 and Theorem 8.2). Beyond the K-cohort case, we prove matching lower and upper bounds of exact logarithmic order $\Theta(1/\log N)$ for general completely monotone memory (Theorem 6.9) and price the resulting disclosure savings against audit-trail budgets.

Several ingredients are known in other settings, and we claim only their closed-loop, passive and financial forms: the δ^{-4} law for telling one component from two closely spaced ones is classical in resolution analysis (Shahram and Milanfar, 2005); logarithmic rates are known for mixtures of exponentials (Rebafka and Roueff, 2015) and for general mixing distributions, by the same moment-matching and polynomial-approximation technique (Wu and Yang, 2020); and labelled samples are known to be far more informative than unlabelled ones in mixture problems (Castelli and Cover, 1996).

Estimating heterogeneous-agent models from prices. A large literature models markets as populations of trader types that switch between forecasting rules (Brock and Hommes, 1998) and estimates such models from price data. Estimates of the parameters that separate the types are often imprecise: Amilon (2008) reports that the estimated model reproduces properties of returns but that its fit is generally poor, Franke and Westerhoff (2012) compare switching mechanisms by the method of simulated moments, and Kukacka and Barunik (2017) find the switching coefficient statistically insignificant by simulated maximum likelihood while the belief parameters are significant. These papers take the identification of the types from prices as given. Our results suggest one reason the type-separating parameters are hard to pin down: around a steady state, prices determine only the merged measure, and separating types whose memories differ by $\Delta\eta$ takes a number of samples of order $(\Delta\eta)^{-(4K-2)}$ for K types (Theorems 3.1 and 6.4). The models these papers estimate are nonlinear, so this is a statement about their linearization, not a claim that their estimates are wrong; it does give a quantity, the separation of the types' memories, that bounds the precision one can expect before estimating.

Closed-loop system identification. Engineering has long studied identification under feedback (Ljung, 1999; Forssell and Ljung, 1999). There the object is the plant, the difficulty is that feedback correlates the input with the noise, and the remedy is external excitation that makes the experiment informative. Our problem is the reverse in two ways. The object is

the feedback law (the traders), not the plant (the price), and the memory rates enter as zeros of the observed transfer function. And external excitation does not remove the difficulty: with a probe the merged measure is identified exactly (Theorem 3.1), so the barrier is not qualitative identifiability but the conditioning of the map from memories to responses, which we measure in samples, quantity and capital. In finance, Leach and Madhavan (1993) show that a market maker can gain by setting prices to learn about demand, so learning through trading has an equilibrium counterpart; our probing observer is a reduced-form version of that idea without the equilibrium. In econometric terms, as $\Delta\eta \to 0$ the cohort parameters pass from strong to weak identification (Andrews and Cheng, 2012), and our rates state that passage in samples.

Inferring trader types from market data. Market microstructure has a tradition of inferring hidden trader populations from public data. The probability of informed trading (Easley et al., 1996) fits a mixture model to daily counts of buys and sells to estimate the share of informed orders, and what it captures has been debated (Duarte and Young, 2009). With regulatory account data, Kirilenko et al. (2017) classify traders by their holding and trading patterns. Our results speak to the case where types differ in how long they remember prices: from prices alone the separation is costly, and account or class records of the kind Kirilenko et al. (2017) use remove most of the cost (Proposition 8.1).

Statistical identification. The identification results use two-point testing arguments (Le Cam, 1986; Tsybakov, 2009). The moment-matching construction of Section 6 is the dynamical analogue of the slow rates in finite-mixture estimation (Manole and Ho, 2022) and of super-resolution of near-colliding sources (Liu and Ammari, 2022; Batenkov et al., 2021); the general-memory result combines best-approximation duality (Bernstein, 1914; Rivlin, 1981) with the moment comparisons used by Kong and Valiant (2017) and Wu and Yang (2020), and its logarithmic rate parallels supersmooth deconvolution (Fan, 1991); recovery of exponentials connects to ESPRIT and matrix-pencil methods (Roy and Kailath, 1989; Hua and

Table 3: Comparison with finite-mixture estimation and open-loop identification

Feature	Finite-mixture estimation (Heinrich and Kahn, 2018; Wu and Yang, 2020)	Open-loop exponential fitting (Roy and Kailath, 1989; Batenkov et al., 2021)	Closed-loop market identification (this paper)
Where rates live	Component locations / moments of density	Poles of impulse response	Transmission zeros of closed-loop $H(z)$ (Thm 3.1)
Active vs. passive data	i.i.d. draws only	Designed input vs. output spectrum	At least the same $(\Delta\eta)^{-(4K-2)}$ order, uniformly, via the minimum-phase Whittle gap (Thm 6.11)
Economic unit of cost	Sample size N	Signal-to-noise ratio	Quantity $\Theta(\delta^{-2})$, trading loss and capital $\Theta(\delta^{-4})$ via $\frac{1}{2}v^\top T_h v$ (Thms 4.1–5.4)
Effect of class tagging	Observed component labels ($\Theta(1)$)	Separate sensor per mode	Breaks closed loop cohort-by-cohort; saves $\Theta((\Delta\eta)^{-(4K-2)})$ at fixed accuracy, $\Theta((\Delta\eta)^{-(4K-4)})$ for rates (Thm 8.2)

Sarkar, 1990). Minimum-phase spectral factorization and the Whittle likelihood are classical (Whittle, 1953). What is new here is the mapping from these statistical rates into executed quantity and money through a financially specified account, the passive equivalence, and the value of class records; that the relevant rates are zeros, not poles, of the observable transfer function matters for which estimator is valid, not for the rate.

Non-normal dynamics and supervision. Transient growth of stable non-normal systems is classical (Trefethen and Embree, 2005; Kreiss, 1962). We use it to show that supervision through eigenvalues and second moments of prices is blind to transient risk (Theorem 9.1), which motivates the identification-free control of the companion paper. On the regulatory side, pre-trade risk controls for broker-dealers with market access are governed by SEC Rule 15c3-5 (SEC, 2010), and participant-level audit trails are the natural disclosure regime for Proposition 8.1. A large literature studies how pre- and post-trade transparency affects liquidity and trading costs (Pagano and Röell, 1996; Boehmer et al., 2005; Bessem-

Table 4: Closest results and how this paper differs

Existing result	Observations	Target	Conclusion	Difference from this paper
Mixture estimation (Heinrich and Kahn, 2018; Wu and Yang, 2020)	i.i.d. draws from a K-component mixture	Mixing measure	Rates that slow sharply as components merge; $n^{-1/(4K-2)}$ for Gaussian location mixtures	Same exponent; here the measure is observed through a feedback loop, as zeros of H, actively or passively, and the cost is also priced in quantity and capital
Super-resolution and Prony (Moitra, 2015; Batenkov et al., 2021; Katz et al., 2023)	Noisy Fourier or exponential samples	Point-source locations	Stability governed by Vandermonde conditioning and node separation; Prony attains the min-max bound for clustered exponents	Real nodes extracted from a closed loop; our matching upper bound (Theorem 6.5) adapts this known result
Closed-loop system identification (Ljung, 1999; Forssell and Ljung, 1999)	Input–output data under feedback	The plant	Consistency requires informative (externally excited) experiments	Target is the feedback law (the traders); excitation gives exact identification of M, and the barrier is conditioning, measured in samples and capital
Heterogeneous-agent estimation (Brock and Hommes, 1998; Franke and Westerhoff, 2012; Kukacka and Barunik, 2017)	Prices and returns	Type and switching parameters	Type-separating parameters often imprecise	Bounds that precision by the types' memory separation (in the linearization)
Probability of informed trading (Easley et al., 1996)	Daily buy and sell counts	Share of informed orders	Mixture estimate of trader population	Types that differ in memory; price-only versus class-tagged data
Data informativity (van Waarde et al., 2020)	Finite data, often noiseless	Whether a property is determined	Qualitative conditions	Quantitative minimax costs under noise

binder et al., 2006); we ask a different question, what disclosure lets a supervisor learn about the algorithms that trade, and give the answer in samples. Crypto venues, whose market structure is studied by Makarov and Schoar (2020), are a natural laboratory for this question because they publish class-level positioning every five minutes, while the CFTC publishes it weekly for regulated futures.

Partial fractions, Gaussian testing, convolution inequalities and positive-system gains are established tools. Our contribution is to connect them into a quantitative chain from memory resolution, through probing resources, to money, with matching orders, explicit constants and a supervisory interpretation, and then to measure the relevant quantities in data.

1.2 Outline

Section 2 sets out the model. Section 3 gives the identification boundary. Section 4 links memory resolution to quantity and money for a known pair. Section 5 computes account-level capital. Section 6 treats unknown mixtures and passive data. Section 7 checks the results against stylized nonlinear and adaptive trading rules. Section 8 studies disclosure and compares its cost in dollars with the U.S. Consolidated Audit Trail. Section 9 gives the spectral blind spot. Section 10 reports the empirical evidence. Section 12 concludes. Proofs are in Appendix A; Appendix C collects additional results, and Appendix D maps every number to its source.

2 Model and motivating examples

The model in words. The market has one price and K groups of traders. Each period three things happen. First, the price moves: it drifts back toward its fundamental value, and it is pushed by the groups' trades, by any test trade and by news. Second, each group updates its view by mixing the newest price into its running average, with weight η_i on the newest price and $1 - \eta_i$ on its old view. Third, each group trades in proportion to its view,

with intensity q_i. A group with a large η_i has a short memory and reacts quickly, like a short-horizon trend follower; a group with a small η_i has a long memory and reacts slowly. The equations below write this down.

Time is discrete. Let p_t be the price, p^* the fundamental value and $x_t = p_t - p^*$ the pricing error. A reduced-form dealer sector absorbs net order flow with price impact $\kappa > 0$, an imposed primitive rather than an equilibrium quote, and arbitrage pulls the error toward zero with persistence $a \in [0, 1)$:

$$x_t = a x_{t-1} + \kappa \sum_{i=1}^{K} q_i y_{i,t-1} + \kappa v_t + w_t, \tag{1}$$

where v_t is an executed probe quantity and w_t a price innovation. Cohort i holds a state $y_{i,t}$ with exponential memory and trades $q_i y_{i,t}$:

$$y_{i,t} = (1 - \eta_i)\, y_{i,t-1} + \eta_i x_t, \qquad \eta_i \in (0, 1). \tag{2}$$

The ordering is *price first*: the price responds to the probe, then cohorts update. Write $r_i = 1 - \eta_i$ for the memory rate and $w_i = q_i \eta_i$ for the merged weight. The feedback kernel and its transform are

$$k_m = \sum_i w_i r_i^m, \qquad G(z) = \sum_i \frac{w_i z}{z - r_i}, \tag{3}$$

and the transfer function from probes to prices, $H(z) = \sum_{j \geq 0} h_j z^{-j}$, is

$$H(z) = \frac{\kappa z}{z - a - \kappa G(z)} = \frac{\kappa z D(z)}{(z - a) D(z) - \kappa N(z)}, \qquad G = \frac{N}{D}, \quad D(z) = \prod_i (z - r_i), \tag{4}$$

with impulse response $h_0 = H(\infty) = \kappa, h_1, h_2, \ldots$; in particular $H(1) = \kappa/(1 - S) = \sum_j h_j$. Stacking $z_t = (x_t, y_{1,t}, \ldots, y_{K,t})^\top$ gives $z_t = J z_{t-1} + b v_t$, where, because cohorts update on the same period's price, the probe enters every state on impact through

$$b = \kappa\, (1, \eta_1, \ldots, \eta_K)^\top, \tag{5}$$

and

$$J = \begin{pmatrix} a & \kappa q_1 & \cdots & \kappa q_K \\ \eta_1 a & 1-\eta_1+\eta_1\kappa q_1 & \cdots & \eta_1\kappa q_K \\ \vdots & & \ddots & \vdots \\ \eta_K a & \eta_K\kappa q_1 & \cdots & 1-\eta_K+\eta_K\kappa q_K \end{pmatrix}, \qquad S \equiv a + \kappa \sum_i q_i. \tag{6}$$

In the *positive economy* ($q_i \geq 0$), J is entrywise nonnegative, its price row sums to S, and the economy is Schur stable iff $S < 1$.

Lemma 2.1 (Two motivating examples: filtering and learning)**.** *(i)* Kalman memory. *A desk that tracks a latent random-walk trend $v_{t+1} = v_t + \zeta_{t+1}$, $\zeta \sim N(0, \sigma_\zeta^2)$, from signals $r_t = v_t + \nu_t$, $\nu \sim N(0, \sigma_\nu^2)$, using the steady-state Kalman filter, updates as $\hat{v}_{t+1} = (1-\eta)\hat{v}_t + \eta r_t$ with*

$$\eta = \frac{-\sigma_\zeta^2 + \sqrt{\sigma_\zeta^4 + 4\sigma_\zeta^2\sigma_\nu^2}}{2\sigma_\nu^2} \in (0,1). \tag{7}$$

(ii) Gradient memory. *Online gradient descent with L_2 regularization $\gamma \in (0,1)$, step size $\alpha > 0$ and loss $\frac{1}{2}H(x_t - \theta_t)^2$ gives $\theta_{t+1} = (1-\eta)\theta_t + \eta q x_t$ with $\eta = \gamma + \alpha H \in (0,1)$ and $q = \alpha H/\eta$. A tabular value update with learning rate η at a fixed state has the same form; general reinforcement-learning policies need not.*

Remark 2.2 (Why exponential memory). Lemma 2.1 shows that two different design philosophies, Bayesian filtering and regularized online learning, produce the same observable reduced-form belief state $y_{i,t}$. When desk i chooses order flow $u_{i,t}$ subject to $I_{i,t} = (1-\delta_{\text{inv}})I_{i,t-1} + u_{i,t}$ and an inventory cap $|I_{i,t}| \leq \bar{I}_i$ to maximize expected continuation return $\beta_i y_{i,t} u_{i,t}$ net of quadratic execution cost $\frac{1}{2}\lambda_{\text{tc},i}u_{i,t}^2$ and quadratic inventory holding cost $\frac{1}{2}\gamma_{\text{inv},i}I_{i,t}^2$, its interior first-order condition gives $u_{i,t}^* = q_i y_{i,t} - \psi_{\text{inv},i} I_{i,t-1}$ with market-response weight $q_i = \beta_i/(\lambda_{\text{tc},i} + \gamma_{\text{inv},i})$ and inventory mean-reversion $\psi_{\text{inv},i} = \gamma_{\text{inv},i}(1-\delta_{\text{inv}})/(\lambda_{\text{tc},i} + \gamma_{\text{inv},i})$, which shrinks toward zero as trading costs or inventory penalties tighten and destabilizes the market when the closed-loop spectral radius reaches

one (`microfoundation_and_baselines.py`). This filtering-and-control recursion specifies an individual agent's decision rule, not a complete economic equilibrium (dealer–trader equilibrium is analysed in the companion paper).

Two further simplifications matter for reading the results: the observer is taken to know the fundamental value p^*, so the pricing error x_t is observed, and the dealer sector is a fixed linear impact κ rather than an equilibrium quote. The paragraph before Section 8.4 states what the data tests lose when neither holds. Assuming that the observer knows p^* gives it more information than any real observer has, so the lower bounds hold a fortiori when p^* is unknown; Corollary 6.6 proves that baseline demeaning eliminates unknown p^* without affecting the upper-bound rate.

The model applies these rules to the pricing error x_t: a desk that cannot separate fundamental news from transitory deviations treats recent deviations as a signal of a persistent trend, which is the exponential extrapolation of Barberis et al. (2015, 2018). With $q_i > 0$ such a desk buys after rises, as trend followers do, and nothing in the model says this is profitable; the lemma shows only that the update has the exponential form. In these two illustrative families of algorithms the memory rate is a known function of a tuning parameter, the desk's signal-to-noise ratio or its regularization. That is the kind of private information a supervisor would like to recover, although real trading algorithms need not belong to either family. Section 3 shows why prices reveal it only in merged form.

Lemma 2.3 (Positive price-response geometry)**.** *In the positive, stable economy the impulse response is a nonnegative mixture of decaying exponentials,*

$$h_j = \sum_{\ell} \omega_\ell \lambda_\ell^j, \qquad \omega_\ell \geq 0, \quad \sum_{\ell} \omega_\ell = \kappa, \quad 0 \leq \lambda_\ell < 1.$$

It is completely monotone and nonincreasing; its peak is the direct impact κ, even though J is non-normal. In price/memory coordinates the price-only gain is $\sup_t \|e_x^\top J^t\|_2 = 1$.

As an illustration of the lemma (not part of its proof), $(-1)^n \Delta^n h_j \geq 0$ for $n \leq 4$ held

in all 892 stable random positive economies tested. Lemma 2.3 matters twice: it makes probing costs positive (Section 4), and it shows that in this economy transient amplification of the price response cannot arise from cohort feedback alone. The companion paper locates amplification in order-book depth instead.

2.1 Discussion of the assumptions

Linearity. The model is linear in prices, flows and memory states. Linearity is what makes the merged-measure equivalence exact and the testing problems Gaussian. Nonlinear strategies (for example threshold rules or position limits) can break the equivalence class, but they do so in regions of the state space that a small probe does not visit; for small probes the linearization around the steady state is the relevant object.

Price-first timing. We order the price update before the cohort update. Under the alternative lagged timing the transition matrix becomes $\left(\begin{smallmatrix} a & 1 \\ \eta Q & 1-\eta \end{smallmatrix}\right)$ and its eigenvalues differ, so stability statements must be computed under the timing actually used. The identification results depend on the transfer function and survive either convention.

Exponential memory. Lemma 2.1 justifies exponential memory for filtering and regularized learning. Neural and reinforcement-learning policies are not exponential filters. Rules whose orders are smooth functions of an exponential-memory state are exponential filters to first order, so the results hold for small probes; Section 7 checks this in stylized simulations and finds where it fails, and how closely production algorithms follow such rules remains an empirical question we do not settle. Mixtures of exponentials can approximate any completely monotone kernel, so a power-law memory corresponds to a continuum of cohorts, for which Theorem 6.9 shows that accuracy improves only logarithmically in the data.

Algorithms that react to being probed. The probing results assume that the cohorts keep their memories while they are probed. An algorithm that detects a probe could pause,

shorten its memory or trade against the probe. Four observations limit what this does to our conclusions. First, a reaction that all cohorts share does not help separate them, and an algorithm that mimics another group's memory rate only brings the hypotheses closer, so neither lowers the costs we derive. A reaction that differs across cohorts is a new signature, but one the observer must specify and estimate, so it shifts the identification problem rather than solving it. Second, the passive route needs no probe: Theorem 6.11 shows that natural price variation faces a lower bound of the same polynomial order, so an observer who avoids detection by not trading does not escape the barrier; it pays in time instead of money. We do not show that passive data attain that order. Third, detection is itself a testing problem against the same noise: the generalized-eigenvector probe of Section 5 trades at most 1.25 par units in any period, and a probe sized within normal flow variation is hard to tell from noise, so the threat bites mainly for large single pulses, which our ledger shows are not optimal anyway. Fourth, what we cannot rule out is an algorithm designed to defeat surveillance, which turns identification into a game between observer and observed; our results are the benchmark against which such a game would be measured, and we leave it to future work.

Known primitives. The resource theorems take κ, a and the noise law as known. Unknown primitives can only make separation harder, so the lower bounds remain valid. Assuming known primitives simplifies the baseline statements and strengthens the lower bounds; Corollary 6.6 (and its Lean 4 formalization) proves that estimating (p^*, κ, a) at the parametric $O_P(N^{-1/2})$ rate and plugging them into the closed-loop deconvolution leaves the $\Theta(\sigma^2(\Delta\eta)^{-(4K-2)})$ upper bound unchanged. In an equilibrium in which dealers set the impact κ from the order flow they see, as in Kyle (1985), a prober's unexpected trades would also move κ; we do not model this, and expect it to raise the cost of probing rather than lower it.

Observation contracts. Each result names what the observer sees. A price-only observer, an observer with executed quantities, one with noisy statements, one with audited equity

and one with tagged flow face different experiments. The paper's central message is that these differences, not the observer's statistical skill, determine what can be learned.

2.2 The observer, its options and its costs

The results are stated from the side of one observer outside the trader groups: a supervisor, an exchange's surveillance unit or a large desk. Its *target* is attribution, recovering the merged weight and memory rate of each group to a stated accuracy, or deciding between two candidate populations at a stated error. Before it observes, it knows the price persistence a, the impact κ and the law of the observation noise, and, for the K-group results, the number of groups; it does not know their rates or weights. These are the conditions under which the lower bounds are statements about the data rather than about the observer's ignorance; unknown primitives can only enlarge the set of candidate populations.

The observer has three options, and each result prices one of them.

(a) *Wait.* Record stationary prices driven by natural innovations, without trading. The cost is time: the number of observations needed (Theorems 6.10 and 6.11).

(b) *Probe.* Execute a sequence v_t of trades against the groups and record the price response. Probes execute at the average price of a linear book after the uncontrolled state update; the prober bears the execution loss $\frac{1}{2}v^\top T_h v$ plus fees (Lemma 4.2) and must fund margin throughout. The counted resource is the number of independent experiments, each starting from the reset state, or the executed quantity and capital of a single experiment (Theorems 4.1–5.2).

(c) *Acquire disclosure.* Receive each class's flow or position, observed with noise, at an operating cost per supervisory cycle and a compliance cost per market (Proposition 8.1, Theorem 8.2, Proposition 8.4).

Within each period the order is: the probe (if any) executes and the price updates, the groups update their states on the new price, and the observer records the price with noise.

The groups do not change their rules in response to being observed; Section 2.1 discusses what happens if they do. Probing experiments start from the reset state. A probe sent into a running market adds the free response of the current state, which by linearity superposes with the forced response; if the state is unknown it acts as a nuisance term, which can only make separation harder, so the lower bounds still hold, while the upper bounds as stated need the reset state or a known one. Independent resets are a benchmark experiment that no outside participant can run in a live market; we use them because they make the information content of the data exact, and the economic comparisons say whenever they rely on them. A numerical check without resets, in which a known random probe is sent into one long running market whose groups' state is unknown, keeps the $(\Delta\eta)^{-6}$ order for $K = 2$ and raises the constant by a factor of about 1.4 to 6 (Appendix C.8). Regulators do not, in practice, trade to learn; the probing option describes an outside participant, or prices the information a supervisor would otherwise have to buy by trading.

Three money quantities appear and are kept apart: the *capital required*, the equity an account must hold so that its margin constraint is never breached while it probes (Theorem 4.3 and Lemma 4.5); the *round-trip trading loss*, the execution loss of a probe that ends flat, which is what the account actually spends (Theorems 5.2, 5.3 and 5.4); and the *annualised loss* of renewing that probe weekly, used for the comparison with audit-trail budgets (Section 8.2). The lower bounds hold for the stated probe classes, deterministic probes and, in Theorem 5.3, non-anticipating sequential probes; they are not lower bounds on information obtained through other channels, which is the point of Section 8. Two rates recur and measure different things: δ^{-4} is the capital (and the average-power horizon) for a *known* symmetric pair whose rates differ by 2δ, while $(\Delta\eta)^{-(4K-2)}$, which is $(\Delta\eta)^{-6}$ for $K = 2$, is the number of experiments for K *unknown* groups with spacing $\Delta\eta$.

3 The identification boundary

A central question in algorithmic market surveillance is whether an observer of executions and prices can identify the individual strategies active in the market.

Theorem 3.1 (Merged rate–weight equivalence)**.** *Let $\kappa > 0$ and $a \in [0, 1)$ be known. Use zero initial states and a common probe contract, and observe the modeled price state x_t.*

(i) Equivalence class. *Two economies $\mathcal{E}^A = \{(q_i^A, \eta_i^A)\}$ and $\mathcal{E}^B = \{(q_j^B, \eta_j^B)\}$ produce identical price responses for every probe sequence if and only if their merged measures coincide:*

$$\sum_{i:\eta_i^A=\eta} q_i^A \eta \;=\; \sum_{j:\eta_j^B=\eta} q_j^B \eta \qquad \forall \eta \in (0, 1]. \tag{8}$$

Write $a' = a + \kappa M(\{0\})$, so that cohorts with $\eta = 1$ (rate 0) act as part of the anchor. For $a' \neq 0$ the minimal realization has dimension $1 + |\mathrm{supp}(M) \cap (0, 1)|$, counting distinct positive rates with nonzero merged weight after signed cancellation; for $a' = 0$ the price carries no state of its own and the dimension is $|\mathrm{supp}(M) \cap (0, 1)|$. A cohort with $\eta = 1$ merges into the anchor a without adding a pole.

(ii) Separation complexity under average power. *Consider testing a single cohort with rate $\bar{\eta}$ against two cohorts with rates $\bar{\eta} \pm \delta$ and the same total weight, on a compact family that is uniformly Schur stable with all rates in $[\epsilon, 1 - \epsilon]$ for some $\epsilon > 0$. Observe $\tilde{x}_t = x_t + \nu_t$ with $\nu_t \sim N(0, \sigma_\nu^2)$ and use deterministic probes with $\frac{1}{T}\sum_{t<T} v_t^2 \leq P$. Any test with worst-case error at most $\alpha \in (0, 1/4)$ needs*

$$T \;\geq\; \frac{4\sigma_\nu^2(1 - 2\alpha)^2}{C_{\mathcal{F}}^2 P\, \delta^4} = \Omega(\delta^{-4}), \tag{9}$$

where $C_{\mathcal{F}}$ depends only on the family, through ϵ, the bound on $|H|$ implied by uniform stability, and the total weight q.

In words. If two trading desks update on prices with the same memory, no amount of price data can tell them apart or say how much each trades: prices reveal only how much trading happens at each memory speed. Desks with different memories can be told apart in principle, and the rest of the paper prices how hard that is.

The memory rates are poles of G and *zeros* of H, not closed-loop poles. This has a practical consequence we verify in Section 6: pole-fitting methods applied directly to the price impulse response recover closed-loop eigenvalues (0.4256 and 0.8224 for $K = 1$, $\eta = 0.3$, $a = 0.5$, $\kappa = 0.8$, $q = 0.2$) and miss the memory rate 0.70.

Remark 3.2 (Scope). The result concerns the stated model and observation contract. Cohorts sharing an exact rate enter prices through their merged weight; this does not determine identities or wealth allocations. Nearby rates can be distinguishable with enough information. The δ^{-4} statement in (ii) is a lower bound on a declared family, with family-dependent constants. Individual algorithms may choose their rates strategically, but any unilateral change maps into the same observer-facing merged measure, which reinforces rather than circumvents the boundary.

Proposition 3.3 (Endogenous dealer pricing and equilibrium invariance)**.** *Let the price-impact parameter κ be determined endogenously by a competitive or risk-averse dealer who sets prices given observed aggregate anonymous order flow $y_t = \sum_{k=1}^{K} u_{k,t} + v_t + \xi_t$, where $\xi_t \sim$ i.i.d. $N(0, \sigma_\xi^2)$ is uninformative liquidity demand.*

(i) Equilibrium invariance on equivalence classes. *Any equilibrium pricing functional $\kappa = \mathcal{K}(\Gamma_y)$ that depends on the economy only through the aggregate observable order-flow second moments $\Gamma_y(k) = \mathrm{Cov}(y_t, y_{t-k})$ satisfies*

$$\kappa^*(\mathcal{E}^A) = \kappa^*(\mathcal{E}^B) \tag{10}$$

for all economies $\mathcal{E}^A, \mathcal{E}^B$ sharing the same merged rate–weight measure M. Observational equivalence in Theorem 3.1(i) is therefore strictly invariant to the dealer's

endogenous pricing response.

(ii) Moment-matched stability and sample complexity. *For any two K-cohort economies $\mathcal{E}^A, \mathcal{E}^B$ whose first $2K-1$ moments match ($|m_j(\mathcal{E}^A) - m_j(\mathcal{E}^B)| = O((\Delta\eta)^{2K-1})$), the equilibrium price impacts satisfy*

$$|\kappa^*(\mathcal{E}^A) - \kappa^*(\mathcal{E}^B)| = O((\Delta\eta)^{2K-1}). \tag{11}$$

Consequently, endogenizing dealer pricing does not relax the lower bound: distinguishing $\mathcal{E}^A$ from $\mathcal{E}^B$ from price and aggregate flow observations continues to require sample complexity $\Omega((\Delta\eta)^{-(4K-2)})$.

Defeating the Lucas critique. A standard objection in market microstructure is the Lucas critique: if a supervisor attempts surveillance or intervention, an endogenous dealer might alter their pricing slope κ, generating identifiable price-path divergences. Proposition 3.3 proves that endogenous dealer pricing cannot circumvent the identification boundary. Because a competitive or risk-averse dealer observes only anonymous aggregate net order flow, the dealer's filtering and quote-setting problem depends strictly on the autocovariance sequence of aggregate flow. Since observationally equivalent economies generate identical aggregate flow autocovariances, the dealer faces the exact same informational blindness as the supervisor, producing identical equilibrium pricing coefficients κ^*. For near-colliding cohorts, any equilibrium adjustment in κ is confined to the infinitesimal order $O((\Delta\eta)^{2K-1})$, preserving the super-resolution sample complexity barrier $\Omega((\Delta\eta)^{-(4K-2)})$. Full proof is in Supplement S1.

4 Probing resources, memory resolution and money

Observation length alone cannot buy identification when the lifetime quantity of controlled executions is fixed. This section connects the merged-memory experiment to a financially

specified input/output contract. Throughout, Φ is the standard normal distribution function.

4.1 Quantity

Theorem 4.1 (Sharp probing-resource orders for a known symmetric pair)**.** *Fix* $0 \le a < 1$, $\kappa, q > 0$, $S = a + \kappa q < 1$ *and* $\eta \in (0,1)$. *Compare a single cohort of weight* q *and rate* η *with two cohorts of weight* $q/2$ *and rates* $\eta \pm \delta$, *both in* $[\epsilon, 1-\epsilon]$. *All parameters are known. The price receives executed quantity through* κv_t, *starts from zero, has no process noise, and is observed with independent* $N(0,\sigma^2)$ *noise. Probes are deterministic and common to both alternatives. Let* $d_T = (h_j^\delta - h_j^0)_{j<T}$ *and* $K_* = 2\kappa^2 q/[\epsilon^3(1-S)^2]$.

(i) For every probe with $\sum_t |v_t| \le V$,

$$D_{\mathrm{KL}}(P_0^v \| P_\delta^v) \le \frac{K_*^2 V^2 \delta^4}{2\sigma^2}, \qquad e_{\mathrm{worst}} \ge \frac{1}{2}\Big(1 - \frac{K_* V \delta^2}{2\sigma}\Big)_+. \tag{12}$$

(ii) The optimal worst-case error over these probes and all classifiers is exactly

$$e_T^*(V) = \Phi\left(-\frac{V\|d_T\|_2}{2\sigma}\right), \tag{13}$$

attained by a pulse at the first input and the likelihood-ratio test. Waiting adds information but $\|d_T\|_2 \uparrow \|d_\infty\|_2 < \infty$, *so it cannot make the error arbitrarily small at fixed* V. *Error at most* $\alpha < 1/2$ *requires* $V \ge 2\sigma(1-2\alpha)/(K_*\delta^2)$ *independently of* T.

(iii) The first two taps coincide and $d_2 = -\kappa^2 q \delta^2$. *A single pulse of size*

$$A = \frac{2\sigma\Phi^{-1}(1-\alpha)}{\kappa^2 q \delta^2} \tag{14}$$

and the third observation achieve error exactly α. *The necessary and sufficient quantity order is* δ^{-2}.

(iv) Under average power $\sum_{t<T} v_t^2 \le PT$, *the necessary horizon order is* δ^{-4}, *attained by*

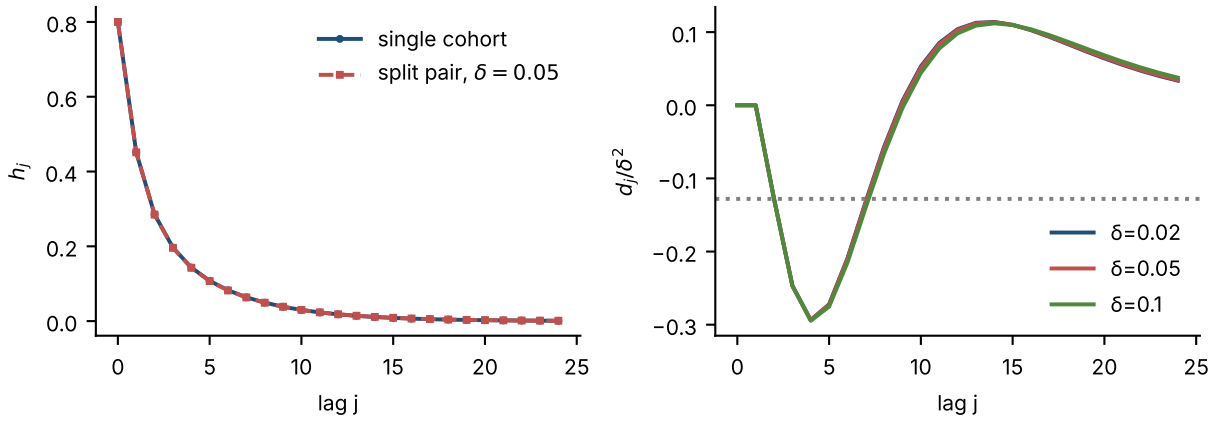


Figure 2: Left: impulse responses of the single cohort and the split pair ($\delta = 0.05$) are nearly indistinguishable. Right: the difference scaled by δ^2 collapses onto one curve, with first nonzero value $-\kappa^2 q$ at lag 2 (dotted). Calibration $a = 0.5$, $\kappa = 0.8$, $q = 0.2$, $\eta = 0.4$.

the pulse whenever $T \geq \max\{3, \lceil A^2/P \rceil\}$.

(v) For a fixed monetary price exposure $\lambda \neq 0$*, additive reserve* L *and cost envelope* $c \sum |v_t|$*, the sharp robust quantity allowance under budget* $B > L$ *is* $V_B = (B - L)/(|\lambda|\kappa + c)$.

In words. To tell one desk from two desks whose memories differ slightly, the observer must trade against them, and the quantity it must trade grows like $1/\delta^2$: halving the difference quadruples the trading needed. Waiting longer after a fixed trade does not help, because the trade's echo in the price dies out.

The identity in (iii) follows by direct substitution, and code checks it at $\delta \in \{0.02, 0.05, 0.10\}$: the difference is -1.2×10^{-4}, -7.5×10^{-4} and -3.0×10^{-3} at $\kappa = 1$, $q = 0.3$, and -3.2×10^{-4} at the calibration $\kappa = 0.8$, $q = 0.2$, $\delta = 0.05$. The theorem does not solve estimation of an unknown mixture (Section 6).

A Monte Carlo check of the exact error. Formula (13) is exact for the likelihood-ratio test. We simulated 200,000 noisy paths per hypothesis at the calibration ($\sigma = 0.01$, $\delta = 0.05$) for several pulse sizes and horizons (Table 6); simulated errors match the formula within Monte Carlo standard error under both hypotheses.

Figure 2 shows why the problem is hard. The two economies have the same direct impact and the same first-lag response; they first differ at lag two, by an amount proportional to δ^2, and the whole difference sequence scales with δ^2. A test therefore needs the signal-to-noise ratio $V\|d\|/\sigma$ to grow like δ^{-2}, which is the quantity order, and because the adverse mark of a probe grows with the square of its size, capital grows like δ^{-4}.

4.2 Execution cost

A probe changes inventory as well as price. We therefore specify cash and execution timing. Price has units of currency per asset, and κ currency per squared quantity.

Lemma 4.2 (Fair execution)**.** *For a quote schedule refreshed after the uncontrolled state update, execution at the average price of a linear book costs, relative to the constant fundamental mark,*

$$C(v) = \tfrac{1}{2} v^\top T_h v + c\|v\|_1, \qquad (T_h)_{ij} = h_{|i-j|}. \tag{15}$$

In the positive stable economy T_h is positive definite, so no nonzero round trip produces a deterministic gain.

We verify the ledger identity (15) against an explicit cash account to 5.7×10^{-13} on 1,000 random probes, and the smallest eigenvalue of T_h over 1,000 random stable positive economies is 0.141. A direct-only proxy $\sum \frac{1}{2}\kappa v_t^2$ omits cross-time impact: for $v = (+1, -1)$ it gives 0.80 against a true Toeplitz loss of $h_0 - h_1 = 0.3488$.

4.3 Capital supplied by a counterparty

Theorem 4.3 (Sharp robust quantity cap from cash and margin)**.** *Use the positive, stable, reset economy with constant fundamental price $P_* > 0$ and no process noise. A designated dealer fills only nonnegative probes $v_t \geq 0$ at the average price $P^0_{t+1} + \kappa v_t/2$ of a linear schedule, holds inventory $I_{t+1} = I_t - v_t$ and cash $K_{t+1} = K_t + v_t(P^0_{t+1} + \kappa v_t/2) - c v_t$, starts*

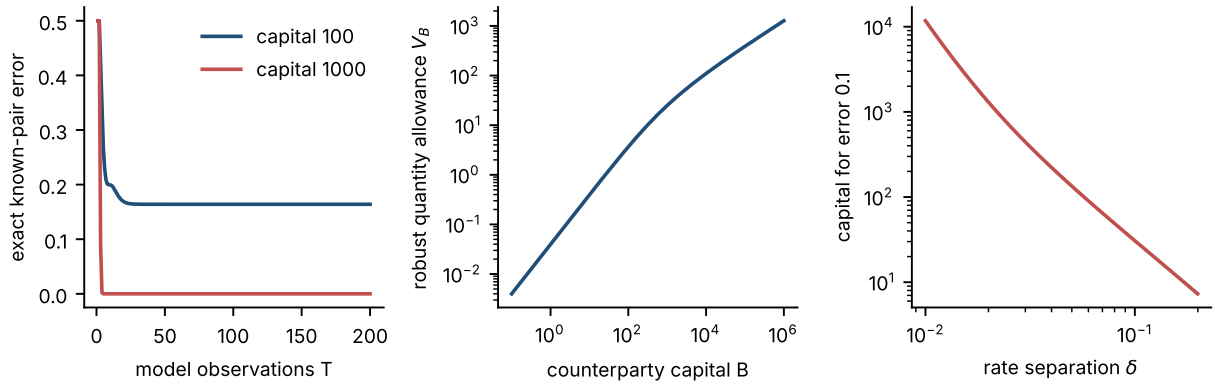


Figure 3: Known-pair experiment funded by the counterparty cash/margin contract. (a) Exact Gaussian error at fixed capital 100 and 1,000, separation $\delta = 0.05$. (b) The sharp robust quantity allowance (16). (c) Capital for error 0.1 with 200 observations. Parameters: $a = 0.5$, $\kappa = 0.8$, $q = 0.2$, $\eta = 0.4$, observation s.d. 0.003, reference price 100, margin 0.25, cost 0.02. No process noise.

with equity $B > 0$, and must satisfy $\mu|I_t|P_t \le E_t$ at every time. For the robust contract admitting every probe with $\sum_t v_t \le V$, the largest allowance guaranteeing margin is

$$V_B = \frac{2B}{\beta + \sqrt{\beta^2 + 4(\mu + \frac{1}{2})\kappa B}}, \qquad \beta = c + \mu P_*, \tag{16}$$

equivalently $(\mu+\frac{1}{2})\kappa V^2 + (c+\mu P_*)V \le B$. The cap holds across all positive cohort representations with $S < 1$, without resolving individual rates. An observer that receives only noisy prices and executed quantities has best error $\Phi[-V_B\|d_T\|_2/(2\sigma)]$ within this allowance; fixed error $\alpha < 1/2$ requires capital of order at least δ^{-4}, and a single positive pulse attains that order.

In words. A probe moves the price against the account that makes it, so the margin it needs grows with the square of the probe. Because the probe grows like $1/\delta^2$, the capital grows like $1/\delta^4$: halving the difference multiplies the capital needed by 16.

Even with zero initial inventory, the probe creates a quadratic adverse mark that a fixed initial price exposure would miss. Figure 3 illustrates the experiment.

Proposition 4.4 (Exact financial disclosure changes the experiment)**.** *In this noiseless-state*

account, exact disclosure of the third equity mark after a single positive pulse $0 < A \le V_B$ *distinguishes the known pair with zero error for every* $\delta > 0$.

In words. If the observer can see an account's exact equity, a single test trade settles the question, however close the two candidate markets are. The hard part of the problem is therefore not the market itself but what the observer is allowed to see.

Exact marked equity is extra information, not a free observation assumed by the price-only theorem. Real account marks have their own precision: with reports rounded to 0.01, a guaranteed change requires a larger trade (Section 8).

4.4 Capital paid by the probe trader

Lemma 4.5 (Time-domain Toeplitz coercivity: capital order for funded probes)**.** *Use the same economy and execution law. The statistical observer receives only noisy prices and executed quantities. The probe trader starts with cash* $B > 0$ *and zero inventory; a finite signed sequence* v *changes inventory by* $+v_t$ *and cash by* $K_{t+1} = K_t - v_t(P^0_{t+1} + \kappa v_t/2) - c|v_t|$, *after which trading stops and the trader must remain solvent as the price returns to* P_*. *Put*

$$r = 1 - \epsilon(1 - S) < 1, \qquad m = \kappa \frac{1 - r}{1 + r} > 0, \qquad K_* = \frac{2\kappa^2 q}{\epsilon^3 (1 - S)^2}.$$

Every solvent strategy satisfies $D_{\mathrm{KL}}(P^v_0 \| P^v_\delta) \le K_*^2 B \delta^4/(\sigma^2 m)$, *so error at most* $\alpha < 1/2$ *requires*

$$B \;\ge\; \frac{2\sigma^2 m (1 - 2\alpha)^2}{K_*^2 \delta^4}, \tag{17}$$

without a quantity cap and at any horizon. Conversely, for any margin fraction $0 < \mu \le 1$, *the pulse* A *in* (14) *achieves error* α *and remains margin-solvent at every later mark with capital* $B_A = \frac{\kappa}{2} A^2 + (c + \mu P_*) A$. *Necessary and sufficient capital orders are* δ^{-4}.

Lemma 4.5 is an order statement: an elementary time-domain coercivity argument for Toeplitz forms that establishes the δ^{-4} rate before the exact frequency-domain optimization

of Section 5. Its explicit constant is very conservative, because it scales as $m/K_*^2 \propto \epsilon^7(1-S)^5$: at the calibrated parameters it is 4.9×10^{-4} par units at $\delta = 0.05$ with $\epsilon = 0.35$, five orders of magnitude below any feasible design. We therefore state it as a lemma; Section 5 computes the sharp constant, within 5% of a feasible design at every separation.

4.5 A worked example: two Kalman desks

To make the orders concrete, consider two desks that track the same trend with Kalman filters (Lemma 2.1(i)) and differ only in their signal-to-noise ratios: $\sigma_\zeta^2/\sigma_\nu^2 = 0.09$ and 0.16. Their gains are $\eta_1 = 0.2584$ and $\eta_2 = 0.3279$, so $\bar\eta = 0.2931$ and $\delta = 0.0348$. With equal weights, the economy is observationally the split pair of Theorem 4.1 around $\bar\eta$, and the response difference first appears at lag two, $d_2 = -\kappa^2 q\delta^2 = -1.55\times10^{-4}$ (verified).

A supervisor who wants to tell this market from a single desk with gain $\bar\eta$ at 5% error using a single pulse and the third observation needs $A = 2\sigma\Phi^{-1}(0.95)/(\kappa^2 q\delta^2) = 212$ units of executed quantity, and, under a 0.5% margin and 4 basis points of fees, single-pulse capital of 18,166 par units. The pulse is very inefficient: the spectral lower bound of Theorem 5.2, evaluated at this pair ($\bar\eta = 0.2931$, $\delta = 0.0348$), gives about 77 par units for the optimal flat probe (rescaling the $\bar\eta = 0.4$ constant by δ^{-4} alone would give 262, because the constant also depends on $\bar\eta$). Converting with the measured Roll noise (factor $(0.40/0.01)^2$), the optimal probe costs roughly 1.2×10^5 USDT and the naive pulse about 3×10^7 USDT. Two desks whose signal-to-noise ratios differ by almost a factor of two are therefore expensive, but not impossible, to tell apart by probing; three or more desks are much harder (Theorem 6.4).

4.6 What prices hide, in economic terms

The worked example states the barrier in signal-to-noise ratios. That is the general picture. Write $s = \sigma_\zeta^2/\sigma_\nu^2$ for a desk's signal-to-noise ratio, so that by (7) its memory rate is $1-\eta(s)$

with

$$\eta(s) = \tfrac{1}{2}\Big(\sqrt{s^2+4s} - s\Big), \qquad \eta'(s) = \tfrac{1}{2}\Big(\frac{s+2}{\sqrt{s^2+4s}} - 1\Big) > 0,$$

and η is increasing and concave. With no inventory penalty, the first-order condition of the Remark on exponential memory gives the trading intensity $q_i = \beta_i/\lambda_{\mathrm{tc},i}$: what the desk expects to earn per unit of belief, over what it pays to trade.

Proposition 4.6 (Identification in economic primitives)**.** *Consider K Kalman desks (Lemma 2.1(i)) with signal-to-noise ratios $s_i \in [\underline{s}, \bar{s}] \subset (0, \infty)$ and intensities $q_i = \beta_i/\lambda_{\mathrm{tc},i}$, and write $c_- = \eta'(\bar{s})$ and $c_+ = \eta'(\underline{s})$.*

(i) Prices and aggregate order flow depend on the desks only through the measure

$$\sum_i q_i\, \eta(s_i)\, \delta_{1-\eta(s_i)}.$$

Desks are identified through their signal-to-noise ratios and intensities only; expected returns and trading costs enter only through their ratio, and two populations with the same intensity-weighted distribution of signal-to-noise ratios are observationally equivalent.

(ii) Fix a centre s_c and let the ratios lie in $s_c + \Delta s\,[-1, 1] \subset [\underline{s}, \bar{s}]$. Recovering the intensity-weighted distribution of signal-to-noise ratios to Wasserstein accuracy of order Δs, or its weights $q_i\eta(s_i)$, from N noisy reset experiments with worst-case error probability below $1/4$ requires $N \geq c\,\sigma^2(\Delta s)^{-(4K-2)}$; on separated configurations, Prony's method of Theorem 6.5 composed with η^{-1} achieves this order. The constants are those of Theorems 6.4 and 6.5 multiplied by powers of c_- and c_+.

In words. What prices reveal about Kalman-filter desks is the mix of how good their information is (their signal-to-noise ratios), weighted by how hard they trade. Expected returns and trading costs matter only through their ratio. Pinning down that mix carries the same steep cost as separating memory rates.

Proof. (i) By Lemma 2.1(i) desk i is a cohort with memory rate $1-\eta(s_i)$ and weight q_i, so Theorem 3.1 applies, and $\beta_i, \lambda_{\mathrm{tc},i}$ enter the economy only through q_i. (ii) On $[\underline{s}, \bar{s}]$, η is a C^1 bijection onto its image with $c_- \le \eta' \le c_+$. Lower bound: take the two economies of Theorem 6.4 with centre $r_c = 1-\eta(s_c)$ and rate spacing $\Delta\eta = c_-\Delta s$. By the mean value theorem their gains are images of ratios in $s_c + \Delta s\,[-1,1]$, so they are two populations of Kalman desks in the class, and the bound $N \ge c\,\sigma^2(\Delta\eta)^{-(4K-2)} = c\,c_-^{-(4K-2)}\sigma^2(\Delta s)^{-(4K-2)}$ applies to recovering their weights. Because η is c_+-Lipschitz, the Wasserstein distance between the two distributions of ratios is at least $1/c_+$ times that between the two distributions of rates, which Theorem 6.4 bounds below by a constant times $\Delta\eta = c_-\Delta s$. Upper bound: ratios in $s_c + \Delta s\,[-1,1]$ with spacing at least $\gamma\Delta s$ have gains in $\eta(s_c) + c_+\Delta s\,[-1,1]$ with spacing at least $c_-\gamma\Delta s$, a class of the form required by Theorem 6.5 with rescaled separation $\gamma c_-/c_+$; Prony's method recovers the rates and weights at that theorem's rate, and η^{-1}, which is $1/c_-$-Lipschitz, maps rate errors to ratio errors. □

The private information a supervisor would most like to know, how good each desk's information is, is exactly what prices hide. In the worked example the two desks' ratios differ by a factor of almost two, and separating them already costs hundreds of thousands of dollars of probing capital. The proposition takes intensities to vary independently of information quality, because they also reflect costs and capital. The next proposition asks what happens when the economics ties them together.

Proposition 4.7 (Intensity tied to information quality)**.** *Let each desk's intensity be a known function of its signal-to-noise ratio, $q = \varphi(s)$ with φ positive and C^2 on $[\underline{s}, \bar{s}]$ (for a mean–variance desk, $\varphi(s)$ increases with s), and let $n_i \ge 0$ denote the mass of desks with ratio s_i (a divisible population, not an integer desk count). Write $s(\eta)$ for the inverse gain map and $w(\eta) = \varphi(s(\eta))\,\eta$ for the weight per desk as a function of its gain.*

(i) Unknown numbers of desks. *The merged measure is $\sum_i n_i\, w(\eta(s_i))\,\delta_{1-\eta(s_i)}$. Every positive measure with finitely many atoms in $[1-\eta(\bar{s}), 1-\eta(\underline{s})]$ is of this form, with*

$n_i = m_i/w(\eta_i)$ for an atom of mass m_i at $1-\eta_i$. So when the numbers of desks of each type are not known, tying intensity to information quality removes no economy, and every lower bound of Proposition 4.6(ii) holds unchanged.

(ii) Known numbers of desks. *If the n_i are also known, the set of economies is smaller, but the single-versus-pair problem keeps its order: one type at gain $\bar\eta$ with $2n$ desks and two types at $\bar\eta \pm \delta$ with n desks each have taps that differ by $n\,\delta^2\,\partial_\eta^2[w(\eta)(1-\eta)^m]_{\bar\eta} + o(\delta^2)$, so separating them needs $N \geq c\,\sigma^2\delta^{-4}$ reset experiments, the order the symmetric pair has when intensities vary freely.*

(iii) Self-confirming configurations. *Suppose the ratio each desk perceives is the one implied by the closed-loop dynamics it trades in, through a map T from configurations $(s_1, \ldots, s_K)$ to perceived ratios that is continuous on a box $\prod_i [\underline{s}, \bar{s}]$ and maps it into itself. Then a configuration with $T(s) = s$ exists, and parts (i)–(ii) and Proposition 4.6 apply at it.*

Proof. (i) Desk type i is a cohort with gain $\eta_i = \eta(s_i)$, memory rate $1-\eta_i$ and total weight $n_i\varphi(s_i)$, so it contributes $n_i\varphi(s_i)\eta_i = n_i w(\eta_i)$ to the atom at $1-\eta_i$. Since $w > 0$, any atom mass $m_i \geq 0$ is attained with $n_i = m_i/w(\eta_i)$. (ii) Write $g_m(\eta) = w(\eta)(1-\eta)^m$; the tap difference is $n[g_m(\bar\eta+\delta) + g_m(\bar\eta-\delta) - 2g_m(\bar\eta)] = n\delta^2 g_m''(\bar\eta) + o(\delta^2)$, with $\sum_m \sup|g_m''|^2 < \infty$ on gains bounded away from zero, so the Euclidean distance between the two tap vectors is $O(n\delta^2)$ and the Kullback–Leibler divergence over N experiments is $O(Nn^2\delta^4/\sigma^2)$; Le Cam's two-point argument gives the bound. (iii) is Brouwer's fixed-point theorem. □

So the barrier does not depend on intensities varying freely: it survives any known link between intensity and information quality as long as the supervisor does not also know how many desks of each kind there are, and even then the simplest separation problem keeps its δ^{-4} order. Whether the full $(\Delta\eta)^{-(4K-2)}$ exponent survives with known numbers of desks is open. Part (iii) shows only that self-confirming configurations exist under continuity; it does not derive the map T from a model of the desks' learning.

4.7 Natural excitation and the cost of ambiguity

Natural price innovations can also separate the pair with zero controlled quantity: an exact finite-sample test on independent three-step blocks needs order δ^{-4} blocks (Internet Appendix, Proposition C.2), the same order as the passive Whittle bound of Section 6. Non-identification also has a price on the quote side: when the relevant set of economies includes costlier ones, a robust dealer widens its quote, which the companion paper analyses (Niroula et al., 2026); cohort non-identification alone does not establish that the worst relevant cost exceeds the informed one.

5 Account-level capital

Lemma 4.5 gives the order but not the constant. This section computes what a trader actually needs under a calibrated linear-impact environment, then converts it to Bitcoin dollars.

Numerical values. Decimal optima and brackets below are numerical evaluations of the stated bounds.

5.1 Observation design and the tri-region partition

An observer distinguishes two candidate mechanisms with rates θ_0 and $\theta_1 = \theta_0 + \delta$ over N periods, submitting probes $V = (V_1, \dots, V_N)^\top$ that generate inventory $I_t = \sum_{s \le t} V_s$. Under contract I_0 (active probing, Gaussian observation noise) $p_t = (h_i * V)_t + \epsilon_t$, $\epsilon_t \sim N(0, \sigma^2)$, so

$$D_{\mathrm{KL}}(P_0^{\otimes N} \| P_1^{\otimes N}) = \frac{1}{2\sigma^2} \sum_{t=1}^{N} \Big((h_1 * V)_t - (h_0 * V)_t \Big)^2. \tag{18}$$

The observer minimizes capital $C(N, V) = \sum_t (\lambda V_t^2 + \phi |V_t|) + \max_t \mu p_0 |I_t|$ subject to (i) separation $D_{\mathrm{KL}} \ge \frac{1}{2}(z_\alpha + z_\beta)^2$, (ii) caps $|V_t| \le V_{\mathrm{eff}}$ and $|I_t| \le I_{\max}$, and (iii) a solvency

budget $\max_t \mathrm{Loss}_t(V) \le B$. Here $V_{\mathrm{eff}} = \min(V_{\max}, D_{\max}/\kappa)$ combines resting depth and a price-displacement limit.

Proposition 5.1 (Tri-region partition and finite-session resolution floor)**.** *Under observation horizon T per pulse, (δ, B) space splits into three regions.*

(i) Feasible single pulse. *For $\delta \ge \delta_{\min}(T)$ and $B \ge C^*(\delta)$ one pulse within capacity suffices. For a simple rate shift ($\eta_1 = \eta_0 + \delta$) the response difference is $\Theta(\delta)$ and $C^* = \Theta(\delta^{-2})$; for symmetric splitting the first-order variation cancels, the difference is $\Theta(\delta^2)$ and $C^* = \Theta(\delta^{-4})$.*

(ii) Sequential bounded probing. *For $\delta < \delta_{\min}(T)$ the required quantity exceeds V_{eff}. Capped pulses each supply divergence $V_{\mathrm{eff}}^2 \|d_T\|^2/(2\sigma^2)$, so*

$$N_{\mathrm{pulses}}(T, \delta) = \left\lceil \frac{4\sigma^2 z_\alpha^2}{V_{\mathrm{eff}}^2 \|d_T\|_2^2} \right\rceil. \tag{19}$$

(iii) Infeasible with isolated pulses. *If the budget cannot finance sequential costs, peak inventory breaches $I_{\max}$, or a session deadline binds, the finite-session floor $\delta_{\min}(T_{\mathrm{session}})$ is the root of $V_{\mathrm{eff}} \|d_{T_{\mathrm{cycle}}}(\delta)\|_2 = 2\sigma z_\alpha / \sqrt{\lfloor T_{\mathrm{session}}/T_{\mathrm{cycle}} \rfloor}$, with small-$\delta$ approximation $\delta_{\min} = \Theta(T_{\mathrm{session}}^{-1/4})$.*

In words. A single test trade works only when the groups are far enough apart and the account has enough capital. Closer groups need many capped test trades, and a longer observation window after each trade makes every test trade more informative.

At the calibration $\kappa = 0.8$, $q = 0.2$, $\sigma = 0.01$, $V_{\max} = 50$, $D_{\max} = 5$ ($V_{\mathrm{eff}} = 6.25$), the single-tap floor is $\delta_{\min}(3) = \sqrt{2\sigma z_\alpha/(V_{\mathrm{eff}}\kappa^2 q)} = 0.2028$; extended observation lowers it to 0.1376 ($T = 4$), 0.0973 ($T = 10$), 0.0914 ($T = 20$) and 0.0905 ($T = 100$). At $\delta = 0.05$ the simplified one-leg benchmark (cost $\frac{1}{2}\kappa V_{\mathrm{eff}}^2 + \phi P_* V_{\mathrm{eff}} = 15.875$ per pulse plus initial margin 3.125) needs 271 pulses and 4,305.25 par units with $T = 3$, 58 pulses and 923.88 with $T = 4$, but only 12 pulses and 193.63 with $T = 20$ and 11 pulses and 177.75 with $T = 100$; optimizing

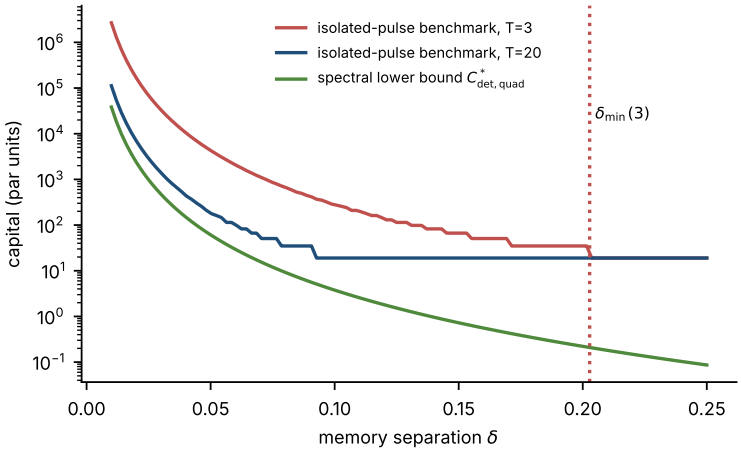


Figure 4: Capital against separation: isolated-pulse benchmarks with three and twenty observations per pulse, and the spectral lower bound of Theorem 5.2. The dotted line is the single-tap resolution floor $\delta_{\min}(3) = 0.2028$. Recomputed in this revision.

the pulse size (6.06 rather than 6.25) lowers the $T = 20$ figure to 182.03. Most of the saving comes from observing the transient response rather than a single tap (Figure 4).

These are costs within an isolated-pulse design class, not minimum capital; the next subsection replaces the formula by a ledger.

5.2 Account-level capital and the gap to the lower bound

We run each candidate design through an explicit synthetic ledger at reference price $P_* = 100$: fills at the average linear-book price, taker fees of 4 basis points, marks at the prevailing price, and 0.5% maintenance margin. Initial capital is $C_0 = \max_t(\text{margin}_t - \text{equity}_t)$, maximized over both hypotheses, and the test uses the full price path including unwind legs.

Table 7 gives three conclusions. First, twelve buys that are never unwound fail the 5% target once the test uses the joint price path (divergence $4.95 < 5.41$) and end with 75 units of inventory. Second, if every unwind leg is itself informative, five or six round trips meet the target with zero terminal inventory. Third, pulses are not optimal: a generalized-eigenvector probe, which maximizes divergence per unit of fee-free round-trip cost, needs about 65 par units with maximum trade 1.25 and maximum displacement 2.23, well inside limits.

For flat fee-paying policies $v \in \mathbb{R}^L$ with $\sum_t v_t = 0$, the capital problem is

$$\min_{v} \max_{i\in\{0,1\}} L_i(v) + \sum_t \phi P_* |v_t| + \max_t \mu P_* |I_t| \quad \text{s.t.} \quad \|(H_1 - H_0)v\|_2 \geq 2\sigma z_{1-\alpha}, \tag{20}$$

where $L_i(v) = \frac{1}{2} v^\top T_{h_i} v$ is the exact Toeplitz round-trip loss. Zero-padding embeds horizon L into $L+1$ at identical cost and information, so the minimum is non-increasing in L.

Theorem 5.2 (Spectral dissipation lower bound)**.** *Let $\hat{H}_i(z) = \sum_{j\geq 0} h_{i,j} z^{-j}$, $\Delta\hat{H} = \hat{H}_1 - \hat{H}_0$, and define the dissipation kernel*

$$g_i(\omega) = \operatorname{Re} \hat{H}_i(e^{i\omega}) - \kappa/2 > 0 \qquad \forall \omega \in [-\pi, \pi]. \tag{21}$$

(i) Toeplitz energy identity. *For any flat round trip, $L_i(v) = \frac{1}{2\pi} \int_{-\pi}^{\pi} g_i(\omega) |V(e^{i\omega})|^2 d\omega$.*

(ii) Lower bound. $\|(H_1 - H_0)v\|_2^2 \leq \sup_\omega \frac{|\Delta\hat{H}(e^{i\omega})|^2}{g_i(\omega)} L_i(v)$, *so every deterministic round trip with error at most α has*

$$L_i(v) \;\geq\; 4\sigma^2 z_{1-\alpha}^2 \inf_\omega \frac{g_i(\omega)}{|\Delta\hat{H}(e^{i\omega})|^2} \;\equiv\; C^*_{\text{det,quad}}(\delta). \tag{22}$$

(iii) Bracket. *At $\delta = 0.05$, $C^*_{\text{det,quad}} = 61.14$ at $\omega^* = 0.2526$; adding the unavoidable fee $\phi P_* \cdot 2\sigma z / \sup_\omega |\Delta\hat{H}|$ gives the numerical lower estimate $C_{\text{det}} \geq 61.41$. A feasible 240-step sequence attains 64.48, so the fee-paying optimum lies in* $[61.41, 64.48]$.

(iv) Process noise. *Let the price also receive the innovations $w_t \sim N(0, \sigma_w^2)$ of (1), independent of observation noise $\nu_t \sim N(0, \sigma_\nu^2)$, so that from the reset state the observer sees $\tilde{x}_t = (h_i * v)_t + \kappa^{-1}(h_i * w)_t + \nu_t$, $t < T$. For the known pair of Theorem 4.1, put $\sigma_{\text{eff}}^2 = \sigma_\nu^2 + \sigma_w^2/4$, $\bar{M} = 1/(1-S)$ and $C_w = 2\sigma_w^4 \bar{M}^2 K_*^2 (\sigma_\nu^2 + \sigma_w^2 \bar{M}^2)/(\kappa^2 \sigma_\nu^2 \sigma_{\text{eff}}^4)$. Every deterministic round trip whose test has errors α and β under the two hypotheses satisfies*

$$L_i(v) \;\geq\; 2\sigma_{\text{eff}}^2 \big[\operatorname{kl}(\alpha, 1-\beta) - C_w T \delta^4\big]_+ \inf_\omega \frac{g_i(\omega)}{|\Delta\hat{H}(e^{i\omega})|^2}, \tag{23}$$

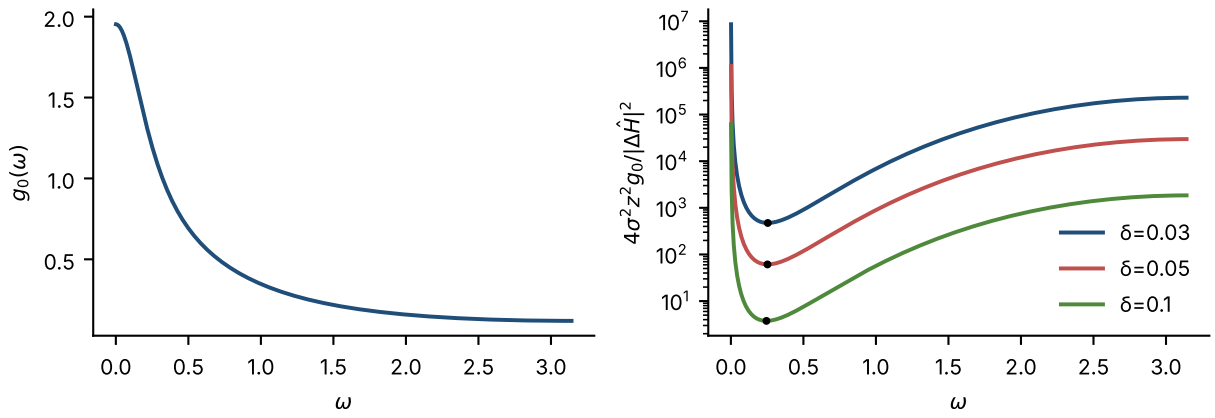


Figure 5: Left: the dissipation kernel $g_0(\omega)$ is strictly positive. Right: the capital ratio in (22); dots mark its minimum, near $\omega^* \approx 0.25$ for every separation. Recomputed in this revision.

with kl *as in Theorem 5.3. In-loop noise therefore leaves the* δ^{-4} *capital order intact for every horizon* $T \le \mathrm{kl}(\alpha, 1-\beta)/(2C_w\delta^4)$*; at horizons of order* δ^{-4} *and longer, the noise itself separates the pair without a probe (Theorem 6.11).*

In words. This turns the order into a dollar floor. No round trip that tells the two candidate markets apart at 5% error can lose less than this amount, and a concrete probe comes within 5% of it.

Positivity of g_i follows from Lemma 2.3: $g_i(\omega) = \frac{1}{2}\sum_\ell \omega_\ell \frac{1-\lambda_\ell^2}{1-2\lambda_\ell\cos\omega+\lambda_\ell^2} > 0$. We compute $C^*_{\mathrm{det,quad}}$ across separations (Table 8): 61.14 at $\delta = 0.05$, with $\delta^4 C^*_{\mathrm{det,quad}}$ between 3.67×10^{-4} and 3.84×10^{-4}, confirming the δ^{-4} law. The fee term, the feasible upper bounds and the adaptive bound of Theorem 5.3 are reported in the same table.

The infimum over ω is evaluated on a fine frequency grid; a local refinement agrees to four digits (61.1377 at $\omega^* = 0.2526$), and the reported constant is the larger of the two hypotheses' bounds. "Certified" here means a numerical evaluation of the stated bound, not an interval-arithmetic enclosure.

Theorem 5.3 (Adaptive and sequential probing)**.** *Let* $(v_t)_{t\le\tau}$ *be any non-anticipating policy adapted to the price filtration, with an almost surely finite stopping time and* $\sum_{t\le\tau} v_t = 0$*,*

followed by a terminal test $\psi_\tau \in \{0,1\}$ *with* $P_0(\psi_\tau = 1) \le \alpha$ *and* $P_1(\psi_\tau = 0) \le \beta$*. Then*

$$\mathbb{E}_0[L_0(v)] \;\ge\; 2\sigma^2 \,\mathrm{kl}(\alpha, 1-\beta) \inf_{\omega} \frac{g_0(\omega)}{|\Delta \hat{H}(e^{i\omega})|^2}, \tag{24}$$

where $\mathrm{kl}(\alpha, 1-\beta) = \alpha \log \frac{\alpha}{1-\beta} + (1-\alpha) \log \frac{1-\alpha}{\beta}$*. Hence adaptation can lower expected fee-free loss by at most the factor* $\mathrm{kl}(\alpha, 1-\beta)/(2z_{1-\alpha}^2) \approx 0.4897$ *at* $\alpha = \beta = 0.05$*, and cannot alter the* δ^{-4} *order.*

In words. Letting the observer adapt its test trades as data arrive, and stop as soon as it is confident, helps only a little: at the calibration it can cut the expected loss by at most about half. The steep cost is a property of the market, not of a clumsy probing plan.

At $\delta = 0.05$ the adaptive bound is 29.94, or 30.22 with the minimal fee. Pairing it with the feasible deterministic 64.48 gives a numerical ratio of at least 0.464. In 20,000 simulated sequential probability-ratio tests, which split the optimal probe into ten copies of equal information, mean loss was 39.5 under H_0, between the two bounds, but losses vary with the stopping time: the 95th and 99th percentiles of required account capital were 91.7 and 128.4, and funding a trader with only the mean loss gives a 37.5% insolvency rate before the test ends. Expected-loss savings are not solvency savings.

Theorem 5.4 (Finite-horizon Rayleigh–Hölder duality)**.** *For a horizon of* $L \ge 4$ *trading periods, let* $E \in \mathbb{R}^{L \times (L-1)}$ *be the first-difference basis of flat probes (*$\sum_t v_t = 0$*),* $Q_i = \frac{1}{2} T_{h_i}$ *(so that* $L_i(v) = v^\top Q_i v$*),* $A_{\mathrm{diff}} = (H_1 - H_0)^\top (H_1 - H_0)$*, and let* $G_i = E^\top Q_i E \succ 0$ *and* $M = E^\top A_{\mathrm{diff}} E \succeq 0$ *be the reduced execution-loss and information matrices, with* G_i*-orthonormal generalized eigenvectors* e_j *and eigenvalues* $0 \le \lambda_1 \le \cdots \le \lambda_{L-3} < \lambda_{L-2} \le \lambda_{L-1}$*; put* $c_{\mathrm{req}} = 4\sigma^2 z_{1-\alpha}^2$ *and* $c_{\mathrm{fee}} = \phi P_*$*. For* $\theta \in [0, \pi)$ *let* $d(\theta) = \cos\theta \, e_{L-2} + \sin\theta \, e_{L-1}$*,* $\lambda(\theta) =$

$\lambda_{L-2}\cos^2\theta + \lambda_{L-1}\sin^2\theta$, $s_\theta = \mathrm{sign}(Ed(\theta))$, *and*

$$F_{\mathrm{dom}}(\theta) = c_{\mathrm{fee}}\sqrt{\frac{c_{\mathrm{req}}}{\lambda(\theta)}}\,\|Ed(\theta)\|_1, \qquad W_\perp(\theta) = c_{\mathrm{fee}}^2 \sum_{j\le L-3} \frac{(s_\theta^\top E e_j)^2}{1-\lambda_j/\lambda(\theta)}, \qquad \gamma(\theta) = \frac{\lambda(\theta)\,\lambda_{L-3}}{(\lambda(\theta)-\lambda_{L-3})\,c_{\mathrm{req}}}. \tag{25}$$

Then every flat probe meeting the separation constraint of (20) *satisfies*

$$L_i(v) + c_{\mathrm{fee}}\|v\|_1 \;\ge\; \min_\theta\left\{\frac{c_{\mathrm{req}}}{\lambda(\theta)} + \Psi(\theta)\right\} \;\equiv\; C_{\mathrm{RH}}(L) \;\ge\; \frac{c_{\mathrm{req}}}{\lambda_{L-1}}, \tag{26}$$

$$\Psi(\theta) = \min_{x\ge 0}\left\{x^2 + \left[F_{\mathrm{dom}}(\theta)\big(1-\gamma(\theta)x^2\big)_+ - \sqrt{W_\perp(\theta)}\,x\right]_+\right\} \ge 0,$$

a piecewise-quadratic minimization in closed form; when $\gamma F_{\mathrm{dom}} < 1$ *and the minimizer has* $\gamma x^2 \le 1$, $\Psi \ge \max\big(0, F_{\mathrm{dom}} - W_\perp/(4(1-\gamma F_{\mathrm{dom}}))\big)$, *which is what the bound uses.*

In words. This is a computable lower bound on the cost of a finite-horizon probe, fees included; equality with the optimum is not asserted. It confirms that the dollar floor of Theorem 5.2 is not an artefact of assuming an infinitely long probe.

The bound is computable for any horizon. Table 9 reports its quadratic part $c_{\mathrm{req}}/\lambda_{L-1}$, which decays toward the infinite-horizon value 61.14 at rate $O(L^{-1})$, the full bound $C_{\mathrm{RH}}(L)$, and the best feasible fee-paying probe we find at each horizon. At $L = 80$ the full bound is 64.72 against a best probe of 64.76 on the same loss-plus-fee objective (Table 9 reports 65.02 for ledger capital, which also includes margin), so the dual is essentially tight at short horizons; at longer horizons the fee term of the dual weakens and the bracket widens to about 5%. The fee of unpadded full-support eigenvectors grows as $\Theta(\sqrt{L})$, which is why zero-padded or L_1-localized probes are needed at long horizons.

5.3 Robustness: heteroskedastic impact and probing protocols

Corollary 5.5 (Invariance under impact-induced heteroskedasticity)**.** *Suppose probing inflates innovation variance as* $\sigma_{i,t}^2(v) = \sigma^2(1 + \xi[(h_i * v)_t]^2)$ *and slippage scales as* $\kappa_{\text{eff}} = \kappa(1 + \xi_0|v_t|)$*. For capacity-bounded probes,* $D_{\text{KL}}(P_0^v \| P_1^v) \le \frac{1+O(\xi^2 V_{\text{eff}}^2)}{2\sigma^2} \|(H_1 - H_0)v\|_2^2 = O(\|v\|_2^2 \delta^4)$*, and convex impact only raises dissipation. Required capital remains* $\Theta(\delta^{-4})$*.*

Four probing protocols (unconstrained single pulse, repeated capped pulses, a sequential test with dynamic throttling, and a predetermined generalized-eigenvector trajectory) all show capital slopes near -4 in $\log \delta$: -3.90, -3.84, -3.47 and -3.81 over $\delta \in [0.03, 0.12]$, where the sequential slope is flattened by the 1.7 pulses needed at $\delta = 0.12$ (Appendix Table 34). At equal observation horizon, sequential stopping saves 14–49% of capital relative to repeated capped pulses, but most of the apparent saving from the single-tap baseline comes from observing the transient response, not from adaptivity.

5.4 From par units to Bitcoin dollars

Capital scales as $\sigma^2 \delta^{-4}$, so the dollar magnitude hinges on the observation noise. Treating the 4 USDT ten-second price volatility as observation noise would give 9.8 million USDT at $\delta = 0.05$, but that conflates price volatility with measurement error. We now estimate observation noise from Binance BTCUSDT one-second closes on January 10–20, 2023, with the Roll covariance estimator $\hat{\sigma} = \sqrt{-\text{cov}(\Delta p_t, \Delta p_{t-1})}$ (Roll, 1984). Daily estimates range from 0.09 to 0.77 USDT, median 0.40 USDT.

At realistic noise, separating rates that differ by 0.05 costs about 9.9×10^4 USDT, which is not prohibitive for a well-capitalized desk (Table 10). The barrier becomes large only for finer separation (3.9 million USDT at $\delta = 0.02$) and, through Theorem 6.4, when several cohorts are unknown. Two factors we do not model would raise these figures: walking several levels of a real book costs more than the linear top-of-book impact assumed here, and impact-induced heteroskedasticity raises the noise a probe must overcome. The capital

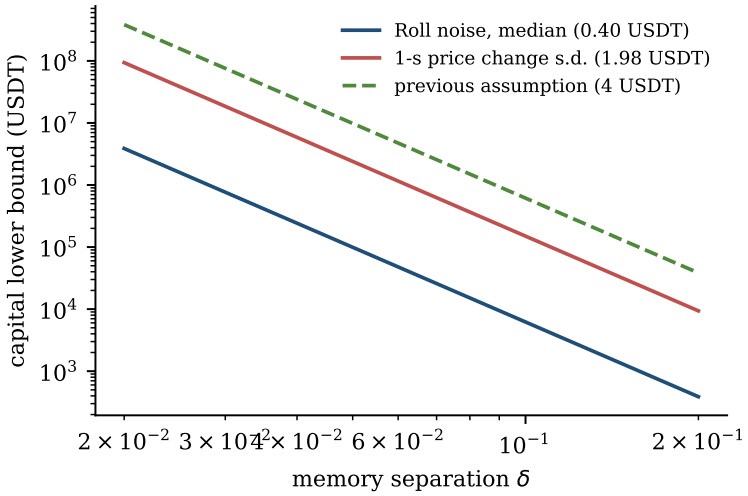


Figure 6: Capital lower bound against memory separation for three observation-noise calibrations.

figures are therefore lower bounds for the stated model.

6 Unknown mixtures and passive data

Theorem 4.1 concerns a *known* pair. A supervisor faces an unknown number of cohorts with unknown rates and weights. This section asks how fast the cost of recovering them grows.

6.1 Why naive pole fitting fails

Because the memory rates are zeros of H (Theorem 3.1), extracting poles directly from the price impulse response by matrix-pencil or Prony methods targets closed-loop eigenvalues. A four-step inversion repairs this: (1) decouple $g_t = h_{t+1} - ah_t$; (2) form a Hankel matrix of g and compute a minimal realization by Kung's eigensystem realization; (3) compute its transmission zeros, which are the roots of the denominator of G; (4) project reconstructed flow onto the recovered memory filters to obtain the weights (Appendix C.3). Without noise this recovers $K = 0$, $K = 1$ and well-separated two-cohort systems exactly (Appendix C.3); with noise, per-cohort weights become unreliable once rates are close. The reason is information-theoretic, as the rest of this section shows.

6.2 Local information: the Cramér–Rao benchmark

Proposition 6.1 (Cramér–Rao benchmark)**.** *Let θ collect the weights $w_1, \dots, w_K$ and rates $\eta_1, \dots, \eta_K$. Suppose the impulse response $h_t(\theta)$, $t \le T$, is observed with i.i.d. Gaussian error from N reset probes of size V_{probe} against innovation noise σ_p, so the error has standard deviation $\sigma_h = \sigma_p/(V_{\text{probe}}\sqrt{N})$. With Jacobian $J_{tj} = \partial h_t/\partial\theta_j$ and $F_{\text{unit}} = J^\top J$, any regular unbiased estimator has $\mathrm{Var}(\hat{w}_k) \ge \sigma_h^2[F_{\text{unit}}^{-1}]_{kk}$, and a 95% interval of half-width* tol *requires*

$$N^* \ge \left(\frac{1.96}{\text{tol}}\right)^2 \left(\frac{\sigma_p}{V_{\text{probe}}}\right)^2 [F_{\text{unit}}^{-1}]_{kk}. \tag{27}$$

For $K = 2$ with spacing $\Delta\eta$, $[F_{\text{unit}}^{-1}]_{kk} = \Theta((\Delta\eta)^{-6})$, so $N^* = \Theta((\Delta\eta)^{-6})$.

In words. This gives the number of test trades any well-behaved estimator needs to pin down one group's trading weight to a stated precision. Each extra group with a nearby memory raises the number sharply, as Table 12 shows.

At the critical noise level maximum likelihood attains the bound: its RMSE for the hardest weight is between 0.88 and 1.22 times the Cramér–Rao value and every weight is within tolerance in 93% to 98% of replications (Table 12), so the critical noise levels, not estimator failure, are what make close cohorts expensive. Under ordinary surveillance budgets ($N \le 50$) unbiased estimators cannot resolve close weights.

6.3 A global two-point bound for any estimator

The Cramér–Rao benchmark of Proposition 6.1 applies only to regular unbiased estimators. Could biased, ridge, shrinkage or Bayesian estimators escape the $(\Delta\eta)^{-6}$ scaling? No.

Theorem 6.2 (Global two-point bound and matching upper bound, $K = 2$)**.** *In this theorem w_k denotes cohort k's trading intensity q_k, not the merged weight $q_k\eta_k$ of Section 2, and $\Delta\eta$ is the actual spacing of the two rates. Fix total intensity $q = 0.20$, mean rate $\bar{\eta} = 0.40$, target*

separation $\tau = 0.01$*, and* $\Delta\eta \in (0, 8/11)$ *(so that* $\bar{\eta} - p_+\Delta\eta > 0$*). Let* $p_\mp = \frac{1}{2} \mp \tau/q$ *and*

$$\theta_0 = (qp_-, qp_+, \bar{\eta} - p_+\Delta\eta, \bar{\eta} + p_-\Delta\eta), \qquad \theta_1 = (qp_+, qp_-, \bar{\eta} - p_-\Delta\eta, \bar{\eta} + p_+\Delta\eta). \tag{28}$$

(i) Label cohorts in increasing order of rate, so that weight vectors are compared rate by rate (as unordered sets the two weight vectors coincide). Both economies have spacing $\Delta\eta$*, weight separation* $\|w(\theta_1) - w(\theta_0)\|_\infty = 2\tau$ *in this labelling, and equal rate moments* $M_m = \sum_k w_k \eta_k^m$ *for* $m \in \{0, 1, 2\}$*, while*

$$M_3(\theta_1) - M_3(\theta_0) = 4\tau p_- p_+ (\Delta\eta)^3. \tag{29}$$

Hence $\|h(\theta_1) - h(\theta_0)\|_2 = \Theta(\tau(\Delta\eta)^3)$*.*

(ii) For any measurable estimator (including biased, ridge, James–Stein or MAP) that reports weights in the same rate order, the supremum being over the class of two-cohort economies with total weight q*, mean rate* $\bar{\eta}$ *and spacing* $\Delta\eta$*,*

$$\sup_\theta P_\theta(\|\hat{w} - w(\theta)\|_\infty \geq \tau) \;\geq\; 1 - \Phi\Big(\frac{\sqrt{N}\|h(\theta_1) - h(\theta_0)\|_2}{2\sigma_h}\Big), \tag{30}$$

so failure probability at most 0.05 *requires* $N \geq 4\sigma_h^2 z_{0.95}^2 / \|h(\theta_1) - h(\theta_0)\|_2^2 = \Omega((\Delta\eta)^{-6})$*. For* $\Delta\eta \geq 2\tau q/(q + 2\tau)$ *(*≈ 0.018 *at the stated values) the same bound holds without any labelling convention, for the permutation-invariant loss*

$$d_\infty(\hat{\theta}, \theta) = \min_{\pi \in S_2} \max_k \Big(|\hat{w}_{\pi(k)} - w_k| + |\hat{\eta}_{\pi(k)} - \eta_k|\Big),$$

because then $d_\infty(\theta_0, \theta_1) = \min\{2\tau(1 + \Delta\eta/q),\, (1 + 2\tau/q)\Delta\eta\} \geq 2\tau$*. For smaller spacings the relevant permutation-invariant statement is the Wasserstein bound of Theorem 6.4.*

(iii) Conversely, the matching upper bound $N = O((\Delta\eta)^{-6})$ *on rich separated two-cohort classes is Theorem 6.5 with* $K = 2$*. On the separated class with weights and rates*

bounded away from the boundary, Toeplitz deconvolution followed by a Prony–Hankel root step has Hankel determinant $g_1 g_3 - g_2^2 = w_1 w_2 \eta_1 \eta_2 r_1 r_2 (\Delta\eta)^2 > 0$ *for the deconvolved kernel* $g_m = \sum_i w_i \eta_i r_i^m$*, and perturbation analysis gives the same order.*

In words. For two unknown groups, no method can tell the two candidate markets apart reliably with fewer than a number of experiments of order $(\Delta\eta)^{-6}$, and a practical method (Prony's method, Theorem 6.5) achieves this order. Halving the spacing multiplies the data needed by 64.

We recomputed the two-point quantities exactly (Table 13): the moment gap equals $4\tau p_- p_+ (\Delta\eta)^3$ to machine precision, and the Le Cam requirement at $\Delta\eta = 0.10$ is 5,927.6 resets. The worst-case failure rates of four re-implemented estimators all lie above the lower bound (Table 13).

Remark 6.3 (The hard directions carry no risk). The two economies of Theorem 6.2 have the same total weight $\sum_i q_i = q$, so their summed price response $\sum_t h_t = \kappa/(1 - a - \kappa \sum_i q_i)$ is identical, 2.353 at the calibration of Table 13: the directions in which prices are least informative about the cohorts are directions in which cumulative price risk does not move. The summed response is a linear functional of the response and is estimated at the parametric rate whatever the spacing of the memories; the companion paper turns this into a certification theorem, a minimax-optimal pre-trade guard and a decision value of disclosure (Niroula et al., 2026). Identification and safety ask different questions of the same data.

6.4 Many unknown cohorts

Theorem 6.4 (K-cohort two-point lower bound)**.** *Fix* $K \geq 2$*, a center* $r_c \in (0, 1)$*, total merged weight* $W > 0$ *and* $0 < \Delta\eta < \min(r_c, 1 - r_c)$*, so that all rates lie in* $(0, 1)$*, and assume* $a + \kappa W/(1 - r_c - \Delta\eta) < 1$*, so that every candidate economy is Schur stable. Each reset experiment observes* $T \geq 2K$ *noisy taps (*$T = \infty$ *allowed); the bound below uses the full response norm, which dominates any finite prefix. There exist two* K*-cohort economies*

with rates in $r_c + \Delta\eta[-1, 1]$*, positive weights, and merged measures whose moments of order* $0, \dots, 2K-2$ *coincide while the moment of order* $2K-1$ *differs by a constant times* $(\Delta\eta)^{2K-1}$*. Their impulse responses differ by* $\|h^A - h^B\|_2 \asymp (\Delta\eta)^{2K-1}$*, so recovering the weights from* N *independent noisy reset experiments with worst-case error probability below* $1/4$ *requires*

$$N \;\geq\; c\,\sigma^2(\Delta\eta)^{-(4K-2)}.$$

The same requirement holds for the permutation-invariant loss $W_1(\hat{M}, M)$ *between estimated and true merged measures at accuracy* $c_K W \Delta\eta$*, where* $2c_K > 0$ *is the Wasserstein distance between the two normalized quadrature rules, so no relabelling of atoms escapes it.*

The two losses have different targets. The weights of the two economies differ by a constant that does not shrink with $\Delta\eta$, so the first statement is about weights at fixed accuracy, which Theorem 6.5 matches; the Wasserstein statement is about the memory distribution at accuracy proportional to $\Delta\eta$, which Theorem 6.7 matches.

The construction uses the K-point Gauss and Gauss–Radau rules of the same weight function: both are positive, have K atoms and integrate polynomials of degree $2K - 2$ exactly. For $K = 2$ it reproduces the order of Theorem 6.2. Table 14 lists the two rules: their low moments agree to machine precision and their moment of order $2K - 1$ differs by 0.222, 0.0533 and 0.0131 for $K = 2, 3, 4$. The shrinking gap explains why constants deteriorate with K on top of the exponent.

What counts as an experiment. In Theorem 6.4 an experiment is a unit impulse from the reset state whose response is observed with noise of variance σ^2 per tap. The exponent is the same for any probe design of bounded energy. For any probe v the mean difference between the two economies is at most $\sup_\omega |\Delta\hat{H}(e^{i\omega})|\,\|v\|_2$, and $\sup_\omega |\Delta\hat{H}| = O((\Delta\eta)^{2K-1})$ because $\Delta\hat{H} = z^{-1}\hat{H}_A\hat{H}_B\,\Delta G$ with both responses bounded on the circle. Experiments with probe energy E each therefore need $NE \gtrsim \sigma^2(\Delta\eta)^{-(4K-2)}$: a larger probe buys the same information with fewer experiments, and Section 5 prices the capital that a larger

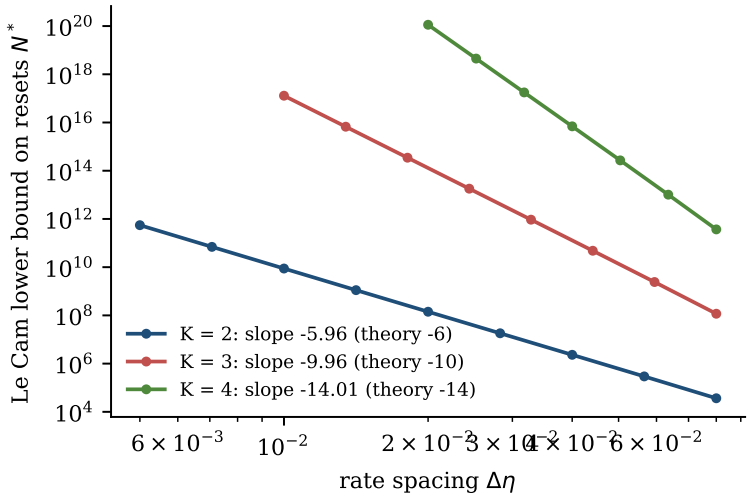


Figure 7: Two-point lower bound on the number of reset experiments needed to separate K-cohort economies, as a function of the rate spacing.

probe requires. The bound does not cover information obtained without trading against the groups, such as class records (Section 8).

Table 15 and Figure 7 confirm the order numerically.

Theorem 6.5 (Matching upper bound for every K)**.** *Fix $K \geq 2$, a compact set $\mathcal{S} \subset \{s \in [-1,1]^K : s_1 < \cdots < s_K\}$ of ordered rescaled rates with $\min_{i \neq j} |s_i - s_j| \geq \gamma > 0$, and merged weights $w_i \in [\omega_{\min}, \omega_{\max}]$ with $\omega_{\min} > 0$. Consider the class of K-cohort economies with rates $r_c + \Delta\eta\, s$, $s \in \mathcal{S}$, and N independent reset experiments whose responses carry Gaussian noise with variance σ^2. There are constants $c_0, C > 0$ depending only on $(K, \mathcal{S}, \gamma, \omega_{\min}, \omega_{\max}, a, \kappa, r_c)$ such that, whenever $\epsilon_N \equiv \sigma N^{-1/2} (\Delta\eta)^{-(2K-1)} \leq c_0 / \sqrt{\log(4K/\beta)}$, Prony's method applied to the first $2K$ estimated kernel values, with its output atoms listed in increasing order of rate, recovers the weights with*

$$\max_i |\hat{w}_i - w_i| \;\leq\; C \sqrt{\log(4K/\beta)}\, \epsilon_N$$

with probability at least $1 - \beta$. Hence $N = C' \sigma^2 (\Delta\eta)^{-(4K-2)}$ experiments suffice for any fixed accuracy and confidence. On any such class that contains the Gauss and Gauss–Radau pair of Theorem 6.4 (a rich class), together with that theorem the minimax sample requirement is

$N^* \asymp \sigma^2(\Delta\eta)^{-(4K-2)}$. *On smaller classes the upper bound need not be tight: if the rates are known, for example, the weights need only order* $\sigma^2(\Delta\eta)^{-2}$ *experiments. The constants* c_0 *and* C *grow quickly with* K*, through the conditioning of the moment map; the theorem gives the order in* $\Delta\eta$ *for fixed* K*, not a usable constant for large* K *(in Supplement Table 36 our double-precision Prony implementation already fails at* $K = 4$ *on the disclosure grid).*

Proof. The kernel values $k_m = \sum_i w_i r_i^m$ are the power moments of the merged measure. The closed-loop recursion $h_t = a h_{t-1} + \kappa \sum_m k_m h_{t-1-m}$ (plus κ at $t = 0$) can be solved for $k_0, \dots, k_{2K-1}$ from $h_1, \dots, h_{2K}$. Because $h_0 = \kappa$ is a known primitive, each k_m is a polynomial in $h_1, \dots, h_{m+1}$ with coefficients depending on (a, κ) (the only division is by $\kappa h_0 = \kappa^2$), so on the bounded set of responses within sup-distance 1 of the class this map is L_h-Lipschitz, with L_h depending on the class but not on $\Delta\eta$. Let E be the event $\max_{1 \le t \le 2K} |\bar{h}_t - h_t| \le \sigma N^{-1/2} \sqrt{2 \log(4K/\beta)}$, where $\bar{h}_t$ averages the N responses; a Gaussian tail bound and a union bound over the $2K$ estimated taps give $P(E) \ge 1 - \beta$. On E, and provided this radius is at most 1, the estimated kernel values satisfy $\max_m |\hat{k}_m - k_m| \le L_h \sigma N^{-1/2} \sqrt{2 \log(4K/\beta)}$. Centring at r_c is a fixed triangular linear map, $\mu_m = \sum_{j \le m} \binom{m}{j} (-r_c)^{m-j} k_j$, and $\mu_m = (\Delta\eta)^m t_m$ with $t_m = \sum_i w_i s_i^m$. The map $F : (s, \omega) \mapsto (t_0, \dots, t_{2K-1})$ does not depend on $\Delta\eta$; its Jacobian is a confluent Vandermonde matrix with determinant $\prod_i \omega_i \prod_{i<j} (s_i - s_j)^4$, bounded away from zero on the compact class. On ordered, γ-separated nodes F is injective, because a measure with K distinct atoms is determined by its first $2K$ moments. A C^1 map that is injective on a compact set Q and has a nonsingular Jacobian there is injective on an open neighbourhood U of Q: otherwise there are $x_n \ne y_n$ with $F(x_n) = F(y_n)$ and distance to Q tending to zero, whose limits $x, y \in Q$ satisfy $F(x) = F(y)$, so $x = y$, contradicting local injectivity at x. Applying the quantitative inverse function theorem uniformly on Q, there are $\rho_0 > 0$ and $L < \infty$, uniform over the class, such that $F(U)$ contains the ball of radius ρ_0 around every point of $F(Q)$ and the single inverse $F^{-1} : F(U) \to U$ is L-Lipschitz on each such ball. An error e in μ_m is an error $e(\Delta\eta)^{-m}$ in t_m, largest at $m = 2K - 1$. Since $\Delta\eta < 1$ and the centring map is fixed, on E we get $\|\hat{t} - t\|_\infty \le$

$C_1\sqrt{\log(4K/\beta)}\,\epsilon_N$ with C_1 equal to $\sqrt{2}L_h$ times the norm of the centring map; and because $\sigma N^{-1/2}\sqrt{2\log(4K/\beta)} \le \sqrt{2}\,\epsilon_N\sqrt{\log(4K/\beta)}$, the radius condition holds whenever $c_0 \le 1/\sqrt{2}$. When this is below ρ_0, $\hat{t}$ lies in the ball where F^{-1} is L-Lipschitz. For an input in such a ball, $x = F^{-1}(\hat{t})$ has K distinct nodes and nonzero weights, so the $K \times K$ Hankel matrix of $\hat{t}$ is nonsingular and Prony's method returns exactly x: the Prony polynomial is determined uniquely by the Hankel system, and since the K real nodes of x reproduce $\hat{t}$ it equals $\prod_i(z-\hat{s}_i)$, so its roots are real and distinct and no complex pair can appear; Prony's method computes the algebraic inverse of the moment map up to the order of the atoms, so, listing its output in increasing order of rate (which the error bound below keeps unambiguous once $C\sqrt{\log(4K/\beta)}\,\epsilon_N < \gamma/2$), its output is $F^{-1}(\hat{t})$ and the stated bound follows with $C = LC_1$ and $c_0 = \min\{\min(\rho_0, \gamma/(2L))/C_1,\, 1/\sqrt{2}\}$, which also keeps $C\sqrt{\log(4K/\beta)}\,\epsilon_N < \gamma/2$. No separate root-perturbation bound is needed: on the ball where F^{-1} is L-Lipschitz, Prony's roots and weights are the coordinates of $F^{-1}(\hat{t})$, so the conditioning of the root step is part of L (for an explicit step-by-step analysis of Prony's method with clustered exponents see Katz et al., 2023). The matching statement requires the class to contain the Gauss and Gauss–Radau pair of Theorem 6.4 (ordered nodes with separation at least the smaller of their two minimum gaps, and weights within $[\omega_{\min}, \omega_{\max}]$ after scaling to total weight W); for any such class the orders match, with constants that depend on the class. □

In words. With K unknown groups whose memories lie within $\Delta\eta$ of each other, the number of experiments needed grows like $(\Delta\eta)^{-(4K-2)}$. For two groups, halving the spread multiplies the data needed by 64; for three groups, by 1,024. Theorem 6.4 shows that no method does better, and this theorem shows that Prony's method achieves it.

Machine check. Lean verifies this theorem for every $K \ge 1$ (`p1_thm_prony_upper_bound`), using only Lean's standard axioms. The formal proof follows the one above: a Lipschitz bound for the closed-loop deconvolution that does not depend on $\Delta\eta$, the $(\Delta\eta)^{-m}$ rescaling, a sub-Gaussian tail bound over the $2K$ taps, injectivity of the Jacobian of the moment map

(proved by Hermite interpolation, without the determinant formula), and a uniform inverse function theorem on the compact class. The formal estimator is the representation with increasing nodes and nonzero weights that matches the $2K$ estimated moments; Lean shows that it is unique and that Prony's linear system has a unique solution whose polynomial has exactly its nodes as roots, so the roots are real. The formal statement allows sub-Gaussian noise and assumes $|r_c| + \Delta\eta \le 1$, so that every memory rate has modulus at most one.

Corollary 6.6 (Plug-in upper bound with unknown primitives (p^*, κ, a))**.** *Suppose the fundamental value p^*, price impact $\kappa \in [\underline{\kappa}, \overline{\kappa}] \subset (0, \infty)$, and persistence $a \in [0, \bar{a}] \subset [0, 1)$ are unknown to the observer. Subtracting the pre-probe price $p_{n,0} = p^* + \nu_{n,-1}$ eliminates p^* (doubling observation noise variance to $2\sigma^2$), and $\hat{\kappa} = \hat{h}_0$ estimates κ from the impact tap $h_0 = \kappa$ with $|\hat{\kappa} - \kappa| \le \sigma N^{-1/2}\sqrt{2\log(8K/\beta)}$. Since the lag-1 response satisfies $h_1 = a\kappa + \kappa^2 k_0$, any $\sqrt{N}$-consistent estimator of a with $|\hat{a} - a| \le C_a \sigma N^{-1/2}\sqrt{2\log(8K/\beta)}$ (for example from lag-0 and lag-1 taps, a pre-probe free-decay step, or regressing price changes on aggregate anonymous order flow under contract I_1) can be plugged in: replacing (a, κ) by $(\hat{a}, \hat{\kappa})$ in the closed-loop deconvolution of Theorem 6.5 yields*

$$\max_{0 \le m \le 2K-1} |\hat{k}_m^{\mathrm{plug}} - \hat{k}_m| \le L_{a,\kappa}\Big(|\hat{a} - a| + |\hat{\kappa} - \kappa|\Big) \le C_{a,\kappa}\sigma N^{-1/2}\sqrt{\log(8K/\beta)}$$

with probability at least $1 - \beta/2$, where $L_{a,\kappa} < \infty$ depends only on $(K, \underline{\kappa}, \overline{\kappa}, \bar{a}, W)$ and not on $\Delta\eta$. Consequently, Prony's method with plug-in $(\hat{a}, \hat{\kappa})$ achieves the exact same minimax sample complexity $N = C''\sigma^2(\Delta\eta)^{-(4K-2)}$ as when (p^*, κ, a) are known.

Table 16 confirms the order numerically. For numerical stability the estimator works in centred and rescaled coordinates (the estimator is the same; only its floating-point evaluation changes). The argument parallels the stability analysis of super-resolution for clustered complex nodes (Batenkov et al., 2021), and Katz et al. (2023) show that Prony's method attains the min-max bound for clustered exponents; here the nodes are real and are first extracted from the closed loop, and it reproduces the ESPRIT slope of -6.15 reported for

$K = 2$ in the earlier version.

6.5 Do the rates survive real market noise?

Theorems 6.4 and 6.5 are proved for observation errors that are independent and Gaussian, and Table 16 checks them with Gaussian noise. Real price changes are different: one-minute returns have heavy tails, although they are almost uncorrelated from one minute to the next. To see whether the rates depend on the Gaussian assumption, we ran a pre-specified test (SHA-256 `596a700da2e4...`, opened once) in which the trader groups are synthetic, with known memory spacing, and only the noise is real. Each reset experiment returns the price response to a unit impulse plus a window of consecutive one-minute log-price changes of Binance BTCUSDT, demeaned and scaled to unit variance. The changes come from 21 days, the 6th of each month from January 2025 to September 2026, that are used nowhere else in this project; their kurtosis is 19.4, against 3 for Gaussian noise. With class tagging, the experiment also returns each group's flow response, observed through independent windows of the same noise. A control arm uses i.i.d. Gaussian noise of the same variance. For each spacing, N^* is the smallest number of experiments at which every recovered weight lies within 25% of the smallest true weight in at least 75% of 200 replications, using Prony's method on prices alone or the per-group regression of Proposition 8.1 with tagging. The noise level and the spacing grids were fixed on Gaussian noise alone, before any real noise was downloaded.

Both pre-specified hypotheses hold, so the test passes (Table 17). With real noise, the price-only requirement for two groups rises from 64 experiments at $\Delta\eta = 0.5$ to 370,728 at $\Delta\eta = 0.125$, a factor of about 5,800, and the fitted slope is -6.01; for three groups it is -9.81. Over the same grid the tagged requirement stays between 256 and 724 experiments, and it falls slightly rather than rises as the spacing shrinks. At every spacing the real-noise requirement is within one step of the grid (a factor $\sqrt{2}$) of the Gaussian control.

Computational exactness across all N. Empirical window averages at any sample size N are computed exactly by drawing allocation counts Multinomial$(N; 1/M, \ldots, 1/M)$ across

the $M = 30{,}219$ historical one-minute return windows. Because matrix multiplication evaluates the exact sample mean $\frac{1}{N}\sum_{j=1}^{M} C_j W_j$ in $O(M)$ time for each replication, exact empirical resampling is maintained throughout the entire grid, including at $N = 370{,}728$, without relying on a central-limit normal approximation. Real one-minute changes exhibit substantial heavy tails (sample kurtosis 19.4 against 3 for Gaussian noise, with lag-1 autocorrelation within $[-0.007, 0.009]$). Across all spacings and cohort configurations, the real heavy-tailed noise preserves the theoretical polynomial scaling (slopes -6.01 and -9.81) without distortion from empirical kurtosis.

In words. With real Bitcoin price noise, the price-only route needs about as many experiments as Gaussian noise, and halving the spacing between two groups multiplies the data requirement by about 64, while class-tagged records keep the requirement flat across all spacings.

6.6 Near-colliding rates: a uniform rate for the merged measure

Theorem 6.5 needs the rates to stay γ-separated and every weight to stay away from zero. Neither restriction can be dropped for weight recovery: two atoms of weights w and $W - w$ at rates r and $r + \Delta$ and two of weights $W - w$ and w at the same rates have merged measures that converge to the same measure as $\Delta \to 0$, so no estimator recovers the individual weights uniformly over near-colliding economies. What can be recovered uniformly is the merged measure itself, in the permutation-invariant Wasserstein distance.

Theorem 6.7 (Uniform rate without separation)**.** *Fix $K \geq 1$, a memory bound $\bar{r} < 1$ and $0 < m_{\min} \leq B$. Let $\mathcal{M}_K$ be the merged measures with at most K atoms in $[0, \bar{r}]$ and total mass $k_0 \in [m_{\min}, B]$, with no separation condition and no lower bound on individual weights, and write $\bar{M} = M/k_0$ for the normalized memory distribution. From N reset experiments as in Theorem 6.5, let $\hat{M}$ be any element of $\mathcal{M}_K$ whose first $2K$ kernel values are closest in the sup norm to the estimated ones. There is C depending only on $(K, \bar{r}, a, \kappa, B)$ such that,*

with probability at least $1-\beta$,

$$W_1(\hat{\bar{M}}, \bar{M}) \;\le\; C\Big(\frac{\sigma\sqrt{\log(4K/\beta)}}{m_{\min}\sqrt{N}}\Big)^{1/(2K-1)} \qquad \textit{uniformly over } \mathcal{M}_K.$$

Hence $N = O(\sigma^2 m_{\min}^{-2}\epsilon^{-(4K-2)})$ *experiments give accuracy* ϵ*, and with the Wasserstein statement of Theorem 6.4 the minimax rate for the merged memory distribution is* $N^{-1/(4K-2)}$*, uniformly over near-colliding economies.*

In words. Even with no lower limit on how close the groups' memories are or how small a group is, the combined memory distribution can still be estimated, and no faster rate is possible. What is lost when groups nearly coincide is the ability to tell them apart, not the ability to describe the combined picture.

Proof. As in the proof of Theorem 6.5, the kernel values $k_0, \dots, k_{2K-1}$ are a Lipschitz function of the first $2K+1$ responses on the bounded set of responses the class produces, so a union bound gives $\max_{t<2K} |\bar{k}_t - k_t| \le e = C_1\sigma\sqrt{\log(4K/\beta)/N}$ with probability at least $1-\beta$ (the Lipschitz constant is for a bounded neighbourhood of the class; when the noise is too large for that, the right-hand side of the theorem exceeds $\bar{r} \ge W_1$ once C is enlarged, so the bound holds trivially). The class is compact and the moment map continuous, so a closest element $\hat{M}$ exists, and $\max_t |k_t(\hat{M}) - k_t(M)| \le 2e$ by the triangle inequality. Since all atoms lie in $[0,1]$, $|k_t| \le k_0$, and the normalized moments $m_t = k_t/k_0$ satisfy $|m_t(\hat{M}) - m_t(M)| \le 2e/k_0 + (|\hat{k}_t|/\hat{k}_0)(2e/k_0) \le 4e/m_{\min}$. The affine map $r \mapsto 2r/\bar{r} - 1$ carries both normalized measures to distributions on $[-1,1]$ with at most K atoms, multiplies W_1 by $2/\bar{r}$ and the differences of moments of order below $2K$ by at most $(1+2/\bar{r})^{2K-1}$. The moment comparison inequality for K-atomic distributions on $[-1,1]$ (Wu and Yang, 2020), $W_1 \le c_K\delta^{1/(2K-1)}$ when the first $2K-1$ moments differ by at most δ, then gives the bound with $C = c_K(\bar{r}/2)(1+2/\bar{r})(4C_1)^{1/(2K-1)}$. □

The theorem separates two questions that Theorem 6.5 answers together. Which mem-

ory distribution do the participants have? That has the same polynomial cost whether or not their rates collide. Which participant has which weight? That becomes arbitrarily expensive as rates approach each other (Theorem 6.4) and impossible when they coincide (Theorem 3.1). A supervisor who needs the first can have it at the rate of Theorem 6.4; one who needs the second has to be told the number of cohorts and a separation, or has to obtain tagged data (Section 8).

6.7 General memory: the minimax rate is logarithmic

Exponential memory is a modelling choice. Suppose instead that the feedback kernel is completely monotone, $k_m = \int_{[0,\rho]} r^m \, d\mu(r)$ with $\rho < 1$, where μ is any nonnegative measure of total mass W: a finite or infinite mixture of exponential memories with merged weight W. Write $\mathcal{M}(\rho, W)$ for this class, assume the closed loop is uniformly stable, $s \equiv a + \kappa W/(1-\rho) < 1$, and measure accuracy by

$$d(\mu, \mu') = \sup\Big\{ \textstyle\int f \, d(\mu - \mu') : \;\; f \text{ 1-Lipschitz on } [0,\rho], \;\; f(0) = 0 \Big\},$$

which equals the Wasserstein-1 distance between measures of equal mass and also compares estimates whose mass is not exactly W. Let $E_L(|x|)$ denote the error of the best uniform approximation of $|x|$ on $[-1,1]$ by polynomials of degree L, and let $c_* = \inf_{L \geq 1} L \, E_L(|x|)$. Bernstein (1914) showed that $L \, E_L(|x|)$ converges to a positive constant (0.2802; Varga and Carpenter, 1985), so $c_* > 0$; numerically $c_* = 1/4$, attained at $L = 2$ (Table 18).

Lemma 6.8 (Moment-matched cohort pairs)**.** *For every $n \geq 2$ there are two probability measures on $[0, \rho]$, each with exactly n atoms, whose moments of orders $0, \dots, 2n-2$ coincide and whose Wasserstein-1 distance is at least $\rho \, E_{2n-2}(|x|) \geq c_* \rho/(2n-2)$.*

The atoms are the points where the best approximation of $|r - \rho/2|$ by polynomials of degree $2n-2$ attains its maximal error with alternating signs, and the weights are those of the divided difference on these points, which annihilates every polynomial of degree $2n-2$

(Appendix A). Each measure is an n-atom merged measure in the class; one atom may sit at rate 0 (a cohort with $\eta = 1$), which the class allows.

Theorem 6.9 (Minimax rate for general memory)**.** *Let N independent reset experiments observe the impulse response with Gaussian noise of variance σ^2 per tap, and let $B(\rho) = (1+\sqrt{2})(1+2/\rho)$ and $C_h = \kappa^2/(1-s)^2$.*

(i) Lower bound. *Let $\Lambda_N = \log(N/\sigma^2) + \log\big(2C_h^2 W^2/((1-\rho)\rho)^2\big)$. Whenever $\Lambda_N > 4\log(1/\rho)$, every estimator $\hat{\mu}$ satisfies*

$$\sup_{\mu \in \mathcal{M}(\rho, W)} P_\mu\Big(d(\hat{\mu}, \mu) \geq \frac{c_*\, \rho W \log(1/\rho)}{\Lambda_N}\Big) \;\geq\; \frac{1}{4e},$$

and the bound already holds over n-cohort economies, for a finite n that grows with N.

(ii) Upper bound. *With a and κ known, invert the averaged response into kernel values $\hat{k}_0, \ldots, \hat{k}_L$, with $L+1 = \lfloor \log(N/\sigma^2)/(4\log B(\rho)) \rfloor$, and let $\hat{\mu}$ be any measure on $[0, \rho]$ minimizing $\max_{j \leq L} |\int r^j d\hat{\mu} - \hat{k}_j|$ (a linear program). For every $\beta \in (0,1)$ there is N_0 depending only on $(\rho, W, a, \kappa, \beta)$ such that for $N/\sigma^2 \geq N_0$,*

$$\sup_{\mu \in \mathcal{M}(\rho, W)} P_\mu\Big(d(\hat{\mu}, \mu) > \frac{3\pi\rho W \log B(\rho)}{\log(N/\sigma^2)}\Big) \;\leq\; \beta.$$

Hence the minimax accuracy is of exact order $1/\log(N/\sigma^2)$, and reaching accuracy w requires and suffices $N/\sigma^2 = \exp(\Theta(1/w))$.

Proof sketch. (i) Put the measures of Lemma 6.8, scaled to mass W, into the model. Their difference annihilates polynomials of degree $2n-2$, so the generating function of the kernel difference is $\int (1-rz)^{-1} d(\mu_A - \mu_B) = z^{2n-1} \int r^{2n-1}(1-rz)^{-1} d(\mu_A - \mu_B)$; on $|z| = 1$ it is at most $2W\rho^{2n-1}/(1-\rho)$, and Parseval bounds $\|k^A - k^B\|_2$ by the same quantity. The response difference is $H_A H_B$ times the shifted kernel difference, so $\|h^A - h^B\|_2 \leq C_h \|k^A - k^B\|_2$. Choosing the smallest n for which the Kullback–Leibler divergence of the N-sample

experiments is at most one, the Bretagnolle–Huber form of Le Cam's lemma gives the bound with half the Wasserstein separation of the lemma. (ii) For a 1-Lipschitz f with $f(0) = 0$, Jackson's theorem gives a polynomial q of degree L with $\sup|f - q| \le \pi\rho/(4(L+1))$, and expanding q in Chebyshev polynomials of $2r/\rho - 1$ bounds the sum of its monomial coefficients by $\rho B(\rho)^{L+1}$. Hence $d(\hat{\mu}, \mu) \le \pi\rho(W + \delta)/(2(L+1)) + 2\delta\rho B(\rho)^{L+1}$, where δ is the largest kernel error. The kernel is recovered from the response by a series reciprocal, $K(z) = (1 - az - \kappa/H(z))/(\kappa z)$, whose error is at most $2((1+s)/\kappa)^2(L+2)$ times the largest response error, so $\delta = O_P(\sigma N^{-1/2} L\sqrt{\log L})$. The choice of L makes $\delta B(\rho)^{L+1} = O_P((\sigma^2/N)^{1/4} L\sqrt{\log L})$, negligible against the first term. Full proofs are in Appendix A. □

Machine check. Lean verifies part (i), for N reset experiments of any length, using only Lean's standard axioms (`p1_thm_general_memory_lower`). The formal proof replaces the pair of Lemma 6.8 by an elementary one: the binomial weights $W\binom{2L+1}{k}/2^{2L}$ on the rates $\rho k/(2L+1)$, odd k against even k, which have $n = L+1$ atoms each and equal moments of orders $0, \dots, 2n-2$. The 1-Lipschitz function $f(r) = (h/\pi)(1-\cos(\pi r/h))$, $h = \rho/(2L+1)$, vanishes at 0 and at even nodes and equals $2h/\pi$ at odd nodes, so the pair is at distance at least $2\rho W/(\pi(2n-1))$, more than the $c_*\rho W/(2n-2)$ of the lemma, without Bernstein's theorem. Because the raw moments agree, the kernels agree below lag $2n-1$ and differ by at most $2W\rho^{2n-1}/(1-\rho)$ in ℓ_1 over all lags; the closed-loop ℓ_1 bounds, the Gaussian divergence and the testing lemma then show that, with $C = 2\kappa^2 W/((1-\rho)(1-s)^2)$, $\Lambda = \log(2NC^2/\sigma^2) > 0$ and $L = \lceil \Lambda/(4\log(1/\rho)) \rceil$, no estimator reaches accuracy $2\rho W \log(1/\rho)/(\pi(\Lambda + 6\log(1/\rho)))$ with probability at least $3/4$ under both economies. This is the order of part (i), with success probability $3/4$ in place of $1/(4e)$ and constants of its own. An estimator $\hat{\mu}$ enters through its integrals $\int f\, d\hat{\mu}$ of 1-Lipschitz functions with $f(0) = 0$, which is how $d(\hat{\mu}, \mu)$ is defined. Lean also verifies part (ii) (`p1_thm_general_memory_rate`). Mathlib has no Jackson theorem, so the formal proof builds one: on $M = n+2$ equally spaced nodes the Fejér–Korovkin kernel $K(u) = |\sum_{k \le n} a_k e^{iku}|^2 / \sum_k a_k^2$, $a_k = \sin(\pi(k+1)/(n+2))$, is nonnegative,

has node mean one and node mean of $K(u)(1-\cos u)$ equal to $1-\cos(\pi/(n+2))$ (from the recurrence of the sines and the orthogonality of roots of unity), so the positive operator it defines maps an L-Lipschitz function on $[-1,1]$ to a polynomial of degree at most n within $L\pi/(n+2)$, by Cauchy–Schwarz on the nodes. Written in Chebyshev polynomials and rescaled to $[0,\rho]$, its coefficients sum in absolute value to at most $\rho(n+1)B(\rho)^n$. With an explicit bound for the deconvolution, $\alpha e(1+\beta')^t$ for kernel value t, and the sub-Gaussian union bound, any measure on $[0,\rho]$ that fits the first $L+1$ deconvolved kernel values at least as well as the truth, the linear-programming estimate among them, is within a finite-sample bound of μ with probability at least $1-\beta$ (`p1_thm_general_memory_upper`); taking $L=\lfloor \log(N/\sigma^2)/(8\log D)\rfloor$, $D=B(\rho)(1+\beta')$, gives $16\pi\rho W\log D/\log(N/\sigma^2)$ for all N/σ^2 beyond a threshold depending only on $(a,\kappa,\rho,s,W,\beta)$. The formal constants are larger than the theorem's 3π and $B(\rho)$, because the formal deconvolution bound grows geometrically in L; the order is the same.

In words. If trader memory is a general mixture of exponentials rather than a few clean ones, accuracy improves only with the logarithm of the data: to halve the error, the data must be squared. The few-groups model of the previous results is therefore the optimistic case.

Theorem 6.9 replaces a polynomial barrier with a logarithmic one. With exponential memory and K known cohorts, halving the error multiplies the required number of experiments by a constant; with general memory, halving the error squares it. The lower bound holds already for finite mixtures, so knowing that the memory is a finite mixture of exponentials does not help unless the number of cohorts is also known (Theorem 6.7). The constants of the two bounds differ, by a factor that grows as $\rho\to 1$, so the dependence on how long memory lasts is not pinned down. The source of the gap is visible in the proof: as $\rho\to 1$ the moment-matched pairs separate their kernels more slowly (ρ^{2n-1} decays slowly), which weakens the lower bound through $\log(1/\rho)$, while the upper bound pays $\log B(\rho)$, which

stays bounded away from zero, for the conditioning of polynomial approximation on $[0, \rho]$. Logarithmic rates of this kind are familiar from deconvolution with supersmooth errors (Fan, 1991) and from moment methods for mixing distributions (Kong and Valiant, 2017; Wu and Yang, 2020); here they arise from the closed loop of a market.

Table 18 checks the construction and the estimator. The pairs of Lemma 6.8 have exactly n atoms, agree in all moments below order $2n - 1$ to 10^{-61}, and remain at Wasserstein distance about $0.4/n$, while the distance between their kernels shrinks by a factor of about 0.2 per atom at $\rho = 0.9$. The linear-programming estimator applied to a continuous memory distribution has error times $\log(N/\sigma^2)$ between 0.18 and 0.28 for N from 10^4 to 10^{24}: a 10^{20}-fold increase in data halves the error. The exponential-memory model of the previous subsections, where the cost is polynomial in $1/\Delta\eta$, is therefore the optimistic case: relaxing the parametric form makes identification from prices much harder, not easier.

6.8 Passive data

Theorem 6.10 (Passive identification: minimum phase)**.** *In a stable economy with $0 < a < 1$, distinct $0 < r_i < 1$ and nonzero merged weights, the zeros of H are the memory rates r_i together with a structural zero at $z = 0$, and H is minimum phase. Degenerate anchor and zero-rate cases require the cancellations described in the merged-equivalence proof. With a single price innovation, the price spectrum identifies H up to scale, hence the memory rates; with κ and a known, as in Theorem 3.1, it identifies the merged measure.*

The result holds for signed weights (mean-reverting cohorts): in 1,733 stable random economies with weights of both signs, the nonstructural zeros equalled the memory rates and lay inside the unit disk. With a second independent shock in the cohorts' memories, the Jacobian of the price spectrum with respect to six model parameters has full rank, with singular values from 34.51 down to 0.0698. Passive noise helps identification rather than hindering it.

Theorem 6.11 (Passive cost)**.** *Suppose the price is observed passively as the stationary output $x_t = H_\theta(L)\xi_t$ of i.i.d. Gaussian innovations $\xi_t \sim N(0, \sigma^2)$, $\sigma^2 > 0$ known, without observation noise, for θ in a compact, uniformly Schur-stable family with rates in $[\eta_{\min}, \eta_{\max}] \subset (0, 1)$, and let $D(\theta_1 \| \theta_0)$ be the Kullback–Leibler divergence rate per observation.*

(i) For a single cohort at $\bar{\eta}$ against the symmetric pair at $\bar{\eta} \pm \delta$ with the same total weight, $D = \Theta(\delta^4)$.

(ii) For the moment-matched two-cohort economies of Theorem 6.2, $D = \Theta(\tau^2(\Delta\eta)^6)$.

(iii) For every $K \geq 2$, for the K-cohort Gauss and Gauss–Radau economies of Theorem 6.4, $D = \Theta((\Delta\eta)^{4K-2})$.

Hence any test based on T passive observations with worst-case error at most $\alpha < 1/4$ needs $T = \Omega(\delta^{-4})$, $T = \Omega(\tau^{-2}(\Delta\eta)^{-6})$ and $T = \Omega((\Delta\eta)^{-(4K-2)})$ respectively, uniformly as the separation shrinks: passive data cost at least the same polynomial order as active probing for every number of cohorts. For each fixed pair, at fixed type-I error the likelihood-ratio test's type-II error decays at exponential rate D (Stein's lemma), and both errors decay at the Chernoff-information rate, which is of the same order in the separation; so the orders are the right ones; we do not claim that this attainment is uniform as the separation shrinks.

In words. Watching the market without trading does not escape the barrier. A passive observer needs at least the same polynomial order of observations as a prober needs experiments; it pays in time rather than in trading losses. Prices recorded with extra measurement noise that is independent of the price are a garbling of the noiseless record, so by the data-processing inequality they carry less information, and the lower bound holds a fortiori.

The proof (Appendix A) combines the Whittle form of the divergence rate with the resolvent identity $H_1 - H_0 = z^{-1} H_1 H_0 (G_1 - G_0)$: the two spectra differ by the leading unmatched moment, of order δ^2 for the pair and $\tau(\Delta\eta)^3$ for the moment-matched economies, and the divergence rate is the square of that difference. Numerically, the fitted slopes of the

required sample length are -4.21 and -5.93, against the theoretical -4 and -6; for part (iii) the fitted slopes of the Whittle divergence rate in $\Delta\eta$ are 5.91, 9.98 and 14.05 for $K = 2, 3, 4$, against 6, 10 and 14 (`hardening/p1_checks.py`).

Passive surveillance is possible in principle but needs at least the same polynomial order of observations as active probing needs experiments. Proposition C.2 is the exact finite-sample version for the known pair.

7 Robustness to nonlinear and adaptive trading rules

The theorems of Sections 3–6 use the linear exponential-memory recursion (2). Real desks cap positions, trade only beyond thresholds, adapt their gains and learn online. A small probe of size A moves the state by $O(A)$, so any rule whose order is a twice-differentiable function of an exponential-memory state is linear to first order and differs from it by $O(A^2)$; the theorems should then hold for small probes and fail for large ones. We check this, and locate where it fails, in simulations of five stylized rules at $a = 0.5$, $\kappa = 0.8$, total weight $q = 0.2$ and mean rate $\bar{\eta} = 0.4$, over 240 periods after a pulse probe.

The rules are: (1) *momentum with a position limit*, the cohort's order a smooth saturating (tanh) function of its memory state $y_{i,t}$, normalized to unit slope at the operating point; (2) a *threshold rule* that trades only when $|y_{i,t}| > \tau = 0.05$, in two versions, an exact dead zone and one averaged over background state variation of standard deviation 0.04; (3) an *adaptive-gain rule* in the form of recursive least squares with forgetting factor $1 - \eta_i$, whose gain falls when the regressor is large; (4) an *online neural readout*, a linear readout of a fixed two-layer tanh feature map with eight hidden units, trained by regularized stochastic gradient steps of size η_i; and (5) a *tabular value rule* that updates a 21-bin table at rate η_i, in two versions, kernel-smoothed across bins and a single-bin lookup. These are stylized: each is built so that its small-signal dynamics are an exponential filter with rate η_i, and they are not calibrated to any production system.

Two checks follow the theory. *Test A (merged measure):* one cohort with weight q and position limit 1 against two cohorts with weights $q/2$ at the same rate and limits 0.55 and 1.65, so the two economies have the same merged measure; the raw response gap $\|X^{(1)}(A) - X^{(2)}(A)\|_2$ should be $O(A^2)$ (slope 2 in $\log A$) and the probe-normalised gap $\|X^{(1)}(A) - X^{(2)}(A)\|_2/A$ should be $O(A)$ (slope 1 in $\log A$) over $A \in [0.002, 0.064]$. *Test B (δ^{-4} law):* one cohort at $\bar{\eta}$ against the symmetric pair at $\bar{\eta} \pm \delta$, $\delta \in [0.015, 0.12]$, at probe size $A = 0.005$; the normalized response difference $\|\Delta h\|_2$ should scale as δ^2 and the capital lower bound of Theorem 5.2 as δ^{-4}. The saturation threshold A_{sat} is the smallest probe at which $\|\Delta h\|_2$ at $\delta = 0.05$ departs by more than 10% from its small-probe value.

To check that these conclusions do not depend on hand-built rules, we pre-specified a second experiment in which online learners are trained from scratch inside the closed-loop market for $T_{\text{burn}} = 5{,}000$ periods of flow and noise innovations; their learned parameters are saved and reloaded into the probing simulator without rescaling. Four learning rules are used: (i) recursive least squares with forgetting factor $\lambda = 0.995$, learning its return-forecast slope; (ii) a two-layer tanh network with eight hidden units whose readout is trained by regularized online stochastic gradient descent on trading profit; (iii) a scalar policy-gradient rule trained on trading profit; and (iv) a tabular Q-learner with 21 states and 21 actions, TD(0) Bellman backups ($\gamma = 0.90$) and Boltzmann action choice. Each learner acts on an exponential-memory state $y_{i,t}$ with rate η_i: training sets its response gain and curvature, not its memory rate, so the experiment asks whether learned policies preserve the merged-measure and δ^{-4} predictions, not whether learning produces exponential memory. The specification fixed four predictions: (a) a single exponential fits each trained learner's small-probe impulse response with $R^2 \geq 0.90$; (b) two trained populations with the same merged measure differ by at most 1% in relative L_2 response at $A = 0.005$; (c) the capital lower bound to separate trained populations whose memory rates differ by 2δ scales with log–log slope in $[-4.15, -3.85]$; and (d) we report where equivalence and the δ^{-4} law fail when learning continues during the probe.

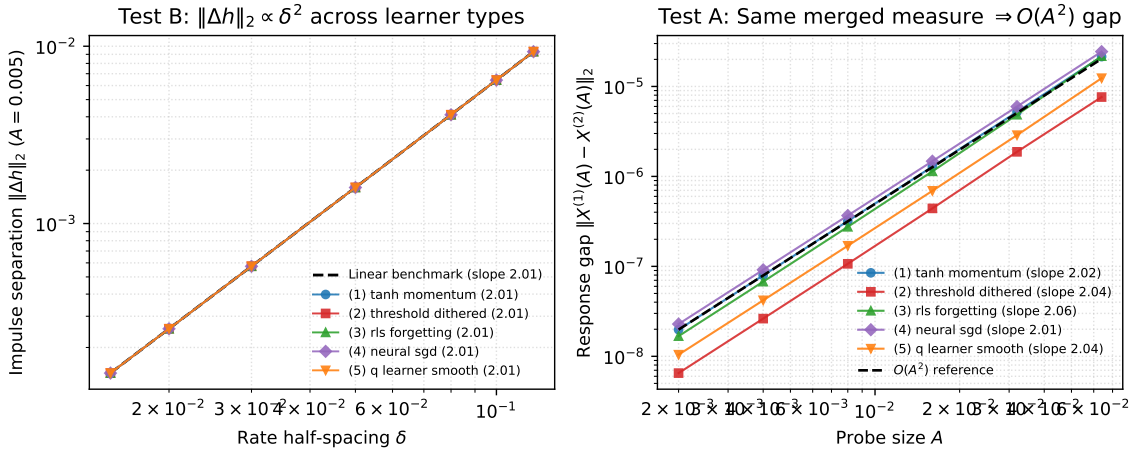


Figure 8: Left: Test B, probe-normalised response difference $\|\Delta h\|_2$ against the rate half-spacing δ at probe size $A = 0.005$, five stylized rules and the linear benchmark. Right: Test A, probe-normalised response gap $\|X^{(1)}(A) - X^{(2)}(A)\|_2/A$ between two economies with the same merged measure against the probe size A, with an $O(A)$ reference line (corresponding to an $O(A^2)$ raw gap).

Table 19 and Figure 8 show both where the linear reduction holds and where it breaks down. In Panel A, for small probes every smooth or dithered rule behaves like the linear model: economies with the same merged measure have raw response gap $O(A^2)$ (slopes 2.011 to 2.064) and probe-normalised gap $O(A)$ (slopes 1.011 to 1.064), while the response difference and capital lower bound scale as δ^2 (slope 2.007) and δ^{-4} (slope -4.001), with capital at $\delta = 0.05$ within 0.1% of the linear 61.42. In Panel B, predictions (a)–(c) pass for all four trained learners: a single exponential explains $R^2 \geq 0.993$ of each small-probe impulse response, populations with the same merged measure differ by at most 10^{-5} in relative response at $A = 0.005$, and capital lower bound scales with slopes -4.014 to -3.973, inside the registered band. The level of capital differs across learners because training sets the response gain.

Panel C shows where this changes. When the two populations are trained on independent innovation streams, prediction (c) fails: the capital slopes are -1.75, -0.97, -1.62 and -0.38. Independently trained populations end with learned gains that differ by sampling error of order $T_{\text{burn}}^{-1/2}$, so they differ in merged weight as well as in memory rate, and a difference in

weight is cheaper to detect than an $O(\delta^2)$ difference in rates. The δ^{-4} law therefore prices the separation of memory rates between populations whose feedback weights match; it is not a claim about telling apart any two trained populations. When learning continues during the probe at rate 0.35, the law survives (slopes -4.02 to -3.98) and the merged-measure gap rises only to 1.1×10^{-3}, so the breakdown we anticipated under prediction (d) did not occur in this design.

The linear theory fails in two further places. First, saturation, adaptive gains and the neural readout leave the small-signal regime once probes exceed $A_{\text{sat}} \in [0.751, 1.271]$ par units, and the dithered threshold rule leaves it already at $A_{\text{sat}} = 0.049$, the threshold itself. Second, rules with hard discontinuities are different in kind: an exact dead zone ($\tau = 0.05$) or a single-bin lookup does not respond at all to probes smaller than 0.139 and 0.149 par units ($\|\Delta h\|_2 = 0$), concealing those cohorts from small probes altogether. Neither smooth nonlinearity nor hard thresholds offers a cheaper route to separating memory rates than the linear benchmark; populations whose learned weights differ are told apart by their weights, not their rates.

8 What disclosure buys

Because active probing is costly and passive data are slow, the natural policy question is what a supervisor gains from disclosure. We compare six observation contracts at $\delta = 0.05$.

- I_0: public prices only.
- I_1: prices and aggregated executed flow.
- I_2: prices and one independent noisy report of dealer marked equity per reset ($\sigma_E = 0.5$).
- I_3: prices and exact audited equity (Proposition 4.4), also with reports rounded to 0.01.

- I_4: algorithm update timestamps without the update values.
- I_5: a supervisory audit registry with access to internal algorithm states.

Table 20 shows that disclosure is not one thing. Timestamps without update values carry no information about δ in this model. Noisy periodic statements barely help. Exact audited equity collapses the requirement to a single trade, limited only by lot sizes and report precision, and an audit registry removes the trading requirement altogether.

The cleanest form of useful disclosure is flow tagged by participant class.

Proposition 8.1 (Tagged flow removes the barrier)**.** *Let the price be driven by i.i.d. innovations of variance σ_ξ^2 in a positive, stable economy, with stationary autocovariances $\gamma_x(s) = \sum_\ell c_\ell \lambda_\ell^{|s|}$ ($c_\ell \geq 0$, $0 \leq \lambda_\ell < 1$ the closed-loop modes). Suppose cohort i has merged weight $w_i > 0$ and the observer receives prices and each cohort's flow $\tilde{u}_{i,t} = u_{i,t} + \varepsilon_{i,t}$, $u_{i,t} = q_i y_{i,t}$, with i.i.d. measurement noise of variance σ_f^2 independent of prices. Then $\tilde{u}_{i,t} = r_i \tilde{u}_{i,t-1} + w_i x_t + e_{i,t}$ with $e_{i,t} = \varepsilon_{i,t} - r_i \varepsilon_{i,t-1}$, and the just-identified IV estimator of $\beta_i = (r_i, w_i)$ with instruments $Z_t = (x_{t-1}, x_t)^\top$ satisfies $\sqrt{T}(\hat{\beta}_i - \beta_i) \to N(0, V_i)$,*

$$V_i = \Sigma_i^{-1} \Omega_i \Sigma_i^{-\top}, \qquad \Sigma_i = \begin{pmatrix} w_i A_0 & \gamma_x(1) \\ w_i A_1 & \gamma_x(0) \end{pmatrix}, \qquad \Omega_i = \sigma_f^2 \big[(1 + r_i^2)\Gamma_0 - r_i(\Gamma_1 + \Gamma_1^\top) \big], \quad (31)$$

where $A_m = \sum_{s \geq 0} r_i^s \gamma_x(s+m)$, $\Gamma_0 = \left(\begin{smallmatrix} \gamma_x(0) & \gamma_x(1) \\ \gamma_x(1) & \gamma_x(0) \end{smallmatrix}\right)$ and $\Gamma_1 = \mathbb{E}[Z_t Z_{t-1}^\top]$. The Jacobian determinant is

$$\det \Sigma_i = w_i \sum_{\ell,m} c_\ell c_m \frac{1 - \lambda_\ell \lambda_m}{1 - r_i \lambda_m} \; > \; 0. \quad (32)$$

V_i depends only on cohort i's own (w_i, r_i), σ_f^2 and the price autocovariances, which are bounded and nondegenerate uniformly on a uniformly stable family with merged weights bounded below, whatever the rate spacing; so the merged weights are estimated at rate $T^{-1/2}$ with a constant that does not depend on $|r_i - r_j|$.

In words. If each group's trading is reported separately, each group becomes an ordinary regression. Its precision depends on the group's own rate and on the price dynamics, but it has no singularity as the groups' rates approach each other.

Each tagged flow follows $u_{i,t} = r_i u_{i,t-1} + w_i x_t$, a regression whose information about w_i does not involve the other cohorts' rates; the closed form (32) matches the simulated Jacobian (0.1330 against 0.1322 in a two-cohort economy with 4×10^5 periods; `hardening/p1_checks.py`). In simulation ($T = 5{,}000$, 60 replications) the root-mean-square error of the recovered weights is flat in $\Delta\eta$ (fitted slope 0.011, error between 5.8×10^{-4} and 6.6×10^{-4}). This is the formal content of a disclosure regime such as a consolidated audit trail: tagging flow by participant class turns a polynomial barrier into an ordinary estimation problem.

8.1 How much disclosure is worth

Proposition 8.1 says that tagging removes the barrier; the next result says by how much.

Theorem 8.2 (Value of class-tagged disclosure)**.** *Take a rich class of Theorem 6.5: rates $r_c + \Delta\eta\, s$ with s in a compact set of γ-separated ordered points that has nonempty interior and contains the Gauss and Gauss–Radau pair of Theorem 6.4, and merged weights in a nondegenerate interval $[\omega_{\min}, \omega_{\max}]$, $0 < \omega_{\min} < \omega_{\max}$. In each reset experiment the observer applies a unit impulse and records the price response with independent $N(0, \sigma^2)$ noise per tap; under tagging it also records each class's flow response $f_{i,t} = w_i \sum_{s \le t} r_i^{t-s} h_s$ with independent $N(0, \sigma_f^2)$ noise per tap. Let N_p and N_t be the numbers of experiments needed, without and with tagging, for a fixed accuracy and confidence. There are constants $0 < c \le C < \infty$, depending on (σ, σ_f) and the class but not on $\Delta\eta$, such that*

(i) to recover the weights to a fixed accuracy smaller than half the weight difference between the Gauss and Gauss–Radau economies (so that recovering the weights decides between them), $c \le N_t \le C$, so $N_p/N_t \asymp (\Delta\eta)^{-(4K-2)}$;

(ii) to recover the rates to accuracy $c_0\Delta\eta$, for a fixed $c_0 > 0$ smaller than the radius of a ball contained in the class of rescaled nodes (so that no fixed guess meets the target), $c(\Delta\eta)^{-2} \le N_t \le C(\Delta\eta)^{-2}$, so $N_p/N_t \asymp (\Delta\eta)^{-(4K-4)}$.

Proof. $N_p \asymp \sigma^2(\Delta\eta)^{-(4K-2)}$ by Theorems 6.4 and 6.5. *Upper bounds.* Average the N price responses to estimate h with error of order $\sigma N^{-1/2}$ per tap. Given h, class i's mean flow is a two-parameter family $F(w_i, r_i; h)$ whose Jacobian columns $\sum_s r_i^{t-s} h_s$ and $w_i \sum_s (t-s) r_i^{t-s-1} h_s$ are linearly independent, with a Gram determinant bounded below uniformly on the class because $w_i \ge \omega_{\min}$, the rates lie in a compact subset of $(0, 1)$ and $h_0 = \kappa > 0$; the bound does not involve the other classes or $\Delta\eta$. Least squares on class i's flows with the plug-in h therefore estimates (w_i, r_i) with error $O_P(N^{-1/2})$, uniformly on the class, which gives the upper bounds in (i) and in (ii) for accuracy $\epsilon\Delta\eta$. For (i) there is also a finite-sample version: in the first period the tagged flow of class i is $w_i\kappa$, so the mean of the N first-period records gives $P(|\hat{w}_i - w_i| \ge \epsilon) \le \sigma_f^2/(N\kappa^2\epsilon^2)$ by Chebyshev's inequality, for every configuration (this step is machine-checked). *Lower bounds.* At a configuration in the interior of the class, perturb one class's weight by ϵ, or its rate by $\epsilon\Delta\eta$; for small ϵ both perturbed configurations remain in the class because the weight interval is nondegenerate and the rate set has nonempty interior. The mean differences of both the price and the flow records are of order ϵ, respectively $\epsilon\Delta\eta$, so the Kullback–Leibler divergence of N experiments is at most $CN\epsilon^2$, respectively $CN\epsilon^2(\Delta\eta)^2$, and Le Cam's two-point lemma gives the lower bounds. □

In words. This is the value of disclosure measured in data. Tagging replaces a requirement that explodes as the groups become similar with one that does not grow at all (for the weights) or grows only like $(\Delta\eta)^{-2}$ (for the rates).

Proposition 8.1 and Theorem 8.2 use different observation arrangements: the proposition concerns passive, stationary data with prices observed without noise, where the price is a valid instrument; the theorem concerns reset experiments with noisy prices. With noisy prices, regressing flow on the observed price is an errors-in-variables problem, which is why

the theorem works from the known impulse instead.

Proposition 8.3 (Robustness to strategic evasion: Sybil splitting, camouflage, and chilling)**.** *Consider the class-tagged disclosure regime of Theorem 8.2 and Proposition 8.1. Let trader group i attempt to evade supervisory identification through one of three strategic responses:*

(i) Sybil account-splitting. *Suppose group i divides its trading volume across $M_i \geq 1$ distinct pseudonymous accounts or sub-entities, each executing a fraction $\alpha_{i,m} > 0$ ($\sum_{m=1}^{M_i} \alpha_{i,m} = 1$) with the same underlying memory rate r_i and flow $u_{i,m,t} = \alpha_{i,m} u_{i,t} + \varepsilon_{i,m,t}$, where $\varepsilon_{i,m,t}$ is i.i.d. reporting noise with variance $\sigma^2_{f,m}$. Then for each sub-account m, the instrumental-variables Jacobian determinant satisfies*

$$\det \Sigma_{i,m} = \alpha_{i,m} w_i \sum_{\ell,k} c_\ell c_k \frac{1 - \lambda_\ell \lambda_k}{1 - r_i \lambda_k} > 0, \tag{33}$$

which is bounded strictly away from zero for any non-vanishing sub-account ($\alpha_{i,m} \geq \underline{\alpha} > 0$). Group i's total weight $w_i = \sum_{m=1}^{M_i} \hat{w}_{i,m}$ and common persistence rate r_i are identified at the standard parametric rate $T^{-1/2}$, with no polynomial degradation in $\Delta\eta$.

(ii) Strategic camouflage and noise injection. *Suppose group i injects uninformative pseudo-random trading noise $\zeta_{i,t}$ with variance σ^2_ζ into its orders to obscure its regression signature. To lower the signal-to-noise ratio sufficiently to restore the price-only super-resolution complexity $(\Delta\eta)^{-(4K-2)}$, group i must inject noise variance satisfying*

$$\sigma^2_\zeta = \Omega\Big((\Delta\eta)^{-(4K-2)}\Big). \tag{34}$$

By Lemma 4.2, trading this camouflage volume against a linear order book incurs a cumulative per-period quadratic execution loss of at least $\frac{1}{2}\kappa\sigma^2_\zeta = \Omega((\Delta\eta)^{-(4K-2)})$, which strictly dominates any finite expected trading profit $\Pi_i < \infty$, rendering strategic camouflage economically unviable.

(iii) Intensity chilling. *If group i strategically curtails its trading intensity by scaling w_i*

down to $\gamma_i w_i$ with $\gamma_i \in [\underline{\gamma}, 1]$ for some $\underline{\gamma} > 0$, the instrumental Jacobian determinant (33) scales proportionally by $\gamma_i \geq \underline{\gamma} > 0$, preserving $T^{-1/2}$ identification. If group i chills its activity to zero ($\gamma_i \to 0$), its contribution to market instability S is eliminated, achieving the supervisor's prudential objective without enforcement.

In words. A firm cannot dodge class disclosure by splitting into shell accounts, because each sub-account still obeys the same regression and their estimated weights add back up. It cannot hide behind injected noise trades, because trading enough random volume to fake the price-only barrier costs far more in execution fees than any algorithm can make. And if it chills its trading to avoid detection, it shrinks its market impact and cures the fragility problem on its own. Full proof is in Supplement S1.

With the same noise and lags for both observers, the Cramér–Rao requirements give fitted slopes of -6.08 for prices and -4.34 for the ratio at $K = 2$ (theory -6 and -4). Calibrated to the class memories estimated below, prices alone would need about 333 times as many days as the published class ratios to separate the top-trader and all-account memories in BTCUSDT (rates 0.9889 and 0.9986 per five minutes; about 6,000 days against 18), and about 19 times as many in ETHUSDT (rates 0.9793 and 0.9972). These rates come from day-demeaned regressions, which bias persistence down (Section 8.5), so the factors are illustrative.

8.2 Disclosure in dollars

This subsection is a hypothetical scale illustration of an outside prober's shadow loss, not a public cost–benefit calculation.

Theorem 8.2 states the value of class-tagged disclosure in samples. Theorems 5.2 and 6.4 let us translate that value into a flow of round-trip trading losses for an outside participant. Two adjustments make the figures exact:

(i) *Exact K-cohort moment-gap constant.* Rather than extrapolating from Cramér–Rao

ratios, we evaluate the spectral round-trip trading loss bound of Theorem 5.2,

$$L^*_{\mathrm{rt}}(K, \Delta\eta) \;=\; \frac{4\sigma^2 z^2_{1-\alpha}}{\sup_\omega\Big(|\Delta\hat{H}_{K,\Delta\eta}(e^{i\omega})|^2/g_0(\omega)\Big)},$$

directly for the K-point Gauss versus Gauss–Radau economies of Theorem 6.4 and Table 14, whose moments below order $2K-1$ coincide and whose $(2K-1)$-th moment gap has the exact constant $\Delta_2 = 0.222222$, $\Delta_3 = 0.053333$, and $\Delta_4 = 0.013061$ at spacing $\Delta\eta$ (with $\Delta\eta = 2\delta$ for the symmetric known pair).

(ii) *Flow against flow.* A single probe loss L^*_{rt} is a one-time certification cost, whereas an audit trail is an ongoing annual expenditure. Because algorithm populations turn over, we report the annualised round-trip trading loss $52 \times L^*_{\mathrm{rt}}$ under weekly renewal (matching the CFTC weekly reporting cycle of Section 8.4). For orientation only, we also record official audit-trail figures of two different kinds; we do not divide one by the other, because a prober's trading loss is a private transfer to other traders calibrated to a Bitcoin perpetual, while an audit trail's budget is a resource cost for equities and options. The first is CAT LLC's own operating budget: $248.8 million for 2025 in the November 2024 filing (SEC Release No. 34-101901, Dec. 12, 2024, p. 4), revised to $228.3 million in May 2025 and reduced further later in 2025, and the revised 2026 Financial and Operating Budget of 31 March 2026, with total expenses of $147,044,869, which we use (CAT LLC, 2026). The second is historical, ex ante estimates of industry-wide compliance costs, which are not operating budgets and cover a different scope: the SEC's 2016 CAT NMS Plan Approval Order estimate of $2.4 billion in initial industry build costs and $1.7 billion per year in ongoing industry costs (SEC, 2016, Release No. 34-79318, Nov. 15, 2016, 81 FR 84696 at 84801, p. 363); and the SEC's Rule 613 proposing and adopting releases estimating $4.0 billion initial and $2.1 billion annual ongoing costs (SEC, 2012, Release No. 34-62174, May 26, 2010, 75 FR 32556 at 32598, 32600, pp. 147, 153; Release No. 34-67457, July 18, 2012, 77 FR 45722 at 45790–45795,

pp. 220–238).

The figures are a shadow cost of learning by trading, the price an outside participant would pay; a supervisor without an audit trail does not probe, and its cost is detection delay, which Proposition 8.4 models. The cost results require a positive economy: with signed weights the execution form $\frac{1}{2}v^\top T_h v$ can be indefinite and a round trip can earn money, so none of the dollar figures applies there. Probing cost is a lower bound on the cost of learning from prices by trading, not on passive learning and not a social cost. It is also a lower bound for deterministic probes only: sequential probing can lower the expected loss by a factor of up to 0.49 (Theorem 5.3). Only the observation noise is estimated from data; the structural parameters are illustrative, and the constant is sensitive to them. At $\delta = 0.05$ the numerically evaluated lower bound of Theorem 5.2 is 358.8, 179.0, 61.1 and 13.9 par units for $a = 0$, 0.25, 0.5 and 0.7, and 3.9, 19.6, 61.1 and 305.2 for $\bar{\eta} = 0.2$, 0.3, 0.4 and 0.6; at $\Delta\eta = 0.20$ instead of 0.10 the unknown-pair figures fall by a factor of about 64. The Gauss–Radau pair is used in the orientation that gives the larger, still valid, lower bound. The known-pair rows span a rate width of $2\delta = \Delta\eta$, whereas the $K = 2$ Gauss–Radau rows on $r_c + \Delta\eta[-1, 1]$ span $\frac{4}{3}\Delta\eta$. Within that interpretation, Table 21 shows how rapidly the polynomial exponent $(\Delta\eta)^{-(4K-2)}$ and the shrinking Gauss–Radau moment-gap constant Δ_K turn the comparison in favour of disclosure. For a known pair at coarse separation $\Delta\eta = 0.20$ ($\delta = 0.10$), weekly renewal of price-only probing costs \$0.32 million per market per year at Roll noise and \$7.79 million at one-second noise. At $\Delta\eta = 0.10$ ($\delta = 0.05$) the unknown $K = 2$ Gauss–Radau pair costs \$14.74 million (Roll) or \$356.44 million (one-second) per market per year; with $K = 3$ unknown cohorts at $\Delta\eta = 0.20$ the figures rise to \$41.83 million and \$1.01 billion, and at $K = 3$, $\Delta\eta = 0.10$ or $K = 4$, $\Delta\eta = 0.20$ they run to billions. For orientation only, CAT LLC's revised 2026 budget is \$147.0 million; as explained above, the two numbers measure different things. The renewal frequency is an assumption: annual figures scale with it, so monthly renewal divides them by 4.3 and renewal on each of 252 trading days multiplies them by 4.8.

8.3 When disclosure is worth acquiring

When is disclosure worth its cost? The dollar illustration of Section 8.2 prices disclosure; it does not say when a supervisor should buy it. Remark 6.3 shows that a supervisor who only needs to bound aggregate price risk does not need disclosure at all: the summed response is identified from public prices at the ordinary rate. Attribution matters when supervision is selective: one algorithm must be stopped or sanctioned while the others keep supplying liquidity, and the law requires attributing the conduct to it at a given error level. A supervisor without an audit trail does not run trading probes; it waits for public data to accumulate, and the cost of waiting is the harm done in the meantime.

Consider M markets, a supervisory cycle of $\bar{T}$ periods, and in each market one algorithm whose conduct does harm $\ell > 0$ per period until it is attributed. With prices alone, attribution at the required error level takes $T_p(\Delta\eta) = c_p(\Delta\eta)^{-(4K-2)}$ periods; with class-tagged records it takes $T_t(\Delta\eta) = c_t(\Delta\eta)^{-2}$ periods, the orders of Theorems 6.11 and 8.2 with constants $c_p, c_t > 0$. We take the attribution times to be of exactly these orders, which is a modelling assumption: Theorem 6.11 gives only a lower bound for prices, and if price-only attribution is slower, G below is larger and disclosure is chosen at least on the band the proposition describes. $T_t \propto (\Delta\eta)^{-2}$ is the rate target; for weights at fixed accuracy T_t does not grow as $\Delta\eta$ shrinks, and the lower edge of the band disappears. Attribution that would take longer than the cycle does not happen within it. An audit trail costs F per cycle to operate and c per market to comply with, including any information leakage. The supervisor's expected loss is therefore $\mathcal{L}_{\text{price}} = M\ell \min\{\bar{T}, T_p\}$ without disclosure and $\mathcal{L}_{\text{disc}} = F + Mc + M\ell \min\{\bar{T}, T_t\}$ with it. Write $\Delta\eta_t = (c_t/\bar{T})^{1/2}$ and $\Delta\eta_p = (c_p/\bar{T})^{1/(4K-2)}$ for the spacings at which tagged and price-only attribution just fit within the cycle.

Proposition 8.4 (When an audit trail is chosen)**.** *Let $K \geq 2$ and $\Delta\eta_t < \Delta\eta_p$, and let $G(\Delta\eta) = \min\{\bar{T}, T_p(\Delta\eta)\} - \min\{\bar{T}, T_t(\Delta\eta)\}$ be the attribution delay that disclosure saves.*

(i) The supervisor strictly prefers disclosure if and only if $F/M + c < \ell\, G(\Delta\eta)$.

(ii) *G vanishes for $\Delta\eta \le \Delta\eta_t$, increases on $(\Delta\eta_t, \Delta\eta_p]$ to its maximum $G^* = \bar{T} - T_t(\Delta\eta_p) > 0$, decreases wherever it is positive beyond $\Delta\eta_p$, and is negative for $\Delta\eta > (c_p/c_t)^{1/(4K-4)}$, where it returns toward zero from below. Hence the spacings at which disclosure is chosen form an interval around $\Delta\eta_p$, possibly empty: disclosure is never chosen for algorithms so similar that even tagged records cannot attribute them within the cycle, nor for algorithms so different that prices attribute them quickly.*

(iii) *The interval is non-empty if and only if the audit trail covers more than $M^* = F/(\ell G^* - c)$ markets (with $M^* = \infty$ if $\ell G^* \le c$), and it widens as M grows.*

Part (i) restates the comparison of losses. Part (ii) says that disclosure is valuable in a middle band, where algorithms are close enough that prices cannot attribute them within the cycle but far enough apart that tagged records can; the polynomial gap between the two delays, $(\Delta\eta)^{-(4K-2)}$ against $(\Delta\eta)^{-2}$, is what makes the band wide. Part (iii) is an economies-of-scale result: an audit trail with a fixed operating cost pays only when it spans enough markets. This illustrates a possible economy of scale for market-wide reporting; the cost function is stylised, not estimated from CAT, and the result neither explains CAT's legal architecture nor shows that it is optimal. A supervisor that can trade has a second option, active probing, whose cost is the frontier of Table 21; because that cost grows at the same polynomial order, it shortens price-only attribution only for well-separated algorithms and leaves the comparison for close ones unchanged. The model takes the harm ℓ as given and holds the algorithms' behaviour fixed across regimes; Section 12 returns to both.

A decision in which attribution matters. A supervisor chooses between a uniform reduction in activity and halting the largest verified class, over a declared supervisory horizon. The target below is the feedback budget $S = a + \kappa \sum_i q_i$, not the spectral radius of the state matrix. In an already stable positive economy, imposing $S \le \bar{s} < 1$ is a tighter prudential constraint, not restoring stability to an unstable market.

Proposition 8.5 (Supervisory separation: market-wide versus targeted intervention)**.** *Assume the positive stable model, known a, κ, and $a < \bar{s} < 1$. Write $q_i = w_i/(1 - r_i)$ and $R = S - \kappa \max_i q_i$ for the post-halt feedback budget.*

(i) A uniform scale θ meets the ceiling exactly when $a + \theta(S - a) \le \bar{s}$. Its largest feasible value is $\theta^* = \min\{1, (\bar{s} - a)/(S - a)\}$. The identity $S = 1 - \kappa/\sum_j h_j$ is algebraic; estimating this infinite sum requires a justified tail bound and an observation/noise model. It alone establishes no unconditional parametric learning rate.

(ii) For $K = 2$, the Gauss and Gauss–Radau pair of Theorem 6.4 has limiting largest normalized weights $1/2$ and $3/4$. Thus $|S^G - S^R| = O((\Delta\eta)^3)$ and $q^R_{\max} - q^G_{\max} = W/(4(1 - r_c)) + O(\Delta\eta)$. For sufficiently small spacing, a nonempty interval of ceilings separates their post-halt decisions. For N independent reset records with common Gaussian noise $\sigma^2 I$, both decision errors at most $1/4$ require

$$N \ge (2\Phi^{-1}(0.75)\sigma/\|h^G - h^R\|_2)^2 \ge c\sigma^2(\Delta\eta)^{-6}.$$

The record horizon and the response norm must be the same on both sides.

(iii) Tagged-flow estimation under Theorem 8.2, a rate bound away from one, and a fixed nonzero decision margin gives a bounded sample requirement as the spacing shrinks. The target class must be mapped to an executable, verified class label; a largest latent atom alone supplies no such identity.

Proof. Part (i) follows by scaling the intensities. For (ii), expand $(1 - r_c - \Delta\eta s)^{-1}$: the first three moments agree, so aggregate budgets differ by order $(\Delta\eta)^3$. Each intensity converges to $W\omega_i/(1-r_c)$, giving the stated dominant-weight gap. This exceeds the aggregate discrepancy at sufficiently small spacing. A ceiling between the two post-halt budgets reduces the action problem to testing two Gaussian mean vectors. Their common-covariance testing error is $\Phi(-\sqrt{N}\|h^G - h^R\|_2/(2\sigma))$; the response bound in Theorem 6.4 gives the lower bound. Part

(iii) follows from continuity of $w/(1-r)$ on the declared compact parameter set and the fixed margin. □

General-K scope. The same reduction gives exponent $4K-2$ only for a pair whose price responses have that moment-matching order *and* whose dominant intensities have a separately established fixed gap. Moment matching alone does not supply the latter. In particular, the symmetric odd/even binomial pair has equal maximal normalized weights. The $K=3$ rows below are numerical examples, not a construction proving this condition for every K.

Formal verification scope, 5 November 2026. This revision builds the pinned Lean 4.34.1/Mathlib project and audits its typed declarations using only standard foundations. The exact two-point Gauss–Radau witness in Table 22 is now covered by `Paper1GaussRadau.lean`: positive normalized weights, moments of orders zero through two, the third-moment gap, a dominant-intensity gap of at least $W/(8u)$ for $0 \le \Delta\eta \le u/3$, the aggregate-budget gap and the actual finite-Gaussian-record testing bound. Thus this example no longer relies solely on the other matched pair used by the earlier supervisory theorem. The finite-horizon bound has constants depending on the declared record length and response envelope. Its proof does not supply an all-K quadrature construction. `Paper1SupervisoryDecision.lean` additionally checks the robust threshold/fallback, the actual unsafe-halt event for Proposition 8.6, the fixed-menu no-data optimum and the acquisition algebra of Proposition 8.7. The new receipt binds the build to hashes of the active Paper 1 sources. Declaration counts include supporting results across the project and are not counts of fully formalized manuscript statements. Independent human review, novelty and empirical market assumptions are not certified by Lean.

Explicit robustness condition. Smooth best responses alone do not preserve separation: the constant smooth map $\Phi(M)=0$ sends both post-halt budgets to a. Instead assume that,

over the stated action and holding horizon, $|\widetilde{R}_j - R_j| \le \varepsilon$ for $j \in \{G, R\}$, with $R_R < R_G$. Then

$$R_R + \varepsilon \le \bar{s} < R_G - \varepsilon \quad \Longrightarrow \quad \widetilde{R}_R \le \bar{s} < \widetilde{R}_G. \tag{35}$$

Indeed, $\widetilde{R}_R \le R_R + \varepsilon$ and $\widetilde{R}_G \ge R_G - \varepsilon$. A nonempty robust window exists if $2\varepsilon < R_G - R_R$. This is a maintained response bound, not an equilibrium result. The same allowance must cover the alternative uniform action. At the $K = 2$, $\Delta\eta = 0.10$ calibration, $R_R = 0.532$, $R_G = 0.569909$, and $\bar{s} = 0.55$ has margin 0.018; an allowance 0.005 fits within it.

A complete conditional cost-and-error model. Let j denote the economy, π_j its declared prior, H the holding horizon, b the surplus per active class per period, C_h the implementation cost of halting and L the penalty per period for violating the ceiling. These are assumed accounting primitives, not estimated social welfare. An envelope $S_{\max}$ supports the robust uniform action

$$\theta_\varepsilon = \min\{1, \max\{0, (\bar{s} - \varepsilon - a)/(S_{\max} - a)\}\}, \qquad L_T = HbK(1 - \theta_\varepsilon).$$

This action is feasible only if its worst-case budget meets the ceiling (in particular $a + \varepsilon \le \bar{s}$). For a halt,

$$L_{h,j} = Hb + C_h + HL\,\mathbf{1}\{R_j + d_j > \bar{s}\}, \quad \mathcal{L} = F + Nc + \sum_j \pi_j\{p_j L_{h,j} + (1 - p_j)L_T\}.$$

Here $F + Nc$ is the information bill, p_j the probability of halting, and $|d_j| \le \varepsilon$ the action response. Require $\Pr_j(\text{halt and violate}) \le \alpha$ separately in each world. Compare this loss with the best feasible no-purchase lottery over the *same two actions*, and allow refusing to buy information. Neither arbitrary alternative policies nor an optimum over sample size are claimed. Active probing additionally incurs its execution loss and capital requirement; these are not automatically the per-record fee c and must not be counted twice.

Existing runs and the remaining synthetic check. The archived `R_P1_targeted_intervention.json` supplies six response/decision rows. Recomputing them independently from the stated state recursion agrees to 2.1×10^{-17} in the response discrepancy. Those rows supply S_j, R_j and Gaussian testing distances; they contain no measured disclosure fees, intervention payoffs or strategic response bounds. The existing GPU and historical diagnostics cannot fill these quantities by relabelling their outputs. The new CPU study `studies/supervisory_release_20261105/run_study.py` therefore evaluates only this missing accounting layer, with exact Gaussian probabilities rather than a new large simulation.

For this check alone, assume independent scalar disclosures $Z_n \sim N(R_j, \sigma_t^2)$ and halt when $\bar{Z} + z_{1-\alpha}\sigma_t/\sqrt{N} + \varepsilon \leq \bar{s}$. This is an illustrative measurement contract, *not* a claim that class positions or tagged flow directly observe post-halt budgets. If a world is unsafe under the response allowance, this gate falsely halts with probability at most α under the stated unbiased Gaussian law. The price comparator observes the full 400-tap Gaussian reset record of the same known pair, with a likelihood-ratio rule constrained by the same per-world error budget. Both information channels have the same assumed monetary fee here; neither represents free passive observation.

With $K = 2$, $N = 32$, $\sigma_t = 0.02$, $\varepsilon = 0.005$, $\alpha = 0.05$, $\pi_R = 1/2$, $H = b = 1$, $C_h = 0.02$, $L = 10$, $F = 0.05$, and $c = 0.0005$, the uniform loss is 1.449219, the information bill is 0.066, and disclosure loss is 1.305132. Its conditional net advantage is 0.144087 surplus units. Halt probabilities are approximately zero in G and 0.978928 in R. Buying the price record yields loss 1.515219 at this sample size; its optimal acquisition decision is to decline and retain the no-data action. An informed oracle before information fees has loss 1.234609.

Across 2,304 deterministic configurations varying fees, noise, horizon, prior, sample size and allowance, paid disclosure improves on the no-data benchmark in 1,207, worsens it in 1,079, and ties in 18. These counts have no population or statistical-significance interpretation. Negative cases remain in the output. A downward measurement bias of 0.03 produces

false-halt probability 0.418776; a future response of 0.025 despite an allowance of 0.005 makes halting unsafe in R with probability 0.978928. The latter also invalidates the uniform-action guarantee, so payoff comparisons under that misspecification are not presented as a fair safety comparison. The experiment establishes conditional accounting and failure modes, not real-market gains, a validated equilibrium response, or the general-K theorem.

Proposition 8.6 (A robust scalar disclosure gate)**.** *Let $Z_1, \dots, Z_N$ be independent $N(R, v)$ readings of the fixed post-halt feedback budget, $N > 0$, and let the realized action response d satisfy $d \le \varepsilon$. For $t > 0$, halt only if $\bar{Z} + t + \varepsilon \le \bar{s}$. Then the probability of halting while the realized budget exceeds the ceiling is at most*

$$\Pr\{\bar{Z} + t + \varepsilon \le \bar{s},\ R + d > \bar{s}\} \le \frac{v}{Nt^2}.$$

Consequently $v \le \alpha N t^2$ gives a per-world error bound α.

Proof. On an unsafe-halt event, $\bar{Z} + t + \varepsilon \le \bar{s} < R + d \le R + \varepsilon$, so $t \le |\bar{Z} - R|$. Chebyshev's inequality for the actual Gaussian sample mean gives the bound. Lean checks the product Gaussian law, the sample statistic, this event inclusion and the probability inequality in `p1_supervisory_gaussian_unsafe_halt`; it is not merely a check of a scalar named "probability bound." The response allowance and the interpretation of Z remain explicit assumptions. □

This conservative gate uses $t = \sigma_t/\sqrt{N\alpha}$. It differs from the sharper Gaussian-quantile gate above, whose exact normal probabilities are numerically evaluated but whose quantile identity is not formalized here. Under the same calibration with $N = 128$, the conservative gate has radius 0.007906, information bill 0.114, loss 1.349034 and conditional net advantage 0.100185. Its halt probabilities are approximately zero in G and 0.998023 in R. Across the same 2,304 sensitivity configurations it improves on the fixed-menu no-data benchmark in 781, worsens it in 1,499 and ties in 24. Both gates' rows remain in the result file: 4,608

configurations in total. The welfare and error interpretation remains conditional on the declared observation and response contracts.

Proposition 8.7 (No-data action and information purchase)**.** *For the fixed halt/throttle menu, let* $\overline{L}_h = \sum_j \pi_j L_{h,j}$ *and let* $p_{\max} = \alpha$ *if a halt is unsafe in at least one world, otherwise* $p_{\max} = 1$*. Assuming the fallback meets the ceiling, the best constant no-data halt probability is* $p = 0$ *when* $\overline{L}_h \geq L_T$*, and* $p = p_{\max}$ *when* $\overline{L}_h < L_T$*. A specified information rule with expected action loss* L_I *is strictly preferred if and only if*

$$F + Nc < \mathcal{L}_0 - L_I.$$

Allowing refusal to purchase gives loss $\min\{\mathcal{L}_0, F + Nc + L_I\} \leq \mathcal{L}_0$.

Proof. The no-data loss is affine in its feasible halt probability: $L_T + p(\overline{L}_h - L_T)$. Its slope determines the minimizing endpoint. Subtracting the information rule's total loss gives the break-even inequality, and the minimum inequality proves the optional-purchase statement. These algebraic decisions are machine checked in `Paper1SupervisoryDecision.lean`. They do not establish an optimum over other actions, sample sizes, or information contracts. □

What the data can and cannot test. The theorems concern experiments that no outside researcher can run: reset initial states, known price persistence and impact, and either controlled probes or stationary data from a known innovation law. Public data meet none of these exactly. Order flow is endogenous, the fundamental value is unobserved and replaced by returns, observations overlap in time, parameters drift, and the class reports are positions aggregated over many accounts rather than tagged flow. The tests below therefore check the theory's observable prediction, that class records restrict persistence which prices leave unrestricted, under a stated mapping from class positions to cohort states. They do not estimate the theorems' rates, and a pass does not show that class persistence is learning memory.

8.4 Mandated disclosure: the CFTC's trader-class reports

The stock–flow inventory bridge. To connect the theoretical flow recursion (2) and inventory rule (Remark 2.2) to the empirical position stocks $n_{c,t}$ reported in regulatory filings, recall from Remark 2.2 that cohort k's inventory $I_{k,t} = (1-\delta_{\text{inv}})I_{k,t-1} + u_{k,t}$ under optimal trading $u_{k,t} = q_k y_{k,t} - \psi_{\text{inv},k} I_{k,t-1}$ and belief update $(1-r_k L)y_{k,t} = \eta_k x_t$ satisfies the coupled recursion

$$I_{k,t} = \rho_{\text{inv},k} I_{k,t-1} + q_k y_{k,t} + \nu_{k,t}, \qquad (1-\rho_{\text{inv},k}L)(1-r_k L)I_{k,t} = w_k x_t + (1-r_k L)\nu_{k,t}, \quad (36)$$

where $\rho_{\text{inv},k} \equiv (1-\delta_{\text{inv}}) - \psi_{\text{inv},k} = \lambda_{\text{tc},k}(1-\delta_{\text{inv}})/(\lambda_{\text{tc},k} + \gamma_{\text{inv},k}) \in [0,1)$ is the inventory-turnover root and $r_k = 1-\eta_k \in (0,1)$ is the belief-memory root. Equation (36) has two immediate empirical implications. First, a class position stock $n_{c,t}$ carries two autoregressive roots: the inventory-adjustment root $\rho_{\text{inv},c}$ (which approaches one when inventory holding costs $\gamma_{\text{inv},c}$ are small relative to execution costs $\lambda_{\text{tc},c}$) and the belief-memory root r_c. A one-lag autoregression $n_{c,t} = \alpha_c + r_c^{\text{pos}} n_{c,t-1} + w_c \,\text{ret}_t$ therefore estimates the dominant persistence of the class's position state—which prices alone cannot bound—while separating the faster belief root r_c from slow turnover $\rho_{\text{inv},c}$ requires either the two-lag specification $(1-\phi_1 L - \phi_2 L^2)n_{c,t}$ or first-differenced class flows $\Delta n_{c,t}$, both of which we test below.

The cleanest test uses disclosure mandated by law. The U.S. Commodity Futures Trading Commission requires large traders in regulated futures markets to report their positions and publishes a weekly snapshot of Tuesday positions, usually on Friday, aggregated by reporting category: leveraged funds, asset managers/institutional investors and dealers/intermediaries in financial futures, and managed money, producers/merchants/processors/users and swap dealers in commodities. The categories reflect each trader's reported business, not the algorithm that traded. These are published, aggregated class positions, not class-tagged flow. We pre-specified a test (SHA-256 `6f1d6f35ae88...`) before downloading any data, in eight markets that span asset classes: CME Bitcoin and Ether, the E-mini S&P 500, 10-year

Treasury notes, the euro, the yen, WTI crude oil and Henry Hub natural gas, using every weekly report through October 2026 (282 to 1,039 weeks per market). For each class, $n_{c,t} = (\text{long} - \text{short})/\text{open interest}$ follows $n_{c,t} = \alpha_c + r_c n_{c,t-1} + w_c \,\text{ret}_t$, the weekly version of the cohort recursion. The price-only alternative regresses each week's return on its own twelve lags and fits a two-exponential kernel $\phi_j = b_1 r_1^{j-1} + b_2 r_2^{j-1}$ to the twelve coefficients by least squares over the grid $0 \le r_1 < r_2 \le 0.99$ in steps of 0.02. Weekly futures returns are close to serially uncorrelated, so the fitted weights are close to zero and the rates are weakly identified nuisance parameters (Andrews and Cheng, 2012); that is part of what the wide price-only intervals show. Returns come from Binance spot (Bitcoin, Ether), the S&P 500 index, euro and yen spot rates, minus the change in the 10-year Treasury yield, and front-month crude-oil and natural-gas futures (Yahoo Finance). Only the last two are futures returns; for the other six the basis between spot and the CME contract enters the return regressor; this matters most for the sign of w_c in markets where leveraged funds run basis trades. The comparison is descriptive, not a race between estimators of the same object: the class route fits the persistence of an observed position series, the price route infers latent kernel rates from returns. Intervals come from a 13-week moving-block bootstrap; near a unit root an autoregressive coefficient is estimated at rate T^{-1} rather than $T^{-1/2}$, so the narrow class intervals partly reflect the near-unit-root persistence of position stocks, such percentile intervals are pointwise and can understate uncertainty, and price-only intervals that fill the admissible range are better read as a sign of weak identification than as calibrated 95% intervals. The pre-specified hypotheses were: (H1, primary) the fast class's memory rate is pinned to an interval narrower than 0.25 while both price-only intervals are wider, in at least 75% of markets; (H2) the fast class trend-follows, $w_c > 0$, in at least 75%; (H3) the fast class has the shorter memory in at least half.

The primary hypothesis holds in all eight markets (Table 23). The class reports pin the fitted persistence rate of each fast class's positions to an interval between 0.013 and 0.106 wide, while the price-only kernel leaves both rates anywhere in the admissible range; the

width ratios run from 9 to 71. The ratio partly reflects how little weekly returns reveal, which is the point of the theory rather than an artefact of it. A fitted persistence of aggregated positions is not a structural learning memory: it also reflects how slowly positions turn over and how accounts enter and leave a class. In the model with inventory costs (the remark after Lemma 2.1) a position accumulates order flow, so it carries an inventory-adjustment root, close to one when inventory costs are small, on top of the belief root r_i, and a one-lag fit to positions mostly picks up the former. The classes differ as the cohort model allows. Leveraged funds and managed money have the less persistent positions in six of eight markets, with half-lives of 6 to 18 weeks against months to years for asset managers and producers; in Treasuries and natural gas the two persistences are indistinguishable.

The trend-following hypothesis fails. The fast class buys after rises in the euro, the yen, crude oil and natural gas, but sells after rises in Bitcoin, Ether and the S&P 500 futures. The first half matches the finding of Moskowitz et al. (2012) that speculators' positions tend to go with past returns. We did not predict this. One candidate explanation is that leveraged-fund positions in financial futures are dominated by the short futures leg of cash–futures basis trades, which a single memory per class cannot represent (Barth and Kahn, 2021). A post-hoc check is consistent with it but not decisive: weekly position changes of leveraged funds and asset managers are strongly negatively correlated in Bitcoin, Ether, the S&P 500 and Treasuries (-0.37 to -0.67) but not in the euro or yen (-0.04 and -0.05), and holding the asset managers' change fixed removes the negative response in Bitcoin and halves it in Ether and the S&P 500. Because class positions sum to zero across all reporting categories, some negative correlation is mechanical, so this is a signature, not a test. The failure sharpens rather than weakens the case for class tagging: class labels are not strategies, and a supervisor who wants to know how a class responds to prices needs its positions, which the price process does not reveal. Negative estimated w_c also lies outside the positive economy of Section 2, so the dollar bounds of Section 8.2 do not apply to these classes; the pre-specified restriction test does not use positivity.

Post hoc robustness. After the registered test we reran the same estimators with three changes, approved before any of them was run and reported without re-scoring the registered result: CME futures prices (BTC, ETH, ES, ZN, 6E and 6J front months) for the six markets that used spot, index or yield prices; the futures version with contract-roll weeks removed by a calendar rule fixed in advance; and 26- and 52-week bootstrap blocks. With futures prices the counts are unchanged (H1 8 of 8, H2 4 of 8, H3 6 of 8), and so is the sign of $\hat{w}_{\text{fast}}$ in every market: the fast class still sells after rises in Bitcoin, Ether and the S&P 500. The spot–futures basis in the registered returns therefore does not produce the negative signs. Longer blocks leave the counts unchanged; Bitcoin's response interval then includes zero. Removing roll weeks leaves H1 and H2 unchanged, lowers H3 to 4 of 8 and turns the Treasury response negative. H3 is lost in Bitcoin and Ether, where the rule removes 46% of the weeks (200 of 436 and 131 of 284): a report is dropped whenever any day of its week touches the five-day expiry window, so with monthly expiries the rule removes about two weeks in four, more than intended, and the fast–slow difference is no longer resolved in the shorter samples. In the equity-index, currency and Treasury markets it removes 15% of the weeks, and in crude oil and natural gas none. The primary result, that class positions pin down a persistence that prices leave open, holds in all seven variants.

Class flows: a pre-specified test that failed. Proposition 8.1 is about tagged flow, while the registered test uses positions. A second pre-specified test (SHA-256 `052a63f20617...`, opened once) reads the same reports as flows, the weekly change in the fast class's net position, and estimates $u_t = c + r\,u_{t-1} + w\,\text{ret}_t$ by the instrumental-variables regression of the proposition, with last week's return as the instrument for last week's flow and the same 13-week bootstrap. Its primary hypothesis, that the 95% interval for r is narrower than 0.5 in at least six markets, fails. It holds in three: the S&P 500, crude oil and natural gas, with widths from 0.36 to 0.44. The intervals are 0.57 and 0.60 wide in the euro and the yen, 1.08 in Ether, 3.4 in Bitcoin and 10.8 in Treasuries. The instrument is strong

(first-stage $F \geq 10$) in seven markets, all but Treasuries ($F = 4.3$), and the flow response w is nonzero in six, with the same signs as in positions: negative in Ether and the S&P 500, positive in the euro, the yen, crude oil and natural gas. Read as flows, then, the CFTC reports identify how each class's flow responds to prices, but not how persistent that flow is, at the precision the test required. In Proposition 8.1 the precision is of order $T^{-1/2}$ with a constant that grows with the noise in the reported flow (σ_f^2 in (31)). Weekly changes in aggregated positions also move when accounts enter or leave a class, which is noise for this purpose, and two markets have short samples (279 and 428 weeks); we did not test these explanations. The failure limits the claim that published class aggregates implement the tagged-flow contract of Proposition 8.1. It does not bear on the proposition itself, which is a statement about the model.

Separating memory from turnover: a pre-specified test that failed. A fitted persistence of positions mixes two things: in the model with inventory costs (the remark after Lemma 2.1) a class position follows a slow inventory or turnover root and, on top of it, the faster belief-memory root r_i. A third pre-specified test (SHA-256 `e66335881e2c...`, opened once) asked whether the CFTC positions show that second root. For each class it fits $n_t = \alpha + \phi_1 n_{t-1} + \phi_2 n_{t-2} + b_0 \mathrm{ret}_t + b_1 \mathrm{ret}_{t-1}$, takes the roots $z_1 \geq z_2$ of $z^2 - \phi_1 z - \phi_2$, and requires, for the fast class in at least six markets, that the 95% interval of z_2 lie inside $(0, 0.9)$ and that $z_1 - z_2$ be clearly positive. It fails: the condition holds in two markets, crude oil and natural gas. The larger root is between 0.90 and 0.99 in every market, as in Table 23. The smaller root is between -0.16 and 0.21; it is negative in Bitcoin, Ether, the S&P 500 and Treasuries, and where it is positive its half-life is under half a week (0.22 to 0.44 weeks), far shorter than any learning memory the weekly reports could reveal. Positions therefore behave as one slow component plus short-lived noise. By the interpretation fixed before the run, this is a failure of the two-root reading of CFTC positions, and it confirms the limitation stated above: these data restrict how persistent class positions are, but they

cannot separate learning memory from the turnover of positions.

Replication on twelve new markets. On 31 October 2026 we registered (SHA-256 `591e71c9...`, opened once) the same test, with the code, the bootstrap and the thresholds unchanged, on twelve CFTC markets it had never been run on: gold, silver, copper, corn, soybeans, wheat, live cattle, coffee, sugar, cotton (disaggregated reports, managed money against producers), the Nasdaq-100 and the British pound (financial reports, leveraged funds against asset managers), with Yahoo Finance front-month or index prices and reports from June 2006 to 27 October 2026 (1,039 to 1,041 usable weeks). The primary hypothesis replicates in all twelve markets: the fast class's rate interval is 0.014 to 0.042 wide, while both price-only intervals are 0.85 to 0.96 wide, a width ratio of 23 to 69. Trend following (H2), which failed on the original eight markets, holds in eleven of the twelve: the managed-money or leveraged class buys after price rises everywhere except the Nasdaq-100, where it sells, as in the S&P 500. We report this as a new test on new markets; it does not undo the earlier failure, and the two results together say that the sign of the speculative class's response depends on the market, with commodities and currencies trend-following and equity indices not. H3 fails: the fast class is less persistent than the slow class in three markets (gold, silver and the Nasdaq-100); in the other nine the two classes' persistence cannot be told apart, since managed money and producers hold positions for similar times in agricultural and industrial commodities. The replication therefore extends the primary result from eight to twenty markets and leaves the limitation unchanged: class positions pin down how persistent positions are, not whether that persistence is learning memory.

A descriptive precision conversion. The recorded interval-width ratios are not an identified equivalent sample size: the two routes estimate different objects, price intervals hit parameter bounds, and near-unit-root positions need not obey uniform root-T scaling. Table 24 reports the earlier arithmetic conversion only as a hypothetical calculation conditional on common root-T scaling. It is not a lower bound on years of price data needed. A

first-half/full-sample interval comparison cannot establish an asymptotic rate or repair the estimand mismatch.

Institutional takeaway: within-class sign cancellation in coarse reporting. The pre-specified test scorecard (Table 1) reports that directional trend following (H2) and flow-based persistence tests failed in several markets, notably in equity indices (where leveraged funds sell into rising prices, opposite to commodities). Far from representing an econometric failure, this divergence provides direct institutional proof of the aggregation pathology inherent in coarse regulatory buckets. The CFTC's "Leveraged Funds" category aggregates heterogeneous market participants with opposing trading mandates: cash–futures basis arbitrageurs (who sell futures when prices rise to capture widened spreads) trade alongside directional CTA momentum funds (who buy as prices rise). Because these sub-entities take opposing positions in response to the same price move Δp_t, their responses offset each other within the aggregated position stock, causing within-class sign cancellation and attenuating net estimated flow persistence. This confirms that macroprudential market surveillance cannot rely on coarse institutional entity snapshots, demonstrating the empirical necessity of granular, algorithm-tagged transaction records (such as MiFID II RTS 24 `algo_id`).

Together with Table 25, the evidence now covers two disclosure regimes, one voluntary and one mandated by law, in crypto and in regulated equity-index, rate, currency and commodity futures. In both, class-tagged positions restrict persistence rates that prices leave unrestricted; whether that persistence is learning memory or slow turnover of positions, these data do not say.

8.5 Voluntary disclosure: Binance's trader-class ratios

Binance publishes voluntarily, for one venue but every five minutes, long/short ratios for its top-trader group, by position and by account, and for all accounts. These are class-level positions, the closest public analogue to the tagged data of Proposition 8.1, but not

homogeneous cohorts: the top-trader position ratio is built from position sizes, the all-account ratio from counts of long and short accounts, the top-trader group is a subset of all accounts, and its membership can change. We treat a class's reported net position as the observable counterpart of its cohort state: it obeys the same one-lag recursion class by class, so the block-diagonal information structure of Proposition 8.1 carries over. The mapping is an assumption, because the reports are positions rather than flows and returns stand in for the unobserved deviation x_t. For each class c we estimate $n_{c,t} = r_c n_{c,t-1} + w_c \,\mathrm{ret}_t$, with $n = (R-1)/(R+1)$ from the ratio R. The price-only alternative regresses signed five-minute taker flow on twelve lagged five-minute returns and fits two exponentials to the kernel; it estimates the memory rates without class information. Intervals come from a day bootstrap with 200 draws.

We begin with the one-hour price kernel of Table 25 and end with the fairer horizon-matched comparison on 30 new markets, which reduces the width ratio to a median of 7.3. Table 25 gives three results. First, tagged class data restrict persistence rates that prices and flow leave essentially unrestricted: with the one-hour price-only kernel, the price-only intervals are 28 to 124 times wider in every one of 30 symbols on fresh days, and cover most of the admissible range. Part of this ratio is mechanical, because a series of positions identifies its own persistence directly while the price-only route must fit two exponentials to a weakly identified kernel. Restricting attention to the five symbols in which the price-only fit does identify the merged response (ACH, API3, OGN, ZEN and DASH) does not change the conclusion: the price-only rate intervals are still 0.64 to 0.92 wide and the ratios run from 27.6 to 62.5 (median 47.5). We therefore read the ratios as a measure of how little prices reveal about class memory, not as a horse race between two estimators of equal standing. The ratios also overstate the gap for a second reason: the price-only kernel spans one hour (twelve five-minute lags), while the tagged half-lives are several hours, so the price-only fit cannot see memories of that length. An exploratory check on the 30 symbols and 105 days of the follow-up test refits the price-only kernel with 48 and 96 lags (four and eight hours)

and compares both routes in one unit, the half-life. With 96 lags the price-only slow half-life (median 5.6 hours) reaches the tagged scale (median 4.2 hours). The tagged intervals remain narrower in all 30 symbols, but by a median factor of 7.5 in log half-life rather than 28 to 124, and the price-only interval still reaches a bound of the fit in 25 of 30. The horizon-matched comparison was then pre-specified as a primary hypothesis of a test on 30 symbols never used in this project, before its data were downloaded, and passed: the tagged interval was narrower in 29 of 29 evaluable symbols (median factor 7.3, smallest 2.6), and the price-only interval reached a bound of the fit in 26. A pre-specified secondary of the same test qualifies what the tagged estimate measures. Day-demeaned regressions bias persistence down; corrected by a split-panel jackknife, the median tagged rate is 1.0008 per five minutes, and the corrected interval lies below one in only 7 of 87 class fits. Within a day the class ratios are close to a random walk, so the half-lives of a few hours reported here are largely produced by removing each day's mean. The evidence therefore supports a narrower claim than identification of class memory: positioning data restrict a persistence parameter that prices cannot bound. Whether that persistence reflects learning memory or slow turnover of positions, which theory's tagged contract (class-tagged flow) would distinguish, is not established by these data. Second, the classes differ where theory says the information should be: in BTCUSDT and ETHUSDT the top-trader group's positioning has a half-life of about three to five hours against about one to two days for all accounts, while in thin symbols the classes are often indistinguishable. Third, the pre-specified test fails its 80% criterion, on the part of the prediction that concerns what prices *do* reveal: the merged response is significantly different from zero in only 5 of 30 symbols. Theorem 6.10 says the merged measure is identified; on 42 days of thin-market data it is not estimated precisely. This is a failure of the prediction as pre-specified, and we report it as one. A post hoc reading, not tested here, is that identification is an asymptotic statement and 42 days of passive thin-market data are far from that regime. Two follow-ups were pre-specified before their data were downloaded, and each ran once. On 105 untouched days in the same 30

markets, price-only intervals were again at least ten times wider than class-tagged ones in all 30 symbols (median ratio 74.9 with the one-hour kernel), but the merged response was significantly different from zero in only 9 symbols, and pinning it down was at least ten times cheaper than separating the memory rates in only 10 of 30 (23 were required): that prediction failed. On 58 further days, the advantage of tagging grew with the sample in 22 of 29 symbols, the number required, but the tagged intervals shrank more slowly than the parametric rate (20 of 29 in the predicted range, 22 required), consistent with drift in class persistence over 21 months. Complete tables for both follow-ups are in the public replication repository.

To test directly whether the memory rates identified by class-tagged positions improve out-of-sample return predictions after large order-flow shocks, we pre-specified a dedicated out-of-sample test before downloading the 22nd-of-month data for January 2025–September 2026 across 19 thin perpetuals (6,137 top-decile five-minute flow shocks on 2024 training days; 4,579 shocks on untouched 2025–2026 test days). On the training sample we estimate each symbol's two memory rates $\hat{r}_1, \hat{r}_2$ either from the class-tagged positions or from the price-only two-exponential fit, construct the filtered memory states $y_{k,t} = \hat{r}_k y_{k,t-1} + (1 - \hat{r}_k)\mathrm{ret}_t$, and forecast the 30-minute post-shock cumulative return $R_{t+1\to t+6}$ out of sample relative to an AR(1) return-and-flow baseline. The primary hypothesis fails: on the untouched test shocks, the price-only memory states yield $\Delta R^2 = -0.199$ percentage points relative to the baseline, the class-tagged memory states yield $\Delta R^2 = -1.288$ points, and the tagged-minus-price-only difference is $\Delta R^2 = -1.089$ points (95% date-block bootstrap CI $[-2.605, +0.376]$, Clark–West $t = 0.96$, $p = 0.1684$). The difference is statistically inconclusive, and there is no evidence of a gain from class data. Likewise, adding the level and one-hour change of the top-trader and all-account positions to the hourly tail-risk controls of the companion paper leaves out-of-sample R^2 unchanged (-0.19 points, CI $[-0.99, 0.83]$).

After this test failed, we pre-specified a follow-up on untouched 2nd-and-9th-of-month data for January 2025–September 2026 across the same 19 thin perpetuals (12,046 top-decile

five-minute flow shocks on 24 training dates in 2024; 9,875 shocks on 42 test dates; test opened once). It asks the narrower question that Theorem 3.1 and Proposition 8.1 answer directly: can prices and flow track the class states themselves? A two-exponential price-and-flow filter ($\alpha_{\text{fast}} = 0.50$, $\alpha_{\text{slow}} = 0.10$ on five-minute returns and signed flow, plus shock return, shock flow and symbol fixed effects) is compared with a filter that adds the current class states and their one-bar changes. Both pre-specified primaries pass by a wide margin (Table 26): adding the class states raises out-of-sample R^2 from -4.8% to 46.0% for the 30-minute divergence between top-trader and all-account positions, and from 4.4% to 54.8% for the top-trader state. As the next paragraph shows, the gain reflects the near-unit-root persistence of position levels, not predictability of future trading.

The gain is almost entirely persistence. In an exploratory decomposition on the same test shocks (`hardened_empirical_audits.py` → `R_hardened_empirical_audits.json`), an AR(1) on the current level alone accounts for 100.5% and 99.1% of the improvement, and when the target is the 30-minute *change* in positions the class data add -0.26 and $+1.41$ points. Because the class filter observes the series it forecasts, this is a weak test: it confirms that prices and flow cannot reconstruct where the classes are positioned, which Theorem 3.1 implies, but it does not show that class data reveal how their positions will move, and neither pre-specified test finds that class data improve return forecasts.

October 2026: feedback and flow, two pre-specified tests that failed. The state-tracking result above passes, but its gain is persistence. On 31 October 2026 we registered (SHA-256 `72cbc8ba...`, opened once) two tests on a window no earlier test had used, 1–30 October 2026, for 30 liquid USD-M perpetuals fixed in advance (all 30 met the data rule), with a power check on synthetic data before registration. N2 asks whether price changes over the last half hour improve out-of-sample forecasts of the change in the top-trader position ratio beyond its own six lags, fitting on 1–20 October and evaluating on 21–30 October (about 2,600 five-minute bars per market) with a one-sided Diebold–Mariano test. It fails: $p < 0.05$

in one market (18 required), and the median out-of-sample R^2 gain is -0.01 percentage points, below the 0.8 to 1.6 points the test could detect. At five-minute frequency, then, this class position does not respond to recent prices beyond its own dynamics by a detectable amount. N3 asks whether the taker-flow imbalance, an untagged aggregate flow, pins down the combined one-hour response $\Psi = \sum_{m=1}^{12} k_m$ to past returns, with a Newey–West interval. Its primary criterion (interval excluding zero and half-width below $|\Psi|/2$) fails, holding in 5 of 30 (24 required); as a description, the interval excludes zero in 17 markets, with negative responses in 8 (among them Bitcoin and Solana) and positive in 9 (among them Dogecoin and Filecoin), and the sign is the same in both halves of the month in 21 (24 required). The one secondary prediction about rates holds in every market: the interval from ten days is 1.33 to 2.45 times as wide as the interval from thirty (median 1.69; root-T shrinkage predicts 1.73). Flow data shrink uncertainty at the rate the theory gives, but the aggregate response is small, noisy and of different signs across markets, so a month of five-minute data does not pin it down at the precision we required.

9 A spectral blind spot

Even without identification, could a supervisor monitor risk through eigenvalues or volatility?

Theorem 9.1 (Spectral blindness)**.** *For any eigenvalues $0 < \lambda_2 < \lambda_1 < 1$, one-step price variance $s_1 > 0$ and stationary price variance V large enough, there is a family of two-state markets $J_\theta = \left(\begin{smallmatrix} \lambda_1 & \theta \\ 0 & \lambda_2 \end{smallmatrix}\right)$ with shock covariance $\mathrm{diag}(s_1, s_2(\theta))$ such that all members share the eigenvalues, the one-step price variance and the stationary price variance, while the worst-case transient gain $\sup_k \|e_1^\top J_\theta^k\|$ grows without bound in θ.*

In words. Two markets can look identical by the usual risk measures, the same eigenvalues and the same price variance, and yet one of them can produce far larger losses in the short run. Monitoring eigenvalues or volatility alone therefore cannot certify that a market is safe.

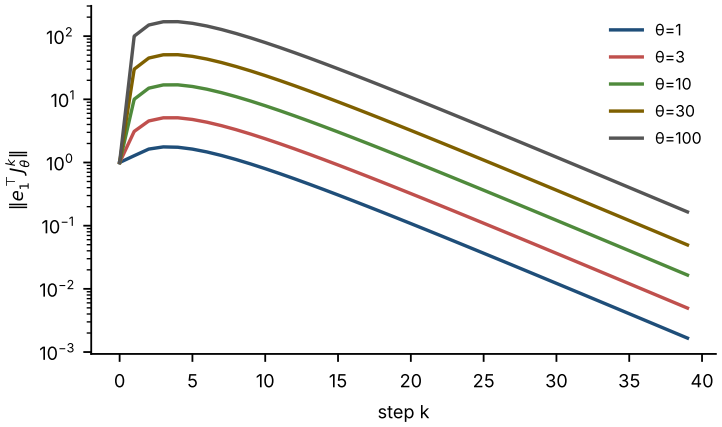


Figure 9: Price-row transient gain of the spectral-blindness family J_θ: eigenvalues $(0.8, 0.7)$, one-step price variance 1 and stationary price variance 5 are the same for every θ.

At $\lambda = (0.8, 0.7)$, $s_1 = 1$ and $V = 5$, the gain is 1.77, 5.10, 16.95, 50.85 and 169.50 for $\theta = 1, 3, 10, 30, 100$, with stationary variance 5.000 in every case; the cumulative gain rises from 18.3 to 1,667.7. No monitor that depends only on eigenvalues and second moments of the price can be uniformly safe (Figure 9).

Two clarifications prevent over-reading. First, the cohort economy of Section 2 is itself non-normal, yet its price output does not amplify (Lemma 2.3). Large weighted gains can be manufactured by output weights alone: with $C_w = (1, 160)$, a price-and-memory functional peaks at 160.003 at $t = 0$, which is the static weight, not dynamic growth (Appendix C.5). Second, in a network of M identical cohort economies coupled by symmetric nonnegative cross-impact, the coupled operator stays symmetric in the coordinates of Lemma 2.3; with cross-impact 0.05, the spectral radius rises from 0.822 ($M = 1$) to 0.973 ($M = 10$) while the Kreiss constant stays at one, as it must for a stable symmetric operator (`recompute_oct20.py`). A stable normal contraction cannot exhibit Euclidean state-norm transient growth. This does not exclude scalar input-output overshoot relative to initial impact, which also depends on the input and output maps. The positive-cohort no-overshoot result is model-specific; the companion paper studies a depth-depletion mechanism.

10 Empirical feedback and its measurement

The theory says what prices can reveal. The empirical work has two roles, and neither is a test of the identification bounds themselves, which concern controlled experiments no outside researcher can run. First, Section 8.5 tests the theory's policy prediction directly: that class-tagged data pin down memories that prices leave unrestricted. Second, this section asks what the reduced-form feedback indices a supervisor would compute from public data actually measure, and how much of them is feedback rather than filtering, timing or trading activity.

10.1 Estimand and filtering mechanics

In theory $x_t = p_t - p_t^*$ is the deviation from an unobserved fundamental. Empirical work replaces p_t^* by a trailing filter. Let p_b be a log price proxy on uniform bars of length Δ and $m_b = (1-\alpha)m_{b-1} + \alpha p_b$ an exponential moving average with $\alpha = 2/(W/\Delta + 1)$. The empirical state is

$$x_b^{\mathrm{EMA}} \equiv p_b - m_b = (1-\alpha)(p_b - m_{b-1}) = \frac{1-\alpha}{1-(1-\alpha)L}\Delta p_b. \tag{37}$$

Because $x_{b+1}^{\mathrm{EMA}} = (1-\alpha)(p_{b+1} - m_b)$, a martingale price gives

$$\mathbb{E}[x_{b+1}^{\mathrm{EMA}} \mid \mathcal{F}_b] = (1-\alpha)x_b^{\mathrm{EMA}} : \tag{38}$$

the filtered state is an AR(1) with mechanical persistence $a_0 = 1-\alpha$, which equals $29/31 \approx 0.9355$ for $W = 300$ seconds and $\Delta = 10$ seconds.

Let f_t be signed flow. The empirical VAR(1)

$$\begin{pmatrix} x_{t+1}^{\mathrm{EMA}} \\ f_{t+1} \end{pmatrix} = \begin{pmatrix} J_{11} & J_{12} \\ J_{21} & J_{22} \end{pmatrix} \begin{pmatrix} x_t^{\mathrm{EMA}} \\ f_t \end{pmatrix} + e_{t+1} \tag{39}$$

has reduced-form coefficients, not authenticated impact or learner parameters. When denominators are nonzero we use the scale-invariant indices

$$\hat{Q} = \frac{\hat{J}_{12}\hat{J}_{21}}{1-\hat{J}_{22}}, \qquad \hat{S} = \hat{J}_{11} + \hat{Q}, \qquad \hat{L} = \frac{\hat{Q}}{1-\hat{J}_{11}}. \tag{40}$$

Equation (38) implies that $\hat{J}_{11}$ must be compared with $1-\alpha$, not with zero: persistence that a naive reading attributes to feedback can be pure filtering.

10.2 Data

Table 27 summarizes the Binance spot and perpetual samples used in this section.

10.3 Clock alignment dominates reduced-form feedback

We estimate (39) on ten-second bars built from one-second Binance klines across a 96-date, eight-pair spot panel (the 10th and 24th of each month, 2021–2024) under three constructions (Table 28), and run the identical pipeline on a placebo random walk with each pair-day's volatility. Pairing the state with *next-bar* flow (a look-ahead) gives $\hat{Q} = 0.0081$; aligning flow with the bar in which it is observed lowers this by more than an order of magnitude, to 0.00037; and replacing the EMA state by raw returns gives an interval that includes zero. Mean $\hat{J}_{11}$ under observed flow is 0.9203, below the mechanical 0.9355, and the placebo reproduces the mechanical value (0.9351) with $\hat{Q}$ indistinguishable from zero.

10.4 Feedback beyond filtering mechanics

We re-estimate feedback on a separate sample and add a placebo. We use Binance BTCUSDT one-second klines on 72 sampled days in 2021–2023. The state x_t is the log price deviation from its EMA (span 30 one-second bins), f_t is next-second signed taker flow, and a VAR(1) gives $\hat{Q}$. Because the filter alone produces persistence, we run the identical pipeline on a placebo: a random walk with the same one-second volatility, independent of the real flow.

Table 30 gives two findings. Feedback, in the sense of a loop gain the filter cannot produce, is detectable: the real $\hat{Q}$ lies far outside the placebo band. It is also small, and the price persistence that a naive reading would attribute to feedback is almost entirely mechanical (0.9364 against 0.9355). Because this construction pairs the state with next-second flow, part of $\hat{Q}$ reflects the timing effect of Table 28; the placebo shows only that the filter does not produce it.

10.5 Order splitting and the estimation of feedback

Persistent order flow makes lagged-flow regressions inconsistent. Suppose flow follows $q_t = \gamma y_t + u_t$ with memory $y_t = (1-\eta)y_{t-1} + \eta\Delta p_{t-1}$, order splitting $u_t = \rho u_{t-1} + \nu_t$, and $\Delta p_t = \kappa q_t + \lambda z_t + e_t$ with an observed exogenous price shock z_t. Lagged returns are correlated with u_{t-1}, so OLS of flow on returns is biased even when $\gamma = 0$. Quasi-differencing ($q_t - \rho q_{t-1} = \gamma(y_t - \rho y_{t-1}) + \nu_t$) or instruments that move y through prices but are independent of u ($z_{t-1}, z_{t-2}, z_{t-3}$) restore consistency. We compare estimators over 500 replications with $\eta = 0.3$ and $\rho = 0.8$ (Table 31).

OLS is badly biased. Cochrane–Orcutt with the true ρ, 2SLS and LIML have correct size and coverage; Cochrane–Orcutt with an estimated ρ under-covers (Cochrane and Orcutt, 1949; Staiger and Stock, 1997). These are properties of estimators in a simulated market. On real data, a 2024 sub-panel of 192 spot units shows why the correction matters little at our bar length: with one VAR lag, the flow-equation residuals have essentially no first-order serial correlation (median $\hat{\rho} = -0.0002$), and Cochrane–Orcutt quasi-differencing moves the mean loop gain from 3.15×10^{-4} to 3.21×10^{-4} (post hoc). Order splitting shows up instead at longer lags, which the 30-lag control of Table 29 addresses. A real-data instrumental-variables estimate needs an instrument whose exclusion can be defended, which we do not claim to have; funding-settlement times, for example, may affect order flow directly.

Controlling for 30 ten-second bars of lagged flow lowers $\hat{Q}$ further (Table 29, column $\hat{Q}_{30}$), consistent with persistent order splitting.

10.6 Feedback across 30 perpetual markets

The companion paper estimates the same loop gain on one-minute bars of 30 Binance USD-M perpetuals spread across the liquidity ranking (2023–2024 training days). Measured feedback rises with depth: its Spearman correlation with log depth within 1% of the mid is +0.70 [0.41, 0.87], from about 0.012 in BTCUSDT and ETHUSDT to 0.0016 on average in the thinnest third, and the placebo pipeline gives $\hat{Q} \approx 0$ for every symbol. Feedback in this reduced-form sense is a feature of deep, algorithmically traded books.

10.7 Two fee events

Binance set spot trading fees on BTC pairs to zero on July 8, 2022 and reinstated them on March 22, 2023. Each event changes the cost of trading BTCUSDT while leaving the seven other pairs unchanged, and the second reverses the first. For every pair and every day within 20 days of each event we compute daily quote volume and the observed-flow $\hat{Q}$, and estimate $Y_{it} = \alpha_i + \gamma_t + \beta\,\mathrm{ZeroFee}_{it} + e_{it}$ with pair and date fixed effects and date-clustered errors, separately for each event and pooled (Table 32).

The fee change moved activity enormously: zero fees multiplied BTCUSDT volume relative to the controls by $e^{0.945} \approx 2.6$ when they began, and their removal divided it by $e^{1.599} \approx 4.9$. Measured feedback did not follow. The two point estimates have opposite signs, neither is significant, and the pooled effect is -0.00016 with an interval of $[-0.00087, 0.00055]$, against a sample mean of 0.00037. A strong first stage with no second-stage response says that $\hat{Q}$ is not a by-product of trading activity, and the reversal design removes the main weakness of a single treated event. With one treated asset the interval remains wide, so this is evidence that $\hat{Q}$ does not scale with activity, not a proof of invariance.

11 Implications for supervision and market design

The results separate two supervisory tasks that are often run together. Monitoring aggregate fragility, for example to set market-wide limits, needs only the combined price response, which public data can identify under the observation and tail conditions of the model (Remark 6.3, Proposition 8.5(i)). Attributing feedback to particular groups, to apply targeted limits or enforcement, faces the $(\Delta\eta)^{-(4K-2)}$ barrier whenever groups react over similar horizons (Proposition 8.5(ii)). Class-tagged records remove that barrier group by group (Proposition 8.1), and in the model persistent pseudonymous class labels deliver the full statistical gain without personal data.

What a supervisor can learn from public data. Prices and anonymous flow reveal the merged measure M and the closed-loop impulse response, hence the transfer function, and nothing finer. A supervisor who wants to know *which* algorithms drive a feedback loop cannot get there from public data at reasonable cost when their memories are close. This is not a failure of econometric technique: Theorem 6.2 holds for every estimator, biased or not.

What probing costs. At realistic observation noise, an exchange or a large desk could separate two well-separated memory rates for about 10^5 USDT of capital. Probing is therefore feasible for coarse questions and infeasible for fine ones. Because capital grows as δ^{-4} and experiments as $(\Delta\eta)^{-(4K-2)}$, halving the separation multiplies the cost by 16 for a known pair and by 2^{4K-2} for K unknown cohorts.

What disclosure buys. Disclosure regimes differ sharply in value (Table 20). Timestamps without update values are worthless for this purpose. Noisy periodic statements are nearly worthless. Audited equity and, above all, flow tagged by participant class remove the barrier entirely (Proposition 8.1). In our attribution experiment, correctly tagged flow is far more informative than the periodic statements considered here. Calibrated to the

fitted day-demeaned persistences of Binance's published class ratios (an illustrative calibration), separating the top-trader group from all accounts in BTCUSDT would take prices alone about 333 times as long (Theorem 8.2). The class reports that the CFTC already requires have the same property in regulated markets (Table 23), so the argument concerns the granularity and frequency of existing disclosure, not a new kind of reporting.

When audit trails already exist. The US Consolidated Audit Trail (SEC Rule 613) is designed to capture specified order and reportable-event information for NMS stocks and listed options across exchanges and broker-dealers, from origination to execution or cancellation, linked to a customer identifier and to the broker-dealer that handled the order; its actual scope is governed by Rule 613, the CAT NMS Plan and exemptions. Since February 2025, under SEC exemptive relief, it no longer requires the names, addresses and years of birth of natural-person customers, who remain linked through an identifier derived from a transformed tax identifier (SEC, 2012, 2016, 2024, 2025). Our results bear on such regimes in three conditional ways. First, for the attribution task modelled here, records that report each class's flow separately remove a barrier that, for an observer who learns by probing, costs round-trip trading loss of order δ^{-4}, and they save samples by a factor of order $(\Delta\eta)^{-(4K-2)}$ at fixed accuracy or $(\Delta\eta)^{-(4K-4)}$ for rates resolved to a fraction of their spacing (Proposition 8.1, Theorem 8.2). This is a measurement value, not a cost–benefit assessment: CAT's budget pays for enforcement, market reconstruction and surveillance, none of which the model measures, and Section 8.2 sets the hypothetical probing loss beside it only as a scale illustration. For estimating class responses in our model a persistent pseudonymous label suffices; this does not settle what enforcement or data-minimization design requires.

Second, an audit trail records legal entities and accounts, not algorithmic decision rules. This institutional design choice directly explains the empirical anomalies and test failures documented in our pre-specified tests (Table 1). Coarse regulatory classifications (such as the CFTC's "Leveraged Funds" category or broad broker-dealer identifiers) inevitably aggregate

heterogeneous, opposing trading strategies under a single umbrella label. In equity-index futures, for example, directional trend-following CTAs (who buy following positive returns) are pooled together with cash–futures basis arbitrageurs (who sell futures as prices rise to capture basis dislocations; Barth and Kahn, 2021). Because these sub-entities execute opposite-signed orders in response to identical price shocks, their positions undergo within-class sign cancellation, driving the aggregate net response coefficient near zero and preventing flow-based regressions from resolving structural persistence. This empirical failure is not a statistical defect: it is direct market evidence of the aggregation pathology of coarse entity snapshots. An observer attempting to attribute feedback from account-level or entity-level records faces our closed-loop identification barrier within each reporting bucket. The policy remedy is direct: macroprudential surveillance cannot rely on coarse institutional classifications (such as Legal Entity Identifiers alone). It requires granular, strategy-specific transaction tags, exemplified by European regulation under MiFID II RTS 24 (Article 2(2) and Table 2, Field 3, requiring venue order records to carry a proprietary `algo_id` identifying the specific trading algorithm responsible for each execution; European Commission, 2017; RTS 6, European Commission, 2017b). Tagging by algorithmic decision rule rather than legal entity eliminates within-class sign cancellation, restoring the full $T^{-1/2}$ statistical benefit of disclosure proved in Proposition 8.1.

Third, the audit trail does not reach the markets we study. Crypto perpetual futures are outside it, futures regulated by the CFTC have separate reporting, and its data are not public: regulators and authorized self-regulatory organizations have controlled access, while outside desks and researchers generally do not. For these observers the price-only and coarse-disclosure contracts of this paper are the relevant ones.

What to monitor instead. Because transient loss can be arbitrarily large at fixed eigenvalues and volatility (Theorem 9.1), supervision that wants to control risk rather than identify participants should monitor quantities that bound transient loss directly: the summed

absolute impulse response and confidence sets for the transfer function, both of which are identified from public data. These are cheap where identification is expensive (Remark 6.3); the companion paper shows that the summed response can be bounded at the parametric rate from controlled experiments, and only under a declared bound on memory length from a finite window (Niroula et al., 2026).

How to measure feedback. Empirical feedback indices should (i) compare persistence with the mechanical value of the filter, (ii) align flow with the bar in which it is observed, (iii) control for order-splitting persistence or use instruments with a defended exclusion, and (iv) report placebo pipelines. Each of these changes the sign or the order of magnitude of published-style estimates in our data.

12 Conclusion and an open problem

In a linear market with exponential-memory cohorts and known impact parameters, prices reveal a merged summary of algorithmic feedback. Recovering the algorithms behind it costs executed quantity of order δ^{-2} and trading capital of order δ^{-4} for a known pair, with a numerically evaluated constant of about 61 par units at the calibration, and a number of experiments of exact order $(\Delta\eta)^{-(4K-2)}$ for K unknown cohorts, and at least that order of passive observations; in simulations of stylized nonlinear rules and of online learners trained in the closed loop the same costs hold when the competing populations' feedback weights match, populations whose learned weights differ are told apart by those weights instead, and hard thresholds hide cohorts from small probes altogether. Without the exponential form, accuracy improves only as $1/\log N$, with matching lower and upper bounds. Class-tagged disclosure removes the barrier under the stated contract, and its value can be stated in samples: a factor of order $(\Delta\eta)^{-(4K-4)}$ for resolving rates (and $(\Delta\eta)^{-(4K-2)}$ at fixed accuracy), about 333 for the trader classes Binance already publishes in BTCUSDT. The CFTC's class reports restrict a fitted position persistence that prices leave unrestricted in all

eight regulated markets we examined; this persistence is not shown to be learning memory, a pre-specified test of what prices do reveal failed, and class-tagged memory states did not improve 30-minute post-shock return forecasts ($\Delta R^2 = -1.09$ points, inconclusive). Read as weekly flows, the same CFTC reports identify how each class's flow responds to prices but failed a pre-specified test of how precisely they pin down its persistence, and a pre-specified test for a second, faster root in positions, which would have been the learning memory, failed as well. With real Bitcoin price noise in place of Gaussian noise, the predicted rates held in a pre-specified test, at the sample sizes where the real noise could be used directly. At realistic noise the capital barrier for a moderate separation is about 10^5 USDT. Measured feedback in crypto markets is real but small, rises with depth, does not move when fees change trading volume severalfold, and is easily confused with filtering and timing artefacts. Finally, eigenvalues and volatility cannot certify safety, because transient loss can be arbitrarily large at fixed spectrum and variance.

Two open questions remain. The first is economic: Proposition 8.5 gives one stylised decision in which knowing the individual groups changes the right action, but turning the information costs here into a value of disclosure needs a supervisory loss function estimated from real interventions. The second is for supervision: *what can be measured and controlled without identifying the algorithms?* The impulse response and the transfer function, and finite-sample confidence sets for them, are identified from public data even when cohorts are not. The structural matrix J is not: economies with the same merged measure can have different state dimensions, so J is identified only up to the equivalence class of minimal realizations of the same transfer function. A well-posed next question is: for a stated class of stable closed-loop markets, given an estimated transfer function and a finite-sample confidence set, which intervention policy of least cost guarantees a stated bound on future price displacement without recovering individual cohorts? The companion paper uses exactly this interface to build an identification-free pre-trade control.

Declarations

Funding. This research received no external funding.

Competing interests. The authors declare no competing interests.

Author contributions. Pratik Niroula: conceptualization, methodology, software, formal analysis, writing (original draft). Prashant Niroula: formal analysis (verification of proofs), validation, writing (review and editing). Sujan Khatiwada: formal analysis (verification of proofs), validation, writing (review and editing). All authors read and approved the final manuscript.

Data availability. All data are public. Binance market data come from the Binance public data archive (`https://data.binance.vision`): spot and USD-margined perpetual futures candles at one-second and one-minute frequency, futures order-book depth, and the published trader-class long/short ratios (top traders by position and by account, and all accounts); every file is verified against Binance's published checksums. The CFTC Commitments of Traders reports (Traders in Financial Futures and Disaggregated, futures only, through 13 October 2026) come from the CFTC public reporting API. Weekly prices for the CFTC markets (stock index, currencies, Treasury yield and CME futures) come from Yahoo Finance. Processed data and all result files are in the public repository below.

Code availability. All code, pre-specified test specifications, result files and test-unlock counters are public at `https://github.com/Pratikn03/Fianince`. The repository's README lists the scripts in run order, and Appendix D maps every number in the paper to the file that produces it.

Ethics approval. Not applicable. The study uses only public market data and involves no human participants or personal data.

Use of AI tools. The research ideas are the authors' own. The authors used AI assistants in carrying them out: Claude (Anthropic) was used to analyse data, read, write and debug analysis code, check derivations and numerical results, and draft and edit text. ChatGPT (OpenAI) and Gemini (Google) were used for planning and to critique drafts; these AI critiques are internal checks, not independent peer review. AI tools are not authors. The authors ran and reviewed all code and results, reviewed all text, and take full responsibility for the content.

References

Alfonsi, A., A. Schied, and A. Slynko (2012). Order book resilience, price manipulation, and the positive portfolio problem. *SIAM Journal on Financial Mathematics* 3. doi:10.1137/110822098.

Alfonsi, A., A. Schied, and F. Klöck (2016). Multivariate transient price impact and matrix-valued positive definite functions. *Mathematics of Operations Research* 41, 914–934.

Almgren, R., and N. Chriss (2001). Optimal execution of portfolio transactions. *Journal of Risk* 3, 5–39.

Amilon, H. (2008). Estimation of an adaptive stock market model with heterogeneous agents. *Journal of Empirical Finance* 15, 342–362.

Andrews, D. W. K., and X. Cheng (2012). Estimation and inference with weak, semi-strong, and strong identification. *Econometrica* 80, 2153–2211.

Banerjee, S., and M. Szydlowski (2026). Trading against algorithms: Price dynamics and risk-sharing in a market with Q-learners. Working paper.

Barberis, N., R. Greenwood, L. Jin, and A. Shleifer (2015). X-CAPM: An extrapolative capital asset pricing model. *Journal of Financial Economics* 115, 1–24.

Barberis, N., R. Greenwood, L. Jin, and A. Shleifer (2018). Extrapolation and bubbles. *Journal of Financial Economics* 129, 203–227.

Barth, D., and R. J. Kahn (2021). Hedge funds and the Treasury cash-futures disconnect. Office of Financial Research Working Paper 21-01.

Batenkov, D., G. Goldman, and Y. Yomdin (2021). Super-resolution of near-colliding point sources. *Information and Inference* 10, 515–572.

Bernstein, S. (1914). Sur la meilleure approximation de $|x|$ par des polynômes de degrés donnés. *Acta Mathematica* 37, 1–57.

Bessembinder, H., W. Maxwell, and K. Venkataraman (2006). Market transparency, liquidity externalities, and institutional trading costs in corporate bonds. *Journal of Financial Economics* 82, 251–288.

Biais, B., T. Foucault, and S. Moinas (2015). Equilibrium fast trading. *Journal of Financial Economics* 116, 292–313.

Boehmer, E., G. Saar, and L. Yu (2005). Lifting the veil: An analysis of pre-trade transparency at the NYSE. *Journal of Finance* 60, 783–815.

Bouchaud, J.-P., Y. Gefen, M. Potters, and M. Wyart (2004). Fluctuations and response in financial markets: The subtle nature of "random" price changes. *Quantitative Finance* 4, 176–190.

Bouchaud, J.-P., J. Bonart, J. Donier, and M. Gould (2018). *Trades, Quotes and Prices: Financial Markets Under the Microscope*. Cambridge University Press.

Brock, W. A., and C. H. Hommes (1998). Heterogeneous beliefs and routes to chaos in a simple asset pricing model. *Journal of Economic Dynamics and Control* 22, 1235–1274.

Brogaard, J., T. Hendershott, and R. Riordan (2014). High-frequency trading and price discovery. *Review of Financial Studies* 27, 2267–2306.

Brunnermeier, M. K., and L. H. Pedersen (2009). Market liquidity and funding liquidity. *Review of Financial Studies* 22, 2201–2238.

Castelli, V., and T. M. Cover (1996). The relative value of labeled and unlabeled samples in pattern recognition with an unknown mixing parameter. *IEEE Transactions on Information Theory* 42, 2102–2117.

Consolidated Audit Trail, LLC (2026). Consolidated Audit Trail, LLC 2026 Financial and Operating Budget, 31 March 2026 (total expenses $147,044,869). https://www.catnmsplan.com/cat-financial-and-operating-budget.

Cespa, G., and X. Vives (2015). The beauty contest and short-term trading. *Journal of Finance*

70, 2099–2154.

Cochrane, D., and G. H. Orcutt (1949). Application of least squares regression to relationships containing auto-correlated error terms. *Journal of the American Statistical Association* 44, 32–61.

Colliard, J.-E., T. Foucault, and S. Lovo (2022). Algorithmic pricing and liquidity in securities markets. CEPR Discussion Paper 17606.

Crouzet, N., I. Dew-Becker, and C. G. Nathanson (2020). On the effects of restricting short-term investment. *Review of Financial Studies* 33, 1–43.

Dou, W. W., I. Goldstein, and Y. Ji (2025). AI-powered trading, algorithmic collusion, and price efficiency. NBER Working Paper 34054.

Duarte, J., and L. Young (2009). Why is PIN priced? *Journal of Financial Economics* 91, 119–138.

Easley, D., N. M. Kiefer, M. O'Hara, and J. B. Paperman (1996). Liquidity, information, and infrequently traded stocks. *Journal of Finance* 51, 1405–1436.

European Securities and Markets Authority (2016). Guidelines on transaction reporting, order record keeping and clock synchronisation under MiFID II. ESMA/2016/1452.

European Commission (2017). Commission Delegated Regulation (EU) 2017/580 of 24 June 2016 supplementing Regulation (EU) No 600/2014 with regard to regulatory technical standards for the maintenance of relevant data relating to orders in financial instruments (RTS 24), Article 2(1)(b)–(c). *Official Journal of the European Union* L 87, 31.3.2017.

European Commission (2017). Commission Delegated Regulation (EU) 2017/589 of 19 July 2016 supplementing Directive 2014/65/EU with regard to regulatory technical standards specifying the organisational requirements of investment firms engaged in algorithmic trading (RTS 6). *Official Journal of the European Union* L 87, 31.3.2017.

Fan, J. (1991). On the optimal rates of convergence for nonparametric deconvolution problems. *Annals of Statistics* 19, 1257–1272.

Forssell, U., and L. Ljung (1999). Closed-loop identification revisited. *Automatica* 35, 1215–1241.

Franke, R., and F. Westerhoff (2012). Structural stochastic volatility in asset pricing dynamics: Estimation and model contest. *Journal of Economic Dynamics and Control* 36, 1193–1211.

Gatheral, J. (2010). No-dynamic-arbitrage and market impact. *Quantitative Finance* 10, 749–759.

Glosten, L. R., and P. R. Milgrom (1985). Bid, ask and transaction prices in a specialist market with heterogeneously informed traders. *Journal of Financial Economics* 14, 71–100.

Hasbrouck, J. (1991). Measuring the information content of stock trades. *Journal of Finance* 46, 179–207.

Heinrich, P., and J. Kahn (2018). Strong identifiability and optimal minimax rates for finite mixture estimation. *Annals of Statistics* 46, 2844–2870.

Hendershott, T., C. M. Jones, and A. J. Menkveld (2011). Does algorithmic trading improve liquidity? *Journal of Finance* 66, 1–33.

Hua, Y., and T. K. Sarkar (1990). Matrix pencil method for estimating parameters of exponentially damped/undamped sinusoids in noise. *IEEE Transactions on Acoustics, Speech, and Signal Processing* 38, 814–824.

Katz, R., N. Diab, and D. Batenkov (2023). On the accuracy of Prony's method for recovery of exponential sums with closely spaced exponents. arXiv preprint arXiv:2302.05883.

Kirilenko, A., A. S. Kyle, M. Samadi, and T. Tuzun (2017). The Flash Crash: High-frequency trading in an electronic market. *Journal of Finance* 72, 967–998.

Kong, W., and G. Valiant (2017). Spectrum estimation from samples. *Annals of Statistics* 45, 2218–2247.

Kreiss, H.-O. (1962). Über die Stabilitätsdefinition für Differenzengleichungen die partielle Differentialgleichungen approximieren. *BIT Numerical Mathematics* 2, 153–181.

Kukacka, J., and J. Barunik (2017). Estimation of financial agent-based models with simulated maximum likelihood. *Journal of Economic Dynamics and Control* 85, 21–45.

Kyle, A. S. (1985). Continuous auctions and insider trading. *Econometrica* 53, 1315–1335.

Le Cam, L. (1986). *Asymptotic Methods in Statistical Decision Theory*. Springer.

Leach, J. C., and A. N. Madhavan (1993). Price experimentation and security market structure. *Review of Financial Studies* 6, 375–404.

Liu, P., and H. Ammari (2022). Super-resolution of positive near-colliding point sources. arXiv preprint arXiv:2212.00536.

Ljung, L. (1999). *System Identification: Theory for the User*, 2nd ed. Prentice Hall, Upper Saddle River, NJ.

Makarov, I., and A. Schoar (2020). Trading and arbitrage in cryptocurrency markets. *Journal of Financial Economics* 135, 293–319.

Manole, T., and N. Ho (2022). Refined convergence rates for maximum likelihood estimation under finite mixture models. *Proceedings of the 39th International Conference on Machine Learning*, PMLR 162.

Manski, C. F. (2003). *Partial Identification of Probability Distributions*. Springer, New York.

Moallemi, C. C., B. Park, and B. Van Roy (2008). Strategic execution in the presence of an uninformed arbitrageur. arXiv preprint arXiv:0801.3001.

Moitra, A. (2015). Super-resolution, extremal functions and the condition number of Vandermonde matrices. In *Proceedings of the 47th Annual ACM Symposium on Theory of Computing (STOC)*, 821–830.

Moskowitz, T. J., Y. H. Ooi, and L. H. Pedersen (2012). Time series momentum. *Journal of Financial Economics* 104, 228–250.

Niroula, P., S. Alkhushayni, Y. Chen, M. Bhavsar, and P. Niroula (2026). From human markets to machine markets: Transient amplification, an exact depth condition, and a maximal pre-trade guard. Working paper, Minnesota State University, Mankato.

Obizhaeva, A. A., and J. Wang (2013). Optimal trading strategy and supply/demand dynamics. *Journal of Financial Markets* 16, 1–32.

Pagano, M., and A. Röell (1996). Transparency and liquidity: A comparison of auction and dealer markets with informed trading. *Journal of Finance* 51, 579–611.

Park, B., and B. Van Roy (2015). Adaptive execution: Exploration and learning of price impact. *Operations Research*, doi:10.1287/opre.2015.1415.

Rebafka, T., and F. Roueff (2015). Nonparametric estimation of the mixing density using polynomials. arXiv preprint arXiv:1002.4516.

Rivlin, T. J. (1981). *An Introduction to the Approximation of Functions*. Dover.

Roll, R. (1984). A simple implicit measure of the effective bid-ask spread in an efficient market.

Journal of Finance 39, 1127–1139.

Roy, R., and T. Kailath (1989). ESPRIT—estimation of signal parameters via rotational invariance techniques. *IEEE Transactions on Acoustics, Speech, and Signal Processing* 37, 984–995.

U.S. Securities and Exchange Commission (2010). Risk management controls for brokers or dealers with market access. Release No. 34-63241.

U.S. Securities and Exchange Commission (2012). Consolidated Audit Trail; Final Rule (Release No. 34-67457, July 18, 2012, 77 FR 45722 at 45790–45795, pp. 220–238; proposing release No. 34-62174, May 26, 2010, 75 FR 32556 at 32598, 32600, pp. 147, 153).

U.S. Securities and Exchange Commission (2016). Joint Industry Plan; Order Approving the National Market System Plan Governing the Consolidated Audit Trail. Release No. 34-79318, November 15, 2016, 81 FR 84696 at 84801, p. 363.

U.S. Securities and Exchange Commission (2024). Notice of Filing of Proposed Rule Change Relating to the 2025 CAT Budget. Release No. 34-101901, December 12, 2024, p. 4 (revised May 2025).

U.S. Securities and Exchange Commission (2025). Order granting exemptive relief from certain requirements of the National Market System Plan governing the Consolidated Audit Trail. Release No. 34-102386, February 10, 2025.

Shahram, M., and P. Milanfar (2005). On the resolvability of sinusoids with nearby frequencies in the presence of noise. *IEEE Transactions on Signal Processing* 53, 2579–2588.

Shleifer, A., and R. W. Vishny (1992). Liquidation values and debt capacity: A market equilibrium approach. *Journal of Finance* 47, 1343–1366.

Sornette, D., S. Lera, J. Lin, and K. Wu (2023). Non-normal interactions create socio-economic bubbles. *Communications Physics* 6, 261.

Staiger, D., and J. H. Stock (1997). Instrumental variables regression with weak instruments. *Econometrica* 65, 557–586.

Trebbi, F., M. B. Zhang, and M. Simkovic (2024). The cost of regulatory compliance in the United States. Working paper.

Trefethen, L. N., and M. Embree (2005). *Spectra and Pseudospectra: The Behavior of Nonnormal*

Matrices and Operators. Princeton University Press.

Tremacoldi-Rossi, P. (2025). Manipulating algorithmic markets. Working paper.

Tsybakov, A. B. (2009). *Introduction to Nonparametric Estimation*. Springer.

van Waarde, H. J., J. Eising, H. L. Trentelman, and M. K. Camlibel (2020). Data informativity: A new perspective on data-driven analysis and control. *IEEE Transactions on Automatic Control* 65, 4753–4768.

Varga, R. S., and A. J. Carpenter (1985). On the Bernstein conjecture in approximation theory. *Constructive Approximation* 1, 333–348.

Wagenmaker, A., and K. Jamieson (2020). Active learning for identification of linear dynamical systems. In *Proceedings of the 33rd Conference on Learning Theory*, PMLR 125, 3487–3582.

Wagenmaker, A. J., M. Simchowitz, and K. Jamieson (2021). Task-optimal exploration in linear dynamical systems. In *Proceedings of the 38th International Conference on Machine Learning*, PMLR 139, 10641–10652.

Wellman, M. P. (2006). Methods for empirical game-theoretic analysis. In *Proceedings of the 21st National Conference on Artificial Intelligence (AAAI-06)*, 1552–1555.

Whittle, P. (1953). Estimation and information in stationary time series. *Arkiv för Matematik* 2, 423–434.

Wu, Y., and P. Yang (2020). Optimal estimation of Gaussian mixtures via denoised method of moments. *Annals of Statistics* 48, 1981–2007.

A Proofs

Lemma 2.1. (i) For $v_{t+1} = v_t + \zeta_{t+1}$ and $r_t = v_t + \nu_t$, the steady-state prior variance P solves $P = P\sigma_\nu^2/(P+\sigma_\nu^2) + \sigma_\zeta^2$, i.e. $P^2 - \sigma_\zeta^2 P - \sigma_\zeta^2\sigma_\nu^2 = 0$, whose positive root is $P = \left(\sigma_\zeta^2 + \sqrt{\sigma_\zeta^4 + 4\sigma_\zeta^2\sigma_\nu^2}\right)/2$. The gain is $\eta = P/(P+\sigma_\nu^2) \in (0,1)$, which simplifies to (7), and the update $\hat{v}_{t+1} = \hat{v}_t + \eta(r_t - \hat{v}_t)$ has the stated form. (ii) $\nabla_\theta \frac{1}{2}H(x_t-\theta)^2 = -H(x_t-\theta)$, so $\theta_{t+1} = (1-\gamma)\theta_t + \alpha H(x_t - \theta_t) = (1-\eta)\theta_t + \eta q x_t$ with $\eta = \gamma + \alpha H$ and $q = \alpha H/\eta$.

Lemma 2.3. Omit zero-weight cohorts. For $a > 0$ put $b_i = \kappa q_i$ and

$$T = \operatorname{diag}\Big(\sqrt{a}, \sqrt{b_1/\eta_1}, \ldots\Big), \qquad u = \Big(\sqrt{a}, \sqrt{b_1\eta_1}, \ldots\Big)^\top.$$

The price-first transition becomes the symmetric matrix $J' = TJT^{-1} = \operatorname{diag}(0, r_1, \ldots) + uu^\top$, the input transforms to κu and $J'e_x = \sqrt{a}\,u$, so $h_j = (\kappa/a)e_x^\top (J')^{j+1} e_x$. The spectral decomposition of J' gives nonnegative weights summing to $h_0 = \kappa$, and Schur stability puts the eigenvalues in $[0, 1)$. The case $a = 0$ follows by continuity. A nonnegative mixture of decaying geometric sequences is completely monotone. Since J is nonnegative with row sums at most one, every price row of J^t has ℓ_1, hence ℓ_2, norm at most one, attained at $t = 0$.

Theorem 3.1. With zero initial states, $y_{i,t} = \sum_{s \le t} \eta_i r_i^{t-s} x_s$, so $\sum_i q_i y_{i,t} = \sum_{s \le t} k_{t-s} x_s$ and every response is determined by (a, κ, k), hence by M. Conversely, equal responses for all probes imply equal H, hence equal G, and partial fractions of G identify the distinct poles with nonzero residues w_i, i.e. M. After cancellation, $H = \kappa z D(z)/[(z-a)D(z) - \kappa N(z)]$ with $D = \prod(z - r_i)$ and $N(z) = \sum_i w_i z \prod_{j \ne i}(z - r_j)$, as in (4); at a retained r_i the denominator equals $-\kappa N(r_i) = -\kappa w_i r_i \prod_{j \ne i}(r_i - r_j) \ne 0$ (memory rates are positive because $\eta_i < 1$), so no root of D cancels. The numerator $\kappa z D(z)$ has no roots other than 0 and the r_i, so these are the only possible cancellations. The remaining numerator factor z cancels if and only if the denominator vanishes at 0; since $N(0) = 0$ the denominator there equals $-aD(0)$, which is nonzero for $a > 0$ and zero for $a = 0$. Hence the McMillan degree, which is the dimension of a minimal realization, is $1 + |\operatorname{supp}(M) \cap (0, 1)|$ for $a' \ne 0$ and $|\operatorname{supp}(M) \cap (0, 1)|$ for $a' = 0$ (an atom at rate 0, a cohort with $\eta = 1$, contributes the constant $M(\{0\})$ to G and no pole, so it shifts a to $a' = a + \kappa M(\{0\})$ and the argument above applies with a') (numerically, the rank of the Hankel matrix of $(h_j)_{j \ge 1}$ is $K + 1$ and K respectively for $K = 1, 2, 3$; `hardening/p1_checks.py`). For (ii), direct algebra gives

$$G_\delta(z) - G_0(z) = -\frac{qz(z-1)\delta^2}{(z - 1 + \bar{\eta})\Big((z - 1 + \bar{\eta})^2 - \delta^2\Big)}.$$

The denominator is the product of the three linear factors $z - r$ with $r \in \{1 - \bar{\eta}, 1 - \bar{\eta} \pm \delta\}$, all in $[\epsilon, 1 - \epsilon]$ by assumption; on $|z| = 1$ each satisfies $|z - r| \geq 1 - r \geq \epsilon$, so $|G_\delta - G_0| \leq 2q\delta^2/\epsilon^3$ there. Uniform stability gives $M_H = \sup |H_\theta| < \infty$, and the identity $H_\delta - H_0 = z^{-1} H_\delta H_0 (G_\delta - G_0)$, whose factor z^{-1} has modulus one on the circle, gives an ℓ_2 gain at most $C_{\mathcal{F}} \delta^2$. The mean difference of a power-constrained probe has squared norm at most $C_{\mathcal{F}}^2 \delta^4 PT$; the equal-noise KL is at most $C_{\mathcal{F}}^2 \delta^4 PT/(2\sigma_\nu^2)$, and Pinsker's inequality with total variation at least $1 - 2\alpha$ gives (9).

Proposition 3.3. (i) In the frequency domain, the closed-loop system is $x(z) = (1 - az^{-1})^{-1}[\kappa y(z) + \xi(z)]$, where $y(z) = \sum_{k=1}^K u_k(z) + v(z) + \xi_0(z)$ is aggregate order flow. Each cohort's flow is $u_k(z) = \frac{w_k}{1 - r_k z^{-1}} x(z) = \frac{w_k z}{z - r_k} x(z)$. Summing across cohorts, the net feedback operator is $F(z) = \sum_{k=1}^K \frac{w_k z}{z - r_k} = \int \frac{z}{z - r} dM(r)$, which depends on the economy only through the merged measure M. The closed-loop map from fundamental innovations and noise to aggregate order flow y_t is therefore an explicit rational function of $F(z)$. For any two economies $\mathcal{E}^A$ and $\mathcal{E}^B$ with identical merged measure M, $F^A(z) \equiv F^B(z)$ outside their poles, so the transfer function and spectral density $S_y(\omega)$ of aggregate anonymous order flow are identical: $S_y^A(\omega) \equiv S_y^B(\omega)$. Consequently, the autocovariance sequence $\Gamma_y(k) = \frac{1}{2\pi} \int_{-\pi}^{\pi} S_y(\omega) e^{ik\omega} d\omega$ coincides for all k. Any pricing rule $\kappa = \mathcal{K}(\Gamma_y)$ set by competitive risk-neutral market makers (minimizing mean squared pricing error, $\kappa^* = \mathrm{Cov}(x_t, y_t)/\mathrm{Var}(y_t)$) or by risk-averse dealers managing inventory variance depends only on Γ_y, so $\kappa^*(\mathcal{E}^A) = \kappa^*(\mathcal{E}^B)$. (ii) For two K-cohort economies whose moments agree up to order $2K - 2$, power-series expansion of the resolvent $(z - r_c - \Delta\eta s)^{-1}$ shows that the transfer functions differ by at most $\|F^A - F^B\|_\infty \leq C(\Delta\eta)^{2K-1}$ on $|z| = 1$, which implies $\|S_y^A - S_y^B\|_\infty = O((\Delta\eta)^{2K-1})$ and $\|\Gamma_y^A - \Gamma_y^B\| = O((\Delta\eta)^{2K-1})$. Since $\mathcal{K}$ is Fréchet differentiable at non-degenerate points, $|\kappa^*(\mathcal{E}^A) - \kappa^*(\mathcal{E}^B)| = O((\Delta\eta)^{2K-1})$. In the testing problem of Theorem 6.4, the perturbed dealer pricing κ^* shifts the closed-loop impulse response h_t by at most $O((\Delta\eta)^{2K-1})$, which is of the same order as $\|h^A - h^B\|_2$. Hence the Kullback–Leibler divergence between the two

economies under endogenous dealer pricing satisfies $D_{\mathrm{KL}}(P^A \| P^B) \le CN(\Delta\eta)^{4K-2}$, and Le Cam's two-point lemma preserves the sample complexity lower bound $N = \Omega((\Delta\eta)^{-(4K-2)})$. □

Theorem 4.1. In either positive economy $J \ge 0$ and $\|J\|_\infty \le 1 - \epsilon(1-S) < 1$, so the response is nonnegative with $\max h = \kappa$ and $\sum h = \kappa/(1-S)$. From the display above, $|G_\delta - G_0| \le 2q\delta^2/\epsilon^3$ on the circle; combined with the resolvent identity and positivity this gives $\|H_\delta - H_0\|_\infty \le K_*\delta^2$, and Parseval bounds $\|h^\delta - h^0\|_2$ by the same constant. Young's inequality gives $\|\mu_\delta - \mu_0\|_2 \le K_*\delta^2\|v\|_1$, and Gaussian KL with Pinsker gives (12). For (ii), the mean difference is a lower-triangular convolution matrix times v; each column is a shifted copy of a prefix of d_T, so $\|\mu_\delta - \mu_0\|_2 \le V\|d_T\|_2$ with equality for a pulse at $t = 0$. For two known equal-covariance Gaussian laws the minimax error is $\Phi(-\|\mu_\delta - \mu_0\|_2/(2\sigma))$, attained by the midpoint likelihood-ratio test. For (iii), the merged measures have equal mass and their first moments differ by $-q\delta^2$, which enters h_2 with factor κ^2 while h_0, h_1 coincide; a pulse A separates the third-observation means by $A\kappa^2 q\delta^2$. Part (iv) follows from the ℓ_2 gain bound, and (v) from superposition and $\max_j h_j = \kappa$, which give $\sup_t |\lambda x_t| \le |\lambda|\kappa V$, attained by a single pulse.

Lemma 4.2. After the uncontrolled update, the pre-probe displacement at step j is $\sum_{i<j} h_{j-i}v_i$; average linear-book execution adds $\kappa v_j/2$. Summing signed execution costs gives $\frac{1}{2}v^\top T_h v$. Each kernel $\lambda^{|i-j|}$ is a positive-definite AR(1) covariance for $0 \le \lambda < 1$, and a nonnegative mixture with positive total weight is positive definite. For a round trip the fundamental cash term cancels.

Theorem 4.3. Let $U_t = \sum_{j<t} v_j$. Positive responses give $0 \le x_t \le \kappa U_t$ and $I_t = -U_t$. Summing cash and marking inventory, $B - E_t = U_t x_t - \sum_{j<t} v_j x^0_{j+1} - \frac{\kappa}{2}\sum_{j<t} v_j^2 + cU_t$. Its quadratic part $v^\top A_t v$ has diagonal $h_{t-1-i} - \kappa/2$ and off-diagonal $\frac{1}{2}(h_{t-1-i} + h_{t-1-j} - h_{j-i})$ for $i < j$; monotonicity of h bounds every entry by $\kappa/2$. Hence $B - E_t \le \kappa U_t^2/2 + cU_t \le \kappa V^2/2 +$

cV, while position value is at most $V(P_* + \kappa V)$; the stated inequality guarantees margin. A single pulse attains both adverse bounds at its first response, so no larger allowance is robust. Solving the quadratic gives (16). The quantity requirement $V_B \gtrsim \delta^{-2}$ then forces $B \gtrsim \delta^{-4}$, and the pulse A funded with $(\mu + \frac{1}{2})\kappa A^2 + (c + \mu P_*)A$ attains it.

Proposition 4.4. The initial cash receipt is common to both alternatives and the dealer then holds $-A$. At time three $E_{3,r} = B + \kappa A^2/2 - cA - A^2 h_{r,2}$, so $E_{3,\delta} - E_{3,0} = A^2\kappa^2 q\delta^2 > 0$.

Lemma 4.5. Marking residual inventory at P_* gives $E_\infty = B - \frac{1}{2}v^\top T_h v - c\sum|v_t| \geq 0$. For each mode $\lambda \leq r$ the Toeplitz kernel has spectral density $(1-\lambda^2)/(1-2\lambda\cos\omega+\lambda^2) \geq (1-r)/(1+r)$, so $v^\top T_h v \geq m\|v\|_2^2$ and $\|v\|_2^2 \leq 2B/m$. The transfer bound gives mean-difference squared norm at most $K_*^2\delta^4\|v\|_2^2$, hence the KL bound; Pinsker gives (17). For one pulse, cash becomes $B_A - A(P_* + \kappa A/2) - cA$ and inventory A; at later marks $E_t = \mu A P_* + A^2 h_{t-1}$ and $\mu|I_t|P_t = \mu A P_* + \mu A^2 h_{t-1}$, so margin holds for $\mu \leq 1$.

Proposition 5.1. Parts (i) and (ii) follow from (18) with $\Delta h = \Theta(\delta)$ for a rate shift and $\Theta(\delta^2)$ for symmetric splitting, and from additivity of KL over independent pulse cycles. Part (iii) solves the session constraint for δ; substituting $d_T(\delta) \approx \delta^2 \tilde{d}_T$ gives the fourth-root approximation.

Theorems 5.2 and 5.3. For a flat round trip, telescoping the cash account gives $L_i(v) = \frac{1}{2}v^\top T_{h_i} v = \frac{1}{2\pi}\int g_i(\omega)|V(e^{i\omega})|^2 d\omega$, with $g_i > 0$ by Lemma 2.3. Parseval gives $\|(H_1 - H_0)v\|_2^2 = \frac{1}{2\pi}\int|\Delta\hat{H}|^2|V|^2 \leq \sup_\omega(|\Delta\hat{H}|^2/g_i)L_i(v)$, and the Neyman–Pearson requirement $\|(H_1 - H_0)v\|_2 \geq 2\sigma z_{1-\alpha}$ gives (22). For adaptive policies, Doob's optional stopping theorem applied to the log-likelihood-ratio martingale and the data-processing inequality give $\mathbb{E}_0[\Lambda_\tau] \geq \mathrm{kl}(\alpha, 1-\beta)$, while $\mathbb{E}_0[\Lambda_\tau] = \mathbb{E}_0\sum_t(\Delta\mu_t)^2/(2\sigma^2) \leq \sup(|\Delta\hat{H}|^2/g_0)\mathbb{E}_0[L_0]/(2\sigma^2)$, which gives (24).

Part (iv). Let G_θ be the $T \times T$ lower-triangular Toeplitz section of H_θ/κ, so the

observation covariance is $\Sigma_\theta = \sigma_w^2 G_\theta G_\theta^\top + \sigma_\nu^2 I$. A finite section of a bounded analytic symbol has norm at most its supremum on the circle, so $\|G_\theta\| \le \sup_\omega |H_\theta|/\kappa \le \bar{M}$ because $|H_\theta(e^{i\omega})| \le \sum_j h_j = \kappa/(1-S)$. The inverse of a lower-triangular Toeplitz section is the section of the reciprocal series; κ/H_θ is analytic on and outside the unit circle because H_θ is minimum phase (Theorem 6.10), and Lemma 2.3 gives $\operatorname{Re} H_\theta(e^{i\omega}) = \sum_\ell \omega_\ell(1-\lambda_\ell\cos\omega)/(1-2\lambda_\ell\cos\omega+\lambda_\ell^2) \ge \sum_\ell \omega_\ell/(1+\lambda_\ell) \ge \kappa/2$, so $\|G_\theta^{-1}\| \le 2$ and $\Sigma_\theta \succeq \sigma_{\mathrm{eff}}^2 I$ (the smallest singular value of G_θ is at least 0.522 in 762 random positive economies and sections; `hardening/p1_checks.py`). The Gaussian divergence splits into a mean and a covariance part. The mean part is $\frac{1}{2}\Delta\mu^\top \Sigma_1^{-1}\Delta\mu \le \|(H_1-H_0)v\|_2^2/(2\sigma_{\mathrm{eff}}^2)$. For the covariance part, the eigenvalues λ_k of $A = \Sigma_1^{-1/2}\Sigma_0\Sigma_1^{-1/2}$ are at least $c = \sigma_\nu^2/(\sigma_\nu^2 + \sigma_w^2\bar{M}^2)$, and $\lambda - 1 - \log\lambda \le (\lambda-1)^2/\lambda$, so it is at most $\|A-I\|_F^2/(2c) \le \|\Sigma_0-\Sigma_1\|_F^2/(2c\,\sigma_{\mathrm{eff}}^4)$. Writing $\Sigma_0 - \Sigma_1 = \sigma_w^2(\Delta G\, G_0^\top + G_1 \Delta G^\top)$ and $\|\Delta G\|_F^2 \le T\|h^\delta - h^0\|_2^2/\kappa^2 \le T K_*^2\delta^4/\kappa^2$ (proof of Theorem 4.1) gives a covariance part of at most $C_w T\delta^4$. The data-processing inequality turns errors (α,β) into $D_{\mathrm{KL}} \ge \mathrm{kl}(\alpha, 1-\beta)$, and part (ii) bounds $\|(H_1-H_0)v\|_2^2$ by $\sup_\omega(|\Delta\hat{H}|^2/g_i)L_i(v)$, which gives (23).

Theorem 5.4. Write a feasible flat probe as $v = Eu$ and expand $u = \sum_j c_j e_j$ in the G_i-orthonormal eigenbasis, so $L_i(v) = \sum_j c_j^2$ and the separation constraint reads $\sum_j \lambda_j c_j^2 \ge c_{\mathrm{req}}$. Write the top-two component as $(c_{L-2}, c_{L-1}) = r(\cos\theta, \sin\theta)$ with $r \ge 0$; replacing θ by $\theta+\pi$ flips d and s_θ and leaves (25) unchanged, so $\theta \in [0,\pi)$ suffices. Put $L_0(\theta) = c_{\mathrm{req}}/\lambda(\theta)$ and $D = L_i(v) - L_0(\theta)$. The constraint gives $r^2\lambda(\theta) \ge c_{\mathrm{req}} - \sum_{j\le L-3}\lambda_j c_j^2$, hence

$$D = r^2 + \sum_{j\le L-3} c_j^2 - L_0(\theta) \;\ge\; \sum_{j\le L-3}\Big(1 - \frac{\lambda_j}{\lambda(\theta)}\Big)c_j^2 \;\ge\; \Big(1 - \frac{\lambda_{L-3}}{\lambda(\theta)}\Big)\sum_{j\le L-3} c_j^2 \;\ge 0,$$

and $r^2 \ge L_0(\theta) - \frac{\lambda_{L-3}}{\lambda(\theta)}\sum_{j\le L-3} c_j^2 \ge L_0(\theta)\big(1 - \gamma(\theta)D\big)$. By Hölder's inequality with the sign vector s_θ, $c_{\mathrm{fee}}\|Eu\|_1 \ge c_{\mathrm{fee}} r\|Ed(\theta)\|_1 + c_{\mathrm{fee}}\sum_{j\le L-3} c_j\, s_\theta^\top E e_j$, and Cauchy–Schwarz with weights $1 - \lambda_j/\lambda(\theta)$ bounds the second term below by $-\sqrt{D\,W_\perp(\theta)}$. Using $r \ge \sqrt{L_0(1-\gamma D)_+} \ge \sqrt{L_0}\,(1-\gamma D)_+$ and the fact that the fee is nonnegative, with

$x = \sqrt{D}$,

$$L_i(v) + c_{\text{fee}}\|v\|_1 \;\geq\; L_0(\theta) + x^2 + \Big[F_{\text{dom}}(\theta)(1-\gamma(\theta)x^2)_+ - \sqrt{W_\perp(\theta)}\,x\Big]_+ \;\geq\; L_0(\theta) + \Psi(\theta).$$

Minimizing over the unknown θ gives (26), and $\Psi \geq 0$, $L_0(\theta) \geq c_{\text{req}}/\lambda_{L-1}$ give the last inequality. On the piece $\gamma x^2 \leq 1$ with $\gamma F_{\text{dom}} < 1$ the bracket is a convex quadratic in x with minimum $F_{\text{dom}} - W_\perp/(4(1-\gamma F_{\text{dom}}))$, which gives the closed form.

Corollary 5.5. Since $\Delta h_t = O(\delta^2)$, the variance difference $\sigma^2_{1,t} - \sigma^2_{0,t} = \sigma^2\xi[(h_1 + h_0) * v]_t[\Delta h * v]_t = O(\delta^2)$, and the variance term of the Gaussian KL, $\frac{1}{2}\sum(x - 1 - \log x)$ with $x = \sigma^2_{0,t}/\sigma^2_{1,t}$, is quadratic in that difference. Convex impact only increases dissipation.

Proposition 6.1. Standard Cramér–Rao for a Gaussian mean model with Jacobian J. For $K = 2$ the columns of J for the two rates become collinear as $\Delta\eta \to 0$; expanding to third order, the smallest singular value of J is of order $(\Delta\eta)^3$, so $[F^{-1}]_{kk} = \Theta((\Delta\eta)^{-6})$.

Theorem 6.2. Expanding $M_m(\theta_1) - M_m(\theta_0)$ under (28) gives zero for $m \leq 2$ and $4\tau p_- p_+(\Delta\eta)^3$ for $m = 3$. The map from moments to impulse responses is analytic with nonsingular derivative at the common limit, so $\|\Delta h\|_2 = \Theta(\tau(\Delta\eta)^3)$. Le Cam's two-point lemma with N independent Gaussian experiments gives (30). For the permutation-invariant loss, the identity matching of θ_0 and θ_1 costs $2\tau + (p_+ - p_-)\Delta\eta = 2\tau(1 + \Delta\eta/q)$ per atom and the swapped matching costs $\max(2p_+, 2p_-)\Delta\eta = (1 + 2\tau/q)\Delta\eta$; when both are at least 2τ, any estimate within τ of one economy in d_∞ is farther than τ from the other by the triangle inequality, and the two-point argument is unchanged. For (iii), the upper bound is Theorem 6.5; the Prony–Hankel determinant identity follows from $g_m = w_1\eta_1 r_1^m + w_2\eta_2 r_2^m$.

Theorem 6.4. Let (s^A, ω^A) and (s^B, ω^B) be the K-point Gauss and Gauss–Radau rules for the uniform weight on $[-1, 1]$, normalized to unit mass. Both are positive and exact for polynomials of degree $2K - 2$, and their moments of degree $2K - 1$ differ by $g_K \neq 0$. With

$r = r_c + \Delta\eta\, s$ and weights $W\omega$, the raw moments $\sum w_i r_i^m$ are polynomials in $\Delta\eta$ that agree below degree $2K-1$, so the kernels differ by $O((\Delta\eta)^{2K-1})$ with nonzero leading coefficient; analyticity of the kernel-to-response map gives $\|h^A - h^B\|_2 \asymp (\Delta\eta)^{2K-1}$, and Le Cam's lemma gives the bound. For the W_1 version, the two merged measures are the images of the fixed normalized rules under $s \mapsto r_c + \Delta\eta\, s$ scaled to mass W, so their distance is $W\Delta\eta$ times the distance $2c_K > 0$ between the rules, and the same two-point argument applies with the metric loss.

Lemma 6.8. Let $L = 2n-2$ and let p^* be the best uniform approximation of $|x|$ on $[-1,1]$ by polynomials of degree L, with error $E = E_L(|x|)$. By Chebyshev's equioscillation theorem (Rivlin, 1981) there are points $t_0 < \cdots < t_{L+1}$ at which $|x| - p^*(x)$ equals $\pm E$ with alternating signs. The divided-difference weights $w_i = 1/\prod_{j\neq i}(t_i - t_j)$ alternate in sign and satisfy $\sum_i w_i p(t_i) = 0$ for every polynomial p of degree at most L; multiplying by -1 if necessary, w_i has the sign of $|t_i| - p^*(t_i)$ for every i. Exactly n weights are positive and n negative, and the positive and negative parts have equal mass because constants are annihilated. Normalizing each part to unit mass and mapping $x \mapsto \rho(x+1)/2$ gives n-atom probability measures π_A (positive part) and π_B (negative part) on $[0,\rho]$ with equal moments of orders $0,\dots,L$. For the 1-Lipschitz function $f(r) = |r - \rho/2| = (\rho/2)|x|$,

$$\int f\, d(\pi_A - \pi_B) = \frac{\rho}{2}\int \big(|x| - p^*(x)\big)\, d(\pi_A - \pi_B) = \frac{\rho}{2}(E + E) = \rho E,$$

because p^* integrates to the same value under both measures, $|x| - p^* = E$ on the support of π_A and $-E$ on that of π_B. By Kantorovich–Rubinstein duality $W_1(\pi_A, \pi_B) \geq \rho E \geq c_* \rho / L$.

Theorem 6.9. Write $K(z) = \sum_m k_m z^m$, $D = 1 - az - \kappa z K$ and $H = \kappa / D$, which is the generating function of the recursion $h_t = a h_{t-1} + \kappa \sum_m k_m h_{t-1-m}$ with $h_0 = \kappa$. Since $k_m \geq 0$ and $\sum_m k_m = \int (1-r)^{-1} d\mu \leq W/(1-\rho)$, the ℓ_1 norm of $az + \kappa z K$ is at most $s < 1$, so $\|H\|_1 \leq \kappa/(1-s)$ and $\|1/H\|_1 = \|D\|_1/\kappa \leq (1+s)/\kappa$.

(i) Let $\mu_A = W\pi_A$ and $\mu_B = W\pi_B$ with π_A, π_B from Lemma 6.8; both lie in $\mathcal{M}(\rho, W)$. With $\nu = \mu_A - \mu_B$, $\sum_m (k_m^A - k_m^B) z^m = \int (1 - rz)^{-1} d\nu(r) = z^{2n-1} \int r^{2n-1} (1 - rz)^{-1} d\nu(r)$, because the first $2n - 1$ terms of the geometric series integrate to zero. On $|z| = 1$ the modulus is at most $2W\rho^{2n-1}/(1-\rho)$, so Parseval gives $\|k^A - k^B\|_2 \le 2W\rho^{2n-1}/(1-\rho)$. Next, $H_A - H_B = z(K_A - K_B)H_A H_B$, so by Young's inequality $\|h^A - h^B\|_2 \le \|H_A\|_1 \|H_B\|_1 \|k^A - k^B\|_2 \le C_h \|k^A - k^B\|_2$. The Kullback–Leibler divergence between the N-sample Gaussian experiments is $N\|h^A - h^B\|_2^2/(2\sigma^2) \le (N/\sigma^2)\, 2C_h^2 W^2 \rho^{2(2n-1)}/(1-\rho)^2$, which is at most one as soon as $2(2n-1)\log(1/\rho) \ge \Lambda_N - 2\log(1/\rho)$. Take n the smallest integer with this property and $n \ge 2$; then $2n - 2 < \Lambda_N/(2\log(1/\rho))$ when $\Lambda_N > 4\log(1/\rho)$. Since d is a norm on signed measures and $d(\mu_A, \mu_B) = W\, W_1(\pi_A, \pi_B) \ge W c_* \rho/(2n-2) > 2c_* \rho W \log(1/\rho)/\Lambda_N$, the Bretagnolle–Huber form of Le Cam's lemma (Tsybakov, 2009, Lemma 2.6 and Theorem 2.2) gives $\max_{j \in \{A,B\}} P_j(d(\hat{\mu}, \mu_j) \ge c_* \rho W \log(1/\rho)/\Lambda_N) \ge \frac{1}{4} e^{-1}$.

(ii) *Approximation.* Let f be 1-Lipschitz on $[0, \rho]$ with $f(0) = 0$, so $|f| \le \rho$. The function $g(x) = f(\rho(x+1)/2)$ is $(\rho/2)$-Lipschitz on $[-1, 1]$, and Jackson's theorem (Rivlin, 1981) gives a polynomial p of degree L with $\sup|g - p| \le (\pi/2)(\rho/2)/(L+1) \le \rho$, hence $\sup|p| \le 2\rho$. Its Chebyshev coefficients satisfy $|b_j| \le 2\sup|p| \le 4\rho$. The polynomial $q(r) = p(2r/\rho - 1)$ approximates f with the same error, and its monomial coefficients satisfy $\sum_j |q_j| \le \sum_{j \le L} |b_j|\, \|T_j(2r/\rho - 1)\|_{\text{coef}} \le 4\rho \sum_{j \le L} B(\rho)^j \le \rho B(\rho)^{L+1}$, where $\|\cdot\|_{\text{coef}}$ is the sum of absolute monomial coefficients, $\|T_j\|_{\text{coef}} \le (1+\sqrt{2})^j$, substituting $y = \alpha r - 1$ multiplies the coefficient sum of y^i by at most $(1+\alpha)^i$, and $B(\rho) - 1 \ge 4$. On the event $\max_{j \le L} |\hat{k}_j - k_j| \le \delta$ the true μ attains objective at most δ, so the minimizer $\hat{\mu}$ satisfies $\max_{j \le L} |\int r^j d\hat{\mu} - k_j| \le 2\delta$ by the triangle inequality. *Moment comparison.* If $\max_{j \le L} |\int r^j d\hat{\mu} - k_j| \le 2\delta$, then $\hat{\mu}$ has mass at most $W + 2\delta$ and

$$\int f\, d(\hat{\mu} - \mu) = \int (f - q)\, d(\hat{\mu} - \mu) + \sum_{j \le L} q_j \Big(\int r^j d\hat{\mu} - k_j \Big) \le \frac{\pi\rho(W + \delta)}{2(L+1)} + 2\delta\rho B(\rho)^{L+1}.$$

Kernel recovery. From $\kappa z K = 1 - az - \kappa/H$, the kernel values $k_0, \ldots, k_L$ are the coeffi-

cients of $(1 - az - \kappa/H)/(\kappa z)$, which depend on $h_0, \dots, h_{L+1}$ only. Work in the algebra of power series modulo z^{L+2} with the ℓ_1 norm, which is submultiplicative. If the averaged response is $\hat{H} = H + E$ and $\|1/H\|_1 \|E\|_1 \le \frac{1}{2}$, then $1/\hat{H} = (1/H)(1 + E/H)^{-1}$ and $\|1/\hat{H} - 1/H\|_1 \le 2\|1/H\|_1^2 \|E\|_1$, so $\max_j |\hat{k}_j - k_j| \le \delta \equiv 2((1+s)/\kappa)^2 (L+2) \max_{t \le L+1} |E_t|$. Each E_t is Gaussian with variance σ^2/N, so with probability at least $1 - \beta$, $\max_t |E_t| \le \sigma\sqrt{2\log(2(L+2)/\beta)/N}$. The true μ is feasible for the linear program with value at most δ, so the minimizer satisfies the moment condition above. *Choice of L.* With $L + 1 = \lfloor \log(N/\sigma^2)/(4 \log B(\rho)) \rfloor$, $B(\rho)^{L+1} \le (N/\sigma^2)^{1/4}$, so $\delta B(\rho)^{L+1} = O((\sigma^2/N)^{1/4} L\sqrt{\log(L/\beta)})$ and $\delta/(L+1)$ are $o(1/\log(N/\sigma^2))$, while $\pi\rho W/(2(L+1)) \le 2\pi\rho W \log B(\rho)/(\log(N/\sigma^2) - 4 \log B(\rho))$. For N/σ^2 larger than a threshold depending on $(\rho, W, a, \kappa, \beta)$ the side condition $\|1/H\|_1 \|E\|_1 \le \frac{1}{2}$ holds on the event and the total is below $3\pi\rho W \log B(\rho)/\log(N/\sigma^2)$, uniformly over $\mathcal{M}(\rho, W)$. The final statement combines (i) and (ii).

Theorem 6.11. For stationary Gaussian processes with spectral densities $f_j \propto |H_j(e^{i\omega})|^2$, the divergence rate is

$$D = \frac{1}{4\pi} \int_{-\pi}^{\pi} \Big(\varrho - 1 - \log \varrho \Big)\, d\omega, \qquad \varrho = f_1/f_0 = |H_1/H_0|^2.$$

By the resolvent identity, $H_1/H_0 - 1 = z^{-1} H_1 \Delta G$ with $\Delta G = G_1 - G_0$, and uniform stability bounds $|H_1|$ on the circle, so $u \equiv \varrho - 1 = 2\,\mathrm{Re}(z^{-1} H_1 \Delta G) + O(|\Delta G|^2)$. Write $G(z) = \sum_k w_k \phi(z, \eta_k)$, with w_k the cohort weights (any factor of η_k is absorbed in ϕ), where ϕ is analytic in η on a neighbourhood of $[\eta_{\min}, \eta_{\max}]$ uniformly for $|z| = 1$ (its pole $z = 1 - \eta$ stays inside the disk). Expanding ϕ around $\bar{\eta}$, $\Delta G = \sum_{j \ge 0} \frac{\partial_\eta^j \phi(z, \bar{\eta})}{j!} \Delta \tilde{M}_j$, where $\tilde{M}_j = \sum_k w_k (\eta_k - \bar{\eta})^j$. For the pair, $\Delta \tilde{M}_0 = \Delta \tilde{M}_1 = 0$ and $\Delta \tilde{M}_2 = q\delta^2$; for the moment-matched economies, $\Delta \tilde{M}_j = 0$ for $j \le 2$ and $\Delta \tilde{M}_3 = 4\tau p_- p_+ (\Delta\eta)^3$ by (29); for the K-cohort Gauss and Gauss–Radau economies with rates $\bar{\eta} + \Delta\eta\, s$ and weights $W\omega$, both rules integrate polynomials of degree $2K - 2$ exactly, so $\Delta \tilde{M}_j = W(\Delta\eta)^j \sum_k (\omega_k^A (s_k^A)^j - \omega_k^B (s_k^B)^j)$ vanishes for $j \le 2K - 2$ and equals $\pm W g_K (\Delta\eta)^{2K-1} \ne 0$ at $j = 2K - 1$. Higher terms are smaller by a factor

of the spacing: $|\Delta\tilde{M}_j| \le 2W(\Delta\eta)^j$, and analyticity of ϕ in η on a disk of radius $\varrho > 0$ around $\bar{\eta}$, uniformly on the circle, gives $|\partial_\eta^j\phi|/j! \le C\varrho^{-j}$, so for $\Delta\eta < \varrho/2$ the tail beyond the first nonzero term is $O((\Delta\eta)^{j_0+1})$. Hence $\sup_\omega |u| = O(\delta^2)$, $O(\tau(\Delta\eta)^3)$, respectively $O((\Delta\eta)^{2K-1})$, and $u = c(\omega)\epsilon + o(\epsilon)$ with $\epsilon = \delta^2$, $\tau(\Delta\eta)^3$ or $(\Delta\eta)^{2K-1}$ and $c(\omega)$ proportional to $2\,\mathrm{Re}\big(z^{-1}H_0\,\partial_\eta^{j_0}\phi(\cdot,\bar{\eta})\big)/j_0!$, $j_0 = 2$, 3 or $2K-1$. The function c does not vanish identically: $f = z^{-1}H_0\partial_\eta^j\phi$ is analytic outside the disk including at infinity, where it vanishes, and is not identically zero; an analytic function on the exterior whose real part vanishes on the circle is an imaginary constant, which here would have to be zero. For $|u| \le \frac{1}{2}$, $u^2/4 \le u - \log(1+u) \le u^2$, so $D \asymp \int c(\omega)^2 d\omega\,\epsilon^2$, which gives (i)–(iii). The sample requirements need a finite-T bound that is uniform in the separation, which we obtain directly. The first T observations are $N(0,\Sigma_j)$ with Σ_j the $T\times T$ Toeplitz matrices of the spectral densities $f_j = \sigma^2|H_j(e^{i\omega})|^2$. On the compact, uniformly stable family the transfer functions have no zeros or poles on the circle (Theorem 6.10), so $0 < \underline{f} \le f_j \le \bar{f} < \infty$ with constants that do not depend on the separation. A Toeplitz matrix has eigenvalues within the range of its symbol, so $\lambda_{\min}(\Sigma_0) \ge \underline{f}$ and $\|\Sigma_1 - \Sigma_0\| \le \sup_\omega |f_1 - f_0| = \sup_\omega f_0|u| \le \bar{f}\sup_\omega |u|$. With $A = \Sigma_0^{-1/2}\Sigma_1\Sigma_0^{-1/2}$, whose eigenvalues lie in $[\underline{f}/\bar{f}, \bar{f}/\underline{f}]$, the inequality $\lambda - 1 - \log\lambda \le (\lambda-1)^2/\lambda$ gives

$$D_{\mathrm{KL}}(P_1^T\|P_0^T) = \tfrac{1}{2}\sum_k(\lambda_k - 1 - \log\lambda_k) \le \frac{\bar{f}}{2\underline{f}}\,\|A - I\|_F^2 \le \frac{\bar{f}}{2\underline{f}}\,T\Big(\frac{\bar{f}}{\underline{f}}\Big)^2 \sup_\omega |u|^2 \le C\,T\epsilon^2,$$

with C independent of T and of the separation, because $\sup_\omega |u| = O(\epsilon)$ uniformly. The Bretagnolle–Huber inequality (Tsybakov, 2009, Lemma 2.6) then gives the stated lower bounds on T for every T, not only asymptotically. For a fixed pair, $T^{-1}D_{\mathrm{KL}}(P_1^T\|P_0^T) \to D$ and Stein's lemma (one error fixed) and the Chernoff bound (both errors vanishing) give the error exponents of the likelihood-ratio test, of the same order in the separation; that statement is pointwise in the pair.

Theorem 6.10. In the delay variable $\lambda = z^{-1}$, $H = \kappa\prod_i(1 - r_i\lambda)/\big[(1-a\lambda)\prod_i(1-r_i\lambda) - \kappa\lambda\sum_i w_i\prod_{j\ne i}(1-r_j\lambda)\big]$, which is (4) divided through by z^{K+1}; its nonstructural zeros are the

memory rates, which do not cancel under the nondegenerate assumptions of Theorem 6.10. In the z-plane there is also the structural timing zero at zero: the numerator is proportional to $z\prod_i(z-r_i)$. Stability puts the poles inside the disk. A rational spectral density $\sigma^2|H(e^{i\omega})|^2$ has a unique minimum-phase factor up to a unimodular constant, which is H; Theorem 3.1 then gives the merged measure.

Proposition 8.1. With tagged flow, $u_{i,t} = r_i u_{i,t-1} + w_i x_t$ holds exactly and involves only cohort i, so $\tilde{u}_{i,t} = r_i\tilde{u}_{i,t-1} + w_i x_t + e_{i,t}$ with the MA(1) error $e_{i,t}$. Prices are independent of the measurement noise, so $\mathbb{E}[Z_t e_{i,t}] = 0$ and, by the ergodic theorem for the stationary linear price process, the IV estimator is consistent when $\Sigma_i = \mathbb{E}[Z_t(\tilde{u}_{i,t-1}, x_t)]$ is nonsingular. Since $u_{i,t-1} = w_i\sum_{s\geq 0} r_i^s x_{t-1-s}$, $\mathbb{E}[x_{t-1}u_{i,t-1}] = w_i A_0$ and $\mathbb{E}[x_t u_{i,t-1}] = w_i A_1$, which gives Σ_i. The price is a nonnegative mixture of geometric responses to its innovation (Lemma 2.3), so $\gamma_x(s) = \sum_\ell c_\ell\lambda_\ell^s$ with $c_\ell = \sigma_\xi^2\kappa^{-2}\omega_\ell\sum_m\omega_m/(1-\lambda_\ell\lambda_m) \geq 0$; substituting, $\det\Sigma_i = w_i(\gamma_x(0)A_0 - \gamma_x(1)A_1) = w_i\sum_{\ell,m} c_\ell c_m(1-\lambda_\ell\lambda_m)/(1-r_i\lambda_m)$, each term positive. By summation by parts, $\sum_{t=1}^T Z_t e_{i,t} = \sum_{t=1}^T(Z_t - r_i Z_{t+1})\varepsilon_{i,t} + O_P(1)$. Since $(\varepsilon_{i,t})$ is i.i.d. with mean zero and variance σ_f^2 and independent of the whole stationary price process, the summands $m_t = (Z_t - r_i Z_{t+1})\varepsilon_{i,t}$ form a stationary ergodic martingale difference sequence with respect to $\mathcal{F}_t = \sigma((x_s)_{s\in\mathbb{Z}}, \varepsilon_{i,s} : s \leq t)$, with $\mathbb{E}[m_t m_t^\top] = \sigma_f^2[(1+r_i^2)\Gamma_0 - r_i(\Gamma_1 + \Gamma_1^\top)] = \Omega_i$, so the Billingsley–Ibragimov martingale central limit theorem gives (31). On a uniformly stable family the modes and weights lie in a compact set with $\lambda_\ell \leq \bar{\lambda} < 1$, so with $w_i \geq w_{\min} > 0$ we get $\det\Sigma_i \geq w_{\min}(1-\bar{\lambda}^2)\gamma_x(0)^2 > 0$ and $\|\Omega_i\|$ is bounded, uniformly in the spacing.

Proposition 8.3. (i) Under Sybil account-splitting, each sub-account $m \in \{1, \ldots, M_i\}$ reports flow $\tilde{u}_{i,m,t} = r_i\tilde{u}_{i,m,t-1} + \alpha_{i,m}w_i x_t + e_{i,m,t}$, where $e_{i,m,t} = \varepsilon_{i,m,t} - r_i\varepsilon_{i,m,t-1}$. Applying the instrumental variables estimator of Proposition 8.1 to each sub-account separately, the instrument covariance is $\Sigma_{i,m} = \mathbb{E}[Z_t(\tilde{u}_{i,m,t-1}, x_t)^\top] = \alpha_{i,m}\Sigma_i$, where Σ_i is the unsplit Jacobian matrix (31). Its determinant is $\det\Sigma_{i,m} = \alpha_{i,m}\det\Sigma_i = \alpha_{i,m}w_i\sum_{\ell,k} c_\ell c_k\frac{1-\lambda_\ell\lambda_k}{1-r_i\lambda_k} > 0$. For any non-vanishing sub-account with fraction $\alpha_{i,m} \geq \underline{\alpha} > 0$, $\det\Sigma_{i,m}$ remains bounded

away from zero independently of $\Delta\eta$. By the Billingsley–Ibragimov CLT, each sub-account estimator satisfies $\sqrt{T}(\hat{\beta}_{i,m} - \beta_{i,m}) \to N(0, V_{i,m})$ with $V_{i,m} = O(1)$. Summing weights across the M_i accounts gives $\hat{w}_i = \sum_m \hat{w}_{i,m} = w_i + O_P(M_i T^{-1/2})$, retaining the standard $T^{-1/2}$ parametric convergence rate without any dependence on rate spacing $\Delta\eta$. (ii) Under noise injection, trader i submits orders $u_{i,t} + \zeta_{i,t}$ with $\zeta_{i,t} \sim$ i.i.d. $N(0, \sigma_\zeta^2)$ independent of the price process. The regression error becomes $e_{i,t} + \zeta_{i,t} - r_i \zeta_{i,t-1}$, so the error variance entering Ω_i in (31) increases by $(1 + r_i^2)\sigma_\zeta^2 \Gamma_0$. The asymptotic variance of the IV estimator becomes $V_i(\sigma_\zeta^2) = \Sigma_i^{-1} \Omega_i(\sigma_\zeta^2) \Sigma_i^{-\top} = \Theta(\sigma_\zeta^2)$. To degrade the estimation error from $O(T^{-1/2})$ to the price-only rate $(\Delta\eta)^{-(2K-1)} T^{-1/2}$ (matching sample complexity $(\Delta\eta)^{-(4K-2)}$), group i must scale noise variance to $\sigma_\zeta^2 = \Omega((\Delta\eta)^{-(4K-2)})$. By Lemma 4.2, submitting zero-mean noise volume through a limit order book with positive impact $\kappa > 0$ yields expected round-trip execution losses $\mathbb{E}[\mathcal{L}_{\text{noise}}] = \frac{1}{2}\kappa\sigma_\zeta^2 = \Omega((\Delta\eta)^{-(4K-2)})$. Because expected trading gains from informational advantage are bounded by available liquidity $\Pi_i \le \bar{\Pi} < \infty$, injecting this volume produces net losses, making intentional noise camouflage strictly dominated. (iii) If group i curtails trading intensity by scaling $w_i \to \gamma_i w_i$ with $\gamma_i \in [\underline{\gamma}, 1]$, the Jacobian determinant scales by γ_i, giving $\det \Sigma_i(\gamma_i) = \gamma_i \det \Sigma_i \ge \underline{\gamma} \det \Sigma_i > 0$. The asymptotic variance V_i scales by $\gamma_i^{-2} \le \underline{\gamma}^{-2} < \infty$, so identification is preserved at rate $T^{-1/2}$. In the limit $\gamma_i \to 0$, group i's contribution to total feedback intensity $\kappa q_i = \kappa\gamma_i w_i/(1 - r_i) \to 0$ vanishes, which strictly decreases market persistence $S = a + \kappa \sum_j q_j$, thereby eliminating the stability threat without supervisory enforcement. □

Theorem 9.1. With J_θ upper triangular the eigenvalues are fixed. The stationary covariance solves $P = J_\theta P J_\theta^\top + \operatorname{diag}(s_1, s_2)$, and P_{11} is affine in s_2 with positive slope, so for each θ a unique $s_2(\theta) \ge 0$ gives $P_{11} = V$ when V exceeds the value at $s_2 = 0$. The one-step price variance is s_1 throughout, and $e_1^\top J_\theta e_2 = \theta$ gives $\sup_k \|e_1^\top J_\theta^k\| \ge |\theta|$.

Proposition 8.4. (i) $\mathcal{L}_{\text{disc}} < \mathcal{L}_{\text{price}}$ is $F + Mc < M\ell G$. (ii) T_p and T_t are decreasing in $\Delta\eta$, with $T_t \ge \bar{T}$ exactly when $\Delta\eta \le \Delta\eta_t$ and $T_p \ge \bar{T}$ exactly when $\Delta\eta \le \Delta\eta_p$. On

$\Delta\eta \le \Delta\eta_t$ both minima equal $\bar{T}$ and $G = 0$. On $(\Delta\eta_t, \Delta\eta_p]$, $G = \bar{T} - c_t(\Delta\eta)^{-2}$, which increases to $G^* = \bar{T} - T_t(\Delta\eta_p) > 0$. On $\Delta\eta > \Delta\eta_p$, $G = c_p(\Delta\eta)^{-(4K-2)} - c_t(\Delta\eta)^{-2} = (\Delta\eta)^{-2}\Big[c_p(\Delta\eta)^{-(4K-4)} - c_t\Big]$; the bracket is positive exactly when $\Delta\eta < \Delta\eta_0 = (c_p/c_t)^{1/(4K-4)}$, and $G'(\Delta\eta) = (\Delta\eta)^{-3}\Big[2c_t - (4K-2)c_p(\Delta\eta)^{-(4K-4)}\Big] < 0$ whenever $c_p(\Delta\eta)^{-(4K-4)} > c_t$, because $4K - 2 > 2$. So G decreases wherever it is positive beyond $\Delta\eta_p$, and is non-positive beyond $\Delta\eta_0$. Every superlevel set of G at a positive level is therefore an interval (the negative tail, where G rises back toward zero from below, never reaches a positive level), and the level $(F/M + c)/\ell > 0$ lies above $G = 0$, so the set where (i) holds is an interval contained in $(\Delta\eta_t, \Delta\eta_0)$ that contains $\Delta\eta_p$ whenever it is non-empty. (iii) The interval is non-empty if and only if the level is below the maximum, $F/M + c < \ell G^*$, that is $M > M^*$; the level decreases in M, so the superlevel set grows. □

B Notation

Symbol	Meaning
x_t	pricing error (price minus fundamental)
a	price persistence (anchor), $0 \le a < 1$
κ	price impact per unit of executed quantity
$q_i,\ \eta_i,\ r_i$	cohort weight, adaptation rate, memory rate $r_i = 1 - \eta_i$
$w_i = q_i\eta_i$	merged weight
$M = \sum_i w_i\delta_{r_i}$	merged rate–weight measure
$k_m,\ G(z)$	feedback kernel and its transform
$H(z),\ h_j$	closed-loop transfer function and impulse response, $h_0 = \kappa$
$J,\ S$	state transition matrix and $S = a + \kappa\sum q_i$
$v_t,\ V$	executed probe quantity and total quantity $\sum \lvert v_t\rvert$
d_T	impulse-response difference between hypotheses, first T taps
σ	observation-noise standard deviation

Symbol	Meaning
$\delta,\ \Delta\eta$	half-separation of a known pair; spacing of unknown cohorts
K_*	transfer-difference constant $2\kappa^2 q/[\epsilon^3(1-S)^2]$
T_h	Toeplitz matrix $(h_{\lvert i-j\rvert})$ of execution cost
$g_i(\omega)$	dissipation kernel $\mathrm{Re}\,\hat{H}_i(e^{i\omega})-\kappa/2$
$B,\ \mu,\ c,\ \phi$	capital, maintenance-margin fraction, proportional cost, taker fee
V_{eff}	per-trade execution capacity $\min(V_{\max}, D_{\max}/\kappa)$
$\hat{Q},\ \hat{S}$	reduced-form loop gain and persistence indices

C Internet appendix

C.1 A finite-window price-surveillance bound

Theorem C.1 (Finite-window abstention)**.** *Observe* $x_t = sx_{t-1}+w_t$, $w_t \sim N(0,\sigma^2)$, $x_0 = 0$, $t \le T$, *with* $s \in (0,2)$. *Let a possibly randomized rule* Ψ *output Stable, Unstable or Abstain, with* $\sup_{s<1} P_s(\Psi = \text{Unstable}) \le \alpha$ *and* $\sup_{s\ge 1} P_s(\Psi = \text{Stable}) \le \alpha$, $\alpha < 1/2$. *Choose* $0 < c_0 < 1$ *with* $c_0 e^{c_0}/(2\sqrt{2}) < 1-2\alpha$. *On* $B_T = (1-c_0/T, 1+c_0/T)$,

$$\inf_{s\in B_T} P_s(\Psi = \text{Abstain}) \ \ge\ 1 - 2\alpha - \frac{c_0 e^{c_0}}{2\sqrt{2}} > 0,$$

and for every finite T *the minimax worst-case error of binary rules over* $s \in (0,2)$ *is at least* $1/2$.

Proof. Pair $s_0 \in (1-c_0/T, 1)$ with $s_1 = s_0 + c_0/T$. Then $D_{\mathrm{KL}}(P_{s_0}\|P_{s_1}) = \frac{(s_1-s_0)^2}{2\sigma^2}\sum_t \mathbb{E}x_{t-1}^2 \le c_0^2 e^{2c_0}/4$, so Pinsker bounds total variation by $c_0 e^{c_0}/(2\sqrt{2})$. Error control gives $P_{s_0}(\text{Stable}) \le \alpha + \text{TV}$ and $P_{s_0}(\text{Unstable}) \le \alpha$, which yields the abstention bound. For binary rules, continuity in total variation and Le Cam's two-point bound give worst-case error tending to $1/2$. □

The interval B_T is an abstention layer that any error-controlled monitor must respect in the scalar Gaussian case; richer disclosure can remove it.

C.2 Passive distinction of a known pair in finite samples

Proposition C.2 (Passive distinction with zero controlled quantity)**.** *Keep the known pair, but add independent uniform innovations $w_t \in [-b, b]$ of variance $s_w^2 = b^2/3$ at each price update, observe $\tilde{x}_t = x_t + \nu_t$ with $\nu_t \sim N(0, \sigma^2)$, use independent three-step reset blocks and no probes, with $P_* > 3b$. Let $b_1 = a + \kappa q\eta$, $b_{2,0} = b_1^2 + \kappa q\eta(1-\eta)$ and $b_{2,\delta} = b_{2,0} - \kappa q\delta^2$. In each block the covariance of the first and third observations differs by $\Delta = \kappa q s_w^2 \delta^2$. With*

$$R = \max_{r \in \{0,\delta\}} \Big\{ (s_w^2 + \sigma^2)\big[s_w^2(1 + b_1^2 + b_{2,r}^2) + \sigma^2\big] + s_w^4 b_{2,r}^2 \Big\},$$

the midpoint test of the average product has worst-case error at most α whenever $n \geq 4R/(\alpha\Delta^2)$ blocks are observed: zero controlled quantity and order δ^{-4} blocks.

The Chebyshev bound is conservative but the order is real. With $b = 0.03$, observation s.d. 0.01 and $\delta = 0.10$, the bound asks for 3.9×10^7 blocks at $\alpha = 0.1$; in simulation the midpoint test still errs 32–38% of the time with 2×10^5 blocks. Independent resets and known innovation laws are extra conditions. The proposition does not show that anonymous public records identify participants; it shows that passive excitation can accumulate memory information while the probe account stays untouched.

Proof. In each block $x_1 = w_1$, $x_2 = b_1 w_1 + w_2$ and $x_3 = b_{2,r} w_1 + b_1 w_2 + w_3$, so $|x_t| \leq 3b$. Observation noise changes variances but not the covariance $s_w^2 b_{2,r}$. The product variance is the product of marginal variances plus the squared covariance plus a fourth-cumulant term $b_{2,r}^2(\mathbb{E}w^4 - 3s_w^4) \leq 0$ for uniform innovations, so it is at most R. Chebyshev bounds each midpoint error by $4R/(n\Delta^2)$. □

C.3 Blind inversion of the closed loop

1. *Impulse decoupling:* $g_t = h_{t+1} - ah_t$.
2. *Minimal realization:* form the Hankel matrix $H_{ij} = g_{i+j-1}$ and run Kung's eigensystem realization with a truncated SVD.
3. *Transmission zeros:* compute the zeros of the realization, which under known (κ, a) isolate the roots of the denominator of G.
4. *Weights:* with recovered rates, project $u_t = (x_t - ax_{t-1})/\kappa$ on the reconstructed filters $y_{k,t-1}$.

Without noise each step is exact: the realization has the order of the true system, its nonstructural transmission zeros are the memory rates (Theorem 3.1); the known timing zero is discarded, and the projection returns the weights, so the procedure removes the pole-shift bias of fitting poles directly to the price response. With noise, per-cohort recovery degrades as rates move together at the rate given by Theorems 6.4 and 6.5, which is why we do not report it as an estimator of individual cohorts.

C.4 Probing protocols

Table 34 compares four probing protocols at observation horizon $T = 4$ per cycle.

C.5 Coordinate-weighted output

For $J = \left(\begin{smallmatrix} a & \kappa q \\ \eta a & 1-\eta+\eta\kappa q \end{smallmatrix}\right)$ with $a = 0.5$, $\kappa = 0.8$, $q = 0.2$, $\eta = 0.3$ ($\rho(J) = 0.8224$, non-normal), the price output has $\sup_t \|e_x^\top J^t\| = 1$. The weighted output $C_w = (1, \lambda/\mathrm{MMR}) = (1, 160)$ with $\lambda = 0.8$ and $\mathrm{MMR} = 0.005$ has $\sup_t \|C_w J^t\| = \|C_w\| = 160.003$, attained at $t = 0$, then 122.32, 98.32, 80.79 [RECOMPUTED IN CLOSED FORM]. The factor 160 is a static coordinate weight; $C_w z$ is not a margin quantity because z contains no inventory. Margin must be computed on a cash-and-inventory ledger.

C.6 Real-noise injection: requirements at each spacing

Table 35 gives the requirements behind Table 17. N^* is read off the grid $2^{j/2}$, so neighbouring values differ by a factor $\sqrt{2}$, and a ratio of $\sqrt{2}$ or $1/\sqrt{2}$ between the two arms is one grid step. Entries at all N use exact empirical resampling via the multinomial count identity described in Section 6.5.

C.7 Numerical stress tests of results that are not machine-checked

Four results were proved on paper but not in Lean when these tests were run (Section 6 and Proposition 8.1); two of them, Theorem 6.5 and the general-memory theorem (Theorem 6.9, whose lower bound Lean proves with a different moment-matched pair than the one checked here), have since been machine-checked in Lean. We stress-tested each across randomly drawn parameter settings. These are numerical checks, not proofs and not pre-specified tests; no prediction failed (Table 36).

C.8 Probing a running market without resets

The probing results count independent experiments that each start from the reset state, which no outside participant can arrange in a live market. As a numerical check, not a proof and not a pre-specified test, we drop the resets. The observer sends a known i.i.d. ± 1 probe into one long stationary run of the closed loop of Table 17 ($K = 2$, $a = 0.5$, $\kappa = 0.8$, total weight 0.08, $r_c = 0.6$), which is also driven by its own price innovations $w_t \sim N(0, \sigma_w^2)$; prices are observed with noise $\nu_t \sim N(0, \sigma_\nu^2)$, $\sigma_\nu = 10^{-3}$ as in the reset experiments. The groups' state is never reset or known: its free response is part of the noise. The observer estimates the response taps $h_1, \ldots, h_4$ by least squares of the observed price on the current and 200 lagged probe values, then applies the same Prony estimator and success rule as Table 16 (every weight within 25% of the smallest true weight in at least 75% of 200 replications). Because the probe is white and independent of the noise, the least-squares error is asymptotically

Gaussian with covariance Γ_e/T, where Γ_e is the autocovariance of the regression residual; the required number of periods T^* is computed from this law, which we checked against 120 simulated runs of 10^5 periods at $\Delta\eta = 0.5$ (standard deviation of the tap estimates 5.44 to 6.05×10^{-6} against 5.35×10^{-6} predicted).

Table 37 shows that the order survives: the number of periods grows with fitted slopes -6.26 and -6.22 in $\Delta\eta$, against -5.87 for reset experiments and -6 in theory. The constant is larger, because the market's own innovations and the groups' unknown state add noise: with $\sigma_w = \sigma_\nu$ a running market needs 1.4 to 2.8 times as many periods as the reset benchmark needs experiments, and with $\sigma_w = 3\sigma_\nu$ four to six times as many (the grid of T and N has steps of $\sqrt{2}$). Since every reset experiment itself uses at least $2K + 1$ periods, the running market is not slower in calendar time. What the check does not show is that a live prober can trade a ±1 sequence every period at the prices assumed, or that other traders ignore it; those are the strategic-response caveats of Section 2.2.

D Reproducibility: source of every number

The replication package contains numbered Colab notebooks and stand-alone scripts; its README lists them in run order. "Part N" below and in table notes refers to notebook N (Part 1: theorem checks and spot tests; Part 2: BTCUSDT futures depth; Parts 4–13: the cross-section, replications, fresh-day tests, second venues, the CFTC test, the certified guard and post-hoc checks). Result identifiers such as `X6` or `U6` name the JSON file that the notebook writes. "Kind" distinguishes closed-form or simulation results (recomputed), results from public market data (real data), and audits.

Result	Source	Kind
Complete monotonicity, 892 economies	`L26_complete_monotonicity`	recomputed
$d_2 = -\kappa^2 q\delta^2$	`P1.7-T1_third_tap`	recomputed

Result	Source	Kind
Ledger identity; T_h positive definite	`R2-T1`, `R2-T3`	recomputed
Figure 3	`make_figs2.py`	recomputed
$\delta_{\min}(3)$, Table 8 recomputed columns, 61.14, 29.94	`recompute_theory.py` → `R_P1_theory_constants.json`	recomputed
Pulse benchmarks (271, 58, 12, 11 pulses), Figures 2–4	`make_figs3.py` → `R_P1_pulse_benchmarks.json`	recomputed
Table 6, Table 14, worked example, passive simulation	`recompute_extra.py` → `R_P1_extra.json`	recomputed
Daily noise (Table 11)	`P1.3-T1_sigma_obs`	real data
Table 13 recomputed columns	same	recomputed
K-cohort slopes (Table 15)	`T1_K_cohort_rate`	recomputed
ESPRIT slope −6.15	`P1.2-T1_esprit_rate`	recomputed
Minimum phase, 1,733 economies; rank 6 of 6	`T2a`, `T2b`	recomputed
Passive slopes −4.21, −5.93	`T3_passive_rate`	recomputed
Tagged-flow slope 0.011	`T4_disclosure`	recomputed
Spectral blindness gains	`T5_spectral_blindness`	recomputed
Noise and capital grid (Table 10)	`P1.3-T1_sigma_obs`	real data
Loop gain and placebo (Table 30)	`P1.6-T1_ema_placebo`	real data
Estimator table (Table 31)	`P1.5-T1`, `T2`, `T3`	recomputed
Data integrity (111 files)	`R2-T9_data_integrity`	audit
Tables 7, 8 (certified, feasible, adaptive columns)	`recompute_p1_capital.py` → `R_P1_capital_ledger.json`	recomputed
Table 20; resolution floors; Table 34	`recompute_p1_protocols.py` → `R_P1_protocols.json`	recomputed

Result	Source	Kind
SPRT capital figures	`recompute_p1_sprt.py` → `R_P1_sprt.json`	recomputed
Table 12 (N^*, σ_{crit})	`recompute_crb.py` → `R_P1_crb.json`	recomputed
Table 12 (ML columns), Table 13 (estimators)	`recompute_p1_estimators.py`	recomputed
Symmetric multi-asset network (Section 9)	`recompute_oct20.py`	recomputed
Proof audit: K_* bound and d_2 identity (Theorem 4.1), moment gaps (Theorem 6.2), spectral-blindness example (Theorem 9.1)	`recompute_audit.py` → `R_audit_oct2026.json`	recomputed
Theorem 6.5, Table 16	`kcohort_upper.py` → `R_P1_kcohort_upper.json`	recomputed
Lem. 6.8, Thm 6.9, Tab. 18	`general_minimax.py`	recomputed
Theorem 8.2 numbers	`disclosure_value.py` → `R_P1_disclosure_value.json`	recomputed
Tables 28, 29, 32	Part 8: `V2_spot_panel`, `V3_fee_events`	real data
Table 25, top panel	Part 4: `X9_disclosure_metrics`	real data
Table 25, bottom panel; class-position forecasts	Part 6: `Z6_C_identification`, `Z6_secondary`; `P6_fresh_days.json` (65acda36…)	real data

Result	Source	Kind
Table 23	Part 11: `C3_cot_data`, `C4_cot_disclosure`; `P11_cot.json` (6f1d6f35…)	real data
Basis-trade signature; real-data Cochrane–Orcutt	Part 13: `R6`, `R7`	real data (post hoc)
Feedback across 30 perpetuals	Part 4: `X3_T6_cross_section`	real data
Theorem 5.4 full bound (Table 9), direct minimization	`hardening/rh_bound.py`, `rh_check.py` → `R_P1_rh_explicit.json`	recomputed
Minimal dimension at $a = 0$; Toeplitz sections; passive slopes for $K \leq 4$; IV determinant	`hardening/p1_checks.py` → `R_P1_hardening_checks.json`	recomputed
Stylized nonlinear and trained online rules (Table 19, Figure 8)	`recompute_p1_learners.py` → `R_P1_learners.json` (Panel A); `validate_trained_learners_hardened.py` → `R_P1_trained_learners_hardened.json` (Panels B–C); specification `P1_trained_learners.json` (3e8bf30c…)	recomputed
Disclosure in dollars (Table 21)	`recompute_p1_cat_costs.py` → `R_P1_cat_disclosure_dollars.json`	recomputed

Result	Source	Kind
Out-of-sample post-shock return test (Section 8.5)	`recompute_prereg_oct21_22nd.py` → `R_P1_disclosure_oos.json`; `P1_disclosure_oos.json` (f0182ce5…)	real data
Out-of-sample cohort-state trajectory test (Table 26)	`recompute_prereg_oct21_followup.py` → `R_P1_disclosure_state_oos.json`; `P1_disclosure_state_oos.json` (5136cc6b…)	real data
Proposition 8.5, Table 22	`theory/targeted_intervention.py` → `R_P1_targeted_intervention.json`	recomputed
Table 13, column N^*_{CRB}; Table 12, spacing-0.10 three-cohort row	`breakthrough_experiments/results_exp17_minimax_lecam_bound.json`, `results_exp15_cohort_information.json`	recomputed (earlier run)
Table 9, best feasible probes and $L = 320, 480$	`R_P1_capital_ledger.json`, `R_P1_cquadL.json` (reproduced by `hardening/rh_quad.py`)	recomputed
Table 7, peak margins of the eigenvector probes	`breakthrough_experiments/results_exp14_capital_frontier_exact.json`	recomputed (earlier run)
Table 18	`theory/R_P1_general_minimax.json`	recomputed

Result	Source	Kind
Table 13 and Table 12, estimator columns	`R_P1_estimators.json`, `R_P1_estimators_crb.json`	recomputed
Theorem 3.1, numerical check of the equivalence	`R2-T5_merged_invariance`	recomputed
Section 8.5 and Table 1: horizon-matched test on 30 new markets (29 of 29, median 7.3, smallest 2.6) and jackknife-corrected persistence (1.0008)	`thesis/evidence/C2_results.json` (specification `thesis/checks/prereg/C2_margin_vs_identify.json`)	real data
Section 8.5: exploratory 48- and 96-lag refit (median 7.5)	`thesis/evidence/pilot3_horizon_matched_results.json`	real data (exploratory)
Section 8.5 and Table 1: 105-day follow-up (median ratio 74.9; 9 identified; 10 of 30, 23 required)	`thesis/evidence/E1v2_results.json`	real data
Section 8.5 and Table 1: 58-day follow-up (22 of 29; 20 of 29, 22 required)	`thesis/evidence/D1_results.json`	real data
Section 6.5, Tables 17, 35 and 1: real-noise injection (slopes −6.01, −9.81; tagged 256–724)	`p1_injection/injection_test.py`, `p1_injection/results/I1_injection_test.json` (specification `p1_injection/prereg/P1_injection.json`, 596a700d…)	semi-synthetic

Result	Source	Kind
Supplement C.7 and Table 36: numerical stress tests	`p1_stress/stress_theorems.py`, `p1_stress/results/S1_stress_tests.json`	numerical check
Appendix C.8 and Table 37: probing a running market without resets	`p1_running/running_market.py`, `p1_running/results/R1_running_market.json`	numerical check
Theorem 3.1(i), Theorem 6.4 and Proposition 8.1 (machine-checked parts)	`research_extension/current_manuscript_lean_20261008/CurrentProofs/Paper1MergedModel.lean`, `Paper1Identification.lean`, `Paper1MomentGap.lean`, `Paper1PoleCount.lean`, `Paper1FiniteIdentification.lean`, `Paper1Tagged.lean`, `Paper1Testing.lean`, `Paper1LowerBound.lean`, `Paper1MatchedPair.lean`	Lean 4 proof
Theorem 6.5 (machine-checked, every K)	`research_extension/current_manuscript_lean_20261008/CurrentProofs/Paper1PronyUpper.lean`, `Paper1PronyMoment.lean`, `Paper1PronyDeconv.lean`, `Paper1UniformInverse.lean`	Lean 4 proof

Result	Source	Kind
Theorems 3.1(ii), 9.1, 6.10 (zeros), 8.2(ii) (tagged upper bound) and Proposition 8.1 (instrument determinant), machine-checked parts	`research_extension/current_manuscript_lean_20261008/CurrentProofs/Paper1MergedSeparation.lean`, `Paper1Blind.lean`, `Paper1MinPhase.lean`, `Paper1TaggedRates.lean`, `Paper1TaggedIV.lean`	Lean 4 proof
Theorem 6.11(i) and (iii), sample bounds for passive data (machine-checked)	`research_extension/current_manuscript_lean_20261008/CurrentProofs/Paper1PassiveCost.lean`, `Paper1PassiveKCohort.lean`, `Paper1PassiveRecord.lean`, `Paper1GaussianTriangular.lean`	Lean 4 proof
Theorem 6.9, both parts (machine-checked)	`research_extension/current_manuscript_lean_20261008/CurrentProofs/Paper1GeneralMemoryLower.lean`, `Paper1Jackson.lean`, `Paper1JacksonPoly.lean`, `Paper1GeneralMemoryUpper.lean`, `Paper1GeneralMemoryRate.lean`	Lean 4 proof

Result	Source	Kind
Proposition 8.5, supervisory separation (machine-checked for two groups; part (i) for every K)	`research_extension/current_manuscript_lean_20261008/CurrentProofs/Paper1SupervisorySeparation.lean`	Lean 4 proof
Theorem 4.1(i) and Lemma 4.5, quantity and capital lower bounds (machine-checked)	`research_extension/current_manuscript_lean_20261008/CurrentProofs/Paper1QuantityCapital.lean`	Lean 4 proof
Section 8.4: post hoc robustness (futures prices, roll weeks, 26- and 52-week blocks)	`p11/robustness_cot.py`, `p11res/results/C4R_cot_robustness.json`	real data (post hoc)
Section 8.4 and Table 1: two-root test (2 of 8; smaller root -0.16 to 0.21)	`p11/two_root_test.py`, `p11res/results/C6M_cot_two_root.json` (specification `p11/prereg_memory/P11M_two_root.json`, e6633588…)	real data
Section 8.4 and Table 1: replication on 12 new markets (H1 12 of 12, H2 11 of 12, H3 3 of 12)	`p11/replication_cot.py`, result `C6R_cot_replication.json` (specification `p11/prereg_replication/P11R_cot_replication.json`, 591e71c9…)	real data

Result	Source	Kind
Figure 1 (paper in one figure)	papers/paper1/figs/make_fig_summary.py (inputs p1_injection/results/I1_injection_test.json, p11res/results/C7_data_equivalence.json)	registered results and post hoc calculation
Section 8.4, Table 24: price data needed to match the class reports (median factor 1,353; post hoc)	p11/data_equivalence.py → p11res/results/C7_data_equivalence.json (inputs C4_cot_disclosure.json, C6R_extract.json)	post hoc calculation
Section 8.5 and Table 1: October 2026 tests N2 (1 of 30) and N3 (5 of 30; 21 of 30; 30 of 30)	p1_oct/oct_tests.py, p1_oct/results/zip/results/P10_N2_feedback.json, P10_N3_combined.json (specification p1_oct/prereg/P10_october2026.json, 72cbc8ba…; power check P10_power_check.json)	real data
Section 8.4 and Table 1: class-flow test (3 of 8; $F \geq 10$ in 7; $w \neq 0$ in 6)	p11/tagged_flow_test.py, p11res/results/C5F_cot_tagged_flow.json (specification p11/prereg_flow/P11F_cot_tagged_flow.json, 052a63f2…)	real data

Development revision and validation status: 29 October 2026

This is a working draft for author and specialist review. The revision corrects hidden-state drift accounting, distinguishes cumulative displacement from peak mark-to-market loss, and restricts claims to their stated model and observation contracts. It does not declare every theorem independently verified or every research objective complete.

Returned execution record. The author returned a completed full-mode notebook archive using source commit `19805d7e`, seed 20261029, Python 3.13.18 and NumPy 2.1.3. All 10 exported-file hashes and the four analytical source modules were checked. The run completed E1, E2, E3, E4 and E6; its log reports 18 targeted tests passing. This is an execution reproduction of the same seeded code, not independent scientific validation. No real-market CSV was provided, and the GPU, known-source-feedback and prospective studies were not run.

What the checks establish. E1 cross-checks the corrected drift identity against a resolvent and exercises a sufficient invariant-box guard. E2 demonstrates filter-induced persistence under martingale nulls. E3 compares known alternatives with equal aggregate risk; it is an illustrative witness, not a uniform rate theorem. E4 matches baseline capacity on training episodes and reports test capacity differences; the guard and volatility control both have zero observed budget breaches in the three bounded-shock synthetic regimes. This does not establish universal guard superiority or economic welfare.

Coverage and the observation contract. E6 uses 1,000 replications, 32 reset responses and 512 taps. The nominal 95% interval assumes independent Gaussian tap errors and a valid memory envelope. The intervals in the table measure Monte Carlo uncertainty about coverage, not financial risk.

Setting	Coverage	MC 95% interval	Contract holds
Independent Gaussian	94.6%	[93.0, 95.8]	yes
Correlated tap errors	48.3%	[45.2, 51.4]	no
Heavy-tailed readings	95.1%	[93.6, 96.3]	no
Biased readings	25.6%	[23.0, 28.4]	no
Unseen slow memory	0.0%	[0.0, 0.4]	no

Failures under correlated errors, biased readings and a false memory floor expose the contract's limits. The heavy-tail outcome in this configuration is not a general robustness result. All outcomes, including negative ones, are retained.

Remaining obligations. Independent review of the load-bearing identification and certification proofs, continuous numerical extrema, economic primitives, the public-observation bridge and policy counterfactuals remains necessary. The current two-paper name-map audit finds 54 of 57 formal environments with existing declaration mappings, three unmapped statements and seven revised statements requiring semantic review. A label mapping is not a new Lean proof. The experiment design and full results are retained under `studies/economic_review_20261029/`.

Economic interpretation and information-to-decision extension

What this revision adds

The information question is distinct from choosing a preferred amount of risk. For a specified financial decision, the relevant issue is which observations support an exposure constraint, what remains unidentified, and whether obtaining additional information changes the attainable decision after costs. This does not introduce portfolio choice or control without full identification as new ideas. The general identification-versus-decision-value result still

requires independent proof and a comparison with the closest existing results.

Time conventions and the institution

The new application is a price-taking, unlevered own-account investor, not a venue-wide supervisor. It begins with cash, buys once, holds a fixed position for 60 one-minute model periods, and closes fully. Round-trip linear and quadratic charges enter cash and wealth. The account cannot change the market's feedback, and the experiment does not measure customer liquidity or welfare. This is separate from the paper's dealer/venue admission contracts.

Quantity	Interpretation
Sampling interval	One minute in the new synthetic account; other historical studies retain their stated sampling clocks.
Historical lookback	Often six hours in the existing forecasting studies; not a holding period.
Forecast horizon	Often the next hour's adverse excursion; not an executed round-trip trade.
New holding period	Exactly 60 model minutes, with explicit entry and exit charges.
Other theory horizons	The existing finite-horizon theorems retain their declared H; no universal holding period is imposed.

Information and decision design

The new comparison uses controlled independent reset observations of 12 impulse taps, a known plant and reading variance, and a declared finite candidate class. It does not treat public market trading as a reset experiment. Common training and validation paths supply 768 and 384 additional reset episodes per repetition; 512 further paths evaluate the selected actions. These are in addition to the reported tap-reading budget, not evidence of learning from 32 observations in total. Verified disclosure restricts the minimum-memory rate to an audited band. The oracle knows the true response, not future shocks. All certificate arms use the same training objective.

The account selects a position by empirical expected logarithmic terminal wealth, subject to its information-supported exposure restriction. This is a finite-horizon application of logarithmic utility, not a newly proved Kelly growth-optimality theorem. Equal response bounds alone do not imply equal expected returns or identical utility-optimal portfolios.

The new account has initial wealth of 1,000 dollars and an initial price of 100 dollars; its position grid runs from zero to 9.9 units. Its round-trip dollar charge is $c(I) = 0.002I + 0.0002I^2$. For bounded external flow, the constraint adds this charge to position times the largest admissible absolute price response over the holding horizon. The resulting bound concerns loss from initial wealth at each date. It is not peak-to-trough drawdown, protection against arbitrary unbounded shocks, or a rolling deployment guarantee.

Returned execution and proof scope

The returned full-mode archive uses seed 20261029. All 35 manifest entries and eight core source files were checked. The run contains 20 independent repetitions per regime, not 16,800 independent experiments: the 16,800 policy rows reuse tapes across policies, budgets and observation counts. Twelve targeted checks passed. Eight conditional decision-logic lemmas compiled with Lean 4.19.0 and their axiom printouts report no axioms. This is a seeded execution reproduction. It neither completes the older manuscript's Lean coverage nor formalizes the Gaussian confidence argument, market dynamics, general minimax separation, portfolio optimality or welfare.

The displayed slice uses 32 reset episodes and a one-dollar loss budget from initial wealth of 1,000 dollars. It is an illustrative presentation of an exploratory grid, not a primary selected before the results were seen. All regimes, budgets, policies and negative outcomes remain in the result archive. The inputs and complete outputs are retained with the verification record in the new information-decision study directory.

Aggregate information and unresolved cohorts: returned simulation

Regime	Unique candidate (%)	Certified units	Oracle units
No feedback	0.0	9.405	9.900
Moderate feedback	5.0	5.000	5.000
Near boundary	0.0	1.200	1.200

Here "unique" means one remaining candidate in the declared finite class, not identification of arbitrary real-world algorithms. In the moderate and near-boundary configurations, the aggregate certificate reaches the oracle's grid capacity in this slice even though the candidate decomposition is usually unresolved. This is an illustrative decision comparison; it does not prove a uniform rate separation for continuous classes.

Disclosure cost and the negative result

Paid minus aggregate	Mean (bp)	Lower (bp)	Upper (bp)
No feedback	-0.100	-0.100	-0.100
Moderate feedback	-0.126	-0.126	-0.125
Near boundary	-0.100	-0.100	-0.100

The intervals are exploratory percentile bootstraps over repetition-level paired means, not simultaneous intervals across the grid. In this slice the unpaid disclosure does not expand capacity, while the assumed one-cent charge reduces net log wealth. This run therefore does not establish a positive net economic value of disclosure. The price is hypothetical, not a calibrated reporting cost. Stronger disclosure need not be worth purchasing when aggregate information already supports the relevant decision.

Assumption failures and calibration diagnostics

The table reports finite-sample frequencies for all seven regimes, combining the three reading budgets. Each band outcome is counted once per repetition and reading budget (60 outcomes

per regime); repeating the same outcome across four financial loss budgets would not create further observations. The three reading budgets within a repetition are not an independent replication.

Regime	Contract valid	Tap-band hit (%)	Empty set (%)
No feedback	Yes	96.7	3.3
Moderate feedback	Yes	93.3	5.0
Near boundary	Yes	96.7	3.3
Hidden slow memory	No	98.3	11.7
Correlated readings	No	98.3	1.7
Biased readings	No	0.0	100.0
Behaviour change	No	95.0	5.0

The nominal one-decision band coverage is 95 percent under the valid contract; these small, dependent-budget frequencies are diagnostics, not a new coverage theorem. A tap-band hit concerns the observed response, not membership in the finite candidate class or stability of future behaviour. In particular, hidden slow memory falls outside the declared class even when its first 12 taps fit the bands. Under behaviour change the readings describe the old response. Biased readings empty the set throughout this run and force the certificate policy to abstain. Correlated readings use an unverified standard-error model. None of these four regimes retains the stated guarantee; a high observed band hit or zero observed financial breaches does not make an invalid contract valid.

Pending evidence: deliberately empty result cells

The following numerical cells remain empty. The returned archive contained no market CSV, dynamic-inventory experiment, institution-level intervention or external replication. The general theorem and independent novelty review also remain open; a table placeholder is not a theorem statement or a proof.

Pending study	Estimate	Uncertainty
Untouched market-data test beyond order-flow imbalance and depth		
Full changing-inventory, execution and funding comparison		
Dealer/venue intervention with endogenous market response		
Net value under empirically calibrated disclosure costs		
Independent implementation and replication		

Further data cannot by themselves prove a general identification or decision-value theorem. Its exact statement, proof and literature priority require separate review. No publication rating follows from the returned run.

Finite-horizon information, cohort ambiguity, and net inventory-service value

Status and contribution boundary. This appendix supplies a restricted analytic separation and its financial consequence. The proof below is written in full; exact rational and numerical checks accompany it. It has not received independent proof review and is not a Lean formalization. Its priority is unconfirmed. Control without unique identification, task-directed exploration, Gaussian functional estimation, moment matching, and the value of information are established ideas (van Waarde et al., 2020; Wagenmaker et al., 2021; Donoho, 1994; Heinrich and Kahn, 2018; Blackwell, 1953). The candidate extension is their explicit combination for a finite-horizon inventory commitment, with a hard stress budget and the information bill deducted from both capital and profit. It does not close the broader open theorem of the preceding drafts.

The observation contract and the financial institution

Fix $H \geq 4$, $0 \leq a < 1$, $\kappa > 0$ and $0 \leq Q \leq \overline{Q} < (1-a)/\kappa$. There are arbitrarily many finite cohorts, with $q_i \geq 0$, $\sum_i q_i = Q$, and $0 < \eta_i \leq 1$:

$$x_{t+1} = ax_t + \kappa\Big(\sum_i q_i y_{i,t} + u_t\Big), \qquad y_{i,t+1} = (1-\eta_i)y_{i,t} + \eta_i x_{t+1}.$$

Initial states are reset to zero. Let $h_0, \ldots, h_{H-1}$ be the price response to a unit impulse, with $h_0 = \kappa$, and put $G = \sum_{t<H} h_t$. Positivity implies $h_t \geq 0$. No positive lower bound on η_i is imposed.

The observer purchases N independent reset experiments, each lasting H model periods and returning all H taps with independent Gaussian errors of known variance σ^2. Thus the tap means are $\overline{h}_t = h_t + \sigma Z_t/\sqrt{N}$, with independent standard normal Z_t. The same observations are available for identifying the cohorts and for choosing the financial position. One experiment is not one scalar reading: it returns H readings and costs c_H, including the declared operating and time charges. Total information cost is $F = Nc_H$. This observation contract is not an ordinary exchange tape; longer horizons cost both time and readings. Plant parameters, the feedback cap, reset, and stationarity are assumptions. Measurements are acquired ahead of a common future service-start deadline. Cash can earn the common reference rate during acquisition; the model does not include earlier service opportunities. The stated measurement price is exogenous, not a calibrated self-impact cost for actual market probes.

The institution is a price-taking, unlevered inventory-warehousing desk. It buys $I \geq 0$ units at P_0, holds them for H periods, and closes. It earns a prequoted service fee pI *after* settlement. Round-trip execution and carry costs are $bI + dI^2$, where $b \geq 0$ and $d > 0$. For example, $b = b_{\text{trade}} + P_0 r_f H/(365 \cdot 24 \cdot 60)$ for one-minute model periods. The service premium is unavailable as collateral until after settlement. Both trading costs and F must be prefunded. Initial cash is W, the stressed-settlement loss budget is B, and $|u_t| \leq \overline{u}$, with

$\overline{u} > 0$. Assume $P_0 > \overline{u}\,\overline{G}$, so the asset price remains positive. This desk does not alter market feedback, depth, or another participant's behaviour. The service fee is exogenous, not an equilibrium price.

For the expected-profit objective, future shocks have conditional mean zero given the observations. Consequently the conditional expected trading price change is zero. The financial objective, including every declared charge, is

$$V(I, F) = g(I) - F, \qquad g(I) = (p - b)I - dI^2.$$

This is expected service profit with a hard stress constraint, not a Kelly growth theorem or a return-prediction result. There are no additional training or validation tapes, no leverage, uncertain fills, or changing inventory.

Lemma D.1 (Sharp prior bounds and exact prefunded capacity)**.** *Write*

$$\underline{G} = \kappa \sum_{t<H} a^t, \qquad \overline{G} = \kappa \sum_{t<H} (a + \kappa \overline{Q})^t.$$

Then $\underline{G} \le G \le \overline{G}$, and both endpoints are attained in the declared class. For a gain bound $C \ge G$ and $0 \le F \le \min(B, W)$, define

$$R(v, r) = \frac{2v}{r + \sqrt{r^2 + 4dv}}, \qquad I_u = \frac{(p - b)_+}{2d},$$
$$I(C, F) = \min\{I_u, R(B - F, \overline{u}C + b), R(W - F, P_0 + b)\}.$$

This is the optimal expected-profit position subject to the stated stress reserve and prefunding constraints. The zero-violation no-observation robust optimum is $I_{\mathrm{nd}} = I(\overline{G}, 0)$, not an arbitrarily conservative benchmark. The response-informed oracle is $I_\circ = I(G, 0)$; it knows no future shocks.

Proof. Let $k_j = \sum_i q_i \eta_i (1 - \eta_i)^j$. The normalized price response is the renewal sequence with defective interarrival probabilities $s_1 = a + \kappa k_0$ and $s_j = \kappa k_{j-1}$ for $j \ge 2$. Their total

mass is $a + \kappa Q < 1$, with the remaining mass placed at infinity. The cumulative interarrival distribution at integer $j \geq 1$ is

$$a + \kappa \sum_i q_i [1 - (1 - \eta_i)^j].$$

It lies between a and $a + \kappa \overline{Q}$. Coupling interarrivals by the same independent uniforms orders every renewal time and hence the number of renewals before H. Taking expectations proves the two bounds. The lower endpoint is the no-feedback economy; the upper is one cohort with $Q = \overline{Q}$ and $\eta = 1$. This also proves sharpness of the no-observation stress bound.

At date $t \leq H$, the price change is $x_t = \sum_{j<t} h_j u_{t-1-j}$, so $|x_t| \leq \overline{u} \sum_{j<t} h_j \leq \overline{u} G$. All costs paid by that date are at most the total prefunded costs. The constant adverse sequence $u_t = -\overline{u}$ attains the total price loss $I\overline{u}G$ at settlement. Before the service fee is credited, worst loss from initial wealth is therefore exactly $I\overline{u}G + bI + dI^2 + F$. The declared robust reserve requirement is $I\overline{u}C + bI + dI^2 + F \leq B$; prefunding requires $P_0 I + bI + dI^2 + F \leq W$. These increasing quadratics yield the two roots. Finally g increases up to I_u and decreases after it, proving the optimizer and the two benchmarks. This exact reserve contract is not a peak-to-trough drawdown or a repeated-deployment guarantee. □

Lemma D.2 (No data at the same error budget)**.** *Let $I_f = \min\{I_u, R(W, P_0 + b)\}$ and allow a data-free rule to randomize, with probability of a stress-budget violation at most γ in every economy of the class. Its greatest expected service profit is*

$$V_{\mathrm{nd},\gamma} = (1 - \gamma) g(I_{\mathrm{nd}}) + \gamma g(I_f).$$

It is attained by the displayed two-position lottery. This is a theoretical benchmark, not a recommended trading policy. The corresponding response- informed same-error-budget oracle has value $(1 - \gamma) g(I_\circ) + \gamma g(I_f)$.

Proof. Truncating an action at I_f weakly improves profit and lowers risk, so restrict to

$0 \le I \le I_f$. At the attainable worst economy with gain $\overline{G}$, every $I > I_{\mathrm{nd}}$ is unsafe unless $I_{\mathrm{nd}} = I_f$, in which case all larger actions are already dominated. Thus a feasible lottery puts at most γ mass above I_{nd}. Below that point profit is at most $g(I_{\mathrm{nd}})$, and above it at most $g(I_f)$. This gives the upper bound. The two-point lottery attains it and is unsafe in any economy with probability at most γ. Replace $\overline{G}$ by the known G for the oracle statement. □

Theorem D.3 (Uniform cost-aware finite-horizon decision)**.** *Fix* $0 < \gamma < 1/2$, $z = \Phi^{-1}(1-\gamma)$, $s_N = \sigma\sqrt{H/N}$, *and*

$$Y = \sum_{t<H} \overline{h}_t, \qquad C_N = \min\{\overline{G}, \max\{\underline{G}, Y + z s_N\}\}.$$

Purchase a fixed, affordable N *and take* $I_N = I(C_N, F)$. *The following hold uniformly over the declared positive-cohort class, without dependence on cohort spacing, cohort count, or a memory-rate floor.*

1. *With probability at least* $1-\gamma$ *over measurements, the position satisfies the stress budget for* every *bounded future shock path.*

2. *Put* $r_0 = \overline{u}\underline{G} + b$, $I_{\max} = W/P_0$,

$$L_G = \frac{\overline{u} I_{\max}}{r_0}, \qquad L_F = \frac{1}{\min(r_0, P_0 + b)},$$
$$A = 2z\sigma\sqrt{H}\,(p-b)_+ L_G, \quad D = 1 + (p-b)_+ L_F.$$

   With probability at least $1 - 2\gamma$,

$$0 \le g(I_\circ) - V(I_N, F) \le A N^{-1/2} + D c_H N.$$

   For sufficiently small $c_H > 0$, *choosing* $N = \lceil (A/(2Dc_H))^{2/3} \rceil$ *is affordable and gives a bound* $3\,2^{-2/3} A^{2/3} (Dc_H)^{1/3} + Dc_H$, *when* $A > 0$. *This is an upper bound from this procedure, not a new minimax optimality claim.*

3. *Let* $I_- = I(G + 2zs_N, F)$. *A sufficient condition for strictly positive expected net information value relative to the strongest no-data rule at the* same *error budget is*

$$\underbrace{(1-2\gamma)g(I_-) - F - V_{\mathrm{nd},\gamma}}_{\mathcal{L}_N} > 0.$$

Precisely, $\mathbb{E}_D[V(I_N, F)] - V_{\mathrm{nd},\gamma} \geq \mathcal{L}_N$. *The conclusion can fail at high data costs or a low service premium.*

Proof. $Y - G$ is Gaussian with variance $H\sigma^2/N$. On $\{Y - G \geq -zs_N\}$, which has probability $1-\gamma$, $C_N \geq G$; clipping preserves this because G lies between the prior bounds. Lemma D.1 proves simultaneous safety for every future bounded path on this single event, including the data-selected position. There is no tail beyond H to estimate, explaining the absence of a memory-floor assumption. This does not certify an infinite horizon.

On $E = \{|Y - G| \leq zs_N\}$, of probability $1 - 2\gamma$, $G \leq C_N \leq G + 2zs_N$. The position function is decreasing in each of C and F, so $I_- \leq I_N \leq I_\circ$ on E. To bound the position difference, implicitly differentiate the stress root: its derivatives are $-\overline{u}I/(\overline{u}C + b + 2dI)$ in C and $-1/(\overline{u}C + b + 2dI)$ in F. Clip this root at $I_{\max}$ without changing the final minimum, since the prefunding root is at most $I_{\max}$. The clipped root is L_G-Lipschitz in C and $1/r_0$-Lipschitz in F. The prefunding root is independent of C and $1/(P_0 + b)$-Lipschitz in F. Minimum with I_u preserves these common Lipschitz bounds. Hence

$$0 \leq I_\circ - I_N \leq 2L_G zs_N + L_F F.$$

On $[0, I_u]$, $0 \leq g' \leq (p - b)_+$. Multiplication and addition of the paid fee give the regret bound. Minimizing $AN^{-1/2} + Dc_H N$ over positive real N gives the stated root and value. Rounding up adds at most Dc_H; affordability follows as $c_H N = O(c_H^{1/3}) \to 0$. The case $A = 0$ has no positive-gain information task and does not need this optimization.

On E, $g(I_N) \geq g(I_-)$, while $g(I_N) \geq 0$ always because $0 \leq I_N \leq I_u$. Thus $\mathbb{E}_D g(I_N) \geq (1-$

$2\gamma)g(I_-)$. Subtracting F and the same-error-budget no-data value proves the last assertion. Conditional mean-zero future flows justify the stated expected-profit interpretation; their distribution is not learned from extra tapes. □

Theorem D.4 (A two-cohort identification obstruction under the same observations)**.** *Fix* $Q > 0$, $a + \kappa Q < 1$, *and* $0 < \eta_0 < 1$. *For sufficiently small* $\delta > 0$, *compare the two economies*

	normalized weights	*adaptation rates*
A_δ	$(1/2, 1/2)$	$(\eta_0 - \delta, \eta_0 + \delta)$
B_δ	$(1/5, 4/5)$	$(\eta_0 - 2\delta, \eta_0 + \delta/2)$.

Actual weights are Q *times those displayed. For fixed* $H \geq 4$, $D_H(\delta) = \|h^{\mathrm{B}}_{0:H-1} - h^{\mathrm{A}}_{0:H-1}\|_2 = \Theta(\delta^3)$. *The exact minimax error for distinguishing these two* known *alternatives from* N *reset experiments is*

$$e_N = \Phi\left(-\frac{\sqrt{N} D_H(\delta)}{2\sigma}\right).$$

Therefore error at most $\beta < 1/2$ *requires* $N \geq [2\sigma\Phi^{-1}(1-\beta)/D_H(\delta)]^2 = \Theta(\delta^{-6})$. *In particular any recovery of the two normalized weights to sup-norm error less than* 0.15 *with probability at least* $1 - \beta$ *in both worlds must satisfy this testing requirement. Both worlds contain exactly two distinct rates. Both have identical infinite summed response, but their finite-horizon gains are not asserted identical.*

Proof. The two rate distributions have equal mass, mean η_0, and variance δ^2. Their third centered moments are respectively 0 and $-3\delta^3/2$. Thus their kernels k_j agree at $j = 0, 1$ and $k_2^{\mathrm{B}} - k_2^{\mathrm{A}} = -3Q\delta^3/2$. The response recursion gives equal h_0, h_1, h_2 and

$$h_3^{\mathrm{B}} - h_3^{\mathrm{A}} = -3\kappa^2 Q\delta^3/2,$$

which is a lower bound for D_H in absolute value.

For an upper bound put $f_t(\eta) = \eta(1-\eta)^t$. Its third derivative is $3t(t-1)(1-\eta)^{t-2} -$

$\eta t(t-1)(t-2)(1-\eta)^{t-3}$, with absent low-degree terms zero. The sum over $t \geq 0$ of its absolute value is at most $12/\eta^3$. For $\delta \leq \eta_0/4$, all intervening rates are at least $v = \eta_0/2$. Taylor's integral remainder in ℓ_1 is therefore at most $2|\eta - \eta_0|^3/v^3$. The matching first two moments cancel the quadratic Taylor terms. The sum of the two third absolute centered moments is $(1 + 17/10)\delta^3 = 27\delta^3/10$, giving

$$\|k^{\mathrm{B}} - k^{\mathrm{A}}\|_1 \leq \frac{27Q}{5v^3}\delta^3.$$

If $K(z) = \sum k_t z^t$ and $H(z) = \sum h_t z^t$, the recursion gives $H(z) = \kappa/[1 - az - \kappa z K(z)]$. Consequently $H_{\mathrm{B}} - H_{\mathrm{A}} = zH_{\mathrm{A}}H_{\mathrm{B}}(K_{\mathrm{B}} - K_{\mathrm{A}})$. The infinite response sums of both economies equal $S = \kappa/(1 - a - \kappa Q)$. Young's inequality yields $D_H \leq S^2\|k^{\mathrm{B}} - k^{\mathrm{A}}\|_1 = O(\delta^3)$. This proves the two-sided order with constants independent of δ.

For equal-covariance Gaussian laws, the likelihood-ratio test projects the sample mean onto the difference of the two known means, and thresholds at their midpoint. The projected normal law gives the displayed error in each world. Its equal-prior Bayes optimality and symmetry establish minimax optimality. A weight estimator accurate within 0.15 distinguishes the two disjoint normalized-weight balls, proving the identification lower bound. Finite-horizon gains need not be equal: only their difference is of cubic order. The test concerns latent cohort recovery, not identification of the input-output response vector, and does not cover alternative probe designs. □

Corollary D.5 (A priced separation)**.** *Fix $c_H > 0$ and an affordable N for which $\mathcal{L}_N > 0$ at the limiting merged economy $\delta = 0$. For all sufficiently small positive δ, the financial rule of Theorem D.3 has a strictly positive expected net information value, with its stated safety probability. Its observation requirement has no dependence on δ. At this same fixed information price, sufficiently small δ makes cohort identification at error β unaffordable within the same hard fee budget B: $c_H[2\sigma\Phi^{-1}(1-\beta)/D_H(\delta)]^2 > B$. Any procedure limited to that many affordable observations is covered, including one that stops early, since it can be viewed as a*

function of the full affordable sample. A rational agent may abandon identification and use the no-data rule. There is no assertion that regulation requires identification.

Proof. Apply the two theorems. The first statement uses the strictly positive lower bound on net value, rather than an assumption that more information is automatically beneficial. The second follows from $D_H = \Theta(\delta^3)$. For a neighbourhood of a fixed δ where the strict value inequality holds, continuity preserves it. To claim the inequality for all sufficiently small δ, it suffices that it hold strictly at the limiting merged economy: finite G_H and I_- are continuous there. The numerical example reports both the primary pair and that limiting check. No finite-class grid is used in the uniform decision guarantee. □

What the result does not establish. The direct gain-sum certificate is exactly ordinary Gaussian projection onto the vector of ones. It is not a new estimator and does not beat that benchmark. The latent-cohort lower bound is not a lower bound for learning the observable response vector; comparing it with FIR model-identification results for $\mathcal{H}_\infty$ norms would conflate different tasks (Tu et al., 2018). No optimal estimation rate over all observation designs is claimed. The data cost is not an empirically measured exchange probing cost. An ordinary public-tape bridge, endogenous quotes and depth, incomplete resets, unknown or misspecified noise, and independent replication remain open. The information purchase in the example is predeclared. Its positive value is pointwise at the displayed economy; no prior over unknown economies or globally optimal Bayesian/robust acquisition policy is claimed. The fixed horizon, independent measurement contract, and conditional mean-zero future flows cannot be removed by the simulation. At another horizon the certificate must be recomputed. A rolling sequence needs simultaneous or time-uniform coverage; a one-decision guarantee is insufficient.

Literature priority. The reviewed sources already establish the general principles used here. The explicit dollar-cost/capital formulation and the combined result above are candidate additions to this project, not verified first-in-literature claims. A specialist must check

whether they are an immediate corollary of existing task-specific or constrained-decision results before publication.

References

H. J. van Waarde, J. Eising, H. L. Trentelman and M. K. Camlibel. Data Informativity: A New Perspective on Data-Driven Analysis and Control. *IEEE TAC* 65, 4753–4768.

A. J. Wagenmaker, M. Simchowitz and K. Jamieson. Task-Optimal Exploration in Linear Dynamical Systems. *ICML*, PMLR 139, 10641–10652.

D. L. Donoho. Statistical Estimation and Optimal Recovery. *Annals of Statistics* 22, 238–270.

P. Heinrich and J. Kahn. Strong Identifiability and Optimal Minimax Rates for Finite Mixture Estimation. *Annals of Statistics* 46, 2844–2870.

D. Blackwell. Equivalent Comparisons of Experiments. *Annals of Mathematical Statistics* 24, 265–272.

S. Tu, R. Boczar and B. Recht. Minimax Lower Bounds for $\mathcal{H}_\infty$-Norm Estimation. *arXiv:1809.10855.*

Checked synthetic financial illustration

This illustration uses the frozen primary design and the stronger comparison recorded in the design amendment. It is exploratory and was developed after the earlier project runs; it is not independent preregistration or replication. The two-cohort primary has $a = 0.4$, $\kappa = 1$, $Q = 0.4$, rates $0.49, 0.51$, and equal weights. The known class cap is $\overline{Q} = 0.55$; $H = 60$ one-minute model periods, $N = 32$, $\sigma = 0.02$, and $\gamma = 0.025$. Initial wealth is 1,000, price 100, stress budget 1, and flow bound 0.04. The service fee is 0.08 per unit. Execution costs are $0.002I + 0.0002I^2$ and carry is charged at 5 percent annualized. The assumed all-in

reset fee is 0.002, so the learning bill is 0.064. These are hypothetical model dollars, not an exchange calibration or returns. The same-error-budget no-data lottery is a theoretical strongest comparator, not a proposed desk practice. The zero-violation comparisons are also shown.

Policy	Mean units	Net expected profit	Unsafe reserve (%)
No data, zero violations	1.306	0.1007	0.00
Best no data, same error budget	1.523	0.1171	2.50
Finite-horizon certificate	4.553	0.2844	2.44
Standard Gaussian projection	4.553	0.2844	2.44
Full-response ellipsoid	4.386	0.2718	0.00
Point estimate (uncertified)	4.601	0.2880	50.41
Exact gain report, fee 0.05	4.669	0.3072	0.00
Response oracle, zero violations	4.914	0.3756	0.00
Response oracle, same error budget	5.041	0.3851	2.50

Profit integrates the future zero-mean flow model and every declared cost; the data-dependent policy values are also integrated over reading noise by one-dimensional Gaussian quadrature. Mean units and empirical reserve violations use 20,000 synthetic reading draws. Oracle and no-data lottery values are analytic. A reserve violation means some bounded future path can exceed the budget, not that a sampled market path actually did so. The 97.5 percent safety statement is over measurement draws for a single future commitment; it is not certainty, market-wide stability, or time-uniform coverage. The point rule buys slightly more expected profit by violating the same risk requirement and is not a valid superiority comparison. The proposed estimator and standard Gaussian projection coincide exactly.

The net information advantage over the strongest same-error-budget no-data rule is 0.1673. The written theorem's sufficient lower bound is 0.1465, strictly positive even after all costs. At the limiting merged economy the lower bound is 0.1465; continuity therefore supports the small-separation corollary. These bounds are evaluated in floating point with

a wide positive margin, not claimed interval-arithmetic certificates.

Separation δ	Known-pair resets required	Implied fee bill
0.0800	478	0.955
0.0400	3.96e+04	79.2
0.0200	2.84e+06	5.67e+03
0.0100	1.91e+08	3.83e+05
0.0050	1.26e+10	2.51e+07
0.0025	8.15e+11	1.63e+09

The known-pair test receives all 60 taps, exactly as the financial rule does. Response differences are calculated in rational arithmetic before conversion to floating point. The sample threshold is the Gaussian testing formula, not a measured runtime or an asserted physical trading-capital requirement. For the primary $\delta = 0.01$, distinguishing the two already-known alternatives requires about 191 million resets at 5 percent error. Even this optimistic information fee is far beyond the unit stress budget. Identifying arbitrary cohorts cannot be easier than this embedded test under the same fixed query contract. Other experimental designs are not covered.

At fixed $N = 32$, the expected net benefit crosses zero at an illustrative reset price of about 0.005825. An exact scalar gain report is worth purchasing only below a fee of about 0.188690 in this example. These are computed break-even prices, not observed reporting costs.

Failures retained. The full output includes all 819 cost, horizon and service-fee sensitivity rows, including negative values and unaffordable information budgets. A low service premium or expensive measurements can make abstaining from data acquisition preferable. Independent checks cover 13 algebra, optimizer, shock-corner, cash-flow and Gaussian-testing properties, including the same-error-budget no-data optimum checked against a separate linear program. An initial low-premium LP check was affected by the solver's absolute objective tolerance; rescaling the objective resolved it, and that failed log is retained. This is

not independent mathematical proof review.

Reading/model contract	Gain coverage (%)	Unsafe reserve (%)
Independent Gaussian, valid	97.36	2.63
Correlated taps, wrong variance	71.16	28.84
Correlated taps, correct variance	97.36	2.63
Negative reading bias, invalid	0.00	100.00
Future feedback changes, invalid	0.00	100.00

The valid Gaussian row has finite Monte Carlo variation around its nominal 97.5 percent coverage. Using the correct correlated variance restores the specified measurement contract at the cost of a larger margin. Bias or future parameter change is not repaired by a larger sample under the wrong model. A signed kernel such as $(1, -0.9)$ has sum 0.1 and absolute sum 1.9, so positivity is indispensable to using the plain sum. The same certificate used repeatedly without correction does not retain its one-decision guarantee.

What changed economically. In this model, the desk can choose a larger inventory commitment while paying for information and retaining the same permitted unsafe-action probability as its strongest no-data comparator. This is an example of positive decision value, not evidence that information always pays, that trading strategies must be identified, or that a real venue's liquidity and welfare improve. Realistic observation costs and a public-data bridge remain requirements for an empirical finance claim.

Table 5: Map of results

Question	Answer	Where
What do prices identify?	The merged rate–weight measure only	Thm 3.1
Quantity to separate a known pair	$\Theta(\delta^{-2})$, exact error formula	Thm 4.1
Capital to separate a known pair	$\Theta(\delta^{-4})$, counterparty or own funds	Thm 4.3, Lem. 4.5
Constant at the calibration	$[61.41, 64.48]$ par units; $\approx 10^5$ USDT	Thms 5.2–5.4, Table 10
Does adaptivity help?	At most by a factor ≈ 0.49 in expected loss	Thm 5.3
Unknown mixtures	$(\Delta\eta)^{-(4K-2)}$, tight on rich classes for every K	Thms 6.4, 6.5
General memory	Accuracy $\asymp 1/\log N$, lower and upper bounds	Lem. 6.8, Thm 6.9
Passive data	Identifies M; at least the same polynomial cost for every K	Thms 6.10, 6.11, Prop. C.2
Nonlinear and adaptive rules?	Settled online learners obey δ^{-4}; fast adaptation or thresholds break it	Table 19
What removes the barrier?	Tagged flow, audited equity	Props 4.4, 8.1
Does attribution change a decision?	Yes in a stylised case: a market-wide limit needs only S; whether a targeted halt suffices needs each q_i	Prop. 8.5
What is tagging worth?	Factor $(\Delta\eta)^{-(4K-2)}$ (fixed accuracy) or $(\Delta\eta)^{-(4K-4)}$ (rates); $\approx 333\times$ for BTCUSDT (illustrative)	Thm 8.2
Hypothetical probing loss?	At least \$15M–\$356M a year per market to re-learn two unknown groups weekly at $\Delta\eta = 0.10$; illustrative, not a social cost	Table 21
Do class data restrict class persistence?	Yes: CFTC 8 of 8; Binance 30 of 30 and 29 of 29 new markets (registered); whether it is memory is open	Tables 23, 25
Do the rates survive real price noise?	Yes: preserved under exact empirical heavy-tailed resampling at all sample sizes	Table 17
Can eigenvalues certify safety?	No	Thm 9.1
What do feedback indices measure?	Mostly filtering and timing; small real feedback	Tables 28, 30
Does feedback track activity?	No: fees move volume 2.6–4.9×, not $\hat{Q}$	Table 32

Table 6: Exact known-pair error: formula versus simulation

Pulse V	Horizon T	$\Phi(-V\|d_T\|_2/2\sigma)$	Simulated error, H_0	Simulated error, H_δ
20	3	0.3745	0.3734	0.3762
40	3	0.2611	0.2599	0.2603
40	10	0.0029	0.0030	0.0027
80	20	1.8×10^{-10}	0	0

200,000 paths per hypothesis; $a = 0.5$, $\kappa = 0.8$, $q = 0.2$, $\eta = 0.4$.

Table 7: Account-level initial capital at $\delta = 0.05$ (par units)

Design	C_0	Error	End inv.	Peak margin	Max $\lvert v_t \rvert$	Trades
12 buys, $T = 20$ apart, never unwound	229.13	0.0578	75.00	39.38	6.25	12
6 buy/sell round trips	189.45	0.0308	0.00	3.28	6.25	12
5 buy/sell round trips	157.89	0.0442	0.00	3.28	6.25	10
12 buys, each unwound after 20 steps	378.82	0.0040	0.00	3.28	6.25	24
Generalized-eigenvector probe ($L = 120$)	65.06	0.0500	0.00	2.74	1.24	120
Generalized-eigenvector probe ($L = 160$)	64.86	0.0500	0.00	2.17	1.09	160
Fee-refined probe ($L = 240$)	64.48	0.0500	0.00	—	1.30	240

The pulse designs place a pulse every 40 steps and close it 20 steps later; the fee-refined probe is re-optimized by sequential quadratic programming over $L = 240$. The numerically evaluated lower bound of Theorem 5.2 is 61.41; the explicit constant of Lemma 4.5 is 4.9×10^{-4}.

Table 8: Capital lower bounds and feasible upper bounds across separations (par units)

δ	ω^*	$C^*_{\mathrm{det,quad}}$	$\delta^4 C^*_{\mathrm{det,quad}}$	Certified C_{det}	Feasible UB	Adaptive LB
0.02	0.2546	2397.85	3.84×10^{-4}	2399.58	2428.25	1176.04
0.03	0.2541	473.19	3.83×10^{-4}	473.97	484.50	232.51
0.05	0.2526	61.14	3.82×10^{-4}	61.41	64.48	30.22
0.08	0.2488	9.26	3.79×10^{-4}	9.37	10.38	4.64
0.10	0.2453	3.77	3.77×10^{-4}	3.83	4.44	1.91
0.15	0.2331	0.73	3.67×10^{-4}	0.76	1.01	0.38

Calibration $a = 0.5$, $\kappa = 0.8$, $q = 0.2$, $\eta = 0.4$, $\sigma = 0.01$, $\alpha = 0.05$. From `recompute_theory.py` and `recompute_p1_capital.py`; the numerical-bound column adds the fee lower bound to the fee-free infimum; the feasible column is the best flat probe found over $L \le 240$; the adaptive column adds the fee lower bound to the adaptive bound.

Table 9: Finite-horizon capital at $\delta = 0.05$ (par units)

Horizon L	Quadratic part $c_{\mathrm{req}}/\lambda_{L-1}$	Full bound $C_{\mathrm{RH}}(L)$	Best feasible probe	Gap to $C_{\mathrm{det}} = 61.41$
80	62.55	64.72	65.02	5.9%
120	61.84	63.39	64.59	5.2%
160	61.54	62.46	64.50	5.0%
240	61.32	61.76	64.48	5.0%
320	61.25	—	≤ 64.48	$\le 5.0\%$
480	61.19	—	≤ 64.48	$\le 5.0\%$

Quadratic bound $c_{\mathrm{req}}/\lambda_{L-1}$ from the generalized eigenproblem of Theorem 5.4 on flat probes (equivalently $c_{\mathrm{req}}/(2\lambda'_L)$ with λ' computed against T_{h_0} rather than $\frac{1}{2}T_{h_0}$); it converges to the infinite-horizon value 61.14. The full bound minimizes (26) over 3,600 values of θ (`hardening/rh_bound.py`); at $L = 80$ and 120 a direct numerical minimization of loss plus fee under the separation constraint gives 64.76 and 64.43, above the bound. Feasible probes: top generalized eigenvector refined by sequential quadratic programming on the full ledger; at $L = 320$ and 480 the $L = 240$ probe padded with zeros is feasible at the same cost, so the best feasible value can only be lower.

Table 10: Capital lower bound under alternative observation-noise calibrations (USDT)

Noise calibration	$\delta = 0.02$	$\delta = 0.05$	$\delta = 0.10$
Roll estimator, median (0.40 USDT)	3,874,696	99,192	6,200
One-second price-change s.d., median (1.98 USDT)	93,688,369	2,398,422	149,901
Earlier assumption (4 USDT, upper scenario)	383,812,500	9,825,600	614,100

$C(\sigma, \delta) = 61.41 \times (\sigma/0.01)^2 \times (0.05/\delta)^4$, with σ the unrounded median. Noise estimated on 11 days of checksum-verified Binance BTCUSDT one-second klines.

Table 11: Daily observation-noise estimates, Binance BTCUSDT one-second closes

Date	Roll $\hat{\sigma}$ (USDT)	One-second change s.d. (USDT)	Mean price (USDT)
2023-01-10	0.152	0.679	17,288
2023-01-11	0.092	0.891	17,456
2023-01-12	0.232	2.292	18,343
2023-01-13	0.348	1.762	19,075
2023-01-14	0.773	3.511	20,818
2023-01-15	0.482	1.790	20,778
2023-01-16	0.629	2.357	21,036
2023-01-17	0.457	1.976	21,194
2023-01-18	0.695	2.541	21,122
2023-01-19	0.402	1.316	20,850
2023-01-20	0.316	2.038	21,306
Median	0.402	1.976	

The Roll estimator measures the negative first-order autocovariance of price changes, which reflects bid-ask bounce and recording noise rather than price volatility.

Table 12: Cramér–Rao sample sizes for multi-cohort weight recovery ($T = 120$ taps)

Scenario	K	N^* probes	σ_{crit}	ML success ($f = 1$)	RMSE/CRB
$\eta = 0.30$	1	0.1	6.05×10^{-3}	94%	1.02
$\Delta\eta = 0.30$	2	47.8	2.31×10^{-4}	94%	1.01
$\Delta\eta = 0.20$	2	597	6.55×10^{-5}	96%	0.98
$\Delta\eta = 0.10$ ($\delta = 0.05$ pair)	2	40,374	7.96×10^{-6}	93%	1.03
$\Delta\eta = 0.05$	2	2,620,320	9.88×10^{-7}	95%	1.03
spacing 0.10	3	17,349,791	3.84×10^{-7}	98%	1.22
spacing 0.20	3	27,038	9.73×10^{-6}	98%	0.88

Tolerance ± 0.01 at 95%. σ_{crit} is the single-probe noise at which one probe suffices; operational single-probe noise is 1.6×10^{-3}. ML columns: least squares over all weights and rates from six generic starts (never the true values), 200 replications per row at σ_{crit}; success means every weight within ± 0.01. Fitted log–log slope of N^* in $\Delta\eta$: -6.04.

Table 13: Two-point sample complexity for two-cohort weight recovery

$\Delta\eta$	ΔM_3	$\|\Delta h\|_2$	N^*_{LeCam} (95%)	N^*_{CRB}
0.30	2.67×10^{-4}	1.98×10^{-3}	7.1	47.8
0.20	7.92×10^{-5}	5.60×10^{-4}	88.2	597.0
0.14	2.72×10^{-5}	1.89×10^{-4}	775.3	5,267.9
0.10	9.90×10^{-6}	6.84×10^{-5}	5,927.6	40,374.2
0.08	5.07×10^{-6}	3.49×10^{-5}	22,740.5	155,048.9
0.06	2.14×10^{-6}	1.47×10^{-5}	128,333.0	875,732.3
0.05	1.24×10^{-6}	8.50×10^{-6}	383,859.1	2,620,321.2
0.04	6.34×10^{-7}	4.35×10^{-6}	1,466,361.5	10,012,662.4

Worst-case failure at $\Delta\eta = 0.10$*,* $N = 50$*, 1,000 trials per hypothesis*: bound 44.1%; ML 94.5%, ridge 47.5%, positive-part shrinkage 100%, heavy prior 73.9%.

$\tau = 0.01$, $\sigma_h = 1.6 \times 10^{-3}$, $a = 0.5$, $\kappa = 0.8$, 121 taps. Recomputed columns from `recompute_theory.py`. Estimators re-implemented in `recompute_p1_estimators.py`: ML is least squares from four starts; ridge and heavy prior are posterior modes with prior standard deviations $(0.05, 0.10)$ and $(0.005, 0.02)$ for (weights, rates) around the symmetric centre of the two hypotheses; positive-part shrinkage pulls the ML weights toward their mean. Failure means some weight is off by at least τ.

Table 14: Gauss and Gauss–Radau rules on $[-1, 1]$ used in Theorem 6.4

K	Gauss nodes / Gauss–Radau nodes	Max gap, orders $< 2K - 1$	Gap at order $2K - 1$
2	± 0.5774 / -1, 0.3333	6×10^{-17}	0.2222
3	0, ± 0.7746 / -1, -0.2899, 0.6899	2×10^{-16}	0.0533
4	± 0.3400, ± 0.8611 / -1, -0.5753, 0.1811, 0.8228	4×10^{-16}	0.0131

All Gauss–Radau weights are positive.

Table 15: Numerical check of Theorem 6.4

	$K = 2$	$K = 3$	$K = 4$
Theoretical exponent of N^*	-6	-10	-14
Fitted log–log slope (small $\Delta\eta$)	-5.96	-9.96	-14.01

Le Cam bound $N^* \propto \sigma^2 / \|h^A - h^B\|_2^2$ with $\sigma = 1$, $a = 0.5$, $\kappa = 1$, total merged weight 0.10, rates centered at 0.70, responses over 400 lags.

Table 16: Sample requirement of Prony's method against the lower-bound order

	$K = 2$	$K = 3$	$K = 4$	$K = 5$
Lower-bound order $-(4K - 2)$	-6	-10	-14	-18
Fitted slope of N^* in $\Delta\eta$ (Prony)	-5.94	-10.00	-13.93	-18.34

Rates $0.7 + \Delta\eta\, s$ with s the K-point Gauss nodes, total weight 0.1, $\Delta\eta \in [0.03, 0.2]$; N^* is the number of experiments for a weight RMSE of 10% of the smallest weight, from 400 simulated estimates per point. At $K = 5$ the fit is numerically sensitive: rerunning the same seeded script with a different numerical library gives slopes -5.94, -10.00, -13.92 and -17.89.

Table 17: Sample requirement with real Bitcoin price noise: fitted slope of $\log N^*$ in $\log \Delta\eta$

	Theory	Real noise	Gaussian control	Pre-specified criterion	Result
Price only, $K = 2$ (H1)	-6	-6.01	-6.22	slope in $[-7.5, -4.5]$	pass
Tagged, $K = 2$ (H2)	bounded	256–724	256–1,024	max/min ≤ 8	pass (2.8)
Price only, $K = 3$ (H3)	-10	-9.81	-9.96	slope in $[-12.5, -7.5]$	holds

Synthetic groups at rates $r_c + \Delta\eta\, s$ (s the K-point Gauss nodes; $r_c = 0.6$ and $\Delta\eta$ from 0.5 to 0.125 for $K = 2$, $r_c = 0.5$ and $\Delta\eta$ from 0.6 to 0.25 for $K = 3$), $a = 0.5$, $\kappa = 0.8$, total weight 0.08, $\sigma = 10^{-3}$, $\sigma_f = 0.02$. The test passes if H1 and H2 hold; H3 is reported either way. The tagged row gives the range of N^* over the spacing grid. Requirements at large N use exact empirical multinomial resampling over historical return windows. Requirements at each spacing are in Supplement Table 35. Specification SHA-256 `596a700da2e4...`; opened once. Kind: semi-synthetic (synthetic groups, real noise).

Table 18: General memory: moment-matched cohort pairs and the moment estimator

A. Best approximation of $\lvert x\rvert$						
L	2	4	6	10	20	30
$L\,E_L(\lvert x\rvert)$	0.2500	0.2705	0.2756	0.2785	0.2797	0.2800
B. Pairs of Lemma 6.8						
Atoms per measure n	2	4	6	8	10	12
$W_1(\mu_A, \mu_B)$, $\rho = 0.9$	0.2250	0.1037	0.0669	0.0494	0.0391	0.0324
$n\,W_1$, $\rho = 0.9$	0.450	0.415	0.402	0.395	0.391	0.388
$\lVert k^A - k^B\rVert_2$, $\rho = 0.9$	4.41×10^{-2}	1.03×10^{-3}	4.08×10^{-5}	1.96×10^{-6}	1.05×10^{-7}	5.95×10^{-9}
$\lVert k^A - k^B\rVert_2$, $\rho = 0.99$	7.41×10^{-2}	5.84×10^{-3}	1.07×10^{-3}	2.76×10^{-4}	8.33×10^{-5}	2.76×10^{-5}
C. Linear-programming estimator, continuous memory distribution, $\rho = 0.9$						
$\log_{10} N$	4	8	12	16	20	24
Moments used $(L + 1)$	2	3	5	6	7	8
Error $d(\hat{\mu}, \mu)$	0.00878	0.00759	0.00432	0.00445	0.00382	0.00399
Error $\times \log(N/\sigma^2)$	0.202	0.245	0.179	0.226	0.229	0.276

A: Remez algorithm; the limit is Bernstein's constant 0.2802. B: moments of orders $0, \ldots, 2n - 2$ agree to 10^{-61} (60-digit arithmetic); every W_1 exceeds $\rho E_{2n-2}(\lvert x\rvert)$. C: true merged measure Beta(2,3) on $[0, 0.9]$ with mass 0.1; kernel values observed with noise $\sigma N^{-1/2}$, $\sigma = 10^{-3}$; L from the rule in Theorem 6.9(ii); measure fitted on a 401-point grid; 20 replications. Kind: recomputed (`general_minimax.py`).

Table 19: Merged-measure equivalence and the δ^{-4} law under stylized nonlinear rules and closed-loop trained online learners

Rule	Test A raw ($O(A^2)$)	Test A norm. ($O(A)$)	Test B slope ($\|\Delta h\|_2$)	Test B slope (capital)	Capital at $\delta = 0.05$	A_{sat} / R^2_{exp}
Panel A. Stylized nonlinear and adaptive rules (fixed operating point)						
Linear benchmark	—	—	2.007	−4.001	61.42	—
1. Momentum (tanh limit)	2.015	1.015	2.007	−4.001	61.47	0.834
2. Threshold rule, averaged	2.043	1.043	2.007	−4.001	61.42	0.049
3. Adaptive gain (RLS)	2.064	1.064	2.007	−4.001	61.46	0.751
4. Online neural readout	2.011	1.011	2.007	−4.001	61.48	0.751
5. Tabular value, smoothed	2.039	1.039	2.007	−4.001	61.44	1.271
Panel B. Learners trained in the closed loop, saved and reloaded (frozen during the probe; shared training stream)						
6. Recursive least squares	3.000	2.000	2.008	−4.014	54.58	$R^2 = 0.993$
7. Two-layer network (SGD)	1.878	0.878	2.007	−4.012	65.19	$R^2 = 0.994$
8. Policy gradient on profit	3.000	2.000	2.004	−4.006	58.05	$R^2 = 0.995$
9. Tabular Q-learner (TD(0))	2.108	1.108	1.994	−3.973	186,821	$R^2 = 1.000$
Panel C. Where the trained-learner results change (range over the four learners)						
Independent training streams	rel. gap $\leq 10^{-5}$		0.13 to 0.92	−1.75 to −0.38	—	(c) fails
Learning during the probe	rel. gap $\leq 1.1 \times 10^{-3}$		1.99 to 2.01	−4.02 to −3.98	—	(c) holds
Exact dead zone ($\tau = 0.05$): $\|\Delta h\|_2 = 0$ at $A = 0.005$; the probe moves the state past τ only for $A > 0.139$.						
Single-bin value lookup: $\|\Delta h\|_2 = 0$ at $A = 0.005$; the probe leaves the central bin only for $A > 0.149$.						

Panel A: deterministic simulations around the operating point; Test A reports the raw response-gap slope in $\log A$ ($O(A^2)$) and the probe-normalised gap slope $\|X^{(1)}(A) - X^{(2)}(A)\|_2 / A$ ($O(A)$, plotted in Figure 8). Panel B: learners trained for 5,000 periods, parameters saved and reloaded without rescaling and frozen during the probe; the two compared populations are trained on the same innovation stream; relative equivalence gaps at $A = 0.005$ are at most 10^{-5}. A raw Test A slope above 2 means the gap is of higher order than A^2. Panel C: the same learners trained on independent innovation streams (first row) or updated at learning rate 0.35 during the probe (second row). Capital in par units, lower bound of Theorem 5.2 (61.42 for the linear model); the Q-learner settles on a small response gain (0.022), so separating it costs far more capital. Kind: recomputed (Panel A: `recompute_p1_learners.py` → `R_P1_learners.json`; Panels B–C: `validate_trained_learners_hardened.py` → `R_P1_trained_learners_hardened.json`).

Table 20: Disclosure contracts at $\delta = 0.05$ (pulse-design capital, par units)

Contract	Channel	Pulses	Benchmark	Saving vs. $T = 20$
I_0: prices, $T = 3$	public prices	271	4,305.25	−2,123.5%
I_0^*: prices, $T = 20$	full price path	12	193.62	0.0%
I_1: prices and signed flow	public flow	12	193.62	0.0%
I_2: noisy equity ($\sigma_E = 0.5$)	statement	12	193.62	0.0%
I_3: exact equity, 5-unit lot	audited equity	1	12.70	93.4%
I_3: exact equity, 0.01 rounding	rounded report	1	15.52	92.0%
I_4: timestamps only	no likelihood	12	193.62	0.0%
I_5: audit registry	internal states	0	0.00	100.0%

Closed-form values. Noisy equity adds information linearly in the number of resets, $D_{\text{KL}} = N\Delta_E^2/(2\sigma_E^2)$, and saves only 1.5% against the $T = 3$ baseline. The equity difference is $\Delta E = \kappa^2 q \delta^2 V^2$; with reports rounded to 0.01, a guaranteed change needs $V \geq 5.59$.

Table 21: Price-only round-trip probing loss per market and annualised weekly-renewal cost (hypothetical outside prober)

	Loss / probe (USDT)		Annual (52×, USDT)	
Economy ($\Delta\eta$)	Roll	1-sec	Roll	1-sec
$K = 2$, known (0.20)	6,200	149,901	322,400	7,794,852
$K = 2$, known (0.10)	99,192	2,398,422	5,157,984	124,717,944
$K = 2$, known (0.04)	3,874,696	93,688,369	201,484,192	4,871,795,188
$K = 2$, Radau (0.20)	4,292	103,788	223,184	5,396,976
$K = 2$, Radau (0.10)	283,488	6,854,612	14,741,376	356,439,824
$K = 3$, Radau (0.20)	804,403	19,450,107	41,828,956	1,011,405,564
$K = 3$, Radau (0.10)	1,137,325,880	27,500,017,407	59,140,945,760	1,430,000,905,164
$K = 4$, Radau (0.20)	156,632,086	3,787,291,895	8,144,868,472	196,939,178,540

Separation variable $\Delta\eta$ throughout ($\Delta\eta = 2\delta$ for the symmetric known pair). Gauss–Radau rows use the exact K-cohort quadrature pair of Theorem 6.4 and Table 14 (moment gaps $\Delta_2 = 0.222222$, $\Delta_3 = 0.053333$, $\Delta_4 = 0.013061$ at $W = 0.08$, $r_c = 0.60$). Loss per probe is the spectral round-trip trading loss bound of Theorem 5.2 at the median BTCUSDT Roll noise (0.40 USDT) and median one-second price-change standard deviation (1.98 USDT); annual loss multiplies by 52 weekly renewals. Probing cost is a lower bound on the cost of learning from prices by trading, not on passive learning; it is a private transfer to other traders, not a social or resource cost, and it is not comparable with audit-trail budgets. Kind: recomputed (`recompute_p1_cat_costs.py` → `R_P1_cat_disclosure_dollars.json`).

Table 22: Targeted versus market-wide intervention: two economies that prices barely separate

K	$\Delta\eta$	$\lvert S^A - S^B \rvert$	$\kappa q^G_{\max}$	$\kappa q^R_{\max}$	Window for $\bar{s}$	N for a correct choice
2	0.20	3.9×10^{-3}	0.112	0.144	[0.527, 0.562]	5
2	0.10	5.0×10^{-4}	0.093	0.131	[0.532, 0.570]	306
2	0.05	6.5×10^{-5}	0.086	0.125	[0.536, 0.575]	18,095
3	0.20	2.8×10^{-4}	0.073	0.092	[0.583, 0.603]	1,092
3	0.10	7.8×10^{-6}	0.071	0.076	[0.587, 0.592]	1.4×10^6

Calibration $a = 0.5$, $\kappa = 0.8$, $W = 0.08$, $r_c = 0.60$, observation noise $\sigma = 0.01$ per tap; G and R are the Gauss and Gauss–Radau economies. Both economies have $S \approx 0.66$–0.68. "Window" is the range of thresholds in which halting the largest group restores $S \le \bar{s}$ in one economy and not in the other. The last column rounds the continuous Gaussian-shift lower bound for a choice correct with probability 3/4 in both; an actual integer sample size must round the unrounded requirement upward. Kind: recomputed (`theory/targeted_intervention.py` → `R_P1_targeted_intervention.json`).

Table 23: Fitted persistence of trader-class positions, CFTC Commitments of Traders, eight regulated futures markets

Market	Weeks	Fast $\hat{r}$ [95% CI]	Slow $\hat{r}$	$\hat{r}_{\text{fast}} - \hat{r}_{\text{slow}}$ [95% CI]	$\hat{w}_{\text{fast}}$	Ratio
Bitcoin (CME)	434	0.951 [0.928, 0.969]	0.986	−0.035 [−0.062, −0.010]	−	23.7
Ether (CME)	282	0.890 [0.826, 0.932]	0.960	−0.070 [−0.143, −0.011]	−	9.1
E-mini S&P 500	1,039	0.914 [0.884, 0.937]	0.954	−0.040 [−0.071, −0.012]	−	18.3
10-year T-note	1,039	0.992 [0.985, 0.998]	0.993	−0.000 [−0.006, +0.007]	(0)	71.4
Euro	1,035	0.961 [0.947, 0.972]	0.996	−0.034 [−0.051, −0.022]	+	37.3
Yen	1,035	0.955 [0.939, 0.968]	0.984	−0.029 [−0.047, −0.013]	+	34.0
WTI crude oil	1,039	0.961 [0.946, 0.974]	0.999	−0.037 [−0.053, −0.023]	+	35.2
Natural gas	1,039	0.981 [0.969, 0.990]	0.980	+0.001 [−0.014, +0.017]	+	44.0

Pre-specified hypotheses: H1 (primary) 8 of 8, pass; H2 4 of 8, fail; H3 6 of 8, pass (two indistinguishable).

"Fast" (designated speculative) class: leveraged funds (financial futures) or managed money (commodities); "slow" (designated comparison) class: asset managers or producers/merchants. The labels are fixed by class role before estimation, not by the estimates; in natural gas the designated fast class has the (insignificantly) higher persistence. $\hat{w}_{\text{fast}}$: sign of the response to returns, + or − when the 95% interval excludes zero, (0) otherwise. Ratio: mean width of the two price-only rate intervals over the width of the fast-class interval; the price-only intervals span 0.95–0.96 of the admissible range $[0, 0.99]$ in every market. Weekly returns: Binance spot (Bitcoin, Ether); S&P 500 index, euro and yen spot, minus the change in the 10-year yield, and front-month crude-oil and natural-gas futures from Yahoo Finance. For the six non-futures series the spot–futures basis is ignored; a post hoc rerun with CME futures prices gives the same counts and the same signs of $\hat{w}_{\text{fast}}$ (see text). Specification SHA-256 `6f1d6f35ae88...`; test opened once. Kind: real data (Part 11).

Table 24: Price data needed to match the precision of the CFTC class reports under standard root-T scaling (post hoc)

Market	Weeks	Ratio	Factor	Years	Market	Weeks	Ratio	Factor	Years
CME Bitcoin	434	23.7	562	4,688	Gold	1041	32.1	1,032	20,652
CME Ether	282	9.1	82	446	Copper	1041	36.3	1,315	26,318
E-mini S&P 500	1039	18.3	336	6,718	Coffee	1041	68.8	4,735	94,796
10-year Treasury	1039	71.4	5,105	102,000	Live cattle	1041	41.0	1,684	33,712
Euro	1035	37.3	1,391	27,684	Nasdaq-100	1041	23.3	541	10,837
Yen	1035	34.0	1,155	22,997	Sugar	1041	52.1	2,719	54,427
Crude oil	1039	35.2	1,239	24,753	Silver	1041	23.1	532	10,658
Natural gas	1039	44.0	1,939	38,750	Corn	1041	48.1	2,313	46,302
Cotton	1041	50.9	2,587	51,793	Soybeans	1041	42.0	1,762	35,266
British pound	1039	29.9	896	17,910	Wheat	1041	44.1	1,943	38,893

Ratio: price-only over tagged width of the 95% interval for the speculative class's persistence, from the registered tests (first eight: original markets; last twelve: the replication). Factor = Ratio^2, the growth of the price sample needed to reach the tagged precision if both intervals shrink as $T^{-1/2}$; Years = Weeks × Factor/52. The price-only intervals are cut off by the admissible range $[0, 0.99]$, so every entry is a lower bound. Median factor 1,353, median 27,001 years. Kind: post hoc calculation from registered results (`p11/data_equivalence.py`).

Table 25: Fitted persistence rates with and without class-tagged data

Symbol	Days	Top position $\hat{r}$ [95% CI]	All accounts $\hat{r}$ [95% CI]	Width ratio	Classes differ
BTCUSDT	48	0.9889 [0.9842, 0.9935]	0.9986 [0.9955, 1.0005]	121.7	yes
ETHUSDT	48	0.9793 [0.9616, 0.9907]	0.9972 [0.9934, 0.9997]	55.6	yes
CELRUSDT	48	0.9884 [0.9830, 0.9922]	0.9882 [0.9804, 0.9914]	81.6	no
CVXUSDT	33	0.9822 [0.9732, 0.9908]	0.9829 [0.9694, 0.9900]	45.9	no

Pre-specified test on fresh days, 30 USD-M perpetuals, 2025–2026	
(i) width ratio > 10	30 of 30 (range 27.6–124.0)
(ii) price-only rate interval wider than 0.25	30 of 30 (range 0.65–0.99)
(iii) merged response $\sum_m k_m$ identified	5 of 30
Classes differ	11 of 30

Top panel: training days 2023–2024; width ratio is the mean price-only rate-interval width over the mean tagged width. Estimates are from day-demeaned regressions, which bias persistence down; corrected by a split-panel jackknife, the median tagged rate is 1.0008 (Section 8.5). The price-only intervals for both rates cover most of $[0, 0.99]$. Bottom panel: specification pre-specified before the fresh days were downloaded (SHA-256 `65acda361861...`); a symbol passes if all three of the first rows hold, and the test required 80% of symbols to pass.

Table 26: Pre-specified out-of-sample cohort-state trajectory test on untouched 2nd-and-9th-of-month days (`P1_disclosure_state_oos`), 2025–2026

30-min post-shock target (19 thin)	R^2_{pf}	R^2_{tag}	ΔR^2 (pts) [95% CI]	CW t (p)	Res.
Prim. 1: Divergence d_{t+6} (9,875 shocks)	−4.766	45.962	+50.728 [+45.291, +55.918]	21.27 ($< 10^{-16}$)	pass
Prim. 2: Top state $n_{\text{top},t+6}$ (9,875 shocks)	4.441	54.845	+50.404 [+47.226, +53.724]	21.85 ($< 10^{-16}$)	pass
Sec.: Divergence d_{t+6}, all (210,144 bars)	−0.586	37.053	+37.639 [+36.232, +39.107]	72.75 ($< 10^{-16}$)	—
Sec.: Top state $n_{\text{top},t+6}$, all (210,144 bars)	0.186	46.779	+46.593 [+45.555, +47.643]	80.44 ($< 10^{-16}$)	—

Out-of-sample R^2 in percent. Trained on 24 dates in 2024 (12,046 shocks, 120,384 bars); evaluated once on 42 untouched test dates (2nd and 9th of each month, Jan 2025–Sep 2026); 2,000-draw date-block bootstrap 95% CIs; Holm $p < 10^{-16}$ across both primaries. Specification SHA-256 `5136cc6b846d...` (`R_P1_disclosure_state_oos.json`).

Table 27: Samples

Sample	Coverage	Size	Source
Binance spot panel, 8 pairs	96 dates, 2021–2024	66.3M one-second bars	Binance archive
Binance spot fee-event windows, 8 pairs	±20 days around 2022-07-08 and 2023-03-22	640 pair-days	Binance archive
Binance USD-M perpetuals, 30 symbols	90 dates, 2023–2026	2,565 symbol-days	Binance archive
Binance BTCUSDT one-second panel	72 dates, 2021–2023	6.2M one-second bars	Binance archive
Binance BTCUSDT noise-calibration days	January 10–20, 2023	11 days	Binance archive

Binance files are verified against the published SHA-256 checksums (every file used). Sampled dates are the 10th and 24th of each month.

Table 28: Multi-year Binance spot panel: reduced-form indices under three timing constructions

Construction	$\hat{Q}$	Date interval	$\hat{S}$	Date interval	$\hat{Q} > 0$
Forward-paired flow (look-ahead)	0.0081424	[0.0072965, 0.0089882]	0.92074	[0.91665, 0.92482]	87.4%
Observed-flow EMA	0.0003733	[0.0002351, 0.0005116]	0.92072	[0.91640, 0.92504]	62.4%
Observed, 30 flow lags	0.0002580	[0.0001415, 0.0003745]	0.91738	[0.91167, 0.92308]	59.1%
Observed raw return	−0.0001555	[−0.0003732, 0.0000622]	−0.04051	[−0.05078, −0.03025]	49.2%
Placebo random walk	0.0000016	[−0.0000020, 0.0000052]	0.93510	[0.93485, 0.93534]	51.0%

768 study units (96 dates × 8 pairs), 6,634,752 ten-second transitions; equal-weight asset means within date, then date means with Student intervals. Mechanical persistence $29/31 = 0.93548$.

Table 29: Asset means under forward-paired and observed flow

Pair	Forward-paired $\hat{Q}$	Observed $\hat{Q}$	$\hat{Q}_{30}$	Observed $\hat{S}$
BTC/USDT	0.0141870	0.0006073	0.0004891	0.93464
ETH/USDT	0.0130885	0.0005675	0.0002549	0.93511
SOL/USDT	0.0087763	0.0002129	0.0001100	0.91966
ADA/USDT	0.0067465	0.0003928	0.0003546	0.91452
DOGE/USDT	0.0058956	0.0002464	0.0001432	0.92496
BNB/USDT	0.0057420	0.0004479	0.0002762	0.89533
XRP/USDT	0.0057034	0.0002249	0.0001508	0.92362
LTC/USDT	0.0049997	0.0002870	0.0002853	0.91792

96 dates per pair. $\hat{Q}_{30}$ controls 30 ten-second bars of lagged flow. Placebo $|\hat{Q}| \leq 1.1 \times 10^{-5}$ for every pair.

Table 30: Feedback beyond filtering mechanics (Binance BTCUSDT, 72 days, 2021–2023)

	Real prices	Placebo (independent random walk)
Mean loop gain $\hat{Q}$	0.0045	0.0000009
95% CI (day bootstrap)	[0.0034, 0.0056]	[−0.0000006, 0.0000023]
Mean price persistence $\hat{J}_{11}$	0.9364	0.9357
Mechanical value $1 - 2/31$		0.9355

Placebo 95% band across days: $[-0.0000095, 0.0000156]$. This construction pairs the state with next-second flow; it is not comparable in level with the observed-flow ten-second index of Table 28.

Table 31: Feedback estimators under order splitting (simulation, 500 replications)

Estimator	Coverage at $\gamma = 0.2$	Rejection at $\gamma = 0$	Bias	RMSE
OLS	0.000	1.000	0.892	0.893
Cochrane–Orcutt, true ρ	0.944	0.054	0.001	0.039
Cochrane–Orcutt, estimated ρ	0.916	0.080	0.002	0.043
2SLS	0.938	0.044	−0.002	0.116
LIML	0.938	0.044	−0.003	0.117

Newey–West standard errors (10 lags), $T = 3{,}000$. Median first-stage $F = 224$; a placebo instrument (future z) has $F = 0.62$. Estimators outside 93–97% coverage give invalid inference.

Table 32: Zero fees on BTCUSDT: activity and measured feedback

	Zero fees introduced (2022-07-08)	Zero fees removed (2023-03-22)	Pooled
Log quote volume, $\hat{\beta}$	0.945 (0.074)	1.599 (0.080)	1.272 (0.084)
Observed-flow $\hat{Q}$, $\hat{\beta}$	0.00058 (0.00032)	−0.00091 (0.00060)	−0.00016 (0.00036)
p-value	0.078	0.140	0.650
Pair-days	320	320	640

$\hat{\beta}$ is the effect of zero fees on BTCUSDT relative to seven control pairs; date-clustered standard errors in parentheses. For the second event ZeroFee equals one before March 22, 2023.

Table 34: Probing protocols at observation horizon $T = 4$ per cycle

Protocol	δ	Executed vol.	Max trade	Error (SE)	Capital
A: single pulse (unconstrained)	0.03	131.5	131.5	0.047 (0.003)	6,985.91
	0.05	47.3	47.3	0.054 (0.004)	921.73
	0.08	18.5	18.5	0.053 (0.004)	146.73
	0.12	8.2	8.2	0.044 (0.003)	31.45
B: repeated capped pulses	0.03	2,768.8	6.2	0.050 (0.003)	7,035.75
	0.05	362.5	6.2	0.054 (0.004)	923.88
	0.08	56.2	6.2	0.047 (0.003)	146.00
	0.12	12.5	6.2	0.038 (0.003)	34.88
C: sequential test, capped	0.03	1,410.2	6.2	0.045 (0.003)	3,585.14
	0.05	195.6	6.2	0.038 (0.003)	500.00
	0.08	35.9	6.2	0.032 (0.003)	94.30
	0.12	10.5	6.2	0.011 (0.002)	29.85
D: generalized eigenvector	0.03	2,359.7	3.7	0.050 (0.003)	2,632.77
	0.05	314.8	3.7	0.048 (0.003)	353.67
	0.08	49.2	3.7	0.046 (0.003)	57.53
	0.12	9.9	3.8	0.045 (0.003)	13.56

2,000 replications per hypothesis and cell. Log–log capital slopes: −3.90, −3.84, −3.47, −3.81. Protocols A and B reproduce the earlier version exactly. C and D are re-implemented because the earlier code was not retained: C is a Wald test with capped pulses; D repeats the top generalized eigenvector of a flat $(T + 1)$-step probe, scaled to a maximum trade of 3.75, and is cheaper than the earlier D. Kind: [RECOMPUTED] (`recompute_p1_protocols.py`).

Table 35: Real-noise injection: number of experiments N^* at each spacing

	$\Delta\eta$	Price, real	Price, Gaussian	Tagged, real	Tagged, Gaussian
$K = 2$	0.500	64	64	724	1,024
	0.400	362	362	362	512
	0.320	1,448	1,024	362	362
	0.250	5,793	5,793	256	256
	0.200	16,384	23,170	256	362
	0.160	65,536	92,682	256	256
	0.125	370,728	370,728	256	256
$K = 3$	0.60	23,170	16,384		
	0.50	65,536	65,536		
	0.40	741,455	741,455		
	0.32	8,388,608	5,931,642		
	0.25	94,906,266	94,906,266		

Real noise: 30,219 one-minute log-price changes of Binance BTCUSDT on 21 days (the 6th of each month, January 2025 to September 2026), kurtosis 19.4. Tagged requirements were computed for $K = 2$ only. Kind: semi-synthetic.

Table 36: Numerical stress tests (numerical checks, not proofs)

Result	What was checked	Outcome
Upper bound, Theorem 6.5	Prony's requirement N^* for random separated configurations (15 for $K = 2$, 12 for $K = 3$); slope in $\Delta\eta$ and whether $N^*(\Delta\eta)^{4K-2}/\sigma^2$ grows as the spacing shrinks	Slopes -6.72 to -5.79 ($K = 2$, theory -6) and -11.01 to -9.29 ($K = 3$, theory -10); the normalized constant changes by a median factor of 1.05 and 0.94 from the widest to the narrowest spacing (at most 2.1); no configuration failed
Passive data, Theorem 6.11	Exact Whittle divergence rate between two moment-matched economies, 10 random settings for each K	Slopes 6.19–6.34, 10.27–10.40 and 14.33–14.54 for $K = 2, 3, 4$ (theory 6, 10, 14): the divergence shrinks at least as fast as the theorem requires
General memory, Lemma 6.8	Alternation pairs on $[0, \rho]$ for 12 random $\rho \in [0.3, 0.99]$ and $n = 2, \ldots, 6$: equal low moments, $W_1 \geq \rho E_{2n-2}$, kernel distance bound	60 of 60 cases pass
Value of disclosure, Theorem 8.2	Tagged and price-only requirements over the spacing grid, $K = 2, 3, 4$	Tagged requirement does not grow as the spacing shrinks (it varies by factors 2.8, 11.3 and 5.7 over the grid); price-only slopes -6.01 and -9.47 for $K = 2, 3$; for $K = 4$ our Prony implementation fails numerically at every spacing, so no price-only comparison

Gaussian observation noise, $\sigma = 10^{-3}$ (prices) and $\sigma_f = 0.02$ (tagged flow), $a = 0.5$, $\kappa = 0.8$, total weight 0.08, 200 replications per point, accuracy: every weight within 25% of the smallest true weight in at least 75% of replications. Moment-matched economies for the passive check: binomial weights on equally spaced nodes (the pair used in the Lean proof of Theorem 6.4). Seed 20261031. Kind: numerical check.

Table 37: Probing with and without resets: required reset experiments N^* and required periods T^* of one running market

$\Delta\eta$	Reset experiments N^*	Running market T^* $\sigma_w = \sigma_\nu$	Running market T^* $\sigma_w = 3\sigma_\nu$
0.5	91	128	362
0.4	362	512	1,448
0.32	1,024	2,048	5,793
0.25	5,793	8,192	32,768
0.2	23,170	32,768	131,072
0.16	65,536	185,364	370,728
Fitted slope in $\Delta\eta$	-5.87	-6.26	-6.22

Numerical check, not pre-specified. $K = 2$ Gauss-node economy of Table 17; i.i.d. ± 1 probe; least squares on 200 probe lags; Prony estimator and success rule as in Table 16.

Returned advanced experiment: evidence and limits

Research question and financial institution. Can information about an aggregate market response improve an exposure decision after information, execution and carry costs, while the hidden trading cohorts remain difficult to recover? The new experiment concerns a price-taking, unlevered inventory-warehousing desk. It buys a fixed inventory, holds it for $H = 10, 60, 240$ one-minute *model* periods and closes. A prequoted service premium arrives after settlement. It cannot change market feedback, quotes or other traders. This is a separate application from a dealer or venue that changes admission and thereby changes the market itself. The objective is expected service profit under a stress reserve and cash constraint, not logarithmic portfolio growth or predicted crypto returns.

Completed execution record. The author returned the Deep Research archive with base seed 20261101. All 611 manifest entries, all 288 checkpoint-to-statistics links and all 15 released source files were verified. The 288 generated economies use 24 seed clusters, 4,096 measurement repetitions per economy and N, $N = 8, 32, 128$, three information-price settings, seven measurement rules and a same-error oracle. The 20,736 policy/cost rows share readings and are not independent market experiments. All 27 implemented tests passed. Numeric policy fields reproduce the local same-seed run to about 1.25×10^{-14}. This is cross-environment reproducibility of the same implementation, not independent replication. The returned runtime used Python 3.13.18 and the pinned numerical packages; the research runner took 52.71 seconds, excluding notebook setup.

What Lean checked. Lean 4.19.0 checked seven concrete rational inequality certificates, with no reported axioms. They include a conservative positive-value numerical witness greater than 0.14 hypothetical model dollars for the original primary contract. They do not formalize the Gaussian analytic tail step, market dynamics, general separation theorem, translation from observations to economics or novelty. The earlier positive illustration is retained as a pointwise example. It must not be read as a positive average-value theorem over all economies or prices.

Economic value depends on the information bill. For the standard Gaussian rule within the independent-Gaussian economy design:

Information price	Mean advantage	95% interval	Positive / total
Free	+0.13047	[0.08458, 0.18420]	99 / 144
Baseline	−0.02953	[−0.07022, 0.01310]	59 / 144
Five times baseline	−0.17212	[−0.21856, −0.12715]	26 / 144

The comparator is the analytically optimal no-data rule with the same permitted unsafe-action probability, $\gamma = 0.025$. Values are hypothetical dollars per inventory-service decision. Baseline cost is $N(0.0005 + 0.000025H)$; its dependence on horizon and N is explicit, not

an exchange calibration. The intervals use 5,000 percentile bootstrap resamples of 24 seed-cluster averages. This price-specific analysis was added after the return and is exploratory, not preregistered confirmation or uncertainty about actual market profit. The baseline-price interval crosses zero: this design does not establish a positive average benefit after costs. The high-price interval is negative under the stated simulator sampling scheme.

Protection and profitability are different outcomes. For purchased iid-Gaussian rules in iid-Gaussian economies, the mean simulated unsafe-action frequency was 0.834%, below the 2.5% allowance. Across 6,535 declared-valid purchased-rule checks, Holm correction produced zero rejections. Individual finite-sample frequencies can exceed the allowance; non-rejection is not a proof. An unsafe action means its reserve can breach the budget under a permitted worst-case shock path. It is not an observed liquidation, default or historical crash. The Gaussian and declared-variance rules coincide exactly under iid readings. Their advantage over a full-response ellipsoid is a task-specific comparison of established methods, not a new estimator.

Failures retained and interpreted.

Purchased Gaussian rule: observation/future setting	Mean unsafe action
Independent Gaussian readings, stationary future	0.834%
Correlated readings, iid variance used	14.695%
Downward-biased readings	31.169%
Future feedback changes	61.033%

These are unweighted descriptive averages of dependent configuration rows, including the zero-price setting. They are not pooled market failure probabilities. Under correlated Gaussian readings, a declared-variance rule reduces the mean to 0.409%, and a projected Student-t rule to 1.153%. Protection can improve while the information purchase remains uneconomic. The Student-t guarantee requires independent Gaussian reset totals; good observed performance on t_3 noise does not extend that guarantee. Median of means needs independent resets and an externally valid variance bound. At insufficient N it uses a paid prior-bound fallback. Neither robust interval repairs mean bias or a change in future dynamics. The full design retains 7,489 negative-advantage rows and 1,743 unaffordable purchase rows. Pooling all methods and scenarios into one success percentage would be misleading.

Information acquisition is an open economic decision. The fixed-N protocol buys information when affordable; it does not optimize the purchase. In 32 of 288 generated economies the known service premium is no greater than known linear execution/carry cost. Positive inventory then has no positive expected service surplus before any information is obtained. In the first indexed case, s000_c000 at $N = 32$ and baseline price, no-data and response-informed oracle profit both equal approximately 0.75902, because response information cannot relax the relevant position limit. Paying 0.024 for information leaves

approximately 0.73500. A useful measurement need not change the financial action. A purchase rule should use only information available before purchase, include the option to abstain, and avoid selecting N or a policy after looking at hidden gains or test outcomes. That rule is still to be proved and validated; it has not been added to this returned run.

Observation boundary and contribution status. There are no historical coins, exchange tapes or crash replays in this run. The quote-data stage is **NOT_RUN**; the unit-test quote fixture is synthetic. Controlled resets, known plant/noise inputs, a feedback cap and stationarity are maintained assumptions. The finite-horizon positive-response certificate needs no adaptation-rate floor when every relevant tap is observed under that reset contract. This does not remove the hidden-tail problem for infinite-horizon or ordinary public-tape certification. New observations do not independently verify the earlier passive, capital, AI or equilibrium claims.

The supported interpretation is a conditional financial separation: aggregate response information can support a specified exposure decision without latent cohort recovery, and positive net value requires additional economic conditions. The written general proof still needs independent review and its distinction from existing identification, data-informativity and decision theory needs specialist assessment. No publication-quality rating follows from execution.

Traceability. The immutable experiment source is repository commit `6a319b0db147`. The return audit is in commit `c0c0c010ec00`. The study directory and its relative audit directory are shown separately for readability:

```
studies/advanced_cost_separation_20261029/
validation/returned_colab_20261030/
```

The uploaded archive is:

```
MM_advanced_deep_seed20261101_v1.zip
```

Its SHA256 begins `e35c10f6b276`; the full hash, configuration, source fingerprint and exploratory price/paired audits are retained in that directory. No original return outcome has been overwritten.

Update of 5 November 2026. The current release adds sixteen kernel-checked supervisory witness and decision theorems, a full-project strict axiom audit and a conservative gate evaluated after information costs. This closes the new core witness and declared scalar-contract verification tasks. It does not mark the all-population construction, empirical market mapping, exhaustive novelty comparison or independent review below as completed.

The authoritative release scope is recorded in `studies/supervisory_release_20261105/RELEASE_REVIEW.md`.

Research work register: 30 October 2026

This register is a dated snapshot, not a claim that the project is complete. *Completed* means the named artifact or check exists. *Written, review open* means an argument is present but has not received independent mathematical sign-off. *Pending* means the evidence or proof is still required. Good outcomes from another experiment do not silently close a task.

Completed records to preserve

Record	Verified completion and boundary
C1	Returned archive, source fingerprint and 288 checkpoint links verified. Same-code reproducibility; no independent implementation.
C2	All 27 implementation checks passed. They do not test every theorem in the papers.
C3	Seven numerical Lean certificates passed. The full general theorem is not formalized.
C4	288-economy synthetic audit completed; every negative and unaffordable row retained. No market-data stage ran.
C5	Pointwise positive-value illustration retained; broad price-sensitive results and assumption failures recorded. No universal profitability claim.
C6	Targeted literature comparison and a written restricted separation exist. Originality and complete proof review remain open.

Core work before stronger contribution claims

R1. Establish the precise literature gap. Status: pending specialist review. Compare the exact assumptions, observation experiment, decision target and conclusion with identification theory, data informativity, task-optimal exploration, functional estimation and constrained financial decisions. Separate latent cohort recovery from learning the observable response. Do not claim that control or allocation without complete identification is new. *Completion evidence:* a source-backed comparison table identifying a specific non-overlap, or a clear statement that the result is an application or corollary; resolved reviewer comments. More simulations cannot complete R1.

R2. Independently audit the load-bearing proofs. Status: written, review open. Review the finite-cohort renewal bounds, observation-to-confidence implication, uniformity in hidden memory, fixed-query testing obstruction, all fee/cash constraints and the same-error no-data optimum. Audit the older broad certification, passive-identification and capital claims separately. Check quantifiers, constants, admissible probes, reset assumptions and counterexamples. *Completion evidence:* theorem-by-theorem assumptions and proof register, resolved objections and independent mathematical review. Record each Lean theorem's actual scope; do not substitute seven scalar facts for a general proof.

R3. Make the information purchase economically explicit. Status: pending. Define what the desk knows before buying observations and whether the acquisition objective is robust, Bayesian or another stated criterion. Include abstention and screen cases where the premium cannot cover known costs or information cannot change the constrained position. Select sample size and price using available information, not hidden gains or test outcomes. *Completion evidence:* a precise acquisition rule, a valid decision-value or regret statement, matched-information/cost comparisons, and a frozen new validation design. Reusing this return to design a rule is exploratory.

R4. Justify the financial institution and measurement contract. Status: partial. The fixed inventory, horizon, post-settlement premium and accounting are explicit. The economic source of the premium, operating cost of reset observations, information reliability and acquisition opportunity cost need justification. An own-account investor, warehousing desk, market maker and venue supervisor have different actions and incentives. Keep their claims separate. *Completion evidence:* a coherent institutional mechanism and documented parameter/cost interpretation, with comparative statics. If logarithmic utility is introduced, rederive the objective rather than relabeling expected profit.

R5. Establish an observation bridge or restrict the empirical claim. Status: pending. Ordinary public quotes are not independent resets. Address initial state, fundamental-value proxies, endogenous order flow, changing parameters, noise covariance, sign changes and unobserved tails. A reduced-form VAR score is not automatically the structural gain in the theorem. *Completion evidence:* a justified identification/partial-identification argument under stated conditions, or a forecasting-only interpretation with certification claims removed from that application. Test the conditions rather than treating good forecasts as proof of them.

R6. Run a protected real-data evaluation. Status: NOT_RUN. Document event-level quote provenance, instruments, units, dates, fees, carry and earlier use. Freeze venues/dates and specifications before test access. Use chronological purging, volatility and OFI/depth baselines, prespecified inference units and all failure outcomes. An exposure replay must account for spread, funding, fill/latency assumptions and capital at risk. Distinguish isolated round trips from a portfolio backtest and from market-wide intervention. *Completion evidence:* licensed/reproducible data description, immutable split and code, held-out results after costs, limitations and external replay. The synthetic unit-test fixture cannot complete R6.

R7. State what happens when certification is unavailable. Status: partial diagnostics. The returned bias and future-change cases demonstrate failure. Specify a defensible response such as abstention or a prior-only restriction; determine what verified disclosure could restore the missing assumption and what it costs. Repeated use requires simultaneous

or time-uniform control, not repeated use of a single-decision confidence level. *Completion evidence:* a stated contract/failure-response rule with a valid guarantee or an explicit impossibility boundary, tested under prespecified violations. Larger samples under a wrong mean do not repair the problem.

R8. Obtain independent replication and reconcile claims. Status: pending. Have another implementation reproduce the central calculation and meaningful stress cases. Reconcile abstracts, conclusions, theorem tags, old result tables, duplicate paper layouts, co-author contributions and AI-use statements. Keep the original evidence and limitations visible. Treat internal "proved" tags as statements about written arguments until independent review is recorded. *Completion evidence:* an external replication record, a claim-to-evidence map, resolved discrepancies and co-author approval for circulation.

Extensions that remain ambitions

Endogenous dealer depth and quotes, venue-wide gating, equilibrium welfare, dynamic inventories and AI-adoption externalities are separate extensions. They should enter the main claim only with a proved mechanism and appropriate evidence. The new fixed-position experiment does not establish those outcomes. A strong core contribution need not contain every extension.

Order of work and maintenance rule

Prioritize R1–R3 together: determine what is new, whether it is true, and when the information purchase changes an economically meaningful decision. R4–R6 determine the strength of the finance application. R7–R8 determine defensible use and reproducibility. Professor review can guide and challenge each item; it is not a substitute for producing the missing proof or evidence.

For each update record the date, task ID, changed claim, immutable code/data reference, result and remaining limitation. Mark an item completed only when its named evidence exists. Update this register, its Markdown/JSON source and all PDF variants together. Publication quality remains a judgment about originality, correctness and usefulness, not a score produced by the notebook.