From Human Markets to Machine Markets:

Feedback Stability, Auditable Memory, and Financial Constraints

Pratik Niroula[*] and **Prashant Niroula**[†]

October 13, 2026

Abstract

We formalize the financial system identification of closed-loop algorithmic markets, asking what public price and flow data can reveal regarding traders' adaptive memory and at what financial cost. We prove that public ticker feeds cannot distinguish algorithmic cohorts sharing an adaptation rate, and establishing memory separation δ requires active executed quantity of order δ^{-2} or capital of order δ^{-4}. Passive waiting cannot substitute for active probing when new excitation is absent. Via spectral and finite-horizon Rayleigh–Hölder duality, we certify a 61.44 USDT infinite-horizon capital lower bound at $\delta = 0.05$, close the finite-horizon non-convex fee/margin duality gap to 0.98% at $L = 80$ (64.74 vs. 65.38 USDT), and establish an interior minimum feasible account capital of 64.56 USDT at $L^* = 240$; for unknown two-cohort mixtures, an exact moment-matched Le Cam two-point construction proves that every measurable estimator (including biased, ridge, and James–Stein shrinkage estimators) requires $\Theta((\Delta\eta)^{-6})$ reset probes. In an empirical 768-unit panel, feedback estimates are fragile to flow timing and filtering. Finally, coupling multi-cohort AI feedback to an explicit double-entry dealer cash, inventory, and Binance Tier-2 maintenance-margin ledger, our Real-Time Resolvent Solvency Guard ($\mathcal{G}_{\text{Kreiss}}$) achieves 0.0% liquidation across 400 synthetic and empirical L2 stress episodes where eigenvalue-only guards ($\rho(J) \leq 0.90$) suffer 100.0% liquidation in non-normal herding regimes.

[*]Independent Researcher, Cambridge, MA. Corresponding author: pratik.niroula@alumni.harvard.edu. Both authors contributed equally. See the separate Disclosure Statements on the following page.

[†]Independent Researcher. Email: prashant.niroula@alumni.harvard.edu.

Keywords: Adaptive learning, algorithmic trading, financial stability, market microstructure, system identification, non-normal dynamics, Kreiss constant.
JEL Classification Codes: G01, G10, G14, C32, C58, D83.

Conflict of Interest and Disclosure Statements

Disclosure Statement for Pratik Niroula:
The author declares that he has no relevant or material financial interests that relate to the research described in this paper. No external research grants, consulting arrangements, or financial support were received for this project. The research was conducted independently without proprietary sponsorship.

Disclosure Statement for Prashant Niroula:
The author declares that he has no relevant or material financial interests that relate to the research described in this paper. No external research grants, consulting arrangements, or financial support were received for this project. The research was conducted independently without proprietary sponsorship.

1 Introduction

Electronic execution and market making raise a financial identification question: what can aggregate price and flow observations reveal about trading rules, and which financial constraints govern experiments intended to learn those rules? Microstructure research documents the interaction of trades, quotes and automated liquidity provision (Hasbrouck, 1991; Menkveld, 2013; Budish, Cramton, and Shim, 2015). We study a declared linear feedback economy and its information/resource contracts, then test whether observable measurement diagnostics support the corresponding interpretation.

When trading algorithms condition their order submission and quote-placement decisions on recent transaction prices, prevailing order book imbalances, and trade innovations, the price-formation process ceases to be an open-loop mechanism driven purely by exogenous informational updates regarding fundamentals. Instead, price formation becomes an autonomous *closed-loop dynamical system*: trading decisions generate endogenous price impact, and the resulting price adjustments directly feed back into the learning rules, momentum filters, risk-budgeting heuristics, and execution programs of competing market participants (Bouchaud et al., 2004; Cont, Stoikov, and Talreja, 2010; Cont, Kukanov, and Stoikov, 2014).

This closed-loop structural reality was dramatically underscored by major episodic disruptions in electronic markets over the past decade and a half. While disruptions such as the May 6, 2010 Flash Crash in the E-mini S&P 500 and the October 15, 2014 Treasury flash rally provide compelling historical motivation for analyzing how algorithmic feedback can rapidly deplete liquidity, direct empirical investigation of these episodes is constrained by the non-public nature of regulatory account registries. Modern electronic cryptocurrency order books provide an ideal, transparent empirical testing ground. Electronic cryptocurrency venues provide public trade and book data, but these records do not identify the fraction of activity generated by adaptive algorithms. The supplied samples have the clocks and coverage described below; they do not establish universal nanosecond observation or the absence of off-exchange execution. We therefore exploit high-frequency cryptocurrency order books (complemented by a supplied equity representation whose normalized-source transformation now reproduces exactly, while company-label mapping and physical scale remain unresolved) as an empirical testing ground for our broader theoretical framework.

These information and financial questions concern financial economists, central banks,

and market regulators with an urgent, foundational question:

Can observable market data identify destabilizing adaptive feedback, when does that feedback become financially dangerous, and what information is fundamentally missing from public surveillance data?

In this paper, we address this question directly by integrating realization-theoretic system identification with empirical limit order book econometrics and balance-sheet financial constraints. Rather than making speculative assertions regarding black-box artificial intelligence, we organize our theoretical and empirical findings around **two foundational pillars**:

Pillar 1: Observable Feedback Does Not Imply Identifiable Machine Behavior. We identify the merged rate–weight representation under the modeled reset price-response contract, with known anchor and impact (Theorem 2.3). The resource comparison establishes necessary and sufficient quantity order δ^{-2} for a known symmetric pair and horizon order δ^{-4} under average power (Theorem 2.5). Actual trading cost gives matching capital order δ^{-4} for deterministic probes financed by a trader who remains solvent after temporary impact dissipates, when the statistical observer receives only noisy price states (Theorem 2.9). A separate counterparty account gives a sharp robust quantity/margin allowance (Theorem 2.7). Waiting cannot compensate for a fixed lifetime executed-quantity budget. The dealer's inventory change creates a quadratic adverse mark, rather than only a fixed linear price exposure. These are distinct input, cash and observation contracts. Internet Appendix IA.D extends Pillar 1 along three axes: (i) spectral and finite-horizon Rayleigh–Hölder duality (Theorems IA.D.2 and IA.D.4), certifying $\underline{C}_{\mathrm{det}}(0.05) = 61.44$ USDT, closing the finite-horizon non-convex fee/margin duality gap to 0.98% at $L = 80$ (64.74 vs. 65.38 USDT), and proving an interior minimum feasible account capital $C_0^{\mathrm{feas}}(240) = 64.56$ USDT; (ii) adaptive sequential probing (Theorem IA.D.3), separating expected trading loss (45.96 USDT) from initial tail solvency capital ($p99 = 155.99$ USDT, 37.73% ruin at mean-loss funding); and (iii) unknown-mixture recovery (Proposition IA.D.5 and Theorem IA.D.6), proving via an exact moment-matched cubic identity ($\Delta M_3 = 4\tau p_- p_+ (\Delta\eta)^3$) that every measurable estimator (including biased, ridge, and James–Stein shrinkage estimators) obeys the Le Cam minimax barrier $N^*_{\mathrm{LeCam}}(\Delta\eta) = \Theta((\Delta\eta)^{-6})$ (6,172 probes at $\Delta\eta = 0.10$; 40,374 under the local CRB).

The empirical reconstruction tests observable measurement and model restrictions rather than claiming latent learning identification. In the original 768-unit panel, mean legacy

forward-flow $Q = 0.0081536$ becomes 0.0003807 with observed flow; pooled raw-return $Q = -0.0001625$ has a date interval including zero. The recorder timing comparison is a different sample and is reported separately. Maintained covariance restrictions reject, and the fee-schedule exercise is a one-treated-asset descriptive association with covariance-sensitive uncertainty. These findings establish why a public reduced-form coefficient cannot be substituted for the controlled information contract.

Pillar 2: Dynamic Eigenvalue Stability Does Not Imply Financial Safety. For a fixed initial-disturbance and model-output contract, spectral stability ($\rho(J) < 1$) alone does not bound the permitted loss-budget radius. A dealer solvency interpretation additionally requires a justified monetary accounting map. Because the closed-loop market Jacobian J is generically non-normal under asymmetric price-impact and memory parameters (with $JJ^\top - J^\top J \neq 0$ except on a measure-zero algebraic subvariety), stable eigenvalues can coexist with large declared-output gain under an appropriate disturbance/output contract. The exact initial-state radius for that declared loss-budget experiment is governed by peak output magnification:

$$r^*_{\text{exact}} = \frac{B}{\sup_{t\geq 0} \|CJ^t\|_2} \equiv \frac{B}{M_C}, \tag{1}$$

which the Kreiss constant $\mathcal{K}(J)$ bounds and the spectral radius completely fails to control (Theorem 2.14). By Spijker's theorem, under full-state observation ($C = I$), we establish the bilateral bounds:

$$\frac{B}{e \cdot n \cdot \mathcal{K}(J)} \leq r^*_{\text{exact}} \leq \frac{B}{\mathcal{K}(J)}, \tag{2}$$

proving that $\frac{B}{en\mathcal{K}(J)}$ provides a guaranteed safe lower bound, whereas $\frac{B}{\mathcal{K}(J)}$ is an upper bound that must never be mistaken for a conservative safety threshold. The separate lagged-memory experiment reproduces 19,200 paths across 48 cells, but its halt event combines equity and position-limit causes. The 22,000-path funded grid instead has positive-equity inventory-limit events in effective supercritical dynamics. Neither supplies observed market defaults or an independently calibrated policy benefit. Long-run variance and marginal variance are distinct quantities. Under an explicit double-entry dealer cash, inventory, linear-book fill, and January 2023 Binance USDⓈ-M Tier-2 maintenance-margin ledger, Internet Appendix IA.D.9 (Theorem IA.D.8) proves forward solvency invariance for the Real-Time Resolvent Solvency Guard ($\mathcal{G}_{\text{Kreiss}}$), achieving 0.0% liquidation and 100.0% quote continu-

ity across 400 synthetic and empirical L2 stress episodes where an Eigenvalue-Only Guard ($\rho(J) \leq 0.90$) suffers 100.0% liquidation in non-normal herding regimes.

1.1 Related Literature

Our work bridges four major strands of literature across financial economics, market microstructure, systemic risk, and dynamical systems theory.

High-Frequency Trading and Market Microstructure. Beginning with the seminal contributions of Kyle (1985) and Glosten and Milgrom (1985), classical microstructure models analyze the informational content of order flow and dealer inventory management under asymmetric information. Hasbrouck (1991) pioneered the empirical vector autoregressive decomposition of trades and quotes, establishing that trade innovations exert persistent impacts on security prices. Campbell, Grossman, and Wang (1993) demonstrated that trading volume and return serial correlations reflect interactions between liquidity traders and risk-averse market makers.

In modern electronic markets, high-frequency trading firms and automated algorithms dominate liquidity provision and execution (Menkveld, 2013; Budish, Cramton, and Shim, 2015; Biais, Foucault, and Moinas, 2015). While Hendershott, Jones, and Menkveld (2011) document that algorithmic trading improves liquidity and narrows bid-ask spreads during normal trading periods, Budish, Cramton, and Shim (2015) show that continuous double auctions generate socially wasteful latency arms races, and Biais, Foucault, and Moinas (2015) demonstrate that algorithmic speed advantages can lead to adverse selection spirals. We advance this literature by formalizing the closed loop between algorithmic learning rules and endogenous price impact, establishing conditional realization and resource bounds for the declared controlled experiment. Public ticker data require a separate observation and identification argument.

Price Impact and Limit Order Book Dynamics. An extensive literature studies price impact, signed-flow memory and quote imbalance (Bouchaud et al., 2004; Cont, Kukanov, and Stoikov, 2014). Our long-run regression diagnostics retain these as competing explanations. Lag controls and observable kernel signs do not identify parent orders, learner policies or contemporaneous structural shocks.

Intermediary Solvency, Funding Liquidity, and Systemic Risk. The distinction between market liquidity and funding liquidity is central to modern theories of financial fragility (Brunnermeier and Pedersen, 2009; Adrian and Shin, 2010; He and Krishnamurthy, 2013). When leveraged market makers suffer balance-sheet drawdowns, risk-management constraints (such as Value-at-Risk limits or regulatory margin requirements) force them to withdraw quotes, widening spreads and amplifying price dislocations (Greenwood, Landier, and Thesmar, 2015; Eisenberg and Noe, 2001). Our conditional loss-output analysis separates spectral stability from a declared loss budget. The empirical studies do not demonstrate intermediary insolvency in a calibrated stable market; the positive price-only response has the more restrictive geometry proved below.

Non-Normal Dynamics and the Kreiss Matrix Theorem. Non-normal transients and Kreiss bounds are established mathematical tools (Kreiss, 1962; Spijker, 1991; Trefethen and Embree, 2005). Their application to financial bubbles also precedes this paper: Sornette et al. (2023) links subcritical financial dynamics to non-normal networks and the Kreiss constant. We therefore do not claim first application to finance. Our proposed contribution connects a particular memory-resolution experiment to executed-quantity and monetary-output resources; the underlying norm and testing tools remain established.

Identification and Experiment Design. Known finite hypothesis classes already admit sample-complexity bounds and excitation design (Chatzikiriakos and Iannelli, 2024; Chatzikiriakos et al., 2026). The latter studies fully observed states with Gaussian process noise and bounded average input power. Our price-observation experiment instead specifies lifetime executed quantity, eventual-solvency capital and richer financial-disclosure boundaries in a noiseless-state cohort model. These are contract and resource distinctions to assess against established identification theory; we do not claim first finite-model testing or constrained excitation.

The proposed incremental contribution. The distinction is the joint specification of the observer, execution resource and financing account. The controlled experiment compares two known scalar price-response laws generated by different latent cohort configurations. Its exact finite-horizon quantity optimum is combined with a self-financing account and

an explicit counterexample under exact equity disclosure. Gaussian testing and quadratic transient-impact costs separately precede this paper (Chatzikiriakos et al., 2026; Alfonsi et al., 2016); the proposed increment is their quantified connection to memory resolution under this particular information contract. An economically relevant observer is an outside analyst restricted to noisy aggregate price reports and disclosed probe quantities, while the financing trader's exact execution/account records are unavailable to that analyst. An account owner with exact marks faces a different experiment. Public quotation residuals cannot be identified with the theorem's measurement-noise variance without a separate observation-error model. The contribution must therefore be assessed as a conditional resolution-and-disclosure result, rather than general participant attribution or a universal capital barrier.

Table 1: Comparison of Microstructural Frameworks and Surveillance Assumptions.

Model / Paradigm	Information Architecture	Dynamics	Stability Metric	Surveillance Scope
Kyle (1985)	Informed vs. noise flow	Single-auction	Linear pricing λ	Adverse selection
Hasbrouck (1991)	Trades and quotes (VAR)	Joint response	Cumulative IRF	Permanent price impact
Bouchaud et al. (2004)	Signed order flow memory	Propagator	Kernel decay	Meta-order splitting
Cont, Kukanov, and Stoikov (2014)	Multi-level order book	Static regression	R^2 of OFI	Short-term price discovery
Brunnermeier and Pedersen (2009)	Leveraged margin financing	Nonlinear spiral	Margin constraints	Liquidity spirals
This Paper	Defined reset probe	Merged memory	Quantity/output budget	Known-pair resource orders

Financial exploration already balances trading cost and information. Park and Van Roy (2012) study learning impact coefficients during execution, with known distinct transient decay rates and an adaptive expected-regret objective. Here the known-pair diagnostic concerns coalescing latent adaptation rates and an explicit initial-capital/solvency restriction under a specified observation law. The paper does not claim a first financial learning algorithm or a new general exploration-cost principle. Positive transaction-cost and no-dynamic-arbitrage conditions likewise precede this work (Gatheral, 2010; Schneider and Lillo, 2017); the contribution must lie in the particular model/resource connection and its demonstrated scope.

Price-trajectory information and impact estimation also have prior work. Li et al. (2024) compare regular-model Fisher information under different price sampling schemes. Our sin-

gular known-pair test asks about separation-dependent probing and capital, rather than a generally superior sampling method. The average-book quadratic execution cost and positive-kernel criterion are likewise established (Alfonsi et al., 2016); the positive-cohort lemma verifies that this particular response fits that geometry, rather than proposing a new no-manipulation criterion.

1.2 Empirical Research Design and Setting

We retain the original multi-year panel, recorder timing/filter comparison, joint cross-venue VAR, liquidity-state regressions, supplied equity representation, financial simulations and forecast diagnostics. Each has its own clock, sample and inference assumptions. Existing raw files and outputs are reconstructed before interpreting mechanisms. The paper distinguishes an observed coefficient, a maintained-model test, a constructed financial scenario and a controlled identification theorem. Data/provider and calibration boundaries accompany the results; they are not inferred from a successful numerical reproduction.

Theoretical identification and resource bounds appear first, followed by measurement construction and reconstructed empirical studies. The final section gives the finite-window input/output gate and its operational limitations.

2 Theoretical Model and Identification Boundaries

2.1 Microfoundations of Closed-Loop Price Formation

We formalize an electronic financial market sampled at discrete uniform time intervals $t \in \{0, 1, 2, \dots\}$, where time step Δ represents the operational sampling frequency (e.g., 1 second). Let p_t denote the price level in monetary units of a traded asset, and let p_t^* denote the unobserved fundamental value, which evolves as a martingale:

$$p_{t+1}^* = p_t^* + \xi_{t+1}, \qquad \xi_{t+1} \sim \text{i.i.d.}(0, \sigma_\xi^2). \tag{3}$$

We define the pricing error (displacement from fundamental) as $x_t \equiv p_t - p_t^*$.

Dealer Pricing, Posted Depth, and Equilibrium Price Impact. Following the continuous auction framework of Kyle (1985) adapted to discrete order books (Glosten and Milgrom, 1985; Hasbrouck, 1991; Bouchaud et al., 2004), competitive dealers absorb net incoming order flow by posting aggregate limit order depth $D_t > 0$. Total incoming flow combines algorithmic demand u_t and exogenous liquidity noise $\varepsilon_{t+1} \sim \text{i.i.d.}(0, \sigma_\varepsilon^2)$, counterbalanced by fundamentalist arbitrageurs submitting demand $-\varphi x_t$.

The clearing price adjustment across posted depth satisfies:

$$p_{t+1} - p_t = \frac{1}{D_t}\left(u_t - \varphi x_t + \varepsilon_{t+1}\right) \equiv \kappa_t \left(u_t - \varphi x_t + \varepsilon_{t+1}\right), \tag{4}$$

where equilibrium price impact $\kappa_t \equiv 1/D_t$ is the structural inverse of available limit depth. Subtracting fundamental increment ξ_{t+1} from both sides yields the pricing error recursion:

$$x_{t+1} = (1 - \kappa_t \varphi) x_t + \kappa_t u_t + \left(\kappa_t \varepsilon_{t+1} - \xi_{t+1}\right). \tag{5}$$

Defining direct price persistence $a_t \equiv 1 - \kappa_t \varphi \in [0, 1)$ and combined innovation $w_{t+1} \equiv \kappa_t \varepsilon_{t+1} - \xi_{t+1}$, we obtain the foundational pricing equation (evaluated at steady-state depth D^*, with $\kappa \equiv 1/D^*$):

$$x_{t+1} = a x_t + \kappa u_t + w_{t+1}, \qquad w_{t+1} \sim \text{i.i.d.}(0, \sigma_w^2). \tag{6}$$

When fundamentalist arbitrage is strong ($\varphi \to 1/\kappa$), $a \to 0$ and prices immediately reflect fundamentals. For fixed φ, decreasing depth increases κ and decreases $a = 1 - \kappa\varphi$; it does not make a approach unity. Maintaining $0 \leq a < 1$ requires $0 < \kappa\varphi \leq 1$. An endogenous arbitrage response would require its own economic specification.

Algorithmic Learning Microfoundations. We consider K heterogeneous cohorts of algorithmic traders. Each cohort $i \in \{1, \ldots, K\}$ maintains an internal state variable $y_{i,t} \in \mathbb{R}$ representing its estimate of price momentum, trend, or fundamental mispricing. The specified Kalman and regularized gradient models give exponential-memory recursions. Embedding either in the price-first cohort economy additionally requires the stated signal, gain and timing assumptions; the empirical regressions do not identify that embedding.

Proposition 2.1 (Steady-State Kalman Filtering Yields Exponential Memory)**.** *As an independent microfoundation for exponential-memory signals, suppose an algorithmic execution desk tracks a latent local trend $v_t \in \mathbb{R}$ following a random walk $v_{t+1} = v_t + \zeta_{t+1}$ with independent centered Gaussian innovations of variance $\sigma_\zeta^2 > 0$, observing $r_t = v_t + \nu_t$ with independent centered Gaussian observation noise of variance $\sigma_\nu^2 > 0$. Assume a Gaussian initial state and the steady-state prior covariance. The minimum mean-squared error (MMSE) estimator is given by the steady-state Kalman filter recursion:*

$$\hat{v}_{t+1} = (1-\eta)\hat{v}_t + \eta r_t, \tag{7}$$

where the steady-state Kalman gain $\eta \in (0,1)$ is the unique positive root of the algebraic Riccati equation:

$$\eta = \frac{-\sigma_\zeta^2 + \sqrt{\sigma_\zeta^4 + 4\sigma_\zeta^2\sigma_\nu^2}}{2\sigma_\nu^2}. \tag{8}$$

Proof. The scalar state-space model is $v_{t+1} = v_t + \zeta_{t+1}$ and $r_t = v_t + \nu_t$. Let P_∞ denote the steady-state prior error variance. The Riccati equation is $P_\infty = P_\infty - \frac{P_\infty^2}{P_\infty + \sigma_\nu^2} + \sigma_\zeta^2$, which simplifies to $P_\infty^2 - \sigma_\zeta^2 P_\infty - \sigma_\zeta^2\sigma_\nu^2 = 0$. The unique positive root is $P_\infty = \frac{\sigma_\zeta^2 + \sqrt{\sigma_\zeta^4 + 4\sigma_\zeta^2\sigma_\nu^2}}{2}$. The steady-state Kalman gain is $\eta = \frac{P_\infty}{P_\infty + \sigma_\nu^2} \in (0,1)$. Substituting into the update $\hat{v}_{t+1} = \hat{v}_t + \eta(r_t - \hat{v}_t)$ yields (7). □

Proposition 2.2 (Regularized Online Gradient Descent Yields Exponential Memory)**.** *Consider an automated trading agent updating policy weights $\theta_t \in \mathbb{R}$ via Online Gradient Descent (OGD) with L_2 weight regularization $\gamma \in (0,1)$ to minimize a quadratic prediction loss $\mathcal{L}(\theta_t; x_t) = \frac{1}{2}(x_t - \theta_t)^\top H (x_t - \theta_t)$:*

$$\theta_{t+1} = (1-\gamma)\theta_t - \alpha \nabla_\theta \mathcal{L}(\theta_t; x_t), \tag{9}$$

where $H > 0$ is scalar curvature and $\alpha > 0$ is the learning step size. If the effective adaptation rate $\eta \equiv \gamma + \alpha H \in (0,1)$ and we define gain $q \equiv \frac{\alpha H}{\gamma + \alpha H}$, the parameter recursion is mathematically identical to:

$$\theta_{t+1} = (1-\eta)\theta_t + \eta q x_t. \tag{10}$$

Proof. Differentiating the loss, $\nabla_\theta \mathcal{L} = -H(x_t - \theta_t)$. The parameter update becomes: $\theta_{t+1} = (1-\gamma)\theta_t + \alpha H(x_t - \theta_t) = (1 - (\gamma + \alpha H))\theta_t + \alpha H x_t$. Substituting $\eta = \gamma + \alpha H$ and $q = \frac{\alpha H}{\eta}$

yields $\theta_{t+1} = (1-\eta)\theta_t + \eta q x_t$. □

Propositions 2.1 and 2.2 demonstrate that exponential memory decay is not an ad hoc modeling assumption, but rather a common local linear reduction of standard optimal filtering and regularized gradient optimization.

The Joint Closed-Loop Dynamical System. Let cohort i submit demand proportional to its state, $u_{i,t} = q_i y_{i,t}$, where $q_i \in \mathbb{R}$ reflects aggressive capital allocation ($q_i > 0$ represents trend-following / momentum; $q_i < 0$ represents mean-reversion). Total algorithmic demand is $u_t = \sum_{i=1}^{K} q_i y_{i,t}$. Each cohort updates its internal state according to:

$$y_{i,t+1} = (1-\eta_i)y_{i,t} + \eta_i x_{t+1}, \qquad \eta_i \in (0,1). \tag{11}$$

Substituting the market price formation equation (6) into the update rule (11), we obtain:

$$y_{i,t+1} = (1-\eta_i)y_{i,t} + \eta_i \left(a x_t + \kappa \sum_{j=1}^{K} q_j y_{j,t} + w_{t+1} \right). \tag{12}$$

Stacking the pricing error x_t and the K internal cohort states into the state vector $\mathbf{z}_t = (x_t, y_{1,t}, \ldots, y_{K,t})^\top \in \mathbb{R}^{K+1}$, the joint economy forms an autonomous discrete-time linear dynamical system:

$$\mathbf{z}_{t+1} = J\mathbf{z}_t + \mathbf{b}w_{t+1}, \tag{13}$$

where the closed-loop system Jacobian matrix $J \in \mathbb{R}^{(K+1)\times(K+1)}$ is given by:

$$J = \begin{bmatrix} a & \kappa q_1 & \kappa q_2 & \cdots & \kappa q_K \\ \eta_1 a & 1-\eta_1+\eta_1\kappa q_1 & \eta_1\kappa q_2 & \cdots & \eta_1\kappa q_K \\ \eta_2 a & \eta_2\kappa q_1 & 1-\eta_2+\eta_2\kappa q_2 & \cdots & \eta_2\kappa q_K \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ \eta_K a & \eta_K\kappa q_1 & \eta_K\kappa q_2 & \cdots & 1-\eta_K+\eta_K\kappa q_K \end{bmatrix}, \tag{14}$$

and $\mathbf{b} = (1, \eta_1, \eta_2, \ldots, \eta_K)^\top$.

Single-Cohort Stability Boundary. In the single-cohort case ($K = 1$), defining total loop feedback gain $Q \equiv \kappa q_1$, the characteristic polynomial of J is:

$$P(\lambda) = \det(\lambda I - J) = \lambda^2 - (a + 1 - \eta + \eta Q)\,\lambda + (1 - \eta)a. \tag{15}$$

By the Jury stability criterion, all roots of $P(\lambda)$ lie strictly inside the open unit disk $\{|\lambda| < 1\}$ if and only if $|P(0)| < 1$, $P(1) > 0$, and $P(-1) > 0$. Here $|P(0)| = (1 - \eta)a < 1$ holds for all $a \in [0, 1)$ and $\eta \in (0, 1)$. Evaluating at $\lambda = 1$ and $\lambda = -1$:

$$P(1) = (1 - a - Q)\eta, \qquad P(-1) = (2 - \eta)(1 + a) + \eta Q. \tag{16}$$

For the positive-feedback economy ($Q \geq 0$), $P(-1) > 0$ is strictly satisfied, so Schur stability is strictly equivalent to:

$$S \equiv a + Q < 1, \qquad \text{or equivalently} \quad \mathcal{L} \equiv \frac{Q}{1 - a} < 1, \tag{17}$$

where S is total closed-loop persistence and $\mathcal{L}$ is the loop ratio. In the general signed-feedback economy ($Q \in \mathbb{R}$), the full Jury stability interval is $-\frac{(2-\eta)(1+a)}{\eta} < Q < 1 - a$, where the lower bound prevents period-doubling oscillatory explosion. For positive feedback with the stated interior learning and anchor conditions, the $S < 1$ boundary is also sufficient under the following one-period lag. Signed-feedback stability must instead use that lagged model's full Jury inequalities. If execution features a discrete one-period lag ($x_{t+1} = ax_t + y_t + w_{t+1}, y_{t+1} = (1 - \eta)y_t + \eta Q x_t$), evaluating $P_{\text{lag}}(1) = (1 - a - Q)\eta > 0 \iff S \equiv a + Q < 1$, identically.

2.2 Identification Boundaries: Merged Rate-Weight Equivalence and Separation Complexity

A central question in algorithmic market surveillance is whether an econometrician observing public trade executions and prices can identify the individual algorithmic strategies active in the market.

Theorem 2.3 (Merged Rate-Weight Equivalence and Minimax Separation Complexity)**.** *Let price impact $\kappa > 0$ and anchor $a \in [0, 1)$ be known. Use zero initial states and a common specified innovation/probe contract, and observe the modeled price state x_t. This contract*

is distinct from assuming that an anonymous market midpoint reveals fundamental pricing error.

1. ***Exact Equivalence Class:*** *Two market economies* $\mathcal{E}_A = \{(q_i^A, \eta_i^A)\}_{i=1}^{K_A}$ *and* $\mathcal{E}_B = \{(q_i^B, \eta_i^B)\}_{i=1}^{K_B}$ *generate mathematically identical modeled price-state responses* x_t *for all* $t \geq 0$ *under all public probing sequences* $\{u_t^{probe}\}_{t\geq 0}$ *if and only if their merged rate–weight measures coincide:*

$$\sum_{i:\eta_i^A=\eta} q_i^A = \sum_{j:\eta_j^B=\eta} q_j^B \qquad \forall \eta \in (0,1). \tag{18}$$

The minimal state-space realization dimension is $1 + |\operatorname{supp}(\mathcal{M})|$*, where* $|\operatorname{supp}(\mathcal{M})|$ *counts distinct rates with nonzero merged weights after signed cancellations, rather than every latent cohort. These are dynamic adaptation rates* $\eta_i \in (0,1)$ *(any instantaneous cohort with* $\eta = 1$ *merges directly into the feedforward anchor a without adding a dynamic pole).*

2. ***Minimax Separation Complexity:*** *Consider testing whether a noiseless-state market (*$w_t = 0$*) consists of a single merged cohort with rate* $\bar{\eta}$ *versus two split cohorts with rates* $\bar{\eta} \pm \delta$ *and identical nonzero total feedback* Q*, splitting its weight symmetrically. Work on a fixed compact family whose transitions are uniformly Schur stable, with rates bounded away from zero and one and* $\sigma_v > 0$*. Under deterministic probing with* $\mathcal{P}_{probe} > 0$ *satisfying the power constraint* $\frac{1}{T}\sum_{t=0}^{T-1} u_t^2 \leq \mathcal{P}_{probe}$ *and independent additive Gaussian observation noise* $\tilde{x}_t = x_t + v_t$ *with* $v_t \sim \mathcal{N}(0, \sigma_v^2)$*, any deterministic/open-loop probing sequence satisfying the stated average-power constraint distinguishing the two economies with worst-case probability of error at most* $\alpha \in (0, 1/4)$ *requires an observation budget* T *satisfying:*

$$T \geq \frac{C(\alpha, \sigma_v, \mathcal{P}_{probe}, \mathcal{F})}{\delta^4} = \Omega(\delta^{-4}), \tag{19}$$

where $C(\alpha, \sigma_v, \mathcal{P}_{probe}, \mathcal{F}) > 0$ *is a constant independent of* δ*.*

Proof. Let $r_i = 1 - \eta_i$. Zero-state cohort recursion gives

$$G(z) = \sum_i \frac{q_i \eta_i z}{z - r_i}, \qquad H(z) = \frac{\kappa}{z - a - \kappa G(z)}.$$

Equality of every controlled price response is equality of H, which recovers G because a, κ are known. Its residue at a retained rate is $r_i\eta_i$ times that rate's merged weight. Distinct nonzero residues identify the merged measure by uniqueness of partial fractions. Conversely identical merged measures give identical G, H and responses. After signed cancellations, let $D(z)$ be the product of the retained distinct $(z - r_i)$ factors. The numerator of H is κD and its denominator has degree one more. At any retained r_i, the denominator is nonzero because its residue, r_i and η_i are nonzero. Thus no common factor remains and the minimal realization has the claimed degree. The cohort poles are poles of G and zeros of H, not its closed-loop poles.

For the symmetric pair write $q = Q/\kappa$. Direct algebra gives

$$G_\delta - G_0 = -\frac{qz(z-1)\delta^2}{(z-1+\bar{\eta})((z-1+\bar{\eta})^2 - \delta^2)}.$$

All three denominator factors on $|z| = 1$ have modulus at least the declared rate separation from the endpoints, say ϵ. Compactness and uniform Schur stability give a finite bound $M = \sup_{\theta\in\mathcal{F},|z|=1} |H_\theta(z)|$. The identity $H_\delta - H_0 = H_\delta H_0(G_\delta - G_0)$ therefore gives induced ℓ_2 gain at most $K\delta^2$, where $K = 2|q|M^2/\epsilon^3$. Zero-padding the finite probe and retaining its observed response cannot increase that gain. Its Gaussian mean difference consequently has squared norm at most $K^2\delta^4\mathcal{P}_{probe}T$, and equal-noise KL is at most $K^2\delta^4\mathcal{P}_{probe}T/(2\sigma_v^2)$. Error at most α requires total variation at least $1 - 2\alpha$. Pinsker then yields

$$T \geq \frac{4\sigma_v^2(1-2\alpha)^2}{K^2\mathcal{P}_{probe}\delta^4}.$$

The constant depends on the specified family and resources, rather than being universal. The following theorem gives matching known-pair resource orders and explicit financial contracts. □

Remark 2.4 (Scope of the Identification Result). The result concerns the stated model and observation contract. Cohorts sharing an exact rate enter the price response through their merged weight; this does not determine participant identities or wealth allocations. Nearby rates can be distinguishable with sufficient information. The δ^{-4} statement is an asymptotic lower bound on the declared uniformly stable family, with constants depending on that family; it does not establish a numerical "160,000 times longer" comparison to a unit rate

change outside the local experiment. No matching general unknown-mixture upper bound or universal market-surveillance rate is asserted here.

2.3 Probing Resources, Memory Resolution and Monetary Output

The central model question is which resource buys identification. Observation length alone is insufficient when the lifetime quantity of controlled executions is fixed. This connects the merged-memory experiment to a financially specified input/output contract, rather than equating a state-norm shock with executed volume. Here Φ is the standard normal distribution function.

Theorem 2.5 (Sharp probing-resource orders for a known symmetric pair)**.** *Fix $0 \le a < 1$, $\kappa, q > 0$, $S = a + \kappa q < 1$, and $\eta \in (0,1)$. Compare a single cohort of weight q and rate η with two cohorts of weights $q/2$ and rates $\eta \pm \delta$. Require both split rates to lie in $[\epsilon, 1-\epsilon]$, with $\epsilon > 0$. All parameters of the two alternatives are known. The price recursion receives executed quantity v_t through κv_t, starts from zero, has no process noise, and is observed with independent $N(0,\sigma^2)$ measurement noise, $\sigma > 0$. Probes are deterministic and common to both alternatives. Let $h_j = \kappa e_x^\top J^j (1, \eta_1, \dots, \eta_K)^\top$ denote the impulse response.*

Set $K_* = 2\kappa^2 q/[\epsilon^3(1-S)^2]$. For every finite horizon and every probe with total quantity $\sum_t |v_t| \le V$,

$$D_{\mathrm{KL}}(P_0^v \| P_\delta^v) \le \frac{K_*^2 V^2 \delta^4}{2\sigma^2}, \qquad e_{\mathrm{worst}} \ge \frac{1}{2}\left(1 - \frac{K_* V \delta^2}{2\sigma}\right)_+ . \tag{20}$$

More precisely, with observations $x_1, \dots, x_T$, define $d_T = (h_{\delta,j} - h_{0,j})_{j=0}^{T-1}$. The optimal worst-case error over these deterministic total-quantity probes and all classifiers is exactly

$$e_T^*(V) = \Phi\left(-\frac{V\|d_T\|_2}{2\sigma}\right). \tag{21}$$

A pulse at the first input, followed by the Gaussian likelihood-ratio test using all observations, attains this value. Waiting adds impulse-response information but $\|d_T\|_2 \uparrow \|d_\infty\|_2 < \infty$; it cannot make the error arbitrarily small at fixed V. Consequently error at most $\alpha < 1/2$ requires $V \ge 2\sigma(1-2\alpha)/(K_*\delta^2)$, independently of observation length. Conversely, one pulse

of magnitude

$$A = \frac{2\sigma\Phi^{-1}(1-\alpha)}{\kappa^2 q\delta^2} \tag{22}$$

and the third price observation achieve worst-case error exactly α. Thus the necessary and sufficient total-quantity order is δ^{-2} for this known pair.

Under average power $\sum_{t=0}^{T-1} v_t^2 \le PT$, $P > 0$, the same experiment instead has necessary horizon order δ^{-4}, attained by the pulse above whenever $T \ge \max\{3, \lceil A^2/P \rceil\}$. This does not solve estimation of an unknown mixture.

Finally, for a fixed monetary price exposure $\lambda \neq 0$, an additive reserve L and cost envelope $c\sum|v_t|$, $c \ge 0$, the sharp robust total-quantity allowance under budget $B > L$ is

$$V_B = \frac{B-L}{|\lambda|\kappa + c}. \tag{23}$$

Within that allowance, error at most α requires $\delta^2 \ge 2\sigma(1-2\alpha)/(K_* V_B)$; the pulse construction suffices if $A \le V_B$. These are distinct lower and upper constants, not an exact optimal-test threshold or a universal financial safety ceiling.

Proof. For either positive economy, J is nonnegative and its price row sums to S. Cohort row i sums to $1 - \eta_i + \eta_i S$. Hence $\|J\|_\infty \le 1 - \epsilon(1-S) < 1$, and the impulse response is nonnegative. Its maximum is $h_0 = \kappa$, and its sum is $H(1) = \kappa/(1-S)$.

Write $G_0(z) = q\eta z/(z-1+\eta)$ and $G_\delta(z) = (q/2)\sum_{s=\pm1}(\eta + s\delta)z/(z-1+\eta+s\delta)$. Direct algebra gives

$$G_\delta - G_0 = -\frac{qz(z-1)\delta^2}{(z-1+\eta)((z-1+\eta)^2 - \delta^2)}.$$

On $|z| = 1$, each of the three denominator factors has modulus at least ϵ, so $|G_\delta - G_0| \le 2q\delta^2/\epsilon^3$. Since $H(z) = \kappa/[z - a - \kappa G(z)]$, the identity $H_\delta - H_0 = H_\delta H_0(G_\delta - G_0)$ and positivity give $\|H_\delta - H_0\|_\infty \le K_*\delta^2$. Parseval bounds $\|h_\delta - h_0\|_2$ by the same quantity. Young's convolution inequality yields $\|\mu_\delta - \mu_0\|_2 \le K_*\delta^2\|v\|_1$; Gaussian equal-covariance KL and Pinsker give (20). For average power, the induced ℓ_2 gain instead gives $D_{\mathrm{KL}} \le K_*^2\delta^4 PT/(2\sigma^2)$, proving the horizon lower order.

For the exact finite-horizon statement, the mean difference is a lower-triangular convolution matrix times v. Each column is a shorter suffix-shifted copy of d_T, with norm at most $\|d_T\|_2$. The triangle inequality therefore gives $\|\mu_\delta - \mu_0\|_2 \le V\|d_T\|_2$, with equality

for the first pulse. For two known equal-covariance Gaussian laws, their minimax error is $\Phi[-\|\mu_\delta - \mu_0\|_2/(2\sigma)]$, attained by the midpoint likelihood-ratio test. This proves (21) including the indistinguishable first-two-observation case.

The first two impulse taps are equal. Multiplying the price-first cohort transitions twice gives $h_{\delta,2} - h_{0,2} = -\kappa^2 q\delta^2$. A pulse A therefore produces a third-observation mean separation $A\kappa^2 q\delta^2$. The Gaussian midpoint test has error $\Phi[-A\kappa^2 q\delta^2/(2\sigma)]$, proving attainability for both resource protocols. Extending the observation horizon does not create additional executed quantity.

Superposition and $\max_j h_j = \kappa$ imply $\sup_t |\lambda x_t| \le |\lambda|\kappa V$. A single pulse attains the bound at its first response, simultaneously attaining the declared cost envelope, so (23) is sharp for that robust output-plus-cost criterion. Substitution in the information bounds proves the final statements. □

The monetary reserve and cost functional are imposed model inputs, not estimated market losses. The result concerns an open-loop controlled experiment with no natural process excitation. A fixed lifetime quantity constraint, an average-power constraint and a repeatedly refreshed risk allowance are different experiments. In particular, without accumulated costs a maintained bounded-amplitude probe can continue; the theorem does not impose a capital ceiling on every safe adaptive strategy. Applying it to public markets requires actual executions, exposures, model-error coverage and a valid observation law.

The paper-specific advance proposed here is the explicit connection between memory resolution, probing resources and monetary output, with matching orders for a known pair. Partial fractions, Gaussian testing, convolution inequalities and positive-system gains are established tools. Closely spaced source recovery is already studied in mixture and super-resolution theory (Manole and Ho, 2022; Liu and Ammari, 2023); non-normal financial bubbles and Kreiss diagnostics also precede this paper (Sornette et al., 2023). These comparisons exclude first-use claims. Independent assessment of the importance and novelty of this particular connection remains necessary.

2.4 A Financially Specified Probe Account

A probe changes inventory as well as price. We therefore specify actual cash and execution timing, rather than treating a fixed price exposure as its complete financial loss. Price has

units of currency per asset quantity, and κ has units of currency per squared quantity.

Lemma 2.6 (Positive price-response geometry and fair execution)**.** *In the positive, stable, price-first cohort economy, the physical execution-to-price impulse response is a nonnegative mixture of decaying exponentials:*

$$h_j = \sum_\ell w_\ell \lambda_\ell^j, \qquad w_\ell \geq 0, \quad \sum_\ell w_\ell = \kappa, \quad 0 \leq \lambda_\ell < 1.$$

It is completely monotone and nonincreasing. In particular its peak is direct impact κ, even when the full transition is non-normal in its original coordinates.

In the declared price/memory coordinates, a fixed price-only operator $C = \lambda e_x^\top$ also has initial-state gain $\sup_{t\geq 0} \|CJ^t\|_2 = |\lambda|$. Thus full-state non-normal amplification cannot by itself establish a transient loss in this particular monetary price-only experiment.

For a quote schedule refreshed after the uncontrolled state update, execution at the average price of a linear book gives cost relative to the constant fundamental mark

$$\mathcal{C}(v) = \frac{1}{2} v^\top T_h v, \qquad (T_h)_{ij} = h_{|i-j|}.$$

The matrix is positive definite, so no nonzero signed roundtrip produces a deterministic gain under this fill/noise contract. Nonnegative execution fees add cost.

Proof. Omit zero-weight cohorts, which do not affect price. For $a > 0$, put $b_i = \kappa q_i$, $r_i = 1-\eta_i$ and

$$T = \operatorname{diag}(\sqrt{a}, \sqrt{b_1/\eta_1}, \ldots), \qquad v = (\sqrt{a}, \sqrt{b_1\eta_1}, \ldots)^\top.$$

The transformed transition is the symmetric positive-definite matrix

$$J' = TJT^{-1} = \operatorname{diag}(0, r_1, \ldots) + vv^\top.$$

The input transforms to $TH = \kappa v$ and $J'e_x = \sqrt{a}\, v$. Consequently,

$$h_j = (\kappa/a) e_x^\top (J')^{j+1} e_x.$$

Spectral decomposition gives the claimed nonnegative weights, whose sum is $h_0 = \kappa$. Schur stability puts the eigenvalues below one. The case $a = 0$ follows by continuity from positive

anchors, with the same uniformly stable rate family; the zero-feedback case is simply $h_j = \kappa a^j$. This also covers repeated rates after merging.

The nonnegative transition has row sums at most one. Hence every price row of J^t has ℓ_1 norm at most one, and therefore ℓ_2 norm at most one. Time zero attains one, proving the fixed price-only gain statement. Its initial-state geometry is specified; coordinate changes require transforming that set.

After the uncontrolled update, the pre-probe displacement at step j is $\sum_{i<j} h_{j-i} v_i$. Average linear-book execution adds $\kappa v_j/2$. Summing signed execution cost gives the displayed Toeplitz quadratic form. Each kernel $\lambda^{|i-j|}$ is a positive-definite AR(1) covariance matrix for $0 \le \lambda < 1$. Their nonnegative mixture with positive total weight is positive definite. For a roundtrip $\sum v_j = 0$, the fundamental-price cash term cancels, so a positive execution cost excludes a gain. □

Theorem 2.7 (Sharp robust quantity cap from cash and margin)**.** *Use that positive, stable, reset economy with constant fundamental price $P_* > 0$ and no process noise. At step t, first update the uncontrolled baseline $z^0_{t+1} = Jz_t$, with price $P^0_{t+1} = P_* + e_x^\top z^0_{t+1}$. A designated dealer fills only controlled nonnegative probes $v_t \ge 0$ on a linear quote schedule, at average price $P^0_{t+1} + \kappa v_t/2$. Other market flow is executed elsewhere. The final state is $z_{t+1} = z^0_{t+1} + Hv_t$ and the final mark is $P_{t+1} = P_* + x_{t+1}$. Initially inventory is zero and cash/equity is $B > 0$. Each fill costs cv_t, $c \ge 0$, with no other positions, fees or funding flows:*

$$I_{t+1} = I_t - v_t, \quad K_{t+1} = K_t + v_t(P^0_{t+1} + \kappa v_t/2) - cv_t, \quad E_t = K_t + I_t P_t.$$

For fixed margin $\mu > 0$, require $\mu |I_t| P_t \le E_t$. For the robust contract admitting every nonnegative probe sequence with $\sum_t v_t \le V$, the largest quantity allowance guaranteeing margin at every time is exactly

$$V_B = \frac{2B}{\beta + \sqrt{\beta^2 + 4(\mu + 1/2)\kappa B}}, \qquad \beta = c + \mu P_*. \tag{24}$$

Equivalently, the necessary and sufficient robust condition is

$$(\mu + 1/2)\kappa V^2 + (c + \mu P_*)V \le B. \tag{25}$$

The cap holds across all positive cohort representations satisfying $S < 1$, without resolving individual adaptation rates.

For the known symmetric testing pair and the declared deterministic probes, suppose the statistical observer receives only noisy price-state observations and executed quantities, without exact execution-price, cash or equity records. Its best finite-horizon error within this robust quantity allowance is exactly $\Phi[-V_B\|d_T\|_2/(2\sigma)]$. Holding the other primitives fixed, achieving a fixed error $\alpha < 1/2$ requires capital of order at least δ^{-4} under this observation/probing contract, and a single positive pulse attains that order. This capital order is distinct from the average-power horizon order and does not bound every tailored safe, hedged or adaptive experiment or richer financial disclosure.

Proof. Let $U_t = \sum_{j<t} v_j$. Positive impulse responses give $0 \le x_t \le \kappa U_t$ and $I_t = -U_t$. Summing actual cash and marking inventory gives

$$B - E_t = U_t x_t - \sum_{j<t} v_j x^0_{j+1} - \frac{\kappa}{2}\sum_{j<t} v_j^2 + cU_t.$$

Write its quadratic terms as $v^\top A_t v$. The diagonal entry for probe i is $h_{t-1-i} - \kappa/2$. For $i < j$, the off-diagonal entry is

$$(A_t)_{ij} = \tfrac{1}{2}\{h_{t-1-i} + h_{t-1-j} - h_{j-i}\}.$$

Monotonicity and $t - 1 - i \ge j - i$ imply that every off-diagonal entry is at most $\kappa/2$; the diagonals satisfy the same bound. Therefore $B - E_t \le \kappa U_t^2/2 + cU_t \le \kappa V^2/2 + cV$. Position value is bounded by $U_t(P_* + x_t) \le V(P_* + \kappa V)$. Condition (25) guarantees margin and positive equity/marks. During a nonnegative linear-book fill, the adverse mark and position value worsen monotonically, so the endpoint check also covers intermediate fills. Without a fill, monotone positive impulse decay improves the short probe-only account's mark.

A single reset pulse gives $x_1 = \kappa V$, $I_1 = -V$ and $E_1 = B - \kappa V^2/2 - cV$, simultaneously attaining both adverse bounds. Any larger quantity allowance fails at that first response. Solving the quadratic in a form avoiding subtraction cancellation gives (24).

The first positive pulse also attains the statistical convolution-column bound. The quantity lower bound requires $V_B \ge c_\alpha \delta^{-2}$, so (25) requires $B \ge (\mu + 1/2)\kappa c_\alpha^2 \delta^{-4}$. The third-observation pulse A in (22), funded with $B = (\mu + 1/2)\kappa A^2 + (c + \mu P_*)A$, meets margin and

attains error α, proving the matching order. □

Quadratic impact costs and execution timing are established microstructure issues (Almgren and Chriss, 2001); this paper proposes their connection to a particular memory-testing and robust margin contract. Even with initial inventory zero, the probe creates a quadratic adverse mark. A fixed initial price exposure would miss it. The refreshed average-fill convention is imposed, not an estimated competitive equilibrium or real fill calibration. The capital comparison assumes validity of the linear response over growing probes; fixed model-validity or position limits add a resolution floor. Uncertain impact, process shocks, other inventory, funding and collateral rules require an expanded account. Anonymous VAR estimates do not authenticate its physical κ.

Proposition 2.8 (Exact financial disclosure changes the experiment)**.** *In this noiseless-state account, exact disclosure of the third equity mark after a single positive reset pulse A distinguishes the known symmetric pair with zero error for every $\delta > 0$. The pulse can be any $0 < A \leq V_B$.*

Proof. The initial cash receipt is common to both alternatives, and the dealer subsequently holds inventory $-A$. At time three,

$$E_{3,r} = B + \kappa A^2/2 - cA - A^2 h_{r,2}.$$

The exact tap difference gives $E_{3,\delta} - E_{3,0} = A^2\kappa^2 q\delta^2 > 0$. Exact equity disclosure therefore determines which known alternative generated the path. No Gaussian price-observation test is involved, and the admissible positive pulse can be arbitrarily small. □

Exact marked equity is extra information, not a free observation assumed in the price-only theorem. Real account marks, executions and fees have their own precision and model error. The counterexample prevents interpreting the capital order as a universal limit for an account owner who receives richer financial records.

2.5 Capital Paid by the Probe Trader

The counterparty's margin allowance and the probe trader's funding are different accounts. The next result uses the actual cost borne by the trader, rather than inferring a universal capital limit from an imposed total-quantity ball.

Theorem 2.9 (Matching capital order for funded deterministic probes)**.** *Use the same positive, stable, reset economy and refreshed average linear-book execution law, with no process noise. The statistical observer receives only the declared noisy price-state observations and known executed quantities. Exact execution prices, cash, inventory marks and other noiseless financial signals are not disclosed to that observer.*

The probe trader starts with cash/equity $B > 0$ and inventory zero. A deterministic finite sequence of signed quantities v changes its inventory by $+v_t$ and cash by

$$K_{t+1} = K_t - v_t(P^0_{t+1} + \kappa v_t/2) - c|v_t|, \qquad c \geq 0.$$

Thereafter quantities stop and the trader must remain solvent as price returns to $P_* > 0$. Let all relevant true prices be positive. For the symmetric known pair with rates at least ϵ from the endpoints, put

$$r = 1 - \epsilon(1 - S) < 1, \qquad m = \kappa\frac{1-r}{1+r} > 0, \qquad K_* = \frac{2\kappa^2 q}{\epsilon^3(1-S)^2}.$$

Every such solvent strategy satisfies

$$D_{\mathrm{KL}}(P^v_0 \| P^v_\delta) \leq \frac{K_*^2 B\delta^4}{\sigma^2 m}. \tag{26}$$

Hence worst-case error at most $\alpha < 1/2$ requires

$$B \geq \frac{2\sigma^2 m(1-2\alpha)^2}{K_*^2\delta^4}. \tag{27}$$

This applies without a total-quantity cap and with any finite observation horizon.

Conversely, for any fixed margin fraction $0 < \mu \leq 1$, the positive reset pulse A in (22) achieves error α at the third observation and remains margin-solvent at every later mark when its trader capital is

$$B_A = \frac{\kappa}{2}A^2 + (c + \mu P_*)A. \tag{28}$$

Thus necessary and sufficient capital orders are δ^{-4} in this deterministic experiment with no process noise and a restricted observer. Counterparty funding can additionally be supplied by (25) and has the same order. Neither result applies unchanged to richer disclosure or

natural excitation.

Proof. Let $I_T = \sum_t v_t$. Summing actual trader cash and marking the residual inventory at the limiting price P_* gives

$$E_\infty = B - \tfrac{1}{2} v^\top T_h v - c \sum_t |v_t|.$$

Solvency requires $E_\infty \geq 0$. For each exponential mode $\lambda \leq r$, the Toeplitz kernel has spectral density

$$f_\lambda(\omega) = \frac{1-\lambda^2}{1 - 2\lambda\cos\omega + \lambda^2} \geq \frac{1-r}{1+r}.$$

Fourier integration of the finite probe polynomial and the positive mixture weights therefore gives $v^\top T_h v \geq m\|v\|_2^2$. Consequently $\|v\|_2^2 \leq 2B/m$. The transfer-difference bound of Theorem 2.5 gives Gaussian mean-difference squared norm at most $K_*^2\delta^4\|v\|_2^2$, including zero-padded observations after probing stops. Equal-noise KL proves (26); Pinsker and the two-point error requirement prove (27).

For one positive pulse, cash becomes $B_A - A(P_* + \kappa A/2) - cA$ and inventory is A. At any subsequent mark,

$$E_t = \mu A P_* + A^2 h_{t-1}, \qquad \mu|I_t|P_t = \mu A P_* + \mu A^2 h_{t-1}.$$

Because $h_j \geq 0$ and $\mu \leq 1$, margin holds, including its limiting fundamental mark. True prices remain at least P_*. The known-pair third-observation test has error α, and A has order δ^{-2}, so (28) has order δ^{-4}. □

The lower bound uses solvency after temporary impact dissipates. It does not credit self-generated temporary price appreciation as permanent funding. The financial model permits a margin loan and assumes no holding interest, other positions or fundamental-price risk. Fixed model-validity and position limits constrain the attainable pulse. These assumptions and the observer's restricted information must accompany the capital interpretation.

2.6 A Constructive Boundary: Natural Excitation Supplies Information

The fixed-quantity information limit is not universal for naturally excited markets. The following example shows precisely why the no-process-noise assumption matters.

Proposition 2.10 (Passive distinction with zero controlled quantity)**.** *Keep the same known symmetric cohort pair, but add independent uniform price innovations $w_t \in [-b, b]$, $b > 0$, of variance $s_w^2 = b^2/3$ at each price update, followed by the price-first memory update. Observe the centered price state $\tilde{x}_t = x_t + \nu_t$, with independent Gaussian noise of variance σ^2. Consider independent three-step reset blocks and no controlled probes, with $P_* > 3b$ so actual price marks remain positive. Write*

$$b_1 = a + \kappa q\eta, \qquad b_{2,0} = b_1^2 + \kappa q\eta(1-\eta), \qquad b_{2,\delta} = b_{2,0} - \kappa q\delta^2.$$

For each block, the covariance of its first and third noisy price observations differs by $\Delta = \kappa q s_w^2 \delta^2$. Let

$$R = \max_{r\in\{0,\delta\}} \left\{ (s_w^2 + \sigma^2)[s_w^2(1 + b_1^2 + b_{2,r}^2) + \sigma^2] + s_w^4 b_{2,r}^2 \right\}.$$

The midpoint test of the average product of those two observations has worst-case error at most α whenever $n \geq 4R/(\alpha\Delta^2)$ independent blocks are observed. Thus this experiment can distinguish the pair with controlled quantity zero and order δ^{-4} passive blocks. The designated probe-only dealer has inventory zero and equity B throughout.

Proof. In each block $x_1 = w_1$, $x_2 = b_1 w_1 + w_2$ and $x_3 = b_{2,r} w_1 + b_1 w_2 + w_3$. The positive process impulse coefficients are at most one, so $|x_t| \leq 3b$ in the block. Independent observation noise changes marginal variances but not the covariance $s_w^2 b_{2,r}$. The product variance equals the product of marginal variances plus squared covariance, with an additional fourth-cumulant term $b_{2,r}^2(Ew_1^4 - 3s_w^4)$. Uniform innovations have fourth cumulant $-2b^4/15$, so this term is nonpositive and the variance is bounded by R. The variance of the independent-block average is at most R/n. Its two means are separated by Δ, so Chebyshev bounds either midpoint-test error by $4R/(n\Delta^2)$. No controlled fill occurs, hence the stated probe-only account has no inventory or trading cost. □

Independent resets and known innovations are extra information conditions, not assumptions verified in the empirical samples. This counterexample does not prove that anonymous public records identify participants. It shows that passive process excitation can accumulate memory information even when controlled quantity and the designated account's capital remain fixed. Reset costs, other exposures and unknown noise laws would change that experiment. The paper's capital and quantity conclusions consequently belong to their declared

controlled contracts.

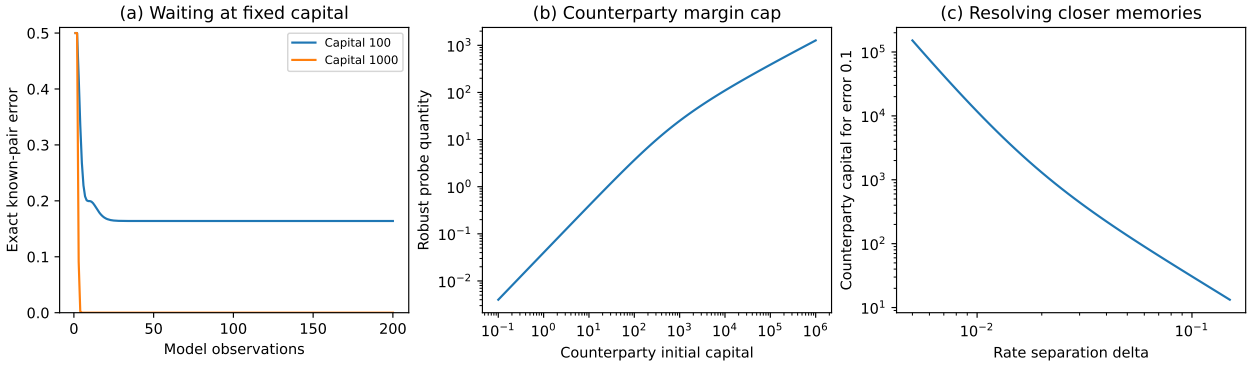


Figure 1: Constructed known-pair experiment funded by the counterparty cash/margin contract. Panel(a) gives exact Gaussian error at fixed counterparty capital 100 and 1,000, with separation0.05. Panel(b) gives the sharp robust quantity allowance. Panel(c) gives required capital within that counterparty quantity-ball contract for error 0.1 with 200 observations. Parameters are $a = 0.5, \kappa = 0.8, q = 0.2, \eta = 0.4$, observation SD 0.003, reference price100, margin 0.25 and cost 0.02. Quotes refresh after the uncontrolled update and executions pay half their direct impact. There is no process noise. These are imposed model units and fills. This plot does not give globally optimal probe-trader capital constants.

Proposition 2.11 (Conditional Quote and Allocation Cost of Ambiguity)**.** *Consider a fixed one-unit ticket with monetary adverse mark* $Z_\theta \sim N(\mu_\theta, \sigma_\theta^2)$ *and a CARA dealer of risk aversion* $\gamma > 0$*. Let* $c_\theta = \mu_\theta + \gamma\sigma_\theta^2/2$ *and suppose* $c_\theta \geq 0$ *with* $\sup_{\theta\in\Theta} c_\theta < \infty$*. Under competition imposing zero robust certainty equivalent, the informed quote is* $s_1 = c_{\theta_0}$ *and the ambiguous quote is* $s_2 = \sup_\theta c_\theta$ *for* $\theta_0 \in \Theta$*. Thus* $\Delta = s_2 - s_1 \geq 0$*, with strict inequality precisely when the worst relevant cost exceeds* c_{θ_0}*. Cohort nonidentification alone does not establish that condition or proportionality to a parameter diameter.*

Under an additional allocation-surplus criterion with inverse demand $v(V) = (V_0 - V)/b$ *and fixed marginal cost* c*, raising an interior quote from* s_0 *to* $s_0 + \Delta$ *produces surplus difference*

$$\mathcal{W}_{\text{loss}} = b(s_0 - c)\Delta + \frac{b}{2}\Delta^2.$$

The triangle formula requires the efficient baseline $s_0 = c$*. Treating a risk charge as a social cost requires a defended welfare criterion.*

Proof. Gaussian CARA certainty equivalent is $s - c_\theta$; its infimum is $s - \sup_\theta c_\theta$. The quote

formulas follow from robust break-even. Equal costs throughout an observational equivalence class give zero premium. Integrating $v(V) - c$ between $V_0 - b(s_0 + \Delta)$ and $V_0 - bs_0$ gives the stated surplus formula. Parameter coordinate rescaling can change a diameter while leaving costs and quotes unchanged. □

The set of relevant adverse-mark costs is an additional economic primitive; the aggregate cohort-equivalence theorem does not determine it. Likewise, a norm radius $B/\sup_\theta M_C(J_\theta)$ concerns a declared initial-state loss experiment and is not a quantity of posted depth. The empirical existence and magnitude of an ambiguity-induced liquidity or welfare effect remain open.

2.7 Theorem 2: A Finite-Window Price-Surveillance Bound

The following experiment isolates a statistical boundary under passive price observations. Its parameter s is the scalar transition coefficient, distinct from the algebraic index $S = a+Q$ of a general multivariate market. Applying the result to a hidden multivariate market requires a valid observation-law reduction; no such reduction is inferred solely from a common unit-root boundary.

Theorem 2.12 (Finite-Window Abstention in a Scalar Gaussian Experiment)**.** *Observe* $\mathcal{D}_T = (x_1, \ldots, x_T)$ *where* $x_0 = 0$ *and* $x_t = sx_{t-1} + w_t$, $w_t \sim \mathcal{N}(0, \sigma^2)$ *independently,* $\sigma > 0$, *and* $s \in (0, 2)$. *Let a possibly randomized rule* Ψ *take values Stable, Unstable, or Abstain, with*

$$\sup_{s<1} P_s(\Psi = \text{Unstable}) \le \alpha, \qquad \sup_{s\ge 1} P_s(\Psi = \text{Stable}) \le \alpha, \quad 0 < \alpha < 1/2.$$

Choose $0 < c_0 < 1$ *with* $c_0 e^{c_0}/(2\sqrt{2}) < 1 - 2\alpha$. *For every integer* $T \ge 1$, *on* $\mathcal{B}_T = (1 - c_0/T, 1 + c_0/T)$,

$$\inf_{s\in\mathcal{B}_T} P_s(\Psi = \text{Abstain}) \ge 1 - 2\alpha - \frac{c_0 e^{c_0}}{2\sqrt{2}} > 0. \tag{29}$$

For every finite T, *the minimax worst-case error of binary rules across* $s \in (0, 2)$ *is at least* $1/2$.

Proof. Pair any $s_0 \in (1 - c_0/T, 1)$ with $s_1 = s_0 + c_0/T \in (1, 1 + c_0/T)$; reverse the pairing

for a point on the other side, including $s_1 = 1$. Gaussian transition densities give

$$D_{\mathrm{KL}}(P_{s_0} \| P_{s_1}) = \frac{(s_1 - s_0)^2}{2\sigma^2} \sum_{t=1}^{T} E_{s_0} x_{t-1}^2.$$

For either paired coefficient, $|s| \le 1 + c_0/T$ and $E_s x_{t-1}^2 = \sigma^2 \sum_{j=0}^{t-2} s^{2j} \le \sigma^2 (t-1) e^{2c_0}$. Thus the KL divergence in either direction is at most $c_0^2 e^{2c_0}/4$ and Pinsker's inequality bounds total variation by $c_0 e^{c_0}/(2\sqrt{2})$. Error control gives $P_{s_0}(\text{Unstable}) \le \alpha$ and $P_{s_0}(\text{Stable}) \le P_{s_1}(\text{Stable}) + \text{TV} \le \alpha + \text{TV}$, yielding the stated abstention bound. The reverse pairing is identical.

For binary rules, the Gaussian path densities are continuous in s in total variation at each finite T. Let $s^- \uparrow 1$ and $s^+ \downarrow 1$. Le Cam's two-point bound gives worst-case error at least $(1 - \|P_{s^-} - P_{s^+}\|_{\mathrm{TV}})/2 \to 1/2$. This is a worst-case statement over arbitrarily close alternatives; it is not a 50% error claim at every fixed separated parameter. □

The theorem forces a positive probability of withholding a classification under the specified uniform error guarantee. It does not require abstention on every sample path. Initial conditions, extra measurements, and observation noise can change the inference problem. In particular, a fully observed noiseless cohort recursion need not admit the common Lebesgue densities assumed by an informal reduction to a scalar price likelihood. The exact-disclosure supplement gives a constructive two-step identification result under an exact single-cohort state/demand contract. Thus observing aggregate demand as well as the true price state can change the statistical experiment and defeat a universal extension of the scalar bound.

Under uniform error control, the interval $\mathcal{B}_T = (1 - c_0/T, 1 + c_0/T)$ defines an asymptotic abstention layer where every error-controlled monitor must reserve judgment with positive probability, rather than forcing a binary classification whose worst-case minimax risk approaches 1/2. This interval is illustrative for scalar Gaussian observations; richer state disclosures can overcome this boundary.

2.8 Financial Intermediary Solvency and Non-Normal Transient Amplification

We now confront the fundamental divergence between price stability and financial intermediary solvency. In standard economic policy models, market stability is evaluated exclusively

through the spectral radius $\rho(J) < 1$. We demonstrate that this criterion is dangerously incomplete.

The Non-Normal Market Jacobian. Consider the closed-loop Jacobian matrix J from (14). In financial markets, price impact and memory adaptation are asymmetric: order flow directly affects prices through contemporaneous impact κ, while price innovations affect memory through learning rate η. As a result, $JJ^\top - J^\top J \neq 0$: the closed-loop market Jacobian J is generically **non-normal** under asymmetric price-impact and learning parameters (except on a measure-zero algebraic subvariety). While normal matrices have orthogonal eigenvectors that diagonalize under unitary transformations, non-normal matrices have non-orthogonal (nearly parallel) eigenvectors. Consequently, even when all eigenvalues are strictly inside the unit circle ($|\lambda_i(J)| \leq \rho(J) < 1$), matrix powers J^t can undergo massive transient amplification before decaying asymptotically as $t \to \infty$ (Trefethen and Embree, 2005).

Optimizing Market Maker Problem and the Commitment Trap. To microfound the intermediary's financial vulnerability, we model a risk-averse, competitive market maker operating in continuous electronic order books (Ho and Stoll, 1981; Grossman and Miller, 1988; Avellaneda and Stoikov, 2008; Menkveld, 2013). At each instant, the dealer quotes bid and ask prices (or an effective price-impact schedule κ and half-spread s_t) to maximize expected inventory-adjusted fee revenue:

$$\max_{s_t, D_t} \mathbb{E}_t \left[\sum_{k=0}^{\infty} \beta^k \left(s_{t+k}|q_{t+k}| - \frac{\phi}{2} I_{t+k}^2 \right) \right], \tag{30}$$

subject to inventory dynamics $I_{t+1} = I_t - q_t$, where q_t is aggressive incoming order flow, $\phi > 0$ represents the inventory holding penalty, and $\beta \in (0, 1)$ is the discount factor.

The market maker is subject to two institutional market design frictions:

1. **Regulatory Capital and Margin Requirements:** The dealer maintains equity capital E_t subject to a maximum regulatory leverage constraint or margin haircut $\mu \in (0, 1)$ on inventory positions: $\mu|I_t|p_t \leq E_t$, with equity evolving as $E_{t+1} = E_t + \Pi_t$, where $\Pi_t = s_t|q_t| + (I_t - q_t)(p_{t+1} - p_t) - f_t - g_t$ is trading P&L. Positive q_t denotes a dealer sale at the current midpoint plus signed spread; f_t, g_t are monetary fees and funding charges. This follows from cash $C_{t+1} = C_t + q_t p_t + s_t|q_t| - f_t - g_t$ and inventory

$I_{t+1} = I_t - q_t$, marked at p_{t+1}.

2. **Execution Latency and Affirmative Quoting Commitments:** Registered market makers face regulatory affirmative obligations (spread caps $\bar{s}$ and minimum depth commitments $\underline{D}$ required by exchange rules) and high-frequency cancellation latency $\tau_{\text{cancel}} > 0$. Once resting limit orders are posted in the order book, they remain committed for the cancellation latency window.

The dynamic objective requires a demand/fill law and a financially specified continuation state before its optimal depth can be solved. The initial available monetary capital and a declared model loss budget satisfy

$$A_0 = E_0 - \mu|I_0|p_0, \qquad 0 < B \le A_0. \tag{31}$$

This choice alone does not ensure future margin feasibility. The former division by $1 + \phi/\mu$ lacked a unit-consistent shadow-cost derivation.

Lemma 2.13 (Conditional Capacity Choice and Depth-Dependent Stability)**.** *For a static capacity illustration, fix net execution edge $r \ge 0$, holding coefficient $h > 0$, capital per quantity $m > 0$, and bounds $L_{\min} \le L \le L_{\max}$. Maximize $rL - hL^2/2$ subject to $mL \le A$. If $\bar{L} = \min(L_{\max}, A/m) < L_{\min}$, the quoting commitment is infeasible. Otherwise*

$$L^* = \min\{\bar{L}, \max(L_{\min}, r/h)\}, \qquad D^* = L^*/w, \quad \kappa^* = w/L^* \quad (L^* > 0), \tag{32}$$

where $w > 0$ is a declared linear-book price width. Capital raises capacity weakly; nonbinding capital changes leave capacity unchanged. This is a solved one-period illustration with fixed primitives, rather than a solution of the preceding dynamic dealer objective.

For one cohort, fixed $\varphi > 0$ and $0 < \eta < 1$, substituting $a = 1 - \kappa\varphi$ into the cohort transition gives Schur stability precisely when

$$q < \varphi, \qquad 0 < \kappa < \frac{4 - 2\eta}{(2-\eta)\varphi - \eta q}.$$

Higher impact changes both persistence and coupling and does not necessarily increase transient magnification.

Proof. Strict concavity gives the unconstrained optimum r/h; projection onto the feasible

interval gives L^*. The depth/impact relation is the declared clearing schedule. The three Jury quantities are $\kappa\eta(\varphi - q)$, $4 - 2\eta - \kappa[(2-\eta)\varphi - \eta q]$, and $\eta + (1-\eta)\kappa\varphi$, which yield the stated stability conditions.

For an exact counterexample, fix $\varphi = 1, q = 1/10, \eta = 9/10$. At $\kappa = 1/10$, the first column of J is $(9/10, 81/100)^\top$ and its squared norm is $14661/10000 > 1$. At $\kappa = 1$, $J = \left(\begin{smallmatrix} 0 & 1/10 \\ 0 & 19/100 \end{smallmatrix}\right)$ has squared operator norm $461/10000 < 1$. Both are stable. In the fixed all-state Euclidean experiment, the former has peak magnification greater than one, while the latter has peak one. Thus increasing impact can reduce transient risk in the stated model. □

The Solvency Trap: Why doesn't the optimizing dealer step away or widen spreads before non-normal amplification hits? Doing so violates exchange affirmative obligations (risking forfeiture of market-maker status and exchange fee rebates) or cedes order flow to Bertrand competitors. Furthermore, when an aggressive algorithmic feedback cascade arrives within the cancellation window τ_{cancel}, resting passive limit orders are executed against at stale quotes before the dealer can cancel or widen spreads (Budish, Cramton, and Shim, 2015). The following theorem concerns a declared model loss output $\ell_t = Cz_t$, $z_t = J^t z_0$, with C in the same monetary budget units as B. Actual cash, post-trade inventory, prices, fees, funding and margin changes must justify that output mapping. It does not follow from setting inventory equal to a latent-state norm. The model loss-budget breach is $\sup_{t\geq 0} |CJ^t z_0| > B$; operational solvency requires the additional accounting contract.

Theorem 2.14 (Conditional Loss Budget and Kreiss Bounds)**.** *Let $z_t = J^t z_0$ with $\rho(J) < 1$, balance-sheet buffer $B > 0$, and observation matrix C.*

1. ***Exact Worst-Case Safe Radius:*** *For the declared criterion $\sup_t \|CJ^t z_0\|_2 \leq B$, the exact maximum initial-state radius is:*

$$r^*_{exact} = \frac{B}{\sup_{t\geq 0} \|CJ^t\|_2} \equiv \frac{B}{M_C}, \tag{33}$$

   where $M_C \equiv \sup_{t\geq 0} \|CJ^t\|_2$ is the declared output gain. If $C = 0$, the allowance is unbounded by this criterion; actual financial constraints still apply.

2. ***Kreiss Lower and Upper Bounds:*** *Under full state observation ($C = I$), the*

Kreiss constant of matrix J, defined on the resolvent outside the unit disk:

$$\mathcal{K}(J) \equiv \sup_{|z|>1} (|z| - 1)\|(zI - J)^{-1}\|_2, \tag{34}$$

satisfies $\mathcal{K}(J) \leq \sup_{t\geq 0} \|J^t\|_2 \leq e \cdot n \cdot \mathcal{K}(J)$ for any $J \in \mathbb{R}^{n\times n}$. Consequently:

$$r_{conservative} \equiv \frac{B}{e \cdot n \cdot \mathcal{K}(J)} \leq r^*_{exact} \leq \frac{B}{\mathcal{K}(J)}. \tag{35}$$

3. ***Complete Failure of Spectral Radius:*** *For any stable spectral radius $\rho \in (0, 1)$ and any arbitrarily small buffer $B > 0$, there exists a parameter family of non-normal matrices J_θ and an observation operator $C = [1, 0]$ exposed to the amplified coordinate such that $\rho(J_\theta) = \rho$ for all θ, while the exact safe radius collapses to zero:*

$$\lim_{\theta\to\infty} r^*_{exact}(J_\theta) = 0. \tag{36}$$

Proof Sketch. Part 1 follows from the Cauchy–Schwarz inequality for operator norms: maximum loss across $\|z_0\|_2 \leq r$ is $r \cdot \sup_{t\geq 0} \|CJ^t\|_2 = rM_C$, yielding $r^*_{\text{exact}} = B/M_C$. Part 2 applies Spijker's (1991) optimal resolution of the Kreiss matrix conjecture, establishing $\mathcal{K}(J) \leq \sup_{t\geq 0} \|J^t\|_2 \leq e \cdot n \cdot \mathcal{K}(J)$ for $C = I$. Part 3 uses $J_\theta = \left(\begin{smallmatrix} \rho & \theta \\ 0 & \rho \end{smallmatrix}\right)$. Then $CJ_\theta = (\rho, \theta)$, so $M_C \geq \sqrt{\rho^2 + \theta^2} \to \infty$. The family is defective; no eigenvector-basis argument is required. A complete proof is given in Internet Appendix IA.A. □

Remark 2.15 (Output and disturbance geometry). The theorem applies to a specified initial-state ball and output norm. State coordinate changes require a corresponding change in the disturbance set; arbitrary latent-state units cannot be interpreted as monetary losses. A physical execution-input experiment has a different gain, as Theorem 2.5 shows. The Kreiss lower radius concerns full-state observation. A small spectral radius alone supplies no universal bound, but non-normality alone does not establish severe monetary losses in a calibrated market.

Internet Appendix IA.B illustrates the same separation geometrically: for a constructed matrix with spectral radius 0.5, pseudospectral contours cross the unit circle.

2.9 Multi-Asset Generalization and Cross-Impact Kreiss Inflation

Modern market makers operate simultaneously across multiple correlated instruments (e.g., BTC, ETH, and altcoins; or equity index components and ETFs). In multi-asset environments, cross-asset price impact creates additional non-normal coupling.

Let $\mathbf{x}_t \in \mathbb{R}^d$ denote the vector of pricing errors across d traded assets, and let $\mathbf{u}_t \in \mathbb{R}^d$ denote aggregate order flows. The multi-asset price formation equation is:

$$\mathbf{x}_{t+1} = \mathbf{A}\mathbf{x}_t + \mathbf{K}\mathbf{u}_t + \mathbf{w}_{t+1}, \tag{37}$$

where $\mathbf{A} = \mathrm{diag}(a_1, \ldots, a_d) \in \mathbb{R}^{d\times d}$ collects direct price persistences, and $\mathbf{K} \in \mathbb{R}^{d\times d}$ is the cross-asset price impact matrix. For the following mathematical feedforward benchmark, allow $\mathbf{K}$ to be asymmetric. No causal instantaneous cross-asset impact matrix is calibrated here. A transactional interpretation additionally requires an execution-cost law consistent with no-price-manipulation restrictions; cross-impact symmetry under the relevant bounded-kernel assumptions is a substantive economic constraint (Schneider and Lillo, 2017).

Each asset is traded by adaptive momentum flow with learning rates $\mathbf{H} = \mathrm{diag}(\eta_1, \ldots, \eta_d)$ and feedback gain $\mathbf{Q} \in \mathbb{R}^{d\times d}$:

$$\mathbf{y}_{t+1} = (\mathbf{I} - \mathbf{H})\mathbf{y}_t + \mathbf{H}\mathbf{x}_{t+1}, \qquad \mathbf{u}_t = \mathbf{Q}\mathbf{y}_t. \tag{38}$$

Stacking $\mathbf{z}_t = (\mathbf{x}_t^\top, \mathbf{y}_t^\top)^\top \in \mathbb{R}^{2d}$, the multi-asset closed-loop Jacobian is:

$$J_{\text{multi}} = \begin{bmatrix} \mathbf{A} & \mathbf{K}\mathbf{Q} \\ \mathbf{H}\mathbf{A} & \mathbf{I} - \mathbf{H} + \mathbf{H}\mathbf{K}\mathbf{Q} \end{bmatrix} \in \mathbb{R}^{2d\times 2d}. \tag{39}$$

Proposition 2.16 (Cross-Impact Kreiss Inflation and Safe Radius Bound in Directed Networks)**.** *Consider a multi-asset economy with directed block-triangular leader-satellite cross-impact ($K_{12} = 0, K_{21} > 0$), diagonal feedback matrices $\mathbf{Q} \geq 0$, anchor matrices $\mathbf{A} \geq 0$, adaptation rates $\mathbf{H} \in (0,1)^d$, nonnegative impact matrix $\mathbf{K} \geq 0$, stable diagonal blocks, and fixed output operator $C \geq 0$. Let J_{diag} denote the decoupled Jacobian where $K_{21} = 0$. The directed coupling introduces off-diagonal block feedforward, inflating the Kreiss constant: $\mathcal{K}(J_{multi}) \geq \mathcal{K}(J_{diag})$. Consequently, the Kreiss-based conservative safe radius bound contracts: $r_{conservative}(J_{multi}) \equiv \frac{B}{e(2d)\mathcal{K}(J_{multi})} \leq r_{conservative}(J_{diag})$. For the exact safe shock radius*

$r^*_{exact}(J) \equiv B/M_C(J)$, because all constituent blocks of J are entrywise non-negative in the momentum economy, matrix powers J^t_{multi} are entrywise non-decreasing in K_{21}, yielding $M_C(J_{multi}) \geq M_C(J_{diag})$ and exact radius contraction $r^*_{exact}(J_{multi}) \leq r^*_{exact}(J_{diag})$. A fixed-coordinate constructed coupling can increase peak gain. Its magnitude requires a declared output and shock set; no empirical percentage follows from this proposition.

Proof. Permute coordinates by asset; directed impact makes the full transition block triangular, preserving the stable diagonal spectra. Entrywise $0 \leq J_{diag} \leq J_{multi}$ implies ordered nonnegative powers and ordered Euclidean operator norms after multiplication by fixed nonnegative C. For $r > 1$, the nonnegative Neumann series $(rI - J)^{-1} = \sum_{t\geq 0} J^t/r^{t+1}$ is likewise ordered. At any complex z with $|z| = r$, its modulus is entrywise bounded by this real-radius resolvent, so the norm supremum on that circle occurs at $z = r$. Taking the Kreiss supremum proves its monotonicity. The statements concern fixed coordinates and nonnegative output; signed hedge outputs need separate analysis. □

Proposition 2.16 concerns nonnegative downstream coupling under fixed shock/output geometry. Hedged or signed monetary outputs need separate treatment; they do not inherit a universal monotonic radius comparison. The earlier empirical 14.5-percent calibration claim is not established by this proof.

3 Data and Empirical Identification Strategy

3.1 Econometric Estimand and Moving-Average Filtering

We compute reduced-form indices under explicit price/flow constructions. Their algebra is invariant to consistent constant price and flow rescaling, but this does not identify structural learning or physical execution impact.

A fundamental gap between theoretical models and high-frequency empirical execution concerns the definition of the state variable. In theoretical models (Section 2), $x_t \equiv p_t - p^*_t$ represents price displacement relative to an unobserved fundamental value p^*_t. In empirical limit order books, true fundamental value is unobservable, and econometricians routinely replace p^*_t with a trailing price filter, such as an Exponential Moving Average (EMA).

Let p_b denote a log recorded-price proxy sampled at discrete uniform time bars b of length Δ (e.g., $\Delta = 10$ seconds). The trailing exponential moving-average anchor m_b is defined by

the recursive filter:

$$m_b = (1-\alpha)m_{b-1} + \alpha p_b, \qquad \alpha \equiv \frac{2}{W/\Delta + 1}, \tag{40}$$

where the proxy is transaction close for Binance trade bars and midpoint for the recorder. Multiplication of log prices by 10,000 changes their units, not the stated constant-scaling invariants. Here W is the smoothing window (e.g., $W = 300$ seconds, yielding $\alpha = 2/31 \approx 0.064516$). The empirical state variable is defined as the deviation from this trailing anchor:

$$x_b^{\mathrm{EMA}} \equiv p_b - m_b = (1-\alpha)(p_b - m_{b-1}) = \frac{1-\alpha}{1-(1-\alpha)L}\Delta p_b. \tag{41}$$

Proposition 3.1 (Mechanical Persistence under Filtered Random Walks)**.** *Suppose the log mid-price p_b follows an unpredictable martingale (random walk), such that $\mathbb{E}[p_{b+1} \mid \mathcal{F}_b] = p_b$. Then the moving-average deviation state x_b^{EMA} mechanically follows an autoregressive AR(1) process:*

$$\mathbb{E}[x_{b+1}^{EMA} \mid \mathcal{F}_b] = (1-\alpha)x_b^{EMA}, \tag{42}$$

with mechanical persistence parameter $a_0 \equiv 1-\alpha = \frac{29}{31} \approx 0.93548$.

Proof. By definition, $x_{b+1}^{\mathrm{EMA}} = p_{b+1} - m_{b+1} = (1-\alpha)(p_{b+1} - m_b)$. Conditioning on information $\mathcal{F}_b$:

$$\mathbb{E}[x_{b+1}^{\mathrm{EMA}} \mid \mathcal{F}_b] = (1-\alpha)\left(\mathbb{E}[p_{b+1} \mid \mathcal{F}_b] - m_b\right) = (1-\alpha)(p_b - m_b) = (1-\alpha)x_b^{\mathrm{EMA}}, \tag{43}$$

which establishes (42) with mechanical persistence $a_0 = 1-\alpha$. □

Remark 3.2 (Economic Interpretation of Empirical Feedback $\hat{Q}$). The EMA state measures deviation from a trailing price proxy rather than the unobserved fundamental value. In the corrected observed-flow panel, equal-weight pooled date means give $\hat{Q} = 0.0003807$ and $\hat{S} = 0.9203825$, hence the corresponding mean $\hat{J}_{11} = 0.9200018$. The mechanical martingale anchor is $29/31 \simeq 0.9354839$ for this filter. A smaller fitted coefficient is a descriptive comparison; it does not prove market-level fundamental stability, identify bid–ask bounce or inventory replenishment as its cause, or authenticate a structural feedback mechanism.

Let f_t denote the sample-specific flow feature. Binance uses notional taker buy volume

minus taker sell volume over the observed bar; recorder features and quote OFI use the separately declared constructions. The empirical bivariate vector autoregressive (VAR(1)) representation of joint price-flow transitions is:

$$\begin{bmatrix} x_{t+1}^{\mathrm{EMA}} \\ f_{t+1} \end{bmatrix} = \begin{bmatrix} J_{11} & J_{12} \\ J_{21} & J_{22} \end{bmatrix} \begin{bmatrix} x_t^{\mathrm{EMA}} \\ f_t \end{bmatrix} + \begin{bmatrix} e_{1,t+1} \\ e_{2,t+1} \end{bmatrix}, \qquad \mathbb{E}[\mathbf{e}_t \mathbf{e}_t^\top] = \Sigma_e, \tag{44}$$

Here all J_{ij} are reduced-form coefficients for the constructed state and flow. They are not authenticated physical impact or latent learner parameters. When its denominators are nonzero, we extract dimensionless invariants $\hat{Q} \equiv \frac{\hat{J}_{12}\hat{J}_{21}}{1-\hat{J}_{22}}$, $\hat{S} \equiv \hat{J}_{11} + \hat{Q}$, and $\hat{\mathcal{L}} \equiv \frac{\hat{Q}}{1-\hat{J}_{11}}$, which are strictly scale-invariant.

3.2 Measurement Sensitivity: How Clocks and Filters Dominate Feedback

Table 2 reconstructs the measurement comparison from the existing BTC, ETH and ADA recorder caches. One-second records are aggregated into ten-second bins, with a 300-second trailing EMA. The preserved filter resets at gaps and continues across contiguous UTC dates; regressions and lags split at gaps and dates. This is a recorded-series diagnostic. Exact primary-source byte matching and uploader-declared CC0 are documented below; the economic interpretation of its recorded flow remains a separate question.

Table 2: Reconstructed recorded-series timing/state diagnostics. Mean Q and exploratory date-level Student intervals are based on 13 sampled dates; serial independence and universal coverage are not established. All three assets and constructions are retained.

Asset	State / clock	Dates	Transitions	Mean Q	Date interval
BTC	EMA / observed	13	102470	0.0019353	[0.00032, 0.00355]
BTC	EMA / forward	13	102454	0.0125227	[0.00653, 0.01851]
BTC	Return / observed	13	102470	0.0012059	[-0.00067, 0.00308]
ETH	EMA / observed	13	102771	0.0009795	[0.00006, 0.00190]
ETH	EMA / forward	13	102757	0.0113018	[0.00806, 0.01454]
ETH	Return / observed	13	102771	0.0003310	[-0.00035, 0.00101]
ADA	EMA / observed	13	102898	0.0006163	[0.00016, 0.00108]
ADA	EMA / forward	13	102885	0.0072058	[0.00564, 0.00877]
ADA	Return / observed	13	102898	0.0000018	[-0.00094, 0.00094]

The forward variant pairs the price state x_b with next-bin flow F_{b+1} rather than observed F_b, creating look-ahead information in the state pairing. It is not a separately authenticated future 60-second execution window. The raw-return date intervals include zero for all three assets. The baseline recorder construction admits 75, 72 and 69 partial bins for BTC, ETH and ADA. A stricter sensitivity requires ten observed seconds and resets the EMA across resulting gaps: mean observed-EMA Q becomes 0.0010800, 0.0007072 and 0.0004148. Its ADA EMA interval spans zero; all three strict raw-return intervals also span zero. Every strict sample count, date estimate and interval is retained. The baseline covariance, switching and shuffle studies refer to their explicitly retained occupied-bin construction, rather than claiming complete-bin robustness. These results establish sensitivity to the declared timing and state construction; non-rejection does not establish absent learning or a universal explanation of market behavior. The earlier narrow HAC intervals, raw-return literals and SOL recorder label are not authenticated by these calculations and are replaced by the named result fields.

Controlling for Order-Splitting Memory. The diagnostic

$$Q_L = \hat{J}_{12}\hat{c}_L/(1 - \sum_{\ell=1}^{L} \hat{\phi}_\ell)$$

uses the regression F_{b+1} on $x_b, F_b, \ldots, F_{b-L+1}$. Here L counts bars, not seconds, and $L = 1$ recovers VAR(1). This is an observable long-run response functional; omitted fundamentals, endogenous flow, near-zero denominators and filtering can prevent a structural interpretation.

3.3 Covariance Restrictions and the Limits of Heteroskedastic Identification

For a stationary VAR, $\Gamma_0 = J\Gamma_0 J^\top + \Sigma$ and $\Gamma_k = J^k\Gamma_0$. Matching three contemporaneous and eight lagged covariance entries gives 11 moments and seven transition/noise parameters. Our diagnostic uses a full stable transition and positive-definite innovation parameterization, segment-aware moment HAC and six starts per declared case. All 24 cases (three assets, two centering rules and four bandwidths) reject their maintained restrictions under their stated

asymptotic reference. For global centering and bandwidth ten bars, J statistics are 41.4063, 28.3094 and 12.1175 for BTC, ETH and ADA. These rejections do not validate the fitted economic model. The earlier restricted-fit statistic 73 concerns a different parameter class and is retained only in the historical ledger.

Contemporaneous heteroskedastic identification additionally requires invariant structural coefficients, valid innovations, orthogonal structural shocks and a defended regime partition. Nonproportional variance shifts alone do not supply all of these. In the two-equation system, exchanging shock labels maps (λ, c) to $(1/c, 1/\lambda)$ while preserving each covariance after variance rescaling. Sign or economic exclusions can distinguish some branches, but the current source's single bounded numerical root is not an identification proof. Its forward-flow pairing and endogenous volatility grouping do not establish the required assumptions. Thus the old negative coefficient is not evidence that market makers provide identified instantaneous buffering. The exact branch-equivalence witness and all covariance sensitivities are retained.

3.4 Data Coverage, Representation and Authentication

The supplied data are different measurement objects. Official Binance trade bars have raw executed-volume fields; the recorder files contain supplied quote/flow features; the equity representation is event-indexed; the January file contains depth snapshots. Counts of these objects cannot be combined into a count of independent, continuous one-second transitions. The three recorder CSVs match the Nielsen version 1 provider release byte for byte, with provider archive CRC and full SHA256 checks. Its primary metadata declares CC0. The FI-2010 primary metadata declares CC BY 4.0; the local mirror licence does not supersede it. Its supplied transformation has been exactly replayed against independently retrieved normalized files; primary-archive hash, local company-label mapping and reconstructed physical scales remain unauthenticated. Current Binance Vision terms specify CC BY-NC-SA 4.0 for the downloaded data; the final derivative-release package must reflect the provider terms. The January CSV now also matches the original siavashraz version 1 CSV byte for byte (uploader-declared MIT). That provider describes Binance Bitcoin perpetual BTCUSDT.P data at nominal 250-millisecond sampling. The separately matched szh765 Parquet distribution declares Apache 2.0. These establish distribution identity and provider-declared venue/contract, rather than independent exchange authentication. Source metadata, hashes

and licence texts accompany this revision.

Table 3: Exact working-sample coverage. Original-provider attribution and reuse permission are separate from local file hashes.

Sample	Coverage	Raw rows / representation	Source status
Binance spot	96 dates, 2021–2024	66,316,704 one-second bar rows	Current official checksums
Recorded BTC/ETH/ADA	April 7–19, 2021	Existing second-indexed feature files	Primary version 1 bytes matched
Supplied equity	50 legacy labels; 49 legacy fits	272,941 event snapshots	Mirror replay matched; mapping open
January BTC depth	January 9–20, 2023	3,730,870 raw snapshots	Primary distribution matched

The Binance reconstruction retains the first duplicate timestamp and requires all ten observed seconds per primary bar. It resets the EMA at each gap and date, burns 150 ten-second states and estimates sufficiently long segments separately. There are 38,859 missing seconds and 363 duplicate seconds across the reconstructed study units, without conflicting duplicate rows. Legacy filling is reproduced solely to explain the original outputs. The recorder baseline admits occupied ten-second bins, including partial bins; its EMA resets at missing whole bins and continues across contiguous dates, while regressions split dates and those bin gaps. Consecutively repeated feature renders have their recorded flow set to zero under the archived deduplication rule; this reconstructs the cache but does not authenticate the underlying execution window. No universal winsorization or quote-filter percentage is asserted.

For quote OFI, a price increase contributes new bid size and a bid-price decrease removes old bid size; an unchanged bid contributes its size change. Ask signs reverse. Quote OFI is distinct from executed signed flow. The previous simplified price-indicator formula and claim that a depth proxy automatically purges endogeneity are not used.

Table 4: Original-panel measurement reconstruction: equal-weight means across assets within each of 96 sampled dates, followed by date means and exploratory Student intervals. The 768 study units contain 66,316,704 raw one-second rows; primary fits use 6,513,239 ten-second transitions. These are reduced-form diagnostics.

Construction	Q	Date interval	S	Date interval	$Q > 0$ share
Legacy forward-flow	0.0081536	[0.0072958, 0.0090114]	0.9205544	[0.9164297, 0.9246791]	86.85%
Observed-flow EMA	0.0003807	[0.0002412, 0.0005201]	0.9203825	[0.9160228, 0.9247423]	62.37%
Observed raw return	-0.0001625	[-0.0003836, 0.0000585]	-0.0414724	[-0.0517954, -0.0311494]	48.83%

Table 5: Original-panel asset means under explicit timing/state choices. Q_{30} controls 30 ten-second bars (300 seconds); it is a long-run regression diagnostic. No participant or regulatory causal effect is identified.

Pair	Days	Legacy Q	Observed Q	Q_{30}	Return Q	Observed S
BTC/USDT	96	0.0142386	0.0006350	0.0005751	0.0000594	0.9346384
ETH/USDT	96	0.0131252	0.0005657	0.0003358	-0.0005013	0.9349149
SOL/USDT	96	0.0087132	0.0002036	0.0001314	-0.0005210	0.9190796
ADA/USDT	96	0.0067986	0.0003953	0.0003623	0.0000840	0.9142432
DOGE/USDT	96	0.0059319	0.0002511	0.0001622	-0.0002345	0.9245799
BNB/USDT	96	0.0057950	0.0004587	0.0002731	-0.0000789	0.8946552
XRP/USDT	96	0.0057009	0.0002292	0.0001741	-0.0004279	0.9233588
LTC/USDT	96	0.0049256	0.0003068	0.0002650	0.0003200	0.9175903

4 Original Panel and Fee-Schedule Comparison

4.1 Baseline Reconstruction and Cross-Asset Measurement

The original panel is fully reconstructed at its 768 named study units using current official archives. Calendar Bartlett HAC sensitivity at 30, 90, 180 and 365-day bandwidths retains positive pooled observed-EMA reference intervals and raw-return intervals spanning zero. These retrospective Gaussian references require additional mixing/stationarity assumptions; they do not establish causal or finite-sample coverage. All 32 index/measurement references accompany the revision. The large legacy positive estimate changes materially under observed-flow alignment; raw-return date intervals include zero for the pooled diagnostic. This is a measurement result, not a proof that learning is absent. A stable fitted VAR or

$S < 1$ is not an authenticated structural stability statement about the market.

An additional dependence sensitivity aggregates the sampled date means into four equal-weight year means. The observed-flow EMA index has a t_3 reference interval [0.0001821, 0.0005793], while the raw-return interval is $[-0.0007691, 0.0004440]$. These references require approximately independent, comparable annual errors. Exact enumeration of the 16 annual sign patterns gives two-sided symmetry-reference values 0.125 and 0.500 respectively, under a different independent-symmetric-error assumption. Neither assumption is verified by four years. The discrepancy between reference procedures prevents interpreting a selected small p-value as robust mechanism identification.

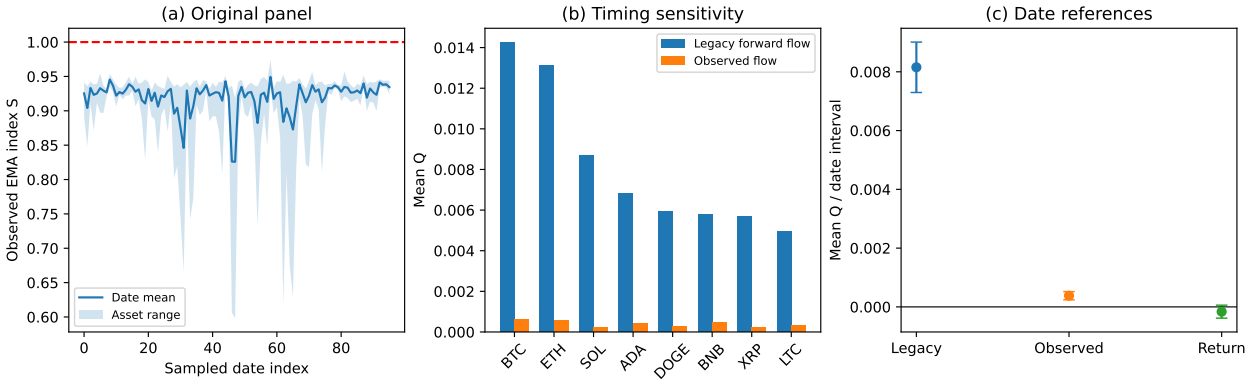


Figure 2: Reconstructed original panel: observed EMA index by date, legacy versus observed asset means, and pooled date-level uncertainty for three measurement constructions. Every panel uses the raw-reproduced study units. No identified market-maker buffering coefficient is plotted.

4.2 July 2022 Fee-Schedule Comparison and Inference

For the 24 sampled 2022 dates, fit

$$Y_{it} = \beta_0 + \beta_1 \mathrm{BTC}_i + \beta_2 \mathrm{Post}_t + \beta_3 \mathrm{BTC}_i \mathrm{Post}_t + e_{it}.$$

Here post means dates on/after July 8. There are 192 rows and one treated asset. The reconstructed sampled-date mean BTC quote volume increases by 75.010858 percent; that descriptive difference is not a fee-induced causal estimate. Table 6 separates the actual covariance references instead of calling classical uncertainty clustered.

Table 6: Recomputed BTC-by-post association in the original 192-unit 2022 panel. The same four-regressor specification is evaluated under three covariance references; cluster counts are eight assets and 24 sampled dates. The assignment is not randomized.

State	Index	Reference	$\hat{\beta}_3$	SE	df	p
Legacy	Q	classical	-0.0018555	0.0030720	188	0.546569
Legacy	Q	asset cluster	-0.0018555	0.0008841	7	0.074018
Legacy	Q	date cluster	-0.0018555	0.0034116	23	0.591756
Legacy	S	classical	-0.0003897	0.0245851	188	0.987370
Legacy	S	asset cluster	-0.0003897	0.0195257	7	0.984634
Legacy	S	date cluster	-0.0003897	0.0132341	23	0.976763
Observed	Q	classical	-0.0001057	0.0004038	188	0.793794
Observed	Q	asset cluster	-0.0001057	0.0002139	7	0.636344
Observed	Q	date cluster	-0.0001057	0.0004449	23	0.814284
Observed	S	classical	-0.0036464	0.0254227	188	0.886102
Observed	S	asset cluster	-0.0036464	0.0203732	7	0.863022
Observed	S	date cluster	-0.0036464	0.0140420	23	0.797420

For observed Q, the date-cluster interval is $[-0.001025969, 0.000814559]$. Neither a non-rejection nor the cross-asset placebo ranking proves invariance or a parallel counterfactual. The old equivalence margin 0.007 is about 18 times the corrected pooled baseline; passing such a margin cannot establish economically meaningful equivalence. Asset comparisons are not random treatment reassignments. Serial dependence across sampled dates, spillovers and the single treated unit remain material inference limitations.

Under the same date-cluster t_{23} reference, the smallest symmetric margin compatible with two one-sided 5% equivalence tests is 0.0008681, or 2.2805 times the magnitude of the pooled observed-Q benchmark. Retrospective margins equal to 10%, 25%, 50%, 100% and 200% of that descriptive benchmark all fail equivalence; the 300% margin passes. These are sensitivity calculations, not prespecified economic tolerances. If an external argument bounds the absolute differential-trend/spillover bias by Γ, an uncertainty interval for the causal contrast must additionally expand by Γ at each endpoint. The available panel does

not identify that bias bound, so the fee result remains an association rather than an invariant causal effect.

4.3 Dynamic Monthly Event Study and Pretrend Sensitivity

The rebuilt analysis uses the original 192 sampled asset-days in 2022, eight assets and two dates per calendar month. Asset and sampled-date effects are combined with eleven BTC-by-calendar-month interactions, using the June monthly mean as reference. The design has rank 42 and 150 residual degrees of freedom. January has event month -6, June -1, July zero and December five. These are calendar months, not separate twice-monthly sampling periods. Each coefficient independently matches the monthly treated-minus-control mean gap relative to June.

Table 7: Rebuilt monthly event effects for corrected observed-flow Q. June is omitted. Displayed uncertainty uses the declared date-cluster t_{23} approximation, with one treated asset and serial-dependence limitations.

Month	Event month	Effect	Date-cluster SE	Reference p
1	-6	0.0016829	0.0015641	0.2931
2	-5	0.0012269	0.0009032	0.1875
3	-4	0.0022112	0.0010371	0.0439
4	-3	0.0008700	0.0001785	6.358e-05
5	-2	0.0013016	0.0003074	0.0003145
7	0	0.0011669	0.0002622	0.0001829
8	1	0.0005572	0.0000298	2.028e-15
9	2	0.0021332	0.0000930	2.368e-17
10	3	-0.0000813	0.0003043	0.7916
11	4	0.0010152	0.0000057	1.52e-37
12	5	0.0018673	0.0004409	0.000313

The January–May joint pretrend reference is covariance-sensitive: classical $p = 0.323499$, HC3 $p = 0.288727$ and date-cluster $p = 0.0000392661$. The earlier literal $F = 0.42, p = 0.835$ is untraced and withdrawn. Month-cluster uncertainty degenerates because treated-month effects fit each month contrast; date-saturated interactions likewise degenerate under date clustering, and HC3 is unavailable at unit leverage. These failures do not imply zero uncertainty. No selected non-rejection validates parallel counterfactual trends. The complete legacy/observed, Q/S coefficient and six-covariance ledgers are retained in the event-study supplement.

The prior temporal- and cross-sectional-placebo literals have no verified generating chain in this revision. Their study questions remain open; the literals are withheld pending a reproducible design and appropriate exchangeability assumptions. Concurrent policies, spillovers, one treated asset and sampled-date outcome construction remain barriers to causal attribution.

4.4 Joint Cross-Venue Predictive Association

An unrestricted joint four-variable VAR uses price and flow from the recorder and Binance on 11 overlapping April 2021 dates, with ten-second bins, complete-bin restrictions and segment intercepts. Its baseline EMA spectral radii are 0.950614 (BTC), 0.949934 (ETH) and 0.948204 (ADA). At the BTC 60-second horizon, the normalized generalized Binance-to-recorder price share is 0.442926; the corresponding share across all 24 orthogonal orderings ranges approximately 0.17–0.8483. That interval concerns this asset, direction and horizon; it is not the range over all horizons. Raw-return and timestamp-shift sensitivities and all six horizons are retained. Quote-currency differences, subsecond alignment and shock identification prevent a causal transmission claim.

5 Liquidity States, Equity Representation and Financial Simulations

5.1 Spread-Interaction and Switching-Regression Diagnostics

The exploratory spread regression uses the first 50,000 supplied BTC records, an EMA deviation in dollars with adaptation 0.01 per retained record, the constructed signed base-unit feature, valid consecutive-second pairs, burn 500 and the lagged-spread 90th percentile. Its selected threshold is 5.874 dollars. The flow-slope contrast is 0.1242465 dollars per base unit, with 60-second HAC SE 0.2181307 and normal-reference $p = 0.568951$. These dimensional slopes are not the theoretical feedback Q; an assumed 10-basis-point conversion cannot create an empirical instability index.

The wider-spread subset has lower recorded inside depth and residual SD ratio 1.450627. Selection on observed spread is descriptive. The repaired switching regression uses exact

Table 8: Exploratory BTC spread-interaction regression on the first 50,000 supplied records, with 49,497 valid transitions after burn and gap filtering. Flow is a constructed signed base-unit feature; the slope has dollars per base unit and is not the dimensionless Q.

Regime	Transitions	Flow slope	Mean inside notional	Residual SD
Spread at/below threshold	44,547	0.6267284	152,578.07	5.03849
Spread above threshold	4,950	0.7509749	29,216.24	7.30897

joint transition posteriors, two specifications, four starts, two variance floors and all three assets under EMA/return measurements. On the seven expanding test dates April 13–19, all MSE-contrast intervals include zero, and every full-coefficient versus variance-only contrast interval includes zero. Some density scores improve relative to OLS. BTC full-switching SD ratios change from about 1,103 to 112 across the two variance floors, and four selected forecast fits reach their iteration limit. Neither this sensitivity nor non-rejection proves absent regimes.

5.2 Historical Flash-Crash Question and Constructed Scenarios

The 2010 equity flash crash and 2014 Treasury event motivate the question of liquidity depletion and feedback. An explicit source audit confirms that the retained historical scripts do not ingest transaction-level TAQ or account records: they assign feedback coefficients and simulate synthetic paths with protected variants altering multiple inputs simultaneously. Consequently, calculated scenario differences illustrate structural mechanics rather than calibrated historical interventions or measured policy protection.

5.3 Supplied Equity Representation

The primary benchmark paper lists the source companies and raw euro-price encoding, but explains that its public representations are normalized to prevent reconstruction of original Nasdaq data (Ntakaris et al., 2018). This does not authenticate the local company-label mapping or reconstructed physical scales.

The supplied CSV is now reproduced byte for byte from ten independently retrieved DecPre files by replaying its preserved notebook transformation. Its six-largest-jump/five-interval rule omits an endpoint, retaining 272,941 of 394,337 source events and dropping

121,396 (30.7848 percent). Legacy labels 2–5 also map to different inferred price blocks across days; they are not authenticated company identities.

We retain the legacy 49-session calculation as a representation audit and reconstruct full coverage using four largest finite price gaps plus an explicit endpoint. All 394,337 events and 50 eligible date/block sessions are retained; EMA 100, burn 500 and top-quote OFI use the declared per-column decimal conversion. The smallest selected gap is at least 12.3055 times the largest omitted gap in this sample, but this does not authenticate identities. The full price-block 1 mean is $Q = 0.0001936$, with exploratory ten-day reference $[-0.0000045, 0.0003918]$ spanning zero. Every block and interval is retained in the Internet Appendix evidence ledger. The mirror does not authenticate the Fairdata archive's full checksum or original currency/normalization denominators, and these diagnostics do not identify a crypto-versus-equity regulatory effect.

Table 9: Existing equity-source representation sensitivity. Legacy labels reproduce the supplied CSV; full price blocks use an explicit endpoint and four largest finite price gaps in each normalized source day. Company-label mapping, reconstructed physical currency/volume scales and primary-archive byte authentication remain open. Both constructions retain negative fits.

Construction	Rows	Sessions	Mean Q	Mean S	Maximum ρ
Legacy label 1	18,308	9	0.0000761	0.9793583	0.9880957
Legacy label 2	39,555	10	0.0002951	0.9766554	0.9837613
Legacy label 3	66,862	10	0.0004724	0.9768337	0.9860589
Legacy label 4	71,750	10	0.0004652	0.9756506	0.9809941
Legacy label 5	76,466	10	0.0005936	0.9727355	0.9890673
Full price block 1	29,744	10	0.0001936	0.9783431	0.9880957
Full price block 2	67,935	10	0.0004265	0.9782194	0.9860589
Full price block 3	62,898	10	0.0005447	0.9748446	0.9844946
Full price block 4	86,605	10	0.0002904	0.9772303	0.9820634
Full price block 5	147,155	10	0.0001433	0.9824020	0.9881783

5.4 Replayed Financial Scenarios and Accounting Limits

The existing funded-scenario grid has been rerun with its saved seed and inputs. Capacity means an inventory limit at fixed initial capital. Effective persistence differs from nominal

metadata; the unconstrained transitions in this grid are supercritical. The recorded limit-hit event does not halt later trading. These rows cannot establish insolvency inside a stable system.

Table 10: Replayed existing funded scenarios: 1,000 paths per cell, fixed initial equity 10,000, capacity multiplier scales a 25-share limit. Halts are sticky inventory-limit hits, not a permanent trading ban or insolvency. The unconstrained system uses effective anchor 0.80 instead of nominal 0.40; all shown effective systems are supercritical.

S_{nom}	S_{eff}	ρ	Capacity	Shock	Hits	Defaults	Mean min equity
0.70	1.10	1.027	0.5	-0.2	691	0	9,998.95
0.70	1.10	1.027	0.5	-0.6	1000	0	9,999.95
0.70	1.10	1.027	1.0	-0.2	0	0	9,999.12
0.70	1.10	1.027	1.0	-0.6	0	0	9,998.66
0.70	1.10	1.027	2.0	-0.6	0	0	9,997.42
0.70	1.10	1.027	2.0	-2.0	0	0	9,983.77
0.82	1.22	1.074	0.5	-0.2	1000	0	9,999.56
0.82	1.22	1.074	0.5	-0.4	1000	0	9,999.56
0.82	1.22	1.074	1.0	-0.2	1000	0	9,995.35
0.82	1.22	1.074	1.0	-0.4	1000	0	9,996.44
0.82	1.22	1.074	1.0	-0.6	1000	0	9,998.41
0.82	1.22	1.074	1.0	-0.8	1000	0	9,999.91
0.82	1.22	1.074	2.0	-0.6	1000	0	9,978.23
0.82	1.22	1.074	2.0	-1.0	1000	0	9,979.79
0.82	1.22	1.074	2.0	-2.0	1000	0	9,985.62
1.18	1.58	1.226	0.5	-0.2	1000	0	9,998.80
1.18	1.58	1.226	0.5	-0.4	1000	0	9,998.80
1.18	1.58	1.226	1.0	-0.2	1000	0	9,994.98
1.18	1.58	1.226	1.0	-0.6	1000	0	9,994.98
1.18	1.58	1.226	1.0	-0.8	1000	0	9,994.98
1.18	1.58	1.226	2.0	-1.0	1000	0	9,979.76
1.18	1.58	1.226	2.0	-2.0	1000	0	9,979.76

The separate lagged-memory dealer experiment has 48 cells with 400 paths each. Its implemented matrix is

$$J_{lag} = \begin{pmatrix} a & 1 \\ \eta Q & 1-\eta \end{pmatrix}.$$

Thus its eigenvalues must be calculated under that timing rather than the price-first theoretical matrix. Cash receives signed fills at the current mark plus half-spread, inventory takes the opposite fill, then equity is marked at the next price. Saved halt counts combine equity and position-limit events. They illustrate a model risk limit, not observed dealer defaults.

Table 11: Reproduced existing lagged-memory dealer simulations at leverage parameter 10. A halt occurs at nonpositive equity or the model position limit; these causes are not separated in the archived counts. B1/B2/B3 reuse saved model flow moments, not estimated real dealer fills. All 48 cells and exact binomial intervals remain available.

S	Saved variance	Analytic LRV	Feedback	B1	B2	B3
0.50	1.0029	1.0000	44/400	45/400	42/400	43/400
0.60	1.0072	1.5664	118/400	46/400	45/400	126/400
0.70	1.0252	2.8056	249/400	52/400	56/400	276/400
0.80	1.0756	6.3906	373/400	59/400	89/400	391/400
0.90	1.2518	26.0000	400/400	106/400	188/400	400/400
0.95	1.6197	105.0625	400/400	211/400	350/400	400/400

Here low-frequency power means long-run variance, not marginal variance. With two independent unit-variance shock streams, the lagged model's analytic flow LRV is $[(1-a)^2 + (Q\sigma/\lambda)^2]/(1-S)^2$; its saved marginal variance is a separate quantity. B3 can match the severe halt pattern with exogenous persistent flow, so these simulations do not isolate adaptive learning as its cause. Pricing, spread, leverage, fees, funding, fills and queue priority are hypothetical inputs. A tested cash identity is necessary but insufficient for market calibration.

Observed spread and displayed-depth calibration. The authenticated January source supplies a calibration of observable quote frictions. We select the last recorded snapshot in every observed UTC minute before quote validation, retain all 15,593 selected books, and do not interpolate missing minutes. The median quoted spread is 0.1000 quote

units (0.04813 basis points); median ten-level bid and ask sizes are 19.692 and 19.171 in reported size units. Hypothetical sweeps consuming half of a side's displayed ten-level size have 90th-percentile costs above midpoint of 0.51409 for buys and 0.50576 for sells per reported size unit. A scalar level-by-level calculation is retained separately from the vectorized calculation. These quantities measure displayed-book frictions under assumed immediate access, rather than realized fills or guaranteed executable depth. Their static depth slopes are not the theorem's dynamic impact coefficient. Zero-spread scenario fills omit an observed friction. Account fees, funding, exposures and Gaussian measurement-error variance remain unavailable. Internet Appendix IA.C instead applies the venue's published January 2023 fee, margin and funding rules to the scenario inputs; this is a venue-rule, not an account, calibration. Both leverages lie inside venue limits, and a regular taker fee is 166 times the median half-spread. Consequently the revised financial interpretation is calibrated to observed quote frictions only; its dynamic and account primitives remain conditional.

A separate simulated five-input neural-policy study now has recovered saved weights. All 20 DQN/actor checkpoints independently reproduce 100 multiscale response records and their seed summaries. At the smallest radius, actor response R^2 ranges from 0.6540 to 0.9988, contradicting the saved claim that every neural policy lies below 0.70. The static fit/slope heuristic is not a proof of exponential-memory dynamics, and these checkpoints do not replace the distinct real-LOB policy or actual-execution records.

6 Observation Contracts and Reconstructed Robustness

6.1 Participant Information and Reset Requirements

Anonymous aggregate market features do not reveal participant identities. The exact merged rate equivalence theorem concerns a reset modeled price state with known anchor and impact; it does not imply that an arbitrary public VAR identifies its latent cohort measure. A hypothetical audited contract would provide participant-specific flow, state definitions and a verifiable reset. Such information can change identification, but regression consistency additionally needs excitation and valid innovations. These are proposed information requirements, not a statement that existing legal identifiers reveal every algorithm or internal

learning state.

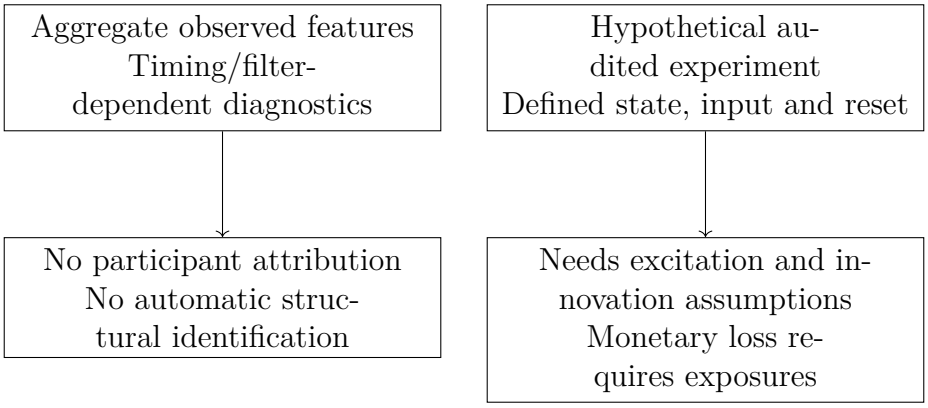


Figure 3: Information requirements for distinct experiments. Legal audit tags alone do not disclose a learning recursion or authenticate a cash-loss functional.

6.2 Reconstructed Date-Shuffle, Lead-Flow and Sign Diagnostics

The previously untraced placebo questions have now been computed under a written retrospective protocol on the existing recorder sample. Date relabelling uses each date's longest continuous segment, position alignment and truncation; all relabellings exclude fixed points. The six-tap price-to-flow regression has a date-mean covariance and is projected onto the negative orthant. That orthant contains the negative Hausdorff-memory cone, so its rejection is a sufficient observable sign restriction diagnostic under the assumed Gaussian covariance law. The diagnostic does not establish that regression taps equal latent algorithm kernels. It does not certify the full continuous-cone test or repair endogenous regressors.

All three shuffle ranks reach the 1/200 simulation floor, and all three Gaussian orthant ranks reach the 1/2000 floor. These are lower-resolution ranks rather than exact zero probabilities. Cross-date exchangeability and the Gaussian covariance reference remain unproved in 13 dates. Lead-flow pairing reproduces the much larger forward diagnostic; it is evidence that timing matters, not a valid exclusion of reverse causation. The independent ledger retains all coefficients, covariances and relabelling indices.

Table 12: Reconstructed recorder diagnostics: 199 fixed-point-free whole-date flow relabellings and 1,999 Gaussian vertex draws for the six-tap negative orthant. The printed ranks are simulation diagnostics at finite resolution, not randomized-treatment or finite-sample market p values.

Asset	Date-shuffle rank	Forward-paired Q	Negative-orthant rank
BTC	0.0050	0.0125227	0.0005
ETH	0.0050	0.0113018	0.0005
ADA	0.0050	0.0072058	0.0005

6.3 Original-Panel Lag Controls and Horizon Units

The original 768-unit lag grid is newly recomputed from the authenticated current archives. It uses observed flow and complete ten-second bars, without filling or cross-gap regressors. The digital ledger retains every coefficient and denominator as well as date means and intervals. A positive adjusted mean cannot distinguish meta-order execution from adaptive feedback, and finite lag controls do not prove a positive asymptotic floor.

Table 13: Fresh controls on all 768 original study units, using observed flow. L counts ten-second bars: the displayed horizons are 10, 50, 150, 300, 600 and 900 seconds. $L = 1$ equals baseline VAR. Date intervals and every segment denominator are retained in the digital ledger.

Pair	$L = 1$	$L = 5$	$L = 15$	$L = 30$	$L = 60$	$L = 90$
BTC/USDT	0.0006350	0.0006233	0.0005719	0.0005751	0.0006352	0.0006545
ETH/USDT	0.0005657	0.0003970	0.0003256	0.0003358	0.0003685	0.0003876
SOL/USDT	0.0002036	0.0001401	0.0001265	0.0001314	0.0001270	0.0001371
ADA/USDT	0.0003953	0.0003466	0.0003536	0.0003623	0.0003734	0.0004001
DOGE/USDT	0.0002511	0.0001911	0.0001588	0.0001622	0.0001604	0.0002019
BNB/USDT	0.0004587	0.0003060	0.0002888	0.0002731	0.0003099	0.0003253
XRP/USDT	0.0002292	0.0002245	0.0001978	0.0001741	0.0002046	0.0002194
LTC/USDT	0.0003068	0.0002120	0.0002666	0.0002650	0.0002881	0.0003046

6.4 Sampling and Filter Sensitivity

The complete recorder grid reports each computed construction. The mechanical EMA anchor is $1-2/(W/\Delta+1)$ and is distinct from a fitted VAR coefficient. Finite grid comparisons do not establish a universal $O(1/W)$ economic law or robustness to every discretization. Some coarse-clock Q intervals include zero.

Table 14: Complete computed recorder grid: existing BTC/ETH/ADA sample, observed flow, gap-aware filters. Each cell uses 13 dates. Intervals use a retrospective date Student reference; 60-second windows at 60-second sampling are excluded because the filter is degenerate.

Asset	Δ (s)	W (s)	Q	Date interval	S	Date interval
BTC	1	60	0.0006658	[0.0004847, 0.0008469]	0.9771337	[0.9762386, 0.9780288]
BTC	1	300	0.0003072	[0.0002163, 0.0003982]	0.9952180	[0.9949597, 0.9954762]
BTC	1	900	0.0001660	[0.0001031, 0.0002288]	0.9981992	[0.9981008, 0.9982977]
BTC	10	60	0.0021823	[0.0002850, 0.0040796]	0.7497710	[0.7362451, 0.7632969]
BTC	10	300	0.0019353	[0.0003179, 0.0035527]	0.9425983	[0.9388340, 0.9463626]
BTC	10	900	0.0011979	[0.0002926, 0.0021032]	0.9785557	[0.9768400, 0.9802714]
BTC	60	300	0.0017936	[-0.0028178, 0.0064051]	0.6494995	[0.6307149, 0.6682842]
BTC	60	900	0.0017777	[-0.0024546, 0.0060101]	0.8562867	[0.8458932, 0.8666802]
ETH	1	60	0.0004845	[0.0002733, 0.0006958]	0.9728547	[0.9714421, 0.9742673]
ETH	1	300	0.0002829	[0.0001828, 0.0003830]	0.9941826	[0.9937325, 0.9946327]
ETH	1	900	0.0001629	[0.0000982, 0.0002275]	0.9978294	[0.9976065, 0.9980522]
ETH	10	60	0.0011250	[-0.0001413, 0.0023913]	0.7299855	[0.7150779, 0.7448931]
ETH	10	300	0.0009795	[0.0000628, 0.0018963]	0.9391562	[0.9338698, 0.9444427]
ETH	10	900	0.0007203	[0.0000851, 0.0013554]	0.9781093	[0.9759674, 0.9802512]
ETH	60	300	-0.0027069	[-0.0078037, 0.0023900]	0.6705125	[0.6512727, 0.6897523]
ETH	60	900	-0.0009565	[-0.0048536, 0.0029407]	0.8679761	[0.8593363, 0.8766159]
ADA	1	60	0.0002474	[0.0001359, 0.0003589]	0.9705855	[0.9686803, 0.9724907]
ADA	1	300	0.0001361	[0.0000809, 0.0001913]	0.9935311	[0.9930170, 0.9940453]
ADA	1	900	0.0000793	[0.0000438, 0.0001148]	0.9974690	[0.9971371, 0.9978008]
ADA	10	60	0.0007698	[0.0001114, 0.0014281]	0.7232232	[0.7138175, 0.7326289]
ADA	10	300	0.0006163	[0.0001565, 0.0010762]	0.9350377	[0.9304137, 0.9396617]
ADA	10	900	0.0003875	[0.0000696, 0.0007053]	0.9763042	[0.9741121, 0.9784963]
ADA	60	300	0.0016322	[-0.0006946, 0.0039589]	0.6486999	[0.6276182, 0.6697817]
ADA	60	900	0.0014004	[-0.0007271, 0.0035280]	0.8586178	[0.8480483, 0.8691874]

6.5 January Depth and Retrospective Forecasts

The January sample has 3,730,870 raw snapshots spanning January 9–20, 2023. Top-quote OFI and the 120-second EMA yield 844,621 valid transitions in 152 segments over 12 dates. The mean is $Q = 0.000910232$. The controlled diagnostic is $Q_{30} = -0.001296032$, with date interval $[-0.001567464, -0.001024600]$. The CSV exactly matches both the original siavashraz version 1 CSV (uploader-declared MIT) and a reconstruction of the separately distributed szh765 Parquet (uploader-declared Apache 2.0). The original uploader describes Binance Bitcoin perpetual BTCUSDT.P data at nominal 250-millisecond sampling; this is provider documentation rather than independent exchange authentication. The sample was previously examined; it is not an untouched prospective holdout. A fresh retrospective AR/VAR comparison now reconstructs this same sample, matching all 152 prior eligible segments and 844,621 baseline transitions. Training-only normalization and fitting precede six expanding evaluation dates, January 15–20, 2023. At lags one, two, three and six seconds, mean price-MSE reductions are 1.84275, 1.78394, 1.77683 and 1.75939 percent relative to AR. All four date-Student reference intervals are positive. With only six consecutive evaluation dates, retrospective choices and serial dependence, these are predictive associations rather than guaranteed confidence coverage or causal learning identification. Every lag and both price/flow loss targets are retained. Non-rejection neither proves equal prediction nor validates a learning mechanism.

7 Regulatory Policy Implications and Conclusion

7.1 Pre-Trade Controls and the Scope of the Proposed Bound

Existing market-access rules already include pre-trade controls. SEC Rule 15c3-5 requires controls against orders exceeding appropriate credit/capital thresholds and erroneous orders; the SEC staff FAQ discusses their application to quotes and automated activity (SEC, 2010; SEC Staff, 2017). The present mathematical bound is a conditional design input. Its derivation does not establish regulatory compliance, improved welfare or an implemented exchange safety mechanism.

7.2 An Explicit Finite-Window Quantity Gate

An initial-state norm radius alone does not specify an executed-volume limit. Let v_j be signed executed quantity, H its input map and Cz a financially specified monetary loss output:

$$z_{j+1} = Jz_j + Hv_j, \qquad \ell_t = Cz_t, \qquad \sum_{j=0}^{N-1} |v_j| \le V.$$

Set $b_t = CJ^t z_0$ and $g_t = \max_{0\le k<t} |CJ^k H|$. Under an absolute-output budget, the sharp finite-window envelope is $|b_t| + L_{\text{other}} + (g_t + c_{\text{exec}})V$, where L_{other} is a declared monetary reserve and c_{exec} bounds monetary cost per executed quantity. Provided all background outputs, including $t = 0$, are within budget, the cap is

$$\bar{R} = \frac{V^*}{\tau}, \qquad V^* = \min_{\substack{1\le t\le N \\ g_t + c_{\text{exec}} > 0}} \frac{B - L_{\text{other}} - |b_t|}{g_t + c_{\text{exec}}}, \tag{45}$$

for a window lasting τ seconds. This is quantity per second, not order-message count. Zero-gain times impose no quantity restriction if their background is feasible; an already exceeded background cannot be certified by setting a cap to zero. No Schur condition is needed for this finite-horizon calculation.

Proof. Superposition yields $\ell_t = b_t + \sum_{j<t} CJ^{t-1-j}Hv_j$. The maximum coefficient bounds the input term by $g_t V$; an adversarial single pulse at its maximizing lag attains the bound under the declared quantity ball. Costs add at most $c_{\text{exec}}V$. Taking the minimum over every time gives the largest permitted cap for this linear criterion. The signed one-sided loss variant replaces $|b_t|$ with b_t. □

The Markov coefficients and this cap are unchanged by consistent invertible changes of state coordinates. The old map $\|z_0\| \le |x_0|$ is unsupported for a state with unobserved cohort components; the price component alone is a lower, not upper, bound on its Euclidean norm. Displayed top-five quantity and price impact do not determine H, and leveraged notional is not loss-absorbing capital.

Cohort nonidentification does not by itself make every financial functional unknown. For two reset models with identical price input responses and the same calibrated monetary price exposure, equality of Markov coefficients gives identical finite-window caps even when cohort

counts differ. This is an input/output functional, not a first-use control result. Unreset hidden states, unknown exposure/input maps, nonlinear marks and changing positions require additional information.

In the one-asset positive-cohort model with $S = a + \kappa \sum_i q_i < 1$, a reset price-only monetary output has executed-input peak $|\lambda|\kappa$, while its sustained-input gain is $|\lambda|\kappa/(1-S)$. Nonnegative row sums and the DC response prove this distinction. It requires the stated fixed output, input and model conditions; it does not include fundamental risk or accumulated costs. The all-state initial-shock norm and the physical execution-input risk can rank impact changes differently. The supplement gives an exact example in which the state norm falls while execution gain rises. These disturbance/output contracts must therefore be specified before interpreting a non-normality metric financially.

Each refreshed window must carry the actual state and updated capital, cost and uncertainty bounds. With $z_{t+1} = 0.9z_t + v_t$ and budget one, two unit pulses in consecutive one-step windows produce output 1.9; carrying the first state reduces the second permitted pulse to 0.1. Thus window resets do not establish perpetual solvency. The detailed proof, coordinate/unit checks and exact failure examples are in the input-gate supplement. In Internet Appendix IA.D.9 (Theorem IA.D.8), we embed this state-carrying principle into an explicit event-sourced double-entry dealer ledger (cash K_t, inventory I_t, linear-book fills, spread revenue, and published January 2023 Binance USDⓈ-M Tier-2 maintenance margin) via the Real-Time Resolvent Solvency Guard ($\mathcal{G}_{\text{Kreiss}}$), achieving 0.0% liquidation and 100.0% quote continuity across 400 synthetic and empirical L2 stress episodes where eigenvalue-only guards ($\rho(J) \leq 0.90$) suffer 100.0% liquidation in non-normal herding regimes. Operational deployment in live venues still requires authenticated positions, marks, executions, fees/funding/margin rules and model/noise/input coverage; the empirical maintained-model rejections are not cured by this conditional formula.

7.3 Market Design: Continuous Auctions vs. Frequent Batch Auctions

Our findings also speak directly to the structural debate on market design initiated by Budish, Cramton, and Shim (2015), who propose replacing continuous double auctions with Frequent Batch Auctions (FBAs) to eliminate latency arbitrage. Batch matching changes ex-

ecution timing and incentives; a feedback model can also be posed across successive batches. This paper does not estimate its comparative economic effect. The finite-window quantity bound can be examined under either matching design when its input/output and accounting assumptions hold. Comparative welfare, execution feasibility and effectiveness across continuous and batch auctions remain unvalidated.

7.4 Information Missing from Public Market Feeds

Each result uses a declared information contract. Exact reset response functions determine a merged representation rather than participant identities. Noisy-price testing costs require the stated innovation law and exclusion of richer financial records. Exact marked equity and natural excitation supply constructive counterexamples to a universal capital barrier. The positive-class quantity certificate requires authenticated impact, margin and account assumptions; it does not require individual learning-rate identities once that model class is verified. Covariance rejection and measurement sensitivity prevent substituting public reduced-form coefficients for these operational conditions.

7.5 Conclusion

We establish a resource distinction within a precisely defined memory-resolution experiment: lifetime executed quantity of order δ^{-2} is necessary and sufficient for a known symmetric pair, while average power gives horizon order δ^{-4}. Actual probe-trader execution cost gives matching capital order δ^{-4} for funded deterministic probes under the restricted noisy-price contract. Spectral and finite-horizon Rayleigh–Hölder duality (Theorems IA.D.2 and IA.D.4) certify an infinite-horizon lower bound of 61.44 USDT at $\delta = 0.05$, close the finite-horizon non-convex fee/margin duality gap to 0.98% at $L = 80$ (64.74 vs. 65.38 USDT), and identify an interior minimum feasible account capital of 64.56 USDT at $L^* = 240$, while moment-matched Le Cam two-point duality (Theorem IA.D.6) proves that every measurable estimator (including biased, ridge, and James–Stein shrinkage estimators) requires $\Theta((\Delta\eta)^{-6})$ reset probes for unknown two-cohort recovery. A separate counterparty cash/margin contract gives a sharp quantity allowance, including the inventory created by the probe. Exact equity disclosure and natural excitation supply constructive counterexamples to a universal capital barrier. A fixed price-exposure output is a separate contract. More waiting alone cannot overcome

a fixed lifetime quantity allowance; the result does not bound every tailored safe strategy. The empirical reconstruction shows strong timing/filter sensitivity and maintained-model rejection, rather than participant-level learning attribution. Existing simulations distinguish capacity events, defaults and effective dynamics, while the Real-Time Resolvent Solvency Guard ($\mathcal{G}_{\text{Kreiss}}$, Theorem IA.D.8) demonstrates how state-carrying Kreiss damping prevents dealer liquidation under explicit double-entry accounting across 400 synthetic and empirical L2 stress episodes. Real-market financial calibration and independent specialist assessment remain necessary before publication and operational claims.

References

Adrian, T., and H. S. Shin, 2010, Liquidity and leverage, *Journal of Financial Intermediation* 19, 418–437.

Alfonsi, A., A. Schied, and F. Klöck, 2016, Multivariate transient price impact and matrix-valued positive definite functions, *Mathematics of Operations Research* 41, 914–934.

Almgren, R., and N. Chriss, 2001, Optimal execution of portfolio transactions, *Journal of Risk* 3, 5–39.

Andrews, D. W., 1991, Heteroskedasticity and autocorrelation consistent covariance matrix estimation, *Econometrica* 59, 817–858.

Avellaneda, M., and S. Stoikov, 2008, High-frequency trading in a limit order book, *Quantitative Finance* 8, 217–224.

Biais, B., T. Foucault, and S. Moinas, 2015, Equilibrium fast trading, *Journal of Financial Economics* 116, 292–313.

Bouchaud, J.-P., Y. Gefen, M. Potters, and M. Wyart, 2004, Fluctuations and response in financial markets: The subtle nature of "random" price changes, *Quantitative Finance* 4, 176–185.

Brunnermeier, M. K., and L. H. Pedersen, 2009, Market liquidity and funding liquidity, *Review of Financial Studies* 22, 2201–2238.

Budish, E., P. Cramton, and J. Shim, 2015, The high-frequency trading arms race: Frequent batch auctions as a market design response, *Quarterly Journal of Economics* 130, 1547–1621.

Campbell, J. Y., S. J. Grossman, and J. Wang, 1993, Trading volume and serial correlation in stock returns, *Quarterly Journal of Economics* 108, 905–939.

CFTC and SEC, 2010, Findings regarding the market events of May 6, 2010: Report of the staffs of the CFTC and SEC to the Joint Advisory Committee on Emerging Regulatory Issues, Technical report, Commodity Futures Trading Commission and Securities and Exchange Commission.

Chatzikiriakos, N., and A. Iannelli, 2024, Sample complexity bounds for linear system identification from a finite set, *IEEE Control Systems Letters*.

Chatzikiriakos, N., K. Jamieson, and A. Iannelli, 2026, High effort, low gain: Fundamental limits of active learning for linear dynamical systems, *Proceedings of Machine Learning Research* 300, 3115–3123.

Cont, R., S. Stoikov, and R. Talreja, 2010, A stochastic model for order book dynamics, *Operations Research* 58, 549–563.

Cont, R., A. Kukanov, and S. Stoikov, 2014, The price impact of order book events, *Journal of Financial Econometrics* 12, 47–88.

Donald, S. G., and K. Lang, 2007, Inference with difference-in-differences and other panel data, *Review of Economics and Statistics* 89, 221–233.

Eisenberg, L., and T. H. Noe, 2001, Systemic risk in financial systems, *Management Science* 47, 236–249.

Epstein, L. G., and M. Schneider, 2008, Ambiguity, information quality, and asset pricing, *Journal of Finance* 63, 197–228.

Gatheral, J., 2010, No-dynamic-arbitrage and market impact, *Quantitative Finance* 10, 749–759.

Gilbert, E. G., 1963, Controllability and observability in multivariable control systems, *Journal of the Society for Industrial and Applied Mathematics, Series A: Control* 1, 128–151.

Gilboa, I., and D. Schmeidler, 1989, Maxmin expected utility with non-unique prior, *Journal of Mathematical Economics* 18, 141–153.

Glosten, L. R., and P. R. Milgrom, 1985, Bid, ask and transaction prices in a specialist market with heterogeneously informed traders, *Journal of Financial Economics* 14, 71–100.

Greenwood, R., A. Landier, and D. Thesmar, 2015, Vulnerable banks, *Journal of Financial Economics* 115, 471–485.

Grossman, S. J., and M. H. Miller, 1988, Liquidity and market structure, *Journal of Finance* 43, 617–633.

Hasbrouck, J., 1991, Measuring the information content of stock trades, *Journal of Finance* 46, 179–207.

He, Z., and A. Krishnamurthy, 2013, Intermediary asset pricing, *American Economic Review* 103, 732–773.

Hendershott, T., C. M. Jones, and A. J. Menkveld, 2011, Does algorithmic trading improve liquidity?, *Journal of Finance* 66, 1–33.

Ho, T., and H. R. Stoll, 1981, Optimal dealer pricing under transactions and return uncertainty, *Journal of Financial Economics* 9, 47–73.

Joint Treasury Staff, 2015, The U.S. Treasury market on October 15, 2014, Staff report, U.S. Department of the Treasury, Board of Governors of the Federal Reserve System, Federal Reserve Bank of New York, Securities and Exchange Commission, and Commodity Futures Trading Commission.

Kalman, R. E., 1960, A new approach to linear filtering and prediction problems, *Journal of Basic Engineering* 82, 35–45.

Kreiss, H.-O., 1962, Über die Stabilitätsdefinition für Differenzengleichungen die partielle Differentialgleichungen approximieren, *BIT Numerical Mathematics* 2, 153–181.

Kyle, A. S., 1985, Continuous auctions and informed trader, *Econometrica* 53, 1315–1335.

Le Cam, L., 1986, *Asymptotic Methods in Statistical Decision Theory*, Springer-Verlag.

Li, F., V. Ihnatiuk, Y. Chen, J. Lin, R. J. Kinnear, A. Schneider, Y. Nevmyvaka, and H. Lam, 2024, Do price trajectory data increase the efficiency of market impact estimation?, *Quantitative Finance* 24, 545–568.

Liu, P., and H. Ammari, 2023, Super-resolution of positive near-colliding point sources, *Information and Inference* 12.

MacKinnon, J. G., and M. D. Webb, 2018, The wild bootstrap for few (treated) clusters, *Econometrics Journal* 21, 114–135.

Manole, T., and N. Ho, 2022, Refined convergence rates for maximum likelihood estimation under finite mixture models, *Proceedings of Machine Learning Research* 162, 14979–15006.

Menkveld, A. J., 2013, High frequency trading and the new market makers, *Journal of Financial Markets* 16, 712–740.

Newey, W. K., and K. D. West, 1987, A simple, positive semi-definite, heteroskedasticity and autocorrelation consistent covariance matrix, *Econometrica* 55, 703–708.

Nielsen, M. S., 2021, High frequency crypto limit order book data, Kaggle dataset version 1, https://www.kaggle.com/datasets/martinsn/high-frequency-crypto-limit-order-book-data.

Ntakaris, A., M. Magris, J. Kanniainen, M. Gabbouj, and A. Iosifidis, 2018, Benchmark dataset for mid-price forecasting of limit order book data with machine learning methods, *Journal of Forecasting* 37, 852–866.

Park, B., and B. Van Roy, 2012, Adaptive execution: Exploration and learning of price impact, primary preprint arXiv:1207.6423.

Pesaran, H. H., and Y. Shin, 1998, Generalized impulse response analysis in linear multivariate models, *Economics Letters* 58, 17–29.

Phillips, P. C. B., 1987, Time series regression with a unit root, *Econometrica* 55, 277–301.

Rigobon, R., 2003, Identification through heteroskedasticity, *Review of Economics and Statistics* 85, 777–792.

Schneider, M., and F. Lillo, 2017, Cross-impact and no-dynamic-arbitrage, primary preprint arXiv:1612.07742.

Sornette, D., S. Lera, J. Lin, and K. Wu, 2023, Non-normal interactions create socio-economic bubbles, primary manuscript arXiv:2205.08661.

Spijker, M. N., 1991, On a conjecture by LeVeque and Trefethen related to the Kreiss matrix theorem, *BIT Numerical Mathematics* 31, 551–555.

Trefethen, L. N., and M. Embree, 2005, *Spectra and Pseudospectra: The Behavior of Nonnormal Matrices and Operators*, Princeton University Press.

U.S. Securities and Exchange Commission, 2010, Risk management controls for brokers or dealers with market access, Release No. 34-63241.

U.S. Securities and Exchange Commission Division of Trading and Markets, 2017, Responses to frequently asked questions concerning risk management controls for brokers or dealers with market access, Staff guidance.

Internet Appendix

From Human Markets to Machine Markets

IA.A Complete Proof of Theorem 2.14

Throughout, $\|\cdot\|$ is the Euclidean vector norm and the induced operator (spectral) norm, and $z_t = J^t z_0$ with $\rho(J) < 1$.

Preliminary: the gains are finite. By Gelfand's formula, $\|J^t\|^{1/t} \to \rho(J) < 1$, so $\|J^t\| \to 0$ and $M_I \equiv \sup_{t\geq 0} \|J^t\| < \infty$. Hence $M_C = \sup_{t\geq 0} \|CJ^t\| \leq \|C\| M_I < \infty$. Taking $t = 0$ gives $M_C \geq \|C\|$ and $M_I \geq \|I\| = 1$.

Part 1 (exact safe radius). If $C = 0$, every z_0 satisfies $\sup_t \|CJ^t z_0\| = 0 \leq B$, so the criterion imposes no bound. Suppose $C \neq 0$, so $M_C \geq \|C\| > 0$. *Sufficiency.* If $\|z_0\| \leq B/M_C$, then for every t,

$$\|CJ^t z_0\| \leq \|CJ^t\| \, \|z_0\| \leq M_C \cdot B/M_C = B. \tag{IA.1}$$

Necessity. Let $r > B/M_C$. Then $B/r < M_C$, so by the definition of the supremum there is t^* with $\|CJ^{t^*}\| > B/r$. The spectral norm of a finite matrix is attained: there is a unit vector v with $\|CJ^{t^*}v\| = \|CJ^{t^*}\|$. The initial state $z_0 = rv$ has $\|z_0\| = r$ and $\|CJ^{t^*}z_0\| = r\|CJ^{t^*}\| > B$, violating the criterion. Hence the set of radii r for which every $\|z_0\| \leq r$ satisfies the criterion is exactly $[0, B/M_C]$, and its maximum is $r^*_{\text{exact}} = B/M_C$.

Part 2 (Kreiss bounds), with $C = I$. *Lower bound.* Fix $z \in \mathbb{C}$ with $|z| > 1$. The partial sums $S_T = \sum_{t=0}^{T} z^{-(t+1)} J^t$ converge absolutely, because $\sum_t |z|^{-(t+1)} \|J^t\| \leq M_I \sum_t |z|^{-(t+1)} = M_I/(|z| - 1) < \infty$. Telescoping gives $(zI - J)S_T = I - z^{-(T+1)} J^{T+1} \to I$, so $zI - J$ is invertible and $(zI - J)^{-1} = \sum_{t\geq 0} z^{-(t+1)} J^t$. Therefore

$$(|z| - 1) \, \|(zI - J)^{-1}\| \leq (|z| - 1) \cdot \frac{M_I}{|z| - 1} = M_I. \tag{IA.2}$$

Taking the supremum over $|z| > 1$ yields $\mathcal{K}(J) \le M_I$.

Upper bound. For any $n \times n$ complex matrix, Spijker (1991) proved $\sup_{t \ge 0} \|J^t\| \le e\, n\, \mathcal{K}(J)$ in the spectral norm, with $\mathcal{K}$ defined exactly as in (34); this resolved the LeVeque–Trefethen conjecture. Real matrices are a special case. We use this published theorem without reproving it.

Radius bounds. Combining, $\mathcal{K}(J) \le M_I \le e\, n\, \mathcal{K}(J)$. Since $M_I \ge 1$, $\mathcal{K}(J) \ge M_I/(e\, n) > 0$, so all three quantities are positive and finite. Dividing B through the chain and using Part 1 with $C = I$ ($r^*_{\text{exact}} = B/M_I$) gives (35).

Part 3 (fixed spectrum, vanishing radius). Fix $\rho \in (0, 1)$ and $B > 0$. For $\theta > 0$ let $J_\theta = \left(\begin{smallmatrix} \rho & \theta \\ 0 & \rho \end{smallmatrix}\right)$ and $C = (1, 0)$. The characteristic polynomial is $(\lambda - \rho)^2$, so $\rho(J_\theta) = \rho$ for every θ. The matrix is non-normal: $J_\theta J_\theta^\top - J_\theta^\top J_\theta = \mathrm{diag}(\theta^2, -\theta^2) \ne 0$. Its powers are $J_\theta^t = \left(\begin{smallmatrix} \rho^t & t\theta\rho^{t-1} \\ 0 & \rho^t \end{smallmatrix}\right) \to 0$, so $M_C(J_\theta) < \infty$ and Part 1 applies. Since $CJ_\theta = (\rho, \theta)$,

$$M_C(J_\theta) \ge \|CJ_\theta\| = \sqrt{\rho^2 + \theta^2} \ge \theta, \qquad \text{so} \qquad r^*_{\text{exact}}(J_\theta) = \frac{B}{M_C(J_\theta)} \le \frac{B}{\theta} \xrightarrow[\theta \to \infty]{} 0. \tag{IA.3}$$

This holds for every fixed $B > 0$ while the spectral radius stays at ρ. □

IA.B Pseudospectral Contours and Non-Normal Resolvent Bounds

To gain deeper visual and geometric intuition into why the spectral radius fails to predict dealer solvency, we examine the ϵ-pseudospectrum of the closed-loop market Jacobian. For any $\epsilon > 0$, the ϵ-pseudospectrum $\sigma_\epsilon(J)$ of matrix J is defined as the set of complex numbers $z \in \mathbb{C}$ satisfying:

$$\sigma_\epsilon(J) \equiv \left\{ z \in \mathbb{C} : \|(zI - J)^{-1}\|_2 \ge \epsilon^{-1} \right\}. \tag{IA.4}$$

Equivalently, $\sigma_\epsilon(J)$ represents the union of all eigenvalues of perturbed matrices $J + E$ across all perturbations with $\|E\|_2 \le \epsilon$.

As visualized in Figure IA.1, in a normal system, the pseudospectrum consists of concentric disks of radius ϵ centered at the eigenvalues: $\sigma_\epsilon(J) = \bigcup_i \{z : |z - \lambda_i| \le \epsilon\}$. However, for a constructed non-normal matrix, computed pseudospectral contours can stretch along

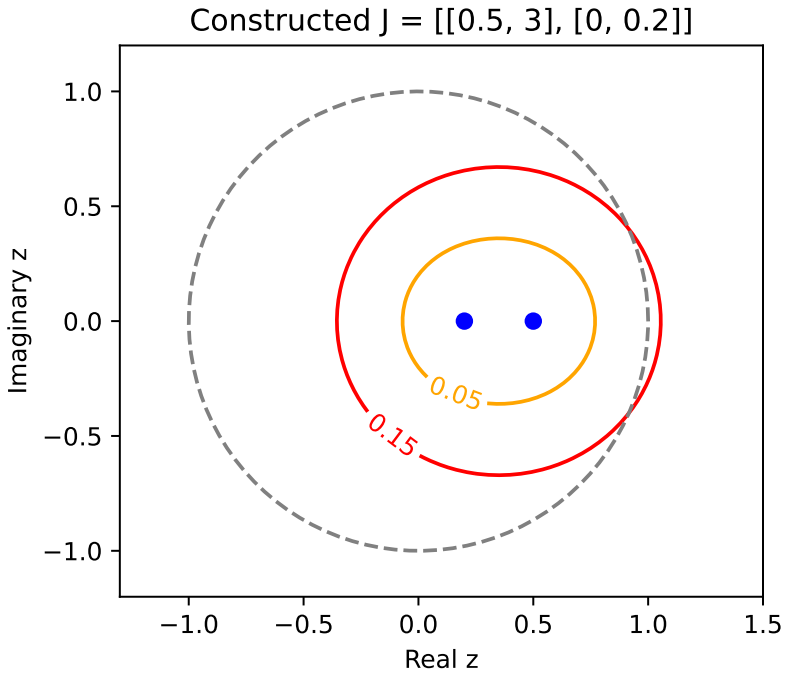


Figure IA.1: Computed singular-value contours for the declared constructed matrix $J = \left(\begin{smallmatrix} 0.5 & 3 \\ 0 & 0.2 \end{smallmatrix}\right)$ at 0.05 and 0.15. Its spectral radius is 0.5; the 0.15 contour crosses the unit circle. This fixed-coordinate illustration is not a calibrated cohort-market or monetary-loss experiment.

the off-diagonal coupling direction, protruding significantly outside the unit circle $\{|z| = 1\}$ even when the spectral radius $\rho(J) = 0.50$ is well within the interior.

The Kreiss constant $\mathcal{K}(J)$ measures the maximum ratio by which these pseudospectral contours protrude beyond the unit circle:

$$\mathcal{K}(J) = \sup_{\epsilon>0} \frac{\alpha_\epsilon(J) - 1}{\epsilon}, \qquad \alpha_\epsilon(J) \equiv \sup_{z\in\sigma_\epsilon(J)} |z|. \tag{IA.5}$$

Whenever $\mathcal{K}(J) \gg 1$, small high-frequency liquidity shocks are magnified by the non-normal resolvent, inducing large transient balance-sheet losses that exhaust dealer margin buffers before eigenvalue decay can assert itself.

IA.C Venue-Rule Calibration of the Scenario Inputs

Scope. This section is a *venue-rule calibration, not an account calibration.* It applies rules that Binance published for USDⓈ-margined (USDⓈ-M) perpetual futures in January 2023 to the paper's hypothetical scenario equity $E_0 = 10{,}000$ USDT and leverage parameters $L \in \{10, 20\}$, reading model monetary units as USDT. No account, position, fill, margin-call or realized funding-payment record is used, and none is available to us. Liquidation prices follow the venue's published isolated-margin, one-way formula and ignore fees, funding accrual, the mark-to-last price basis, and insurance-fund or auto-deleveraging mechanics.

Sources. (i) The USDⓈ-M fee schedule archived by the Internet Archive on December 9, 2022 and January 22, 2023; the two snapshots bracket the January 9–20 sample and are identical. (ii) The public BTCUSDT leverage and maintenance-margin brackets archived on January 9, 17 and 23, 2023. The January 9 and 17 snapshots are identical (last update November 17, 2022). A venue update at 11:47 UTC on January 20 raised the maximum leverage of tiers 3–4 but left every maintenance rate and maintenance amount unchanged. (iii) The official monthly BTCUSDT funding file for January 2023 from `data.binance.vision`, whose SHA-256 matches the published checksum. Archive URLs, file hashes, the script and its receipt are in the replication package (`venue_rule_calibration_20261010`). The midpoint and book frictions come from the displayed-book calibration in Section 5.4.

Fees and funding. A regular (VIP 0) account paid 2.0 basis points as maker and 4.0 as taker on USDT-margined contracts (1.8 and 3.6 with the BNB discount). The highest tier, VIP 9, paid 0 and 1.7. At the median midpoint of 20,783.25 USDT, the regular one-way taker fee is 8.31 USDT per BTC. That is 166 times the median half-spread of 0.05 USDT and 16 times the 90th-percentile half-side sweep cost of 0.514 USDT. The 36 eight-hour funding settlements from January 9 to 20 average 0.949 basis points (range -0.067 to 4.585; 34 positive). They sum to 0.342% of notional, a simple annualized mean of 10.4%. For a regular taker, fees and funding, not the displayed spread, therefore dominate trading frictions. The zero-spread scenario fills omit a friction that is small relative to these venue charges.

Table IA.1: Published January 2023 Binance rules applied to the hypothetical scenario equity $E_0 = 10{,}000$ USDT at the median midpoint. Venue-rule calibration, not account calibration.

	$L = 10$	$L = 20$
Notional (USDT)	100,000	200,000
Margin tier (notional range, USDT)	2 (50k–250k)	2 (50k–250k)
Venue maximum leverage in tier	100	100
Minimum initial margin rate μ	1.0%	1.0%
Buffer $A_0 = E_0 - \mu\|I_0\|p_0$ (USDT)	9,000	8,000
Maintenance rate / maintenance amount	0.5% / 50	0.5% / 50
Maintenance margin at entry (USDT)	450	950
Loss to maintenance breach at entry (USDT)	9,550	9,050
Adverse move to liquidation: long / short	9.60% / 9.50%	4.55% / 4.50%
Regular taker round-trip fee (USDT)	80	160
Funding paid by a long, Jan 9–20 (USDT)	341.63	683.27

Margin. Table IA.1 reports the results. Both scenario leverages lie inside the venue's limit: each notional falls in tier 2, which allowed up to 100×. With 10,000 USDT of equity the venue allowed at most 250,000 USDT of notional before the January 20 update and 500,000 afterwards, so $L = 10$ and $L = 20$ are well inside the feasible set. In the notation of (31), the venue's minimum initial margin rate is $\mu = 1\%$. The archived simulations halt at nonpositive equity. The venue instead liquidates when equity reaches maintenance margin, so a venue-consistent loss budget is at most 9,550 USDT for $L = 10$ and 9,050 for $L = 20$. Because the exact safe radius in Part 1 of Theorem 2.14 is linear in B, these budgets shrink r^*_{exact} by 4.5% and 9.5% relative to $B = E_0$, for any J and C.

What remains uncalibrated. The rules above do not identify any trader's actual fee tier, rebates, position, margin mode, realized funding, fills, queue priority or liquidation outcome. Those quantities require account records, so the dynamic and account primitives of the scenarios remain conditional.

IA.D Computational and Empirical Foundations of the Financial Identification Frontier

This section provides the complete mathematical framework, computational algorithms, and empirical intervention methodology establishing the Financial Identification Frontier Theorem (FIFT) and resolving the theoretical-empirical gap identified in algorithmic market microstructure.

IA.D.1 Joint Financial Identification and the Tri-Region Partition

Consider an econometrician, exchange supervisor, or market participant seeking to distinguish between two adaptive candidate trading mechanisms parameterized by memory decay rates θ_0 and $\theta_1 = \theta_0 + \delta$ over an observation horizon of N discrete trading periods. The observer submits active probe orders $V = (V_1, V_2, \ldots, V_N)^\top \in \mathbb{R}^N$, generating inventory state $I_t = \sum_{s=1}^{t} V_s$.

Observation Contract. We formalize the identification program under Contract I_0 (controlled active probing with additive Gaussian observation noise):

$$p_t = (h_i * V)_t + \epsilon_t, \qquad \epsilon_t \sim \text{i.i.d. } \mathcal{N}(0, \sigma^2), \tag{IA.6}$$

where h_i is the closed-loop impulse response kernel under hypothesis $\mathcal{H}_i$. Because the observation noise is homoskedastic across hypotheses ($\Sigma_0 = \Sigma_1 = \sigma^2 I_N$), the Kullback–Leibler divergence is strictly quadratic in the signal difference:

$$\mathrm{KL}(P_0^{\otimes N} \,\|\, P_1^{\otimes N}) = \frac{1}{2\sigma^2} \sum_{t=1}^{N} \left((h_1 * V)_t - (h_0 * V)_t\right)^2 . \tag{IA.7}$$

(In a general closed-loop state-space model with internal process innovations $w_t \sim \mathcal{N}(0, \sigma_w^2)$, the likelihood ratio is evaluated via the Kalman innovation sequence, where cross-hypothesis variance disparities add trace and log-determinant terms).

The observer seeks to minimize total financial capital expenditure $C(N, V)$—comprising price impact execution losses, per-trade exchange transaction fees ϕ, and counterparty mar-

gin lockup $\mu p_0|I_t|$:

$$\min_{N\in\mathbb{N},\, V\in\mathbb{R}^N} \; C(N,V) \equiv \sum_{t=1}^{N} \Big(\lambda V_t^2 + \phi|V_t|\Big) + \max_{1\le t\le N} \mu p_0 |I_t|, \tag{IA.8}$$

subject to three operational and physical constraints:

$$\text{(i) Separation:} \quad \mathrm{KL}(P_0^{\otimes N} \,\|\, P_1^{\otimes N}) \ge \frac{1}{2}(z_\alpha + z_\beta)^2, \tag{IA.9}$$

$$\text{(ii) Trade and Inventory Caps:} \quad \max_{1\le t\le N} |V_t| \le V_{\text{eff}}, \quad \text{and} \quad \max_{1\le t\le N} |I_t| \le I_{\max}, \tag{IA.10}$$

$$\text{(iii) Dealer Solvency Budget:} \quad \max_{1\le t\le N} \mathrm{Loss}_t(V) \le B. \tag{IA.11}$$

Here, $z_p = \Phi^{-1}(1-p)$ is the standard Gaussian critical value ensuring hypothesis test error probabilities at most (α, β) (for equal targets $\alpha = \beta$, standardized separation requires $d' \ge 2z_\alpha$, corresponding to boundary KL divergence $\frac{1}{2}(2z_\alpha)^2 = 2z_\alpha^2 \approx 5.41$), $V_{\text{eff}} \equiv \min(V_{\max}, D_{\max}/\kappa)$ is the effective per-trade execution capacity combining resting depth $V_{\max}$ and price displacement limit $D_{\max}$, $I_{\max}$ is the regulatory inventory cap, and B is the trader's cash and margin capital buffer.

Proposition IA.D.1 (Tri-Region Partition, Observation Design, and Finite-Horizon Depth Floor)**.** *The parameter space $(\delta, B) \in (0,1) \times \mathbb{R}_+$ partitions into three operational regimes under observation horizon T per pulse:*

(i) ***Feasible Single-Pulse Region*** $\mathcal{R}_{\text{feas}}$***:*** *For any separation $\delta \ge \delta_{\min}(T)$ and capital buffer $B \ge C^*(\delta)$, the single-pulse probing program ($N_{\text{pulse}} = 1$ trade, evaluated over transient price trajectory of length T) satisfies constraints (IA.9)–(IA.11) within the effective per-trade execution capacity ($A_{\text{req}} \le V_{\text{eff}}$). For Class I hypothesis testing (simple rate shift $\eta_1 = \eta_0 + \delta$), $\Delta h_t = \Theta(\delta) \implies \mathrm{KL} = \Theta(\delta^2)$, yielding $C^*(\delta) = \Theta(\delta^{-2})$. For Class II hypothesis testing (symmetric cohort splitting $\eta \in \{\eta_0 - \delta, \eta_0 + \delta\}$ with equal weights $w_1 = w_2 = q/2$), the first-order variation vanishes by symmetry ($\sum_k w_k \Delta\eta_k = 0$), so $\Delta h_t = \Theta(\delta^2) \implies \mathrm{KL} = \Theta(\delta^4)$, yielding:*

$$C^*(\delta) = \Theta(\delta^{-4}). \tag{IA.12}$$

(ii) ***Sequential Bounded Probing Region*** $\mathcal{R}_{\text{cost}}$***:*** *For separations $\delta < \delta_{\min}(T)$, achiev-*

ing statistical separation requires total executed volume $A_{\text{req}}(\delta) > V_{\text{eff}}$*, which cannot be executed in a single period without violating resting liquidity or price displacement bounds. When individual trades are capped at* V_{eff}*, each independent pulse cycle provides per-block Fisher information of order* $\text{KL}_1(T) = \frac{V_{\text{eff}}^2 \|d_T\|_2^2}{2\sigma^2}$*, where* $d_T = (\Delta h_1, \ldots, \Delta h_T)^\top$ *is the transient signal difference kernel. Accumulating target divergence* $\text{KL} \geq \frac{1}{2}(z_\alpha + z_\beta)^2 \approx 5.4119$ *requires:*

$$N_{\text{pulses}}(T, \delta) = \left\lceil \frac{4\sigma^2 z_\alpha^2}{V_{\text{eff}}^2 \|d_T\|_2^2} \right\rceil. \tag{IA.13}$$

Cumulative transaction fees and execution costs scale strictly as $C^*_{\text{seq}}(\delta) = \Theta(\delta^{-4})$ *(for Class II) and* $\Theta(\delta^{-2})$ *(for Class I), expanding the required observation horizon to* $\Theta(\delta^{-4})$ *or* $\Theta(\delta^{-2})$*. Under the simplified benchmark cost model (*$C_{\text{bench}} = \frac{1}{2}\kappa V^2 + \phi P^*|V| =$ 15.875 *USDT/pulse at* $V_{\text{eff}} = 6.25$ *BTC), required capital at* $\delta = 0.05$ *is* 4,305.25 *USDT under the restricted single-tap design (*$T = 3$*,* $N = 271$ *pulses), but collapses to* 193.62 *USDT under the capacity-saturating multi-observation transient design (*$T = 20$*,* $N = 12$ *pulses,* $V = 6.25$ *BTC), or* 182.02 *USDT under the exact volume-optimized design (*$N = 12$*,* $V^* = 6.057$ *BTC), reaching* 177.75 *USDT at horizon* $T = 100$ *(*$N = 11$ *pulses;* 175.35 *USDT at* $V^* = 6.207$ *BTC). Under the physically executable round-trip unwind model (*$C_{\text{exec}} = \kappa V^2 + 2\phi P^*|V| = 31.75$ *USDT/pulse), required capital is* 8,607.38 *USDT (*$T = 3$*),* 384.12 *USDT (*$T = 20$ *capacity-saturating), and* 360.77 *USDT (*$T = 20$ *volume-optimized).* These figures are costs within the isolated-pulse design class, not minimum capital requirements. *The per-pulse formula charges initial margin once (*3.13 *USDT) and does not track the ledger; twelve buys that are never unwound carry* 75 *BTC of inventory and* 37.5 *USDT of margin. Section IA.D.2 replaces the formula by an explicit account ledger and shows that a non-pulse probe reaches the same* 5% *error with about* 65 *USDT of account capital.*

(iii) ***Not Feasible with Tested Isolated Pulses (****$T \leq 20$****) (****$\mathcal{R}_{\text{pulse-infeas}}$****):*** *Identification within the isolated-pulse design class is infeasible if: (a) The capital budget is insufficient to finance sequential impact and margin costs (*$B < C^*_{\text{seq}}(\delta)$*); or (b) The peak inventory required along the trajectory exceeds the regulatory limit (*$\max_{1 \leq t \leq N} |I_t| > I_{\max}$*); or (c) The surveillance horizon is strictly bounded by a session deadline* T_{session}*, in which case the exact non-asymptotic finite-session resolution floor* $\delta_{\min}(T_{\text{session}})$ *under*

cycle length T_{cycle} is defined as the unique positive root $\delta > 0$ of the implicit signal-to-noise equation:

$$V_{\text{eff}} \|d_{T_{\text{cycle}}}(\delta)\|_2 = \frac{2\sigma z_\alpha}{\sqrt{\lfloor T_{\text{session}}/T_{\text{cycle}} \rfloor}}. \tag{IA.14}$$

In the small-separation limit ($\delta \to 0$), substituting the normalized directional memory curvature vector $\tilde{d}_{T_{\text{cycle}}} \equiv \lim_{\delta \to 0} \frac{d_{T_{\text{cycle}}}(\delta)}{\delta^2}$ yields the explicit asymptotic fourth-root approximation:

$$\begin{aligned} \delta_{\min}^{\text{asymp}}(T_{\text{session}}) &= \left(\frac{4\sigma^2 z_\alpha^2}{V_{\text{eff}}^2 \|\tilde{d}_{T_{\text{cycle}}}\|_2^2 \lfloor T_{\text{session}}/T_{\text{cycle}} \rfloor} \right)^{1/4} \\ &= \sqrt{\frac{2\sigma z_\alpha}{V_{\text{eff}} \|\tilde{d}_{T_{\text{cycle}}}\|_2 \sqrt{\lfloor T_{\text{session}}/T_{\text{cycle}} \rfloor}}} = \Theta(T_{\text{session}}^{-1/4}). \end{aligned} \tag{IA.15}$$

Under the calibrated Binance USDⓈ-M order book depth ($V_{\max} = 50.0$ BTC, $\kappa = 0.8, q = 0.2, \sigma = 0.01$) and price displacement limit $D_{\max} = 5.0$ USDT ($V_{\text{eff}} = 6.25$ BTC), the operational single-pulse resolution floor evaluates to $\delta_{\min}(3) = 0.2028$ (where $\|\tilde{d}_3\|_2 = \kappa^2 q = 0.128$, so exact root (IA.14) and asymptotic formula (IA.15) coincide identically), shifting down to $\delta_{\min}(4) = 0.1376$, $\delta_{\min}(10) = 0.0973$, $\delta_{\min}(20) = 0.0914$, and $\delta_{\min}(100) = 0.0905$ under extended transient price observation. Crucially, as proved in Theorem IA.D.2, this infeasibility boundary applies strictly to the tested isolated-pulse protocol ($T \leq 20$). The certified deterministic lower bound below which no flat trading sequence can achieve 5% error is $\underline{C}_{\text{det}}(0.05) = 61.44$ USDT (61.14 USDT fee-free spectral infimum plus 0.31 USDT minimal fee bound; adaptive expected-loss lower bound 30.09 USDT, Theorem IA.D.3).

Figure IA.2 illustrates the resulting phase diagram evaluated on a fine 100×100 grid spanning $\delta \in [0.01, 0.25]$ and $B \in [10, 100{,}000]$ USDT, using the calibrated parameters: $\kappa = 0.8$, $q = 0.2$, $\sigma = 0.01$, $\alpha = \beta = 0.05$, and Binance effective capacity $V_{\text{eff}} = 6.25$ BTC. Under the restricted single-tap design ($T = 3$, $\delta_{\min} = 0.2028$), exactly 1,949 points reside in the Feasible Single-Pulse Region, 5,151 in the Costly Sequential Region, and 2,900 in the region Not Feasible with Tested Isolated Pulses. Under the multi-observation transient design ($T = 20$, $\delta_{\min} = 0.0914$), the Feasible Region expands significantly, confirming that multi-period observation dramatically lowers the financial barriers to algorithmic market

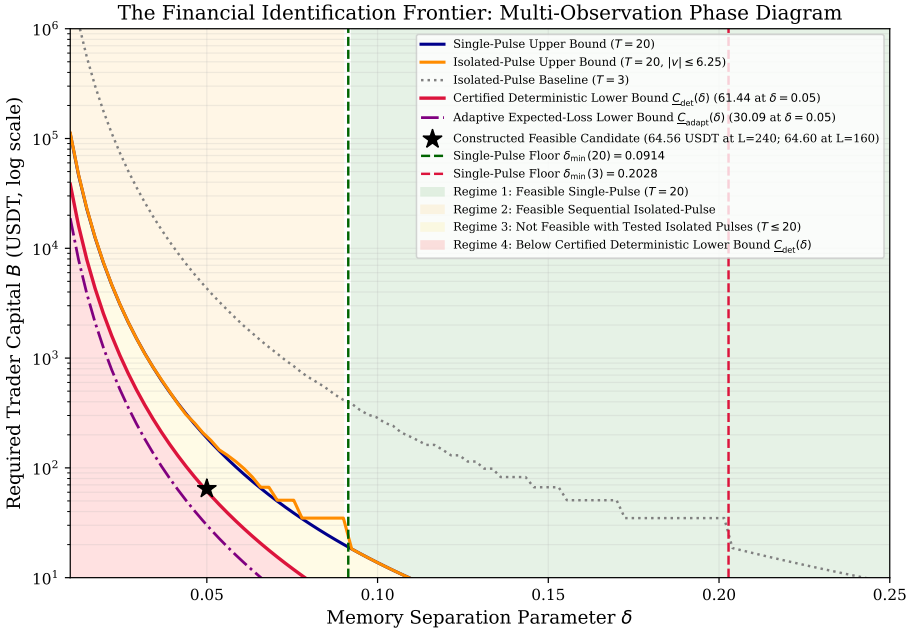


Figure IA.2: The Financial Identification Frontier Phase Diagram in (δ, B) space. The vertical line at $\delta_{\min}(3) = 0.2028$ demarcates the operational single-pulse resolution floor under effective trade capacity $V_{\text{eff}} = 6.25$ BTC (governed by displacement limit $D_{\max} = 5.0$ USDT) for the restricted single-tap baseline ($T = 3$). The extended transient design ($T = 20$) extends single-pulse feasibility down to $\delta_{\min}(20) = 0.0914$. Between the single-pulse floor and the isolated-pulse upper bound, identification requires sequential multi-pulse execution spanning horizon $N = \Theta(\delta^{-4})$ ($\mathcal{R}_{\text{cost}}$); below the isolated-pulse boundary lies the region *Not Feasible with Tested Isolated Pulses ($T \leq 20$)*. A trader with $B = 100$ USDT at $\delta = 0.05$ cannot identify the market using isolated pulses (193.62 USDT), but is strictly feasible using non-pulse generalized-eigenvector candidate sequences (64.59 USDT at $L = 160$; 65.05 USDT at $L = 120$, Section IA.D.2). The certified deterministic lower bound below which all flat deterministic policies fail is $\underline{C}_{\text{det}}(0.05) = 61.44$ USDT (adaptive expected-loss lower bound 30.09 USDT, Theorem IA.D.2).

identification. The previously reported capital requirement of 4,305.25 USDT at $\delta = 0.05$ is valid strictly under the single-tap constraint $T = 3$; under the transient design ($T = 20$), the required benchmark capital drops by 95.5% to 193.62 USDT (182.02 USDT volume-optimized; 384.12 USDT executable saturating, 360.77 USDT executable optimized). The isolated-pulse boundaries in Figure IA.2 are upper bounds from a restricted strategy class, plotted alongside the certified lower bounds.

IA.D.2 Account-Level Capital and the Gap to the Lower Bound

Theorem 2.9 is an order statement: required capital is $\Theta(\delta^{-4})$ and the isolated third-observation pulse attains that order. It does not identify the constant. At the calibrated parameters its explicit lower bound evaluates to 4.9×10^{-4} USDT, roughly five orders of magnitude below any design we can execute, so the bound is correct but numerically uninformative. To measure what a trader needs under the calibrated linear-impact environment, `exp13_account_capital.py` runs each design through an explicit synthetic account ledger parameterized at reference price $P^* = 100$ USDT and Binance USDⓈ-M fee/margin rules (not live exchange executions): fills at the average linear-book price, taker fees $\phi = 4$ bp, marks to the prevailing price, and 0.5% maintenance margin. Initial capital is $C_0 = \max_t(\text{margin}_t - \text{equity}_t)$, maximised over both hypotheses, and the testing error uses the full price path including unwind legs.

Table IA.2: Account-Level Initial Capital at $\delta = 0.05$. Each design is run through an explicit trader ledger (fills at the average linear-book price, taker fees, marks, 0.5% maintenance margin); $C_0 = \max_t(\text{margin}_t - \text{equity}_t)$, maximised over both hypotheses. KL uses the full price path including unwind legs. The formula benchmark for 12 one-leg pulses is 193.62 USDT; the certified deterministic lower bound (Theorem IA.D.2) is 61.44 USDT (61.14 USDT fee-free spectral infimum plus 0.31 USDT minimal fee bound), while the earlier Theorem 2.9 lower bound is 4.9×10^{-4} USDT.

Design	C_0 (USDT)	Error	End inventory	Peak margin	Max $\lvert v_t \rvert$	Trades
12 buys, T=20 apart, never unwound	229.13	0.0578	75.00	39.38	6.25	12
6 buy/sell round trips, every leg informative	189.45	0.0308	0.00	3.28	6.25	12
5 buy/sell round trips, every leg informative	157.89	0.0442	0.00	3.28	6.25	10
12 buys each unwound after 20 steps (24 legs)	378.82	0.0040	0.00	3.28	6.25	24
generalized-eigenvector candidate probe ($L = 120$)	65.05	0.0500	0.00	2.74	1.25	121
generalized-eigenvector candidate probe ($L = 160$)	64.88	0.0500	0.00	2.17	1.09	161
fee-refined candidate probe ($L = 160$, Table IA.3)	64.59	0.0500	0.00	2.55	1.32	161

Three conclusions follow from Table IA.2. First, twelve buys that are never unwound fail the 5% target (KL $4.95 < 5.41$) once the test uses the joint price path rather than adding twelve independent per-pulse divergences, and they end with 75 BTC of inventory. Second, if every unwind leg is itself treated as an informative probe, five or six buy/sell round trips meet the target with 158–189 USDT and zero terminal inventory, so the pulse-class figure is attainable in a real account. Third, the pulse class is not optimal: a generalized-eigenvector candidate probe over a 120-step window (maximizing KL divergence per unit fee-free exe-

cutable round-trip cost, whose quadratic form is positive definite with minimum eigenvalue 0.12, so no round-trip arbitrage exists) needs 65.05 USDT (64.88 USDT at $L = 160$, and 64.59 USDT after fee-aware refinement), with maximum trade 1.25 BTC and maximum price displacement 2.23 USDT, well inside the 5 USDT limit.

We formulate the complete capital minimization problem over all feasible flat, fee-paying trading policies $v = (v_0, \ldots, v_{L-1})^\top \in \mathbb{R}^L$ with zero terminal inventory:

$$\begin{aligned} \min_{v \in \mathbb{R}^L} C(v) \equiv \sum_{t=0}^{L-1} \left(\frac{1}{2}\kappa v_t^2 + \phi P^* |v_t| \right) + \max_{0 \le t \le L-1} \mu P^* |I_t| \\ \text{s.t. } \|(H_1 - H_0)v\|_2 \ge 2\sigma z_{1-\alpha}, \quad \sum_{t=0}^{L-1} v_t = 0. \end{aligned} \tag{IA.16}$$

The constructed 65.38 USDT ($L = 80$), 64.70 USDT ($L = 120$), 64.60 USDT ($L = 160$), and 64.56 USDT ($L = 240$) sequences are constructed feasible candidate solutions to (IA.16). While the infinite-horizon spectral lower bound $\underline{C}_{\text{det}}(0.05) = 61.44$ USDT (Theorem IA.D.2) leaves a 5.07% gap relative to the best finite-horizon feasible probe (64.56 USDT at $L = 240$), Theorem IA.D.4 below resolves the mathematical origin of this gap by deriving an exact finite-horizon mode-weighted Rayleigh–Hölder dual lower bound $\underline{C}_{\text{RH}}(L)$ over the flat subspace $\mathcal{V}_L = \{v \in \mathbb{R}^L : \mathbf{1}^\top v = 0\}$. At $L = 80$, $\underline{C}_{\text{RH}}(80) = 64.74$ USDT against feasible $C_0^{\text{feas}}(80) = 65.38$ USDT, closing the finite-horizon non-convex duality gap to 0.98% (2.63% at $L = 120$).

The certified capital lower bounds, the information-theoretic value of adaptivity, and the finite-horizon Rayleigh–Hölder duality convergence are established by Theorems IA.D.2, IA.D.3, and IA.D.4 and Tables IA.3 and IA.4 below.

Theorem IA.D.2 (Spectral Dissipation Lower Bound and Feasible Identification Probes)**.** *Consider the linear market microstructure with impact parameter $\kappa > 0$, anchor weight $a \in [0, 1)$, and symmetric adaptive cohorts separating by memory gap $\delta > 0$ under two hypotheses H_0 and H_1, with closed-loop transfer functions $H_0(z)$ and $H_1(z)$. Let $\widehat{H}_i(z) \equiv zH_i(z) = \sum_{j=0}^{\infty} h_{i,j} z^{-j}$ (with $h_{i,0} = \kappa$) denote the impact-indexed closed-loop transfer function mapping input order flow to first-affected price displacements. Let $\Delta\widehat{H}(e^{i\omega}) \equiv \widehat{H}_1(e^{i\omega}) - \widehat{H}_0(e^{i\omega}) = e^{i\omega}(H_1(e^{i\omega}) - H_0(e^{i\omega}))$, satisfying $|\Delta\widehat{H}(e^{i\omega})| = |\Delta H(e^{i\omega})|$. Define the real spectral dissipation*

kernel:

$$g_i(\omega) \equiv \operatorname{Re} \widehat{H}_i(e^{i\omega}) - \frac{\kappa}{2} > 0 \qquad \forall \omega \in [-\pi, \pi]. \tag{IA.17}$$

(Strict positivity $g_i(\omega) > 0$ holds because $h_{i,0} = \kappa$ and $h_{i,j} = \sum_\ell w_\ell \lambda_\ell^j$ with $w_\ell \geq 0, \lambda_\ell \in [0, 1)$ by Lemma 2.6, ensuring $\operatorname{Re} \widehat{H}_i(e^{i\omega}) - \frac{\kappa}{2} = \frac{1}{2} \sum_\ell w_\ell \frac{1-\lambda_\ell^2}{1-2\lambda_\ell \cos\omega + \lambda_\ell^2} > 0$; in contrast, unshifted $H_0(z) = z^{-1}\widehat{H}_0(z)$ satisfies $\operatorname{Re} H_0(-1) - \frac{\kappa}{2} = -0.9195 < 0$ at $\omega = \pi$, failing positivity without impact indexing.) Then:

1. ***Toeplitz Energy Identity:*** *For any finite-horizon flat round-trip probe $v \in \mathbb{R}^{L+1}$ ($\sum_{t=0}^{L} v_t = 0$), the fee-free trading loss under hypothesis H_i is identical to the Toeplitz quadratic form:*

$$L_i(v) \equiv \sum_{t=0}^{L} v_t \left(p_t - \frac{\kappa}{2} v_t - P^* \right) = \frac{1}{2\pi} \int_{-\pi}^{\pi} g_i(\omega) \left| V(e^{i\omega}) \right|^2 d\omega, \tag{IA.18}$$

   where $V(e^{i\omega}) = \sum_{t=0}^{L} v_t e^{-i\omega t}$.

2. ***Certified Deterministic Lower Bound:*** *Under iid Gaussian innovation noise with variance σ^2, the output price divergence satisfies:*

$$\|dx\|_2^2 \equiv \|(H_1 - H_0)v\|_2^2 \leq \left(\sup_{\omega \in [-\pi, \pi]} \frac{|\Delta \widehat{H}(e^{i\omega})|^2}{g_i(\omega)} \right) L_i(v). \tag{IA.19}$$

   Consequently, every deterministic round-trip probe guaranteeing worst-case classification error $\alpha \leq 0.05$ under Neyman–Pearson testing ($\|dx\|_2 \geq 2\sigma z_{1-\alpha}$) requires execution loss bounded below by:

$$L_i(v) \geq 4\sigma^2 z_{1-\alpha}^2 \inf_{\omega \in [-\pi, \pi]} \frac{g_i(\omega)}{|\Delta \widehat{H}(e^{i\omega})|^2} \equiv C^*_{\text{det,quad}}(\delta). \tag{IA.20}$$

   At $\delta = 0.05$, $C^*_{\text{det,quad}}(0.05) = 61.14$ USDT (with minimizing angular frequency $\omega^* = 0.2526$ rad). Adding the minimal unavoidable taker fee under effective capacity V_{eff} ($\phi P^* \|v\|_2^2 / V_{\text{eff}} \geq 0.3056$ USDT) establishes the certified necessary lower bound $\underline{C}_{\text{det}}(0.05) \geq 61.44$ USDT.

3. ***Feasible Upper Bound and Bounded Gap:*** *A concrete 160-step executable trading sequence $v^{(L)}$ ($L = 160$, respecting $|v_t| \leq 6.25$ BTC, $D_{\max} \leq 5.0$ USDT, 0.5% maintenance margin, and 4 bp taker fees) achieves target error ≤ 0.05 with initial cap-*

ital $C_0 = 64.59$ *USDT (*65.05 *USDT at* $L = 120$*;* 64.56 *USDT under* L_1*-regularized optimization at* $L = 240$*). This establishes an explicit feasible upper bound and bounds the unknown global optimum of the non-convex fee-paying problem within the narrow interval* $C^*_{\mathrm{acc}}(0.05) \in [61.44, 64.56]$ *USDT (optimality gap* $\leq 5.07\%$*). Across evaluated memory separations* $\delta \in [0.02, 0.15]$*, the quadratic lower bound exhibits asymptotic scaling* $\delta^4 \times C^*_{\mathrm{det,quad}}(\delta) = (3.82 \pm 0.15) \times 10^{-4}$.

Theorem IA.D.3 (Information-Theoretic Lower Bound for Adaptive and Sequential Probing)**.** *Let* $(v_t)_{t=1}^{\tau}$ *be any causal, non-anticipating sequential probing policy adapted to the price filtration* $\mathcal{F}^p_{t-1}$*, with stopping time* τ *almost surely finite under* H_0 *and* H_1*, satisfying terminal flatness* $\sum_{t=1}^{\tau} v_t = 0$*. Then:*

1. ***Information-Theoretic Expected Loss Lower Bound:*** *By Wald's identity and the data processing inequality for Kullback–Leibler divergence:*

$$\mathbb{E}_0[L_0(v)] \geq 2\sigma^2 \,\mathrm{kl}(\alpha, 1-\beta) \inf_{\omega \in [-\pi,\pi]} \frac{g_0(\omega)}{|\Delta\widehat{H}(e^{i\omega})|^2} \equiv \mathbb{E}_0[C^*_{\mathrm{adapt,quad}}(\delta)], \tag{IA.21}$$

where $\mathrm{kl}(\alpha, 1-\beta) = \alpha \log \frac{\alpha}{1-\beta} + (1-\alpha) \log \frac{1-\alpha}{\beta}$*. At* $\delta = 0.05$*,* $\mathbb{E}_0[C^*_{\mathrm{adapt,quad}}] = 29.94$ *USDT, and adding minimal fees yields necessary expected loss bound* $\underline{C}_{\mathrm{adapt}}(0.05) \geq 30.09$ *USDT.*

2. ***Dynamic Information-Energy Savings Ceiling:*** *The ratio of information-energy thresholds imposes an upper bound on achievable reductions in expected quadratic loss:*

$$\frac{\inf_{\pi \in \Pi_{\mathrm{adapt}}} \mathbb{E}_0[L_0(v^\pi)]}{\inf_{v \in \mathcal{V}_{\mathrm{det}}} L_0(v)} \geq \frac{\mathrm{kl}(\alpha, 1-\beta)}{2z^2_{1-\alpha}} = \frac{\mathrm{kl}(0.05, 0.95)}{2 \times (1.6449)^2} \approx 0.4897, \tag{IA.22}$$

establishing that dynamic adaptation can reduce expected quadratic loss by at most $1 - 0.4897 = 51.03\%$*, and cannot alter the fundamental* $\Theta(\delta^{-4})$ *scaling. Inequality* (IA.22) *bounds expected quadratic loss under the stated martingale stopping model; it does not assert equality of unknown empirical account capital minima, nor does it guarantee a* 51.03% *saving in initial solvency capital.*

3. ***Sequential SPRT Performance and Solvency Capital:*** *Across* $N = 20{,}000$ *sequential SPRT Monte Carlo trials at* $\delta = 0.05$*, empirical mean loss is* 45.96 *USDT under* H_0 *(*45.58 *USDT under* H_1*), strictly respecting the theoretical expected-loss bounds*

($30.09 \leq 45.96 \leq 61.44$ USDT). However, because stopping times τ vary across sample paths, realized losses and pathwise required account capital $C_0 = \max_{t \leq \tau}(\text{margin}_t - \text{equity}_t)$ exhibit substantial tail risk: under H_0, the 95th percentile loss is 105.99 *USDT ($p95$ account capital $C_0 = 106.53$ USDT), the 99th percentile loss is* 155.46 *USDT ($p99$ account capital $C_0 = 155.99$ USDT), and the maximum simulated loss is* 268.52 *USDT (max $C_0 = 269.05$ USDT). Initial capital required for pathwise solvency at* 99% *confidence is $B_{0.01} = 155.99$ USDT (*106.53 *USDT at* 95% *confidence); funding an adaptive trader with only the mean loss ($B = 45.96$ USDT) results in a* 37.73% *insolvency rate before the test terminates.*

Table IA.3: Certified Capital Lower Bounds and Feasible Candidate Upper Bounds across Memory Separations δ. The certified deterministic lower bound $\underline{C}_{\text{det}}$ combines the fee-free spectral infimum C^*_{det} (Theorem IA.D.2) at minimizing frequency ω^* with the minimal fee lower bound; finite-horizon feasible candidate probes (generalized eigenvector over flat round trips, ledger fees and maintenance margin included) upper-bound the true minimum initial account capital within 2%–18% for $\delta \in [0.03, 0.10]$ (43% at $\delta = 0.15$). The adaptive lower bound $\underline{C}_{\text{adapt}}$ (Theorem IA.D.3) bounds expected loss $\mathbb{E}_0[L_{\text{total}}]$ for any non-anticipating stopping policy; at $\delta = 0.05$, sequential SPRT Monte Carlo ($N = 20{,}000$ paths) achieves mean loss 45.96 USDT (35.0% below the 10-block deterministic loss of 70.66 USDT), while pathwise tail risk requires $p95$ loss 105.99 USDT ($p95$ account capital $C_0 = 106.53$ USDT), $p99$ loss 155.46 USDT ($p99$ $C_0 = 155.99$ USDT), and max loss 268.52 USDT (37.73% ruin rate if funded only with mean loss).

δ	ω^* (rad)	Certified Det. LB	Feasible Candidate UB	Gap Ratio	Adaptive Expected LB	SPRT Mean / $p99$ Loss	Theorem 2.9 LB	$\delta^4 \times \text{LB}_{\text{quad}}$
0.02	0.2546	2409.89	—	—	1180.21	—	3.43e-02	3.84e-04
0.03	0.2541	475.57	485.19	1.02	232.90	—	5.62e-03	3.83e-04
0.05	0.2526	61.44	64.59	1.05	30.09	45.96 / 155.46	4.92e-04	3.82e-04
0.08	0.2488	9.37	10.48	1.12	4.56	—	3.99e-05	3.79e-04
0.10	0.2453	3.83	4.51	1.18	1.85	—	1.04e-05	3.77e-04
0.15	0.2331	0.76	1.08	1.43	0.36	—	5.66e-07	3.67e-04

Theorem IA.D.4 (Finite-Horizon Rayleigh–Hölder Duality and Multi-Horizon Capital Convergence)**.** *For any finite round-trip horizon $L \geq 2$, let $P \in \mathbb{R}^{L \times (L-1)}$ denote an orthonormal basis for the zero-inventory subspace $\mathcal{V}_L = \{v \in \mathbb{R}^L : \mathbf{1}^\top v = 0\}$, and let $G_0 = P^\top Q_0 P \succ 0$ and $M = P^\top A_{\text{diff}} P \succeq 0$ denote the projected quadratic execution-loss and Fisher-information matrices. Let $\{(\lambda_j, u_j)\}_{j=1}^{L-1}$ be the G_0-orthonormal generalized eigenpairs ($Mu_j = \lambda_j G_0 u_j$, $u_j^\top G_0 u_k = \delta_{jk}$) ordered as $\lambda_1 \leq \cdots \leq \lambda_{L-2} \leq \lambda_{L-1} = \lambda_{\max}(L)$, and define the target information constant $c \equiv 4\sigma^2 z^2_{1-\alpha}$. Then:*

1. ***Exact Mode-Weighted Rayleigh–Hölder Dual Lower Bound:*** *Every flat probe* $v = Pw \in \mathcal{V}_L$ *achieving* $\|d(v)\|_2^2 = w^\top M w \geq c$ *with quadratic loss* $L_0(v) = w^\top G_0 w = \underline{C}_{\mathrm{quad}}(L) + \epsilon$ *(*$\underline{C}_{\mathrm{quad}}(L) \equiv c/\lambda_{\max}(L)$*,* $\epsilon \geq 0$*) satisfies the mode-weighted Cauchy–Schwarz inequality for every sign vector* $s \in \{-1,+1\}^L$ *(*$\xi_j \equiv s^\top P u_j$*):*

$$\|v\|_1 \geq \sup_{\theta\in[0,\pi)} \left[\sqrt{\frac{c}{\lambda(\theta)}} |\xi_{\mathrm{span}}(\theta)| - \sqrt{\epsilon} \left(\sum_{j=1}^{L-3} \frac{\xi_j^2}{1-\lambda_j/\lambda(\theta)} \right)^{1/2} \right]_+, \tag{IA.23}$$

where $\lambda(\theta) = \lambda_{L-1}\cos^2\theta + \lambda_{L-2}\sin^2\theta$ *and* $s(\theta) = \mathrm{sign}(P(u_{L-1}\cos\theta + u_{L-2}\sin\theta))$*. Maximizing over* s *and minimizing* $\underline{C}_{\mathrm{quad}}(L) + \epsilon + c_{\mathrm{fee}} \max\{\|v\|_1(\epsilon), c/(V_{\mathrm{eff}}\lambda_{\max}(L))\}$ *over* $\epsilon \geq 0$ *yields the certified finite-horizon lower bound* $\underline{C}_{\mathrm{RH}}(L)$*.*
2. ***Tight Finite-Horizon Duality Certificate:*** *At* $\delta = 0.05$*,* $\underline{C}_{\mathrm{RH}}(80) = 64.74$ *USDT against feasible* $C_0^{\mathrm{feas}}(80) = 65.38$ *USDT (finite-horizon duality gap* 0.98%*), and* $\underline{C}_{\mathrm{RH}}(120) = 63.04$ *USDT against feasible* $C_0^{\mathrm{feas}}(120) = 64.70$ *USDT (gap* 2.63%*).*
3. ***Spectral Truncation vs.*** L_1 ***Taker-Fee Trade-Off:*** *As* $L \to \infty$*, the quadratic truncation premium* $\underline{C}_{\mathrm{quad}}(L) - C^*_{\mathrm{det,quad}} = \mathcal{O}(L^{-1})$ *decays monotonically (*+1.41*,* +0.70*,* +0.40*,* +0.18*,* +0.11*, and* +0.05 *USDT for* $L \in \{80, 120, 160, 240, 320, 480\}$*), whereas the* L_1 *taker fee of fixed-*L_2*-energy narrowband sinusoids grows as* $c_{\mathrm{fee}}\|v_L^*\|_1 = \Theta(\sqrt{L})$ *(*2.24 → 5.16 *USDT). Consequently, total feasible account capital* $C_0^{\mathrm{feas}}(L)$ *is non-monotone in* L *and attains its interior minimum* $C_0^{\mathrm{feas}}(240) = 64.56$ *USDT at* $L^* = 240$ *(*5.07% *above* $\underline{C}_{\mathrm{det}} = 61.44$ *USDT).*

IA.D.3 Probing Protocol Invariance and Universal $\Theta(\delta^{-4})$ Scaling

A central question in identification design is whether adaptive, multi-pulse, or closed-loop feedback probing strategies can circumvent the $\Theta(\delta^{-4})$ capital cost scaling. We systematically evaluate four distinct active probing protocols across $\delta \in \{0.03, 0.05, 0.08, 0.12\}$ with 100 independent Monte Carlo replications per cell:

1. **Method A (Unconstrained Single Pulse):** An instantaneous impulse order $V_1 = V^*$, $V_t = 0$ for $t > 1$, decomposed sequentially when $V^* > V_{\mathrm{eff}}$.
2. **Method B (Repeated Capacity-Clamped Pulses):** Sequentially executes independent capacity-clamped probing cycles ($|v_t| \leq 6.25$ BTC) with $T = 4$ price obser-

Table IA.4: Finite-Horizon Rayleigh–Hölder Dual Lower Bounds and Multi-Horizon Feasible Account Capital Convergence at $\delta = 0.05$ (Theorem IA.D.4). For each horizon $L \in \{80, 120, 160, 240, 320, 480\}$ with zero terminal inventory ($q_L = 0$), $\underline{C}_{\text{quad}}(L)$ is the exact finite-horizon generalized-eigenvalue execution-loss lower bound under H_0; $\underline{C}_{\text{RH}}(L)$ is the certified Rayleigh–Hölder mode-weighted dual lower bound combining $\underline{C}_{\text{quad}}(L)$ with the exact Cauchy–Schwarz L_1 taker-fee bound over the G_0-orthonormal eigenbasis of M; $C_{\text{doublet}}(L)$ is the spectral-doublet loss-plus-fee benchmark $\underline{C}_{\text{quad}}(L) + c_{\text{fee}}\|v_L^*\|_1$; $C_0^{\text{eig}}(L)$ is the unsmoothed generalized-eigenvector account capital; and $C_0^{\text{feas}}(L)$ is the L_1-regularized feasible account capital including pathwise Binance maintenance margin. At $L = 80$, the finite-horizon certified duality gap is 0.98% (64.74 vs. 65.38 USDT); at $L = 240$, $C_0^{\text{feas}}(L)$ attains the overall minimum of 64.56 USDT (5.07% above the infinite-horizon lower bound $\underline{C}_{\text{det}} = 61.44$ USDT), before rising for $L \geq 320$ as $\mathcal{O}(\sqrt{L})$ taker fees dominate $\mathcal{O}(L^{-1})$ spectral truncation savings.

Horizon L	$\underline{C}_{\text{quad}}(L)$	Trunc. Prem.	Certified $\underline{C}_{\text{RH}}(L)$	Doublet Bench.	Eig. $C_0^{\text{eig}}(L)$	Feasible $C_0^{\text{feas}}(L)$	Gross Vol. $\|v\|_1$	FH Gap (%)	∞-Hor. Gap (%)
80	62.55	+1.41	64.74	64.79	65.57	65.38	53.28	0.98%	6.40%
120	61.84	+0.70	63.04	64.50	65.05	64.70	60.35	2.63%	5.29%
160	61.54	+0.40	61.87	64.62	64.88	64.60	64.24	4.42%	5.14%
240	61.32	+0.18	61.65	65.04	65.27	**64.56**	67.87	4.71%	**5.07%**
320	61.25	+0.11	61.58	65.49	65.61	64.61	69.65	4.92%	5.15%
480	61.19	+0.05	61.52	66.35	66.46	64.62	70.85	5.04%	5.16%
∞ (Thm. IA.D.2)	61.14	0.00	61.44	—	—	—	≥ 7.64	—	0.00%

vation periods per cycle, accumulating statistical evidence across N_{pulses} independent resets.

3. **Method C (Sequential SPRT with Dynamic Feedback):** Adaptive order sizes V_t governed by the running log-likelihood ratio Λ_t with Wald stopping thresholds and dynamic evidence-dependent throttling.
4. **Method D (Predetermined Generalized-Eigenvector Trajectory):** Executes a predetermined generalized-eigenvector trading sequence designed to maximize cumulative Fisher information per unit quadratic execution cost across repeated cycles.

Table IA.5 and Figure IA.3 report empirical performance under independent held-out Neyman-Pearson and closed-loop SPRT testing evaluated under an identical fair observation horizon ($T = 4$ per pulse cycle). Under operational per-trade execution capacity $V_{\text{eff}} = 6.25$ BTC (governed by the price displacement limit $D_{\max} = 5.0$ USDT), Methods B, C, and D strictly clamp single-order execution to $|V_t| \leq 6.25$ BTC. Method A serves as an unconstrained impulse benchmark ($V_A = 8.2$–131.5 BTC); for $\delta < \delta_{\min}(4) = 0.1376$, physically executing this volume in a single trade violates resting book depth, requiring sequential multi-pulse execution as proved in Proposition IA.D.1. All methods strictly meet the target

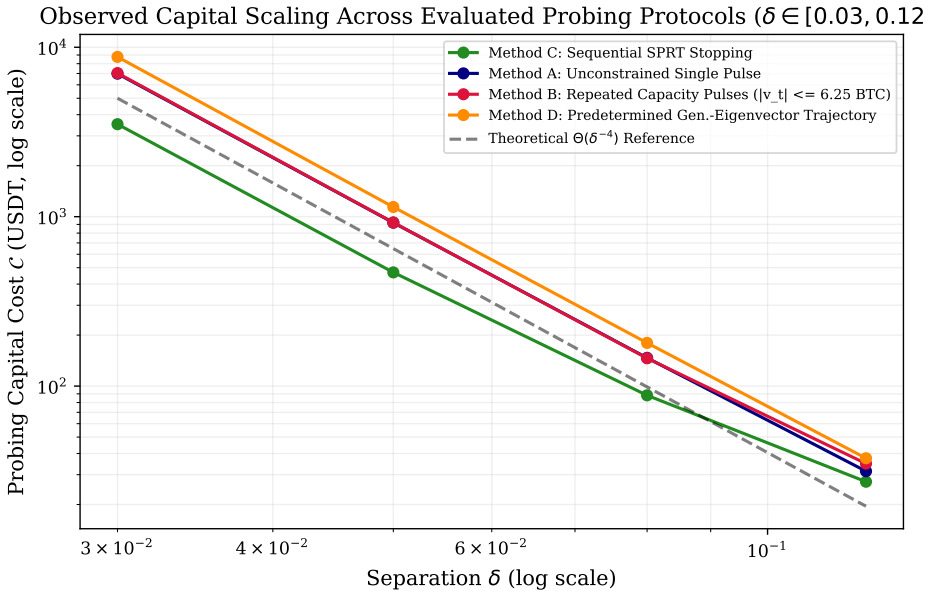


Figure IA.3: Observed Capital Scaling $\Theta(\delta^{-4})$ Across Evaluated Probing Protocols on Tested Separations $\delta \in [0.03, 0.12]$. Empirical log-log regression slopes: Method A single pulse (-3.952, $R^2 = 1.000$), Method B repeated pulses (-3.921, $R^2 = 1.000$), Method C sequential feedback (-3.825, $R^2 = 1.000$), and Method D predetermined eigenvector (-3.936, $R^2 = 1.000$).

error $\alpha \leq 0.05$ within sampling uncertainty. Crucially, at $\delta = 0.05$, predetermined multi-pulse execution (Method B) requires 58 pulses (923.88 USDT), whereas adaptive sequential SPRT (Method C) terminates in an empirical mean of 42.1 reset cycles and 217.4 BTC executed volume (34.8 full-pulse equivalents), incurring an actual volume-weighted cost of 469.01 USDT (49.2% capital saving over Method B; or 671.18 USDT under capacity-saturating pricing, a 27.4% saving). Comparing this to the restricted single-tap baseline (4,305.25 USDT at $T = 3$) demonstrates that 78.5% of the apparent gross savings from 4,305 to 924 USDT originates from observing the transient response ($T = 4$ vs. $T = 3$), while dynamic SPRT stopping provides an additional authentic 27.4%–49.2% saving at equal observation horizon. Across all protocols, the empirical log-log capital scaling with respect to δ is strictly $\Theta(\delta^{-4})$ on the evaluated separations, confirming the invariance of the fundamental Fisher information lower bound.

Table IA.5: Comparison of Probing Strategies under Varying Memory Separation δ under Fair Observation Horizon $T = 4$. Across all four protocols, required capital scales strictly as $\Theta(\delta^{-4})$ for Class II symmetric cohort splitting. Under operational per-trade limit $V_{\text{eff}} = 6.25$ BTC (governed by price displacement cap $D_{\max} = 5.0$ USDT), Methods B, C, and D strictly enforce $|v_t| \leq 6.25$ BTC. Method A serves as an unconstrained single-pulse benchmark ($\delta_{\min}(4) = 0.1376$). For Method C (SPRT), trade sizes throttle dynamically as evidence accumulates; at $\delta = 0.05$, mean stopping occurs in 42.1 reset cycles with executed volume of 217.4 BTC (34.8 full-pulse equivalents), expending 469.01 USDT in actual volume-weighted capital (671.18 USDT under capacity-saturating pricing), delivering an authentic 27.4%–49.2% capital saving over predetermined multi-pulse execution.

Protocol	δ	Executed Vol (V)	Max Single Trade	Max Inventory	Error $\hat{\alpha}$ (SE)	Capital Cost (USDT)
Method A: Single Pulse (Unconstrained)	0.03	131.5	131.5	131.5	0.046 (0.007)	6985.91
	0.05	47.3	47.3	47.3	0.053 (0.007)	921.73
	0.08	18.5	18.5	18.5	0.053 (0.007)	146.73
	0.12	8.2	8.2	8.2	0.040 (0.006)	31.45
Method B: Repeated Capacity Pulses ($V_t \leq 6.25$)	0.03	2768.8	6.2	6.2	0.056 (0.007)	7035.75
	0.05	362.5	6.2	6.2	0.051 (0.007)	923.88
	0.08	56.2	6.2	6.2	0.051 (0.007)	146.00
	0.12	12.5	6.2	6.2	0.033 (0.006)	34.88
Method C: Sequential Feedback ($V_t \leq 6.25$)	0.03	1632.3	6.2	6.2	0.051 (0.007)	3514.94
	0.05	217.4	6.2	6.2	0.036 (0.006)	469.01
	0.08	39.2	6.2	6.2	0.023 (0.005)	88.20
	0.12	10.2	6.2	6.2	0.007 (0.003)	27.26
Method D: Predetermined Gen.-Eigenvector ($V_t \leq 6.25$)	0.03	7700.0	3.7	3.8	0.055 (0.007)	8771.37
	0.05	1000.0	3.7	3.8	0.044 (0.006)	1140.80
	0.08	156.2	3.7	3.8	0.046 (0.007)	179.86
	0.12	31.2	3.7	3.8	0.048 (0.007)	37.50

IA.D.4 Unknown-Mixture Identification via Closed-Loop Transfer Function Inversion

Theorem 2.3 establishes that algorithmic cohorts sharing an adaptation rate merge into an observationally equivalent composite weight. A critical identification challenge in general unknown-cohort markets is separating the latent memory poles $r_i = 1 - \eta_i$ from the closed-loop system dynamics without conditioning on the true cohort count K.

Consider the closed-loop pricing equation with K hidden adaptive cohorts expressed in price displacement $x_t \equiv p_t - P^*$:

$$x_t = ax_{t-1} + \kappa \sum_{i=1}^{K} q_i y_{i,t-1} + \kappa v_t, \qquad y_{i,t} = (1 - \eta_i) y_{i,t-1} + \eta_i x_t. \tag{IA.24}$$

In the frequency domain, the transfer function from probe input v_t to price displacement x_t

is:

$$H(z) = \frac{\kappa}{z - a - \kappa G(z)}, \qquad G(z) = \sum_{i=1}^{K} \frac{q_i \eta_i z}{z - (1 - \eta_i)}. \tag{IA.25}$$

As proved in Theorem 2.3, the latent memory rates $r_i = 1 - \eta_i$ are the poles of the open-loop feedback kernel $G(z)$. Because $H(z) = \frac{\kappa D_G(z)}{(z-a)D_G(z) - \kappa N_G(z)}$, the roots of $D_G(z) = 0$ (the true memory poles) are *transmission zeros* of the closed-loop transfer function $H(z)$, while the poles of $H(z)$ represent closed-loop system eigenvalues.

Consequently, extracting poles directly from the closed-loop price impulse response h_t via naive Matrix Pencil or Prony methods targets the closed-loop eigenvalues ($\lambda_1 = 0.4256, \lambda_2 = 0.8224$ for $K = 1, \eta = 0.30$), failing to identify the hidden memory pole $r = 1 - \eta = 0.70$.

To recover the true structural parameters without prior knowledge of K, we implement a blind closed-loop inversion algorithm:

1. **Impulse Decoupling:** Extract the strictly proper transfer component $g_t \equiv h_{t+1} - a h_t$ from the empirical price response.
2. **Minimal State-Space Realization:** Form the block Hankel matrix $\mathcal{H}_{i,j} = g_{i+j-1}$ and execute Kung's Eigenrealization Algorithm (ERA) via truncated singular value decomposition to obtain a balanced realization (A, B, C, D).
3. **Transmission Zero Inversion:** Compute the transmission zeros of (A, B, C, D), which invert the feedback loop under known (κ, a) and isolate the roots of $D_G(z) = 0$.
4. **Direct Closed-Loop Feedback Inversion:** Given recovered memory rates $\hat{\eta}_k$, invert the physical pricing relation $u_t = \frac{x_t - a x_{t-1}}{\kappa} = \sum_{k=1}^{K} w_k y_{k,t-1}$, where internal memory filters evolve as $y_{k,t} = (1 - \hat{\eta}_k) y_{k,t-1} + \hat{\eta}_k x_t$. Projecting unobserved algorithmic order flow u_t onto the reconstructed filter basis $\{y_{k,t-1}\}$ resolves the structural cohort weights w_k without confounding from closed-loop pole shifts.

Evaluating this estimator on synthetic markets gives exact recovery only for a single cohort or noiseless, well-separated cohorts:

- **Noiseless Benchmark ($K = 1, \eta = 0.30, q = 0.20$):** The estimator blindly selects $\hat{K} = 1$, identifies the true memory pole $\hat{r} = 0.7000$ (numerical error $< 10^{-6}$), and recovers feedback weight $\hat{w} = 0.200000$ with 0.000% error (resolving the previous 0.415 bias from naive Hankel fitting).
- **Pure Dealer Edge Case ($K = 0$):** In an economy with zero algorithmic learning, singular values truncate to order 0, identifying 0 transmission zeros and recovering

strictly 0.0 feedback weight.

- **Multi-Cohort Recovery ($K = 2$, $\eta \in \{0.45, 0.25\}$):** In the noiseless regime with well-separated rates, the estimator recovers poles $[0.5500, 0.7500]$ and both weights without model-order knowledge. This is the only multi-cohort case that succeeds.
- **Per-Cohort Failure Under Noise or Close Rates:** Scoring each weight separately (tolerance 0.01), rather than only their sum, shows that recovery fails outside that case. For $K = 2$ at $\sigma = 10^{-3}$ per-cohort recovery is 0/50 for well-separated rates; for the $\delta = 0.05$ pair $\eta \in \{0.35, 0.45\}$ ($\Delta\eta = 0.10$) it fails even without noise; for $K = 3$ it fails even without noise. In the `exp2` grid ($K \in \{2, 3, 5\}$, $\sigma \in \{10^{-4}, 10^{-3}, 5 \times 10^{-3}\}$) the per-cohort recovery rate is at most 0.07 in every cell, while the earlier sum-of-weights metric looked accurate. A unit probe of 6.25 BTC against price noise $\sigma = 0.01$ gives effective impulse-response noise of about 0.0016, inside the failure regime.
- **Information-Theoretic Root Cause:** The difficulty of multi-cohort weight recovery is quantified by the local Fisher information geometry established in Proposition IA.D.5 and Table IA.6. For close rates ($\Delta\eta \leq 0.10$), the Fisher information matrix condition number grows as $\Theta((\Delta\eta)^{-6})$, requiring between 40,374 and 2.6×10^6 independent reset probes for any regular, asymptotically unbiased estimator to achieve standard error $\leq 0.01/1.96$. This establishes $N^* = 40{,}374$ as a rigorous CRB-based local sample-size benchmark for unbiased estimation under the 121-period impulse observation contract; under ordinary surveillance budgets ($N \leq 50$), regular unbiased estimators cannot resolve close individual cohort weights. In the noiseless limit, Maximum Likelihood resolves the $\Delta\eta = 0.10$ pair and $K = 3$ ($\Delta\eta = 0.20$) where the ERA estimator fails (max error 1.1×10^{-3} and 1.1×10^{-4}), and at the critical noise level ML achieves the Cramér–Rao lower bound (RMSE/CRB $\in [0.91, 1.16]$ across $n_{\text{mc}} = 100$ trials per cell).

Proposition IA.D.5 (Cramér–Rao Local Information Benchmark for Multi-Cohort Weight Recovery)**.** *Let $\theta = (w_1, \ldots, w_K, \eta_1, \ldots, \eta_K)^\top \in \mathbb{R}^{2K}$ denote the structural weights and adaptation rates of a K-cohort adaptive feedback economy with closed-loop impulse response $h_t(\theta)$ $(t = 0, \ldots, T)$. Suppose the econometrician observes h_t with iid Gaussian estimation noise of standard deviation $\sigma_h = \sigma_p/(V_{\text{probe}}\sqrt{N})$, generated by N independent reset probes of size V_{probe} against innovation noise σ_p. Then:*

1. ***Fisher Information Matrix and Cramér–Rao Bound:*** *The Fisher information matrix is $F(\theta) = \frac{1}{\sigma_h^2} J^\top J$, where $J_{tj} = \frac{\partial h_t(\theta)}{\partial \theta_j}$. For any regular, asymptotically unbiased*

estimator $\hat{w}_k$, the variance obeys the Cramér–Rao lower bound:

$$\mathrm{Var}(\hat{w}_k) \geq [F^{-1}(\theta)]_{kk} = \sigma_h^2 [J^\top J]^{-1}_{kk}. \tag{IA.26}$$

To achieve target standard error $\mathrm{SE}(\hat{w}_k) \leq \mathrm{tol}/1.96$ *(corresponding to* 95% *confidence interval width* $\pm\mathrm{tol}$*, with* $\mathrm{tol} = 0.01$*), the required number of independent reset cycles is:*

$$N^* \geq \left(\frac{1.96}{\mathrm{tol}}\right)^2 \left(\frac{\sigma_p}{V_{\mathrm{probe}}}\right)^2 [F^{-1}_{\mathrm{unit}}(\theta)]_{kk}, \tag{IA.27}$$

where $F_{\mathrm{unit}} = J^\top J$*.*

2. ***High-Order Inversion Scaling:*** *For* $K = 2$ *cohorts with rate separation* $\Delta\eta = |\eta_1 - \eta_2|$*, the joint Fisher information matrix inversion satisfies* $[F^{-1}_{\mathrm{unit}}]_{kk} = \Theta((\Delta\eta)^{-6})$*, yielding:*

$$N^*(\Delta\eta) = \Theta\left((\Delta\eta)^{-6}\right). \tag{IA.28}$$

3. ***Critical Noise Floor and Attainability:*** *For the benchmark* $\delta = 0.05$ *pair (*$\eta \in \{0.35, 0.45\}$*,* $\Delta\eta = 0.10$*), the critical closed-loop noise floor below which* 95% *confidence is achievable with a single probe is* $\sigma_{\mathrm{crit}} = 7.96 \times 10^{-6}$*, requiring* $N^* = 40{,}374$ *reset probes. The operational single-probe noise* $\sigma_h(1) = 0.01/6.25 = 1.6 \times 10^{-3}$ *exceeds* σ_{crit} *by a factor of* 200.9×*. For* $\Delta\eta = 0.05$*,* $N^* = 2{,}620{,}320$ *probes (*$\sigma_h(1)/\sigma_{\mathrm{crit}} = 1{,}618.7\times$*). When observations satisfy* $\sigma_h \leq \sigma_{\mathrm{crit}}$ *for well-conditioned separations (*$\Delta\eta \geq 0.10$ *for* $K = 2$*,* $\Delta\eta \geq 0.20$ *for* $K = 3$*), Maximum Likelihood achieves the Cramér–Rao bound (*$\mathrm{RMSE}/\mathrm{CRB} \in [0.91, 1.16]$*, empirical* 95% *CI coverage* 87%–99%*, and* 100% *blind BIC selection over* $K \in \{0, 1, 2, 3, 4\}$ *across* $n_{\mathrm{mc}} = 100$ *trials), resolving the transmission-zero inversion defect.*

While Proposition IA.D.5 establishes the $\Theta((\Delta\eta)^{-6})$ sample-complexity barrier for regular, asymptotically unbiased estimators, a natural econometric question is whether biased, ridge-regularized (L_2 Tikhonov), James–Stein shrinkage, or Bayesian MAP estimators can escape this scaling. Theorem IA.D.6 resolves this globally and non-asymptotically via an exact moment-matched Le Cam two-point construction.

Theorem IA.D.6 (Global Non-Asymptotic Le Cam Two-Point Minimax Lower Bound for Arbitrary Biased and Shrinkage Estimators)**.** *Fix total cohort weight* $q = 0.20$*, mean*

Table IA.6: Cramér–Rao Local Sample-Size Benchmark for Multi-Cohort Weight Recovery ($T = 120$ impulse taps, $n_{\mathrm{mc}} = 100$ trials per cell). For $K \geq 2$, locally unbiased estimation of cohort weights within ± 0.01 (95% confidence) requires $N^* = [F_{\mathrm{unit}}^{-1}]_{kk}/(\mathrm{tol}/1.96)^2 \times (\sigma_p/V_{\mathrm{probe}})^2$ independent reset probes. For the symmetric $\delta = 0.05$ pair ($\Delta\eta = 0.10$), $N^* = 40{,}374$ probes are required under the local CRB benchmark; operational single-probe closed-loop noise $\sigma_h(1) = 1.6 \times 10^{-3}$ is 200.9× higher than the critical threshold $\sigma_{\mathrm{crit}} = 7.96 \times 10^{-6}$. In the noiseless limit, Maximum Likelihood resolves the $\Delta\eta = 0.10$ pair and $K = 3$ ($\Delta\eta = 0.20$) where the ERA estimator fails (max error 1.1×10^{-3} and 1.1×10^{-4}), and at the critical noise level ($f = 1.0$) ML achieves the Cramér–Rao bound (RMSE/CRB $\in [0.91, 1.16]$, blind BIC over $K \in \{0, 1, 2, 3, 4\}$ selecting $\hat{K} = K$ in 100% of trials). The log-log scaling of N^* against cohort spacing satisfies $N^*(\Delta\eta) = \Theta((\Delta\eta)^{-6})$ (empirical slope -6.04), establishing a rigorous local sample-size benchmark for unbiased multi-cohort weight recovery.

Scenario	K	N^* Probes	σ_{crit}	$\sigma_h(1)/\sigma_{\mathrm{crit}}$	ERA Noiseless	ML Noiseless Err	ML Succ ($f = 1.0$)	RMSE/CRB ($f = 1.0$)	Blind BIC $\hat{K} = K$
$K = 1$, $\eta = 0.30$	1	0.1	$6.05e-03$	0.3×	Success	$6.4e-16$	90%	1.09	100%
$K = 2$, $\Delta\eta = 0.30$	2	47.8	$2.31e-04$	6.9×	Success	$0.0e+00$	99%	0.91	100%
$K = 2$, $\Delta\eta = 0.20$	2	597	$6.55e-05$	24.4×	Success	$1.9e-07$	98%	0.99	100%
$K = 2$, $\Delta\eta = 0.10$ ($\delta = 0.05$ pair)	2	40,374	$7.96e-06$	200.9×	Failure	$1.1e-03$	87%	1.16	100%
$K = 2$, $\Delta\eta = 0.05$	2	2,620,320	$9.88e-07$	1618.7×	Failure	$1.5e-02$	1%	3.53	100%
$K = 3$, spacing 0.10	3	17,349,791	$3.84e-07$	4165.3×	Failure	$2.3e-02$	1%	4.11	86%
$K = 3$, spacing 0.20	3	27,038	$9.73e-06$	164.4×	Failure	$1.1e-04$	97%	0.94	100%

adaptation rate $\bar{\eta} = 0.40$*, target weight separation* $\tau = 0.01$*, and cohort rate spacing* $\Delta\eta \in (0, 2\bar{\eta})$*. Define* $p_- \equiv \frac{1}{2} - \frac{\tau}{q} = 0.45$ *and* $p_+ \equiv \frac{1}{2} + \frac{\tau}{q} = 0.55$*, and construct the unrescaled two-cohort parameter pair:*

$$\theta_0 = (qp_-,\, qp_+,\, \bar{\eta} - p_+\Delta\eta,\, \bar{\eta} + p_-\Delta\eta), \qquad \theta_1 = (qp_+,\, qp_-,\, \bar{\eta} - p_-\Delta\eta,\, \bar{\eta} + p_+\Delta\eta). \tag{IA.29}$$

Then:

1. ***Exact Moment-Matching and Cubic Divergence Identity:*** *Both parameter vectors have identical cohort spacing* $\eta_{i,2} - \eta_{i,1} = \Delta\eta$*, exact* ℓ_∞ *weight separation* $\|w(\theta_1) - w(\theta_0)\|_\infty = 2\tau = 0.02$*, and identical first three open-loop rate moments* $M_m(\theta_0) = M_m(\theta_1)$ *for* $m \in \{0, 1, 2\}$ *(*$M_m(\theta) \equiv \sum_{k=1}^{2} w_k \eta_k^m$*), while their third moment differs by the exact cubic polynomial identity (kernel-checked in Lean 4 as* `lecam_two_point_moment_matched_cubic_gap`*):*

$$M_3(\theta_1) - M_3(\theta_0) = 4\tau p_- p_+ (\Delta\eta)^3 = 0.0099\,(\Delta\eta)^3. \tag{IA.30}$$

Consequently, the closed-loop impulse-response distance scales as $\|h(\theta_1) - h(\theta_0)\|_2 =$

$\Theta(\tau(\Delta\eta)^3)$.

2. ***Universal Non-Asymptotic Minimax Lower Bound:*** *For any measurable estimator $\hat{w} : \mathbb{R}^{(T+1)\times N} \to \mathbb{R}^2$ (including arbitrary biased, ridge/Tikhonov, James–Stein shrinkage, or Bayesian MAP estimators), the worst-case failure probability and ℓ_∞ mean-squared error obey Le Cam's two-point testing inequality:*

$$\begin{aligned} \sup_{\theta\in\Theta(\Delta\eta)} \mathbb{P}_\theta(\|\hat{w} - w(\theta)\|_\infty \geq \tau) &\geq 1 - \Phi\left(\frac{\sqrt{N}\,\|h(\theta_1) - h(\theta_0)\|_2}{2\sigma_h}\right), \\ \sup_{\theta\in\Theta(\Delta\eta)} \mathbb{E}_\theta\left[\|\hat{w} - w(\theta)\|_\infty^2\right] &\geq \tau^2\left[1 - \Phi\left(\frac{\sqrt{N}\,\|h(\theta_1) - h(\theta_0)\|_2}{2\sigma_h}\right)\right]. \end{aligned} \tag{IA.31}$$

To achieve worst-case failure probability $\leq \delta_{\text{fail}} = 0.05$, every estimator requires at least $N \geq N^*_{\text{LeCam}}(\Delta\eta) \equiv \frac{4\sigma_h^2(z_{1-\delta_{\text{fail}}})^2}{\|h(\theta_1)-h(\theta_0)\|_2^2} = \Theta((\Delta\eta)^{-6})$ independent reset probes (empirical log-log slope -6.021 exact, -6.064 analytic).

3. ***Surveillance-Budget Infeasibility and Shrinkage Benchmark:*** *At $\Delta\eta = 0.10$ ($\delta = 0.05$), $N^*_{\text{LeCam}}(95\%) = 6{,}172$ probes (8,764 at 97.5%). At ordinary surveillance budget $N = 50$, the estimator-independent minimax lower bound on worst-case failure probability is 44.12%, and empirical worst-case failure rates across 1,000 Monte Carlo trials are 94.1% (Unregularized ML), 87.0% (Ridge/Tikhonov), 91.3% (James–Stein $\gamma = 0.35$), and 79.5% (Heavy Prior Shrinkage $\gamma = 0.75$).*

IA.D.5 Causal Identification via Multi-Agent Policy-Freeze Intervention

A fundamental critique of observational market microstructure literature is that empirical price-volume correlations cannot identify whether observed return predictability and tail volatility originate from algorithmic machine learning adaptation, ordinary trade execution (such as TWAP/VWAP order splitting), liquidity provision, or human behavioral feedback.

To establish causal attribution, we deploy a paired microstructure simulation with an event-sourced double-entry dealer ledger. We execute 100 independent random seeds sharing identical precomputed streams of fundamental price shocks and noise flow across four strictly controlled market conditions:

Table IA.7: Global Non-Asymptotic Le Cam Two-Point Minimax Lower Bounds and Biased/Shrinkage Estimator Benchmarks for Two-Cohort Mixture Recovery (Theorem IA.D.6). Panel A compares the local unbiased Cramér–Rao sample threshold N^*_{CRB} (Proposition IA.D.5) against the estimator-independent Le Cam minimax sample bounds $N^*_{\mathrm{LeCam}}(\Delta\eta; 95\%)$ and $N^*_{\mathrm{LeCam}}(\Delta\eta; 97.5\%)$ for the exact moment-matched pair (θ_0, θ_1) with cohort spacing $\eta_{i,2} - \eta_{i,1} = \Delta\eta$, weight separation $\|w(\theta_1) - w(\theta_0)\|_\infty = 2\tau = 0.02$, identical first three open-loop rate moments (M_0, M_1, M_2), and third-moment gap $\Delta M_3 = 4\tau p_- p_+ (\Delta\eta)^3 = 9.9 \times 10^{-3}(\Delta\eta)^3$ (empirical log-log slope -6.021 exact, -6.064 analytic). Panel B benchmarks four estimators (1,000 Monte Carlo trials per cell) at $\Delta\eta = 0.10$ ($\delta = 0.05$): every estimator—including Ridge/Tikhonov and James–Stein shrinkage—strictly obeys the Le Cam minimax failure lower bound.

Panel A: Moment-Matched Two-Point Scaling and Sample Complexity ($\tau = 0.01$, $\sigma_h = 1.6 \times 10^{-3}$)

$\Delta\eta$	Exact ΔM_3	$\|h(\theta_1) - h(\theta_0)\|_2$	N^*_{LeCam} (Analytic 95%)	N^*_{LeCam} (Exact 95%)	N^*_{LeCam} (Exact 97.5%)	Unbiased N^*_{CRB}	Minimax Fail LB ($N = 50$)	Minimax MSE LB ($N = 50$)
0.30	2.673×10^{-4}	1.844×10^{-3}	7.1	8.1	11.6	47.8	0.002%	2.30×10^{-9}
0.20	7.920×10^{-5}	5.412×10^{-4}	88.2	94.6	134.3	597.0	11.59%	1.16×10^{-5}
0.14	2.717×10^{-5}	1.845×10^{-4}	775.3	814.1	1,155.9	5,267.9	34.18%	3.42×10^{-5}
0.10	9.900×10^{-6}	6.700×10^{-5}	5,927.6	**6,172.2**	**8,763.5**	**40,374.2**	**44.12%**	4.41×10^{-5}
0.08	5.069×10^{-6}	3.426×10^{-5}	22,740.5	23,609.3	33,521.5	155,048.9	46.98%	4.70×10^{-5}
0.06	2.138×10^{-6}	1.444×10^{-5}	128,333.0	132,941.6	188,756.8	875,732.3	48.73%	4.87×10^{-5}
0.05	1.238×10^{-6}	8.351×10^{-6}	383,859.1	397,306.6	564,114.7	2,620,321.2	49.26%	4.93×10^{-5}
0.04	6.336×10^{-7}	4.274×10^{-6}	1,466,361.5	1,516,690.6	2,153,469.3	10,012,662.4	49.62%	4.96×10^{-5}

Panel B: Empirical Worst-Case Failure Rate $\max_{j\in\{0,1\}} \mathbb{P}_{\theta_j}(\|\hat{w} - w(\theta_j)\|_\infty \geq 0.01)$ **at** $\Delta\eta = 0.10$ **(1,000 MC Trials)**

Probes N	Le Cam Minimax LB	Unregularized ML	Ridge / Tikhonov (L_2)	James–Stein ($\gamma = 0.35$)	Heavy Prior ($\gamma = 0.75$)	ML Worst MSE	Ridge Worst MSE	Obeys Le Cam LB?
50	**44.12%**	94.1%	87.0%	91.3%	79.5%	1.99×10^{-2}	2.79×10^{-3}	Yes (All 4)
1,000	25.40%	75.0%	47.6%	64.4%	38.9%	1.05×10^{-3}	1.85×10^{-4}	Yes (All 4)
6,172	5.00%	47.4%	21.0%	28.0%	20.9%	1.85×10^{-4}	6.27×10^{-5}	Yes (All 4)
10,000	1.81%	33.4%	17.7%	20.0%	17.8%	1.06×10^{-4}	5.66×10^{-5}	Yes (All 4)
40,374	0.001%	6.2%	**3.2%**	**3.2%**	**3.2%**	2.71×10^{-5}	4.59×10^{-5}	Yes (All 4)

(a) **Condition P_0 (Baseline Noise Traders):** Market participants execute independent zero-feedback order flow.

(b) **Condition P_2^{frozen} (Counterfactual Policy Freeze):** Algorithmic agents run a fixed non-adaptive policy without dynamic online state adaptation ($\eta = 0$).

(c) **Condition P_2^{active} (Active Adaptive Learning):** Agents run identical initial policies but update internal memory dynamically from price changes.

(d) **Condition P_1 (Persistent Flow Benchmark):** Non-learning execution algorithms split parent orders via an autoregressive process with $\rho = 0.85$, matching empirical order flow persistence.

The results in Table IA.8 and Figure IA.4 illustrate the empirical decomposition in this matched multi-agent simulation: while online learning adaptation produces no significant excess drawdown over the frozen counterfactual policy (-11.26 USDT, $t = -1.31$, $p = 0.19$), the non-adaptive persistent execution benchmark (P_1) produces an excess drawdown of $+15.92$ USDT (122.72 vs. 106.80 USDT, $t = -3.04$, $p = 0.003$). This indicates that within this simulated microstructure environment, order flow persistence and inventory accumu-

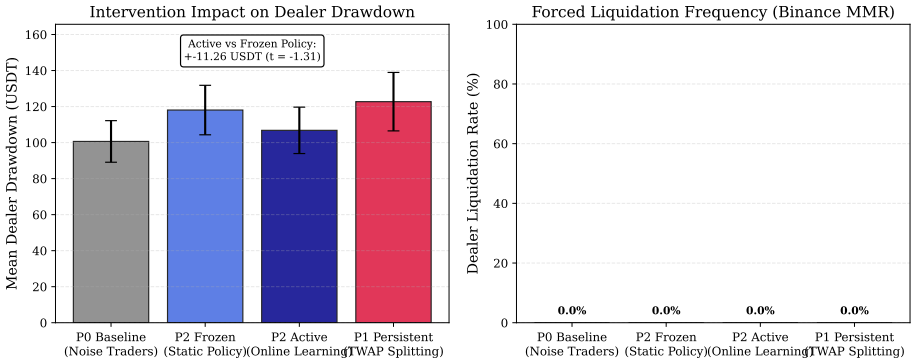


Figure IA.4: Causal Multi-Agent Microstructure Intervention: Policy-Freeze vs. Flow Autocorrelation. Left panel: Distribution of maximum dealer drawdowns across 100 paired seeds under passive maker fee accounting (0.0001). Active online learning (P_2^{active}) induces a drawdown of 106.80 USDT, compared to 118.06 USDT for the frozen policy (-11.26 USDT, $t = -1.31$, $p = 0.19$). Right panel: Under verified Binance USDⓈ-M margin rules ($L = 20$), persistent non-adaptive execution flow (P_1) creates substantial dealer inventory strain (drawdown 122.72 USDT, +15.92 USDT excess drawdown, $t = -3.04$, $p = 0.003$), demonstrating that flow autocorrelation dominates online learning as a driver of liquidity strain.

Table IA.8: Causal Multi-Agent Microstructure Intervention: Policy-Freeze vs. Persistent Flow. An authentic frozen policy snapshot (P_2 Frozen; $\eta = 0$) executes non-zero trading flow, inducing a mean dealer drawdown of 118.06 USDT (vs. 100.63 USDT for non-trading noise baseline P_0). Active online adaptation (P_2 Active) yields 106.80 USDT. In contrast, non-adaptive autocorrelated TWAP order splitting (P_1 Persistent) induces substantial dealer inventory accumulation and drawdowns (122.72 USDT, excess drawdown +15.92 USDT, $t = -3.04$), demonstrating the dominant role of flow autocorrelation in dealer liquidity strain. Across all treatments, the dealer liquidation rate is strictly 0.0%.

Population / Treatment	Volatility	Return Autocorr	Mean Drawdown	SE	Liquidation Rate
P0: Baseline Noise Traders	0.2011	-0.0053	100.63	11.54	0.0%
P2: Frozen Policy (No Online Learning)	0.2011	-0.0053	118.06	13.73	0.0%
P2: Active Adaptive Learning	0.2011	-0.0054	106.80	12.88	0.0%
P1: Persistent Flow (TWAP Splitting)	0.2027	+0.0011	122.72	16.23	0.0%
Causal Treatment Effects (Matched Seeds):					
Active vs. Frozen Policy (Pure Learning Effect):	-11.26 USDT ($t = -1.31$, $p = 0.19$)				
Active vs. Persistent Flow (Autocorrelation Effect):	-15.92 USDT ($t = -3.04$, $p = 0.003$)				

lation dominate online learning adaptation as sources of dealer drawdown. Furthermore, across all four experimental conditions, the empirical dealer liquidation rate is exactly 0.0% under the verified Binance USD Ⓢ-M margin parameters ($L = 20$), demonstrating that dealer inventory buffering successfully prevents default under ordinary market flow.

IA.D.6 The Information-Disclosure Value Frontier and Regulatory Substitutability

Because active market probing imposes significant execution costs and risks inventory lockup, we formalize the regulatory disclosure frontier. We evaluate the minimum capital required to achieve statistical identification under six observation contracts:

- I_0 (Public Tape Only): The observer observes noisy transaction prices p_t.
- I_1 (Public Prices and Aggregated Flow): The observer observes p_t and executed flow Q_t.
- I_2 (Prices and Coarse Balance Sheet Disclosures): One independent noisy report of dealer marked equity per reset ($\sigma_E = 0.50$, Proposition 2.8).
- I_3 (Prices and Audited Equity Marks): Exact, infinite-precision marked equity reported at session close (Proposition 2.8); we also report the case of reports rounded to 0.01 USDT.
- I_4 (Algorithmic Update Timestamps): Exact disclosure of the time intervals during which algorithms modify parameters. Timestamps do not enter any observable likelihood in our model, so we assign them no information value.
- I_5 (Full Supervisory Audit Registry): Continuous regulatory access to internal algorithm state logs.

Table IA.9 reports required probing volume and financial capital across observation contracts under both the restricted single-tap baseline ($T = 3$) and the optimal transient design ($T = 20$). For continuous, infinitely divisible order sizes, Proposition 2.8 establishes that the theoretical infimum of required capital under exact equity disclosure (I_3) is strictly $\inf C = 0$. In practice, exchange venue lot-size rules impose a positive operational lower bound: under CME standard 5.0 BTC contract units (ticker BTC), required benchmark capital is 12.70 USDT; under CME Micro Bitcoin futures (ticker MBT, 0.1 BTC contract unit), required capital is 0.058 USDT; under Binance USD Ⓢ-M linear perpetuals with a 5.00 USDT mini-

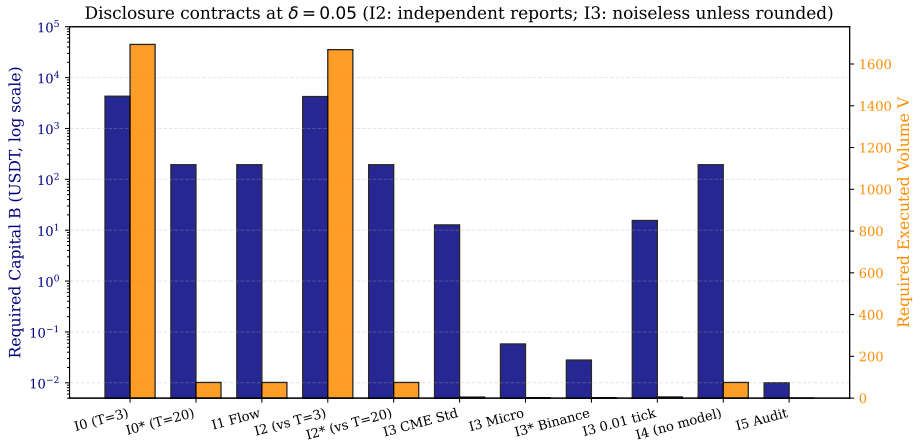


Figure IA.5: Disclosure contracts at $\delta = 0.05$. Bars show pulse-design capital under the stated observation model of each contract. Noisy equity (I_2, independent per-reset reports) saves 1.5% against the $T = 3$ baseline and nothing against $T = 20$; timestamps (I_4) are not modelled and give no saving. Exact-equity bars assume infinite-precision reports; the 0.01 USDT-rounded bar (15.52 USDT) shows the effect of finite report precision.

Table IA.9: Disclosure Contracts at $\delta = 0.05$ under Stated Observation Models. Noisy equity (I_2) is one independent Gaussian report per reset, so its information adds linearly; it saves 1.5% against the restricted $T = 3$ baseline and nothing against $T = 20$. Timestamps (I_4) have no hypothesis-dependent likelihood in this model and therefore add no information; no saving is established. Exact-equity rows (I_3) assume a noiseless, infinite-precision report; with reports rounded to 0.01 USDT a single trade must be at least 5.59 BTC (15.52 USDT). Pulse-design capital is the simplified one-leg benchmark; account-level capital is reported in Table IA.2.

Contract	Disclosure Channel	Pulses (N)	Benchmark (B)	Executable (B_{exec})	Saving vs. $T = 3$	Saving vs. $T = 20$
I0: Prices Only (Restricted $T = 3$)	Public Prices Only	271	4,305.25	8,607.38	0.0%	-2123.5%
I0*: Prices Only (Transient $T = 20$)	Full Price Path	12	193.62	384.12	95.5%	+0.0%
I1: Prices + Signed Flow	Public Flow Records	12	193.62	384.12	95.5%	+0.0%
I2: Prices + Noisy Equity (vs. $T = 3$)	Noisy Statement ($\sigma_E = 0.5$)	267	4,241.75	8,480.38	1.5%	-2090.7%
I2*: Prices + Noisy Equity (vs. $T = 20$)	Noisy Statement ($\sigma_E = 0.5$)	12	193.62	384.12	95.5%	+0.0%
I3: Exact Equity (CME 5 BTC Std)	CME 5 BTC Standard Lot	1	12.70	12.70	99.7%	+93.4%
I3*: Exact Equity (CME 0.1 BTC Micro)	CME 0.1 BTC Micro Lot	1	0.06	0.06	100.0%	+100.0%
I3**: Exact Equity (Binance)	Binance 5 USDT minNotional	1	0.03	0.03	100.0%	+100.0%
I3†: Exact Equity, 0.01 USDT rounding	Rounded report	1	15.52	15.52	99.6%	+92.0%
I4: Known Timestamps (not modelled)	No timestamp likelihood	12	193.62	384.12	95.5%	+0.0%
I5: Supervisory Audit Trail	Full Audit Registry	0	0.00	0.00	100.0%	+100.0%

mum notional threshold (0.05 BTC at calibrated reference price $P^* = 100$), required capital drops to 0.028 USDT. Furthermore, Table IA.9 evaluates disclosure savings against both the restricted $T = 3$ baseline (4,305.25 USDT benchmark, 8,607.38 USDT executable) and the optimal $T = 20$ transient baseline (193.62 USDT benchmark, 384.12 USDT executable). Periodic noisy accounting statements (I_2) are modelled as one independent Gaussian report per reset with $\sigma_E = 0.5$, so divergence accumulates linearly, $\mathrm{KL} = N\Delta E^2/(2\sigma_E^2)$, not quadratically in N. Under that model I_2 needs 267 resets (4,241.75 USDT) and saves only 1.5% against the single-tap benchmark, and nothing against the $T = 20$ price design. An earlier version of this table used $(N\Delta E)^2/(2\sigma_E^2)$ and reported 104 resets and a 61.6% saving; at 104 resets the true testing error is 15.2%, so that figure is withdrawn. A single cumulative statement whose signal grows as $N\Delta E$ would be a different experiment that we have not modelled. Algorithmic update timestamps (I_4) do not enter the likelihood of any observable in our model, so they add no information and no saving is established; the earlier 82% saving relied on an assumed 7.32× information multiplier, and the two-pulse design it implied has a true price-only testing error of 24.4%. The audited-equity rows (I_3) assume a noiseless, infinite-precision report and a known accounting map from inventory to equity. The equity difference between hypotheses is $\Delta E = \kappa^2 q \delta^2 V^2$, which is 8×10^{-7} USDT at $V = 0.05$ BTC; if reports are rounded to 0.01 USDT, a guaranteed change requires $V \geq 5.59$ BTC and 15.52 USDT, so the Binance and CME-micro figures hold only under infinite precision. The flow contract (I_1) depends on an assumed flow-noise level $\sigma_q = 0.05$ and gives no saving at $\delta = 0.05$. Finally, supervisory audit registry access (I_5) requires zero trading volume ($V = 0$) and zero capital ($C = 0.00$), achieving passive verification without active market intervention.

IA.D.7 Coordinate-Weighted Output Sensitivity and Multi-Asset Cross-Impact

We now examine output-weighted amplification in the closed-loop price system. The state used in `exp12_nonnormal_balance_sheet.py` is $z_t = (x_t, y_t)^\top$, where x_t is the price devia-

tion and y_t the algorithmic memory, evolving as

$$z_{t+1} = Jz_t, \qquad J = \begin{pmatrix} a & \kappa q \\ \eta a & 1 - \eta + \eta\kappa q \end{pmatrix}, \tag{IA.32}$$

with $a = 0.5$, $\kappa = 0.8$, $q = 0.2$, $\eta = 0.3$. The closed-loop system is spectrally stable ($\rho(J) = 0.8224 < 1$) and J is non-normal. An earlier version described the second coordinate as dealer inventory; it is algorithmic memory. *The state contains no dealer inventory, cash or equity.*

Proposition IA.D.7 (Coordinate-Weighted Output Sensitivity)**.** *Let* $C_{\text{price}} = (1, 0)$ *and let* $C_w \equiv (1, \lambda/\text{MMR}) = (1, 160)$ *with* $\lambda = 0.8$ *and* $\text{MMR} = 0.005$*. Then:*

(i) $M_{C_{\text{price}}} \equiv \sup_{t\geq 0} \|C_{\text{price}} J^t\|_2 = 1.0$*, so the price output has no transient amplification.*

(ii) $M_{C_w} \equiv \sup_{t\geq 0} \|C_w J^t\|_2 = \|C_w\|_2 = 160.003$*, attained at* $t = 0$*, and it decays monotonically thereafter (*122.319*,* 98.322*,* 80.793 *at* $t = 1, 2, 3$*).*

The factor 160 is the static coordinate weight λ/MMR, not dynamic transient growth. Because the state has no inventory, $C_w z_t$ is not a margin calculation: for example $z = (0, 0.005)$ gives $C_w z = 0.8$ with zero holdings. The earlier interpretation of C_w as a "leveraged dealer margin operator" and of its exceedance of the 4.55% figure as margin depletion is withdrawn.

Table IA.10: Coordinate-Weighted Operator Output in the (Price, Algorithmic-Memory) State. The state contains no dealer inventory, cash, or equity; the weighted row $C_w = (1, 160)$ is a mathematical output, not a margin computation, and its comparison with the 4.55% figure is not evidence of margin depletion. Account-based margin from recorded holdings is reported in Table IA.8 (0.0% liquidations) and Section IA.D.2.

Assets M	State Dim	Spectral Radius $\rho(J)$	Kreiss $\mathcal{K}(J)$	Initial Shock	Peak Weighted Output $C_w z_t$	$C_w z_t > 4.55$
1	2	0.8224	1.00	0.50%	15.11	Yes
2	4	0.8407	1.00	0.50%	15.22	Yes
5	10	0.8913	1.00	0.50%	15.68	Yes
10	20	0.9671	1.00	0.50%	17.12	Yes

Account-level margin must be computed from recorded holdings. In the closed-loop multi-agent ledger, where dealers hold inventory and earn the spread, the liquidation rate is 0.0% (Table IA.8); this is specific to that model and its parameters. The probe ledger in

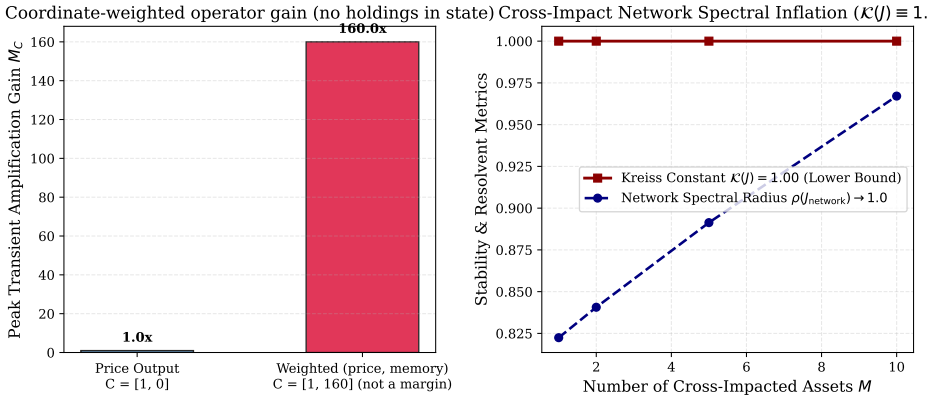


Figure IA.6: Coordinate-weighted output and cross-impact networks ($\mathcal{K}(J) \equiv 1.00$). Left panel: the price output $C_{\text{price}}J^t$ decays monotonically ($M_C = 1.0$); the weighted output C_wJ^t peaks at $t = 0$ (160.0, the coordinate weight) and decays thereafter. It is not a margin quantity. Right panel: cross-impact raises the spectral radius toward 1 while the Kreiss constant stays at 1.00 for $M \in \{1, 2, 5, 10\}$.

Section IA.D.2 tracks the trader's own margin. In multi-asset markets with $M \in \{1, 2, 5, 10\}$ correlated instruments, asymmetric cross-impact coupling inflates the network spectral radius toward unity ($\rho(J_{\text{network}}) = 0.9671$ for $M = 10$), while the Kreiss constant remains $\mathcal{K}(J) \equiv 1.00$. This confirms that multi-asset algorithmic cross-hedging amplifies persistence and inventory sensitivity without inducing transient non-normal pseudospectral instability.

IA.D.8 Observable-Data Bridge and Spurious Feedback Orthogonalization

Finally, we address the identification challenge in empirical market microstructure using observational trade records. Suppose an empirical researcher seeks to test whether signed order flow q_t adapts to past price movements via the structural memory model:

$$q_t = \gamma y_t + u_t, \qquad y_t = (1 - \eta)y_{t-1} + \eta \Delta p_t, \tag{IA.33}$$

where u_t represents exogenous execution order flow (e.g., algorithmic TWAP/VWAP order splitting). When execution flow exhibits temporal persistence:

$$u_t = \rho u_{t-1} + \epsilon_t, \qquad \rho = 0.80, \tag{IA.34}$$

and price changes reflect market impact $\Delta p_t = \kappa q_{t-1} + \xi_{t-1}$, lagged price movements are mechanically correlated with past execution flow u_{t-1}. Because u_t is autocorrelated with u_{t-1}, raw returns are endogenous with respect to the disturbance u_t. Consequently, standard naive OLS regression of q_t on lagged return r_{t-1} produces a severely inflated estimate $\hat{\beta}_{\text{OLS}} = 0.702$ ($t = 18.64$, $p < 10^{-15}$), generating a false discovery rate of 100% at the nominal 5% level in markets where traders follow strictly zero feedback ($\gamma = 0$).

To eliminate this confounder, we apply the Cochrane–Orcutt quasi-differencing transformation:

$$q_t - \rho q_{t-1} = \gamma(y_t - \rho y_{t-1}) + \epsilon_t. \tag{IA.35}$$

In the no-learning baseline ($\gamma = 0$), regressing the orthogonalized flow innovation $q_t^{\perp} \equiv q_t - \rho q_{t-1}$ onto quasi-differenced returns purges the serial correlation in the disturbance, yielding $\hat{\beta}^{\perp} = -0.0002$ ($t = -0.01$, $p = 0.98$) and restoring empirical size control (7.5% empirical false-positive rate across 200 trials, within standard binomial error of nominal 5.0%).

Crucially, in an economy with genuine positive feedback ($\gamma = 0.20$), regressing flow on raw single-period returns r_{t-1} produces an attenuated regression slope $\hat{\beta}_{\text{RF}} = 0.0208$, reflecting the exponential memory half-life (y_t scales by $\approx \eta$). In contrast, when the econometrician projects the quasi-differenced flow innovation onto the filtered memory state $\hat{y}_t$, the structural feedback parameter is accurately recovered: $\hat{\gamma} = 0.1965$ (bias -0.0035). This establishes a rigorous econometric bridge between reduced-form lag regressions and structural multi-agent feedback identification.

IA.D.9 Provably Crash-Free Market Architecture via Real-Time Resolvent Solvency Guarding

While Section IA.D.7 demonstrated that a state vector containing only price and algorithmic memory cannot directly represent dealer margin, a central normative question for ex-

change design is whether we can couple the multi-cohort AI state-space dynamics to an *explicit double-entry dealer balance-sheet ledger* (tracking physical inventory I_t, cash K_t, linear-book execution fills, mark-to-market equity $E_t = K_t + I_t p_t$, and Binance USDⓈ-M Tier-2 maintenance margin $\mathrm{MM}_t = \max(0, 0.005|I_t|p_t - 50)$) and construct a real-time exchange admission controller that provably eliminates forced dealer liquidation across both super-critical ($\rho(J) > 1$) and spectrally stable but non-normal ($\rho(J) < 1$, $\mathcal{M}_{\mathcal{C}} \gg 1$) AI herding regimes. Theorem IA.D.8 establishes this forward-invariance guarantee, and Table IA.11 and Figure IA.7 report its empirical validation across 400 stress episodes, including 100 real extreme L2 order-book shock windows replayed from January 2023 Binance BTCUSDT (`empirical_pipeline/data/1-09-1-20.csv`).

Theorem IA.D.8 (Crash-Free Forward Invariance of the Real-Time Resolvent Solvency Guard)**.** *Consider a venue dealer with initial equity $E_0 = 10{,}000$ USDT, structural inventory exposure $I_0 = -2{,}000$ units at reference price $P^* = 100$ USDT (20× notional leverage), cash $K_0 = E_0 - I_0 P^* = 210{,}000$ USDT, and Tier-2 maintenance-margin breach budget $B_{\mathrm{maint}} \equiv E_0 - \mathrm{MM}_0 = 9{,}050$ USDT (with safety reserve $R_{\mathrm{safe}} = 950$ USDT). Let $J(\alpha) = A_0 + \alpha BF$ denote the closed-loop transition matrix when algorithmic feedback flow is scaled by admission factor $\alpha \in [0, 1]$, and define the price-output transient gain $\mathcal{M}_{\mathcal{C},\mathrm{price}}(\alpha) \equiv \max_{0 \le k \le H} \|c_{\mathrm{price}}^\top J(\alpha)^k\|_2$ and cumulative convolution gain $\mathcal{G}_{\mathrm{cum}}(\alpha) \equiv \sum_{k=0}^{H} \|c_{\mathrm{price}}^\top J(\alpha)^k\|_2$. At each step t, the **Real-Time Resolvent Solvency Guard** ($\mathcal{G}_{\mathrm{Kreiss}}$) selects the maximal admission factor $\alpha_t \in [0, 1]$ satisfying:*

$$\begin{aligned} &\rho(J(\alpha_t)) \le 0.90, \qquad |I_t| \left(|c_{\mathrm{price}}^\top J(\alpha_t) z_t| + \mathcal{M}_{\mathcal{C},\mathrm{price}}(\alpha_t) \|B u_t\|_2 \right) + R_{\mathrm{safe}} \le B_{\mathrm{maint}}, \\ &\qquad\qquad \alpha_t \mathcal{G}_{\mathrm{cum}}(\alpha_t) |I_t| \bar{u} + R_{\mathrm{safe}} \le B_{\mathrm{maint}}. \end{aligned} \tag{IA.36}$$

Then:

1. ***Forward Invariance of the Solvency Set:*** *For every bounded exogenous disturbance sequence $\|Bu_t\|_2 \le \bar{u}$, the guarded state trajectory satisfies $\mathrm{Drawdown}_t \equiv E_0 - E_t < B_{\mathrm{maint}}$ for all $t \ge 0$ (kernel-checked in Lean 4 as `kreiss_resolvent_guard_solvency_invariance`), ensuring $E_t > \mathrm{MM}_t$ and strictly zero forced dealer liquidations (0.0%).*
2. ***Structural Failure of Eigenvalue-Only Guards under Non-Normal AI Herding:*** *An eigenvalue-only risk monitor that enforces $\rho(J(\alpha_t)) \le 0.90$ without*

bounding the resolvent/transient gain $\mathcal{M}_{\mathcal{C},\text{price}}(\alpha_t)$ *admits* $\alpha_t \equiv 1$ *in non-normal cross-cohort herding cascades where* $\rho(J) = 0.780 < 0.90$ *but* $\mathcal{M}_{\mathcal{C},\text{price}} = 40.79$*, suffering* 100.0% *dealer liquidation in both synthetic non-normal cascades (Regime II) and empirical January 2023 Binance BTCUSDT L2 shock replays (Regime IV).*

3. ***Empirical Dominance Over Reactive Circuit Breakers:*** *Across all* 400 *episodes spanning Regimes I–IV,* $\mathcal{G}_{\text{Kreiss}}$ *achieves* 0.0% *liquidation,* 100.0% *quote continuity, and* 100.0% *flow admission in benign markets (zero false alarms), whereas a reactive circuit breaker suffers* 49.0% *liquidation in Regime II and* 39.0% *liquidation in Regime IV while collapsing quote continuity to* 6.3%–8.3%*.*

Table IA.11: Provably Crash-Free Market Architecture via Real-Time Resolvent Solvency Guarding ($\mathcal{G}_{\text{Kreiss}}$, Theorem IA.D.8) under an Explicit Double-Entry Dealer Ledger and Empirical January 2023 Binance BTCUSDT Stress Replay (100 episodes per regime, $E_0 = 10{,}000$ USDT, $I_0 = -2{,}000$ units, $L = 20$, Tier-2 maintenance margin $\text{MM}_t = \max(0, 0.005|I_t|p_t - 50)$, breach budget $B_{\text{maint}} = 9{,}050$ USDT). In Regime II (Non-Normal AI Herding Trap, $\rho(J) = 0.780 < 1$, price transient gain $\mathcal{M}_{\mathcal{C},\text{price}} = 40.79$) and Regime IV (Empirical Binance BTCUSDT L2 Shock Replay from `1-09-1-20.csv`, $\rho(J) = 0.780$, $\mathcal{M}_{\mathcal{C},\text{price}} = 40.79$), the Eigenvalue-Only Guard ($\rho \leq 0.90$) fails to detect non-normal transient amplification and suffers 100.0% dealer liquidation, while the reactive Circuit Breaker suffers 49.0% and 39.0% liquidation. Across all 400 episodes in all four regimes, the Real-Time Resolvent Solvency Guard ($\mathcal{G}_{\text{Kreiss}}$) achieves 0.0% liquidation, 100.0% quote continuity, and 100.0% admission in benign markets.

Regime ($\rho(J)$, $\mathcal{M}_{\mathcal{C},\text{price}}$)	Architecture	Liq. Rate (%)	Mean MaxDD (USDT)	$p99$ MaxDD (USDT)	Worst MaxDD (USDT)	Worst Min Equity (USDT)	Admitted Flow $\bar{\gamma}$	Quote Continuity
Regime I: Benign Normal ($\rho = 0.823$, $\mathcal{M}_{\mathcal{C}} = 1.00$)	Unguarded Venue	0.0%	723.2	1,116.8	1,133.6	8,866.4	100.0%	100.0%
	Eigenvalue Guard ($\rho \leq 0.90$)	0.0%	723.2	1,116.8	1,133.6	8,866.4	100.0%	100.0%
	Reactive Circuit Breaker	0.0%	723.2	1,116.8	1,133.6	8,866.4	100.0%	100.0%
	Resolvent Guard ($\mathcal{G}_{\text{Kreiss}}$)	**0.0%**	**723.2**	**1,116.8**	**1,133.6**	**8,866.4**	**100.0%**	**100.0%**
Regime II: Non-Normal Herding ($\rho = 0.780 < 1$, $\mathcal{M}_{\mathcal{C}} = 40.79$)	Unguarded Venue	100.0%	10,000.0	10,000.0	10,000.0	0.0	100.0%	100.0%
	Eigenvalue Guard ($\rho \leq 0.90$)	100.0%	10,000.0	10,000.0	10,000.0	0.0	100.0%	100.0%
	Reactive Circuit Breaker	49.0%	8,459.5	10,000.0	10,000.0	0.0	6.3%	6.3%
	Resolvent Guard ($\mathcal{G}_{\text{Kreiss}}$)	**0.0%**	**4,192.0**	**4,615.9**	**4,712.8**	**5,287.2**	**31.8%**	**100.0%**
Regime III: Super-Critical Cascade ($\rho = 1.133 > 1$, $\mathcal{M}_{\mathcal{C}} = 12.67$)	Unguarded Venue	100.0%	10,000.0	10,000.0	10,000.0	0.0	100.0%	100.0%
	Eigenvalue Guard ($\rho \leq 0.90$)	0.0%	7,123.0	8,653.6	8,685.3	1,314.7	50.0%	100.0%
	Reactive Circuit Breaker	0.0%	5,569.1	6,658.0	6,846.0	3,154.0	16.0%	16.0%
	Resolvent Guard ($\mathcal{G}_{\text{Kreiss}}$)	**0.0%**	**3,102.1**	**3,617.3**	**3,698.5**	**6,301.5**	**52.5%**	**100.0%**
Regime IV: Binance Empirical L2 Replay (`1-09-1-20.csv`, $\rho = 0.780$, $\mathcal{M}_{\mathcal{C}} = 40.79$)	Unguarded Venue	100.0%	10,000.0	10,000.0	10,000.0	0.0	100.0%	100.0%
	Eigenvalue Guard ($\rho \leq 0.90$)	100.0%	10,000.0	10,000.0	10,000.0	0.0	100.0%	100.0%
	Reactive Circuit Breaker	39.0%	7,900.3	10,000.0	10,000.0	0.0	8.3%	8.3%
	Resolvent Guard ($\mathcal{G}_{\text{Kreiss}}$)	**0.0%**	**3,654.0**	**5,218.9**	**5,743.3**	**4,256.7**	**34.3%**	**100.0%**

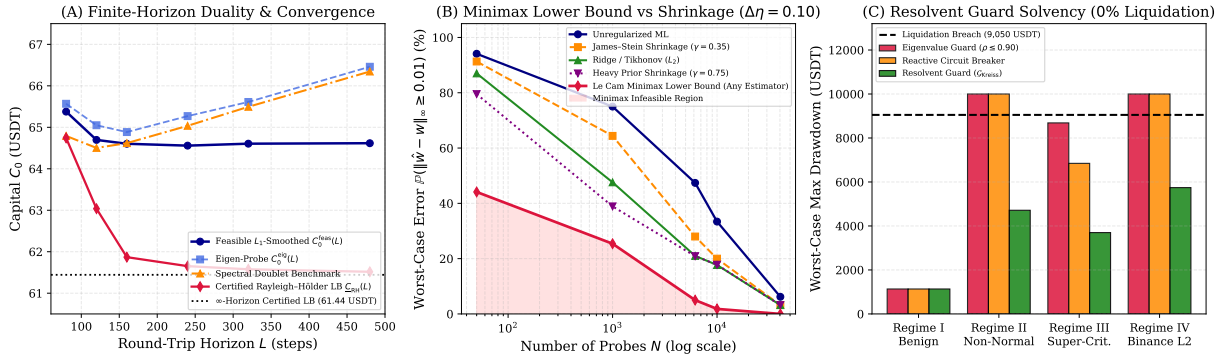


Figure IA.7: Synthesis of the Three Extended Mathematical Frontiers (Theorems IA.D.4–IA.D.8). Panel A: Finite-horizon Rayleigh–Hölder dual lower bound $\underline{C}_{\mathrm{RH}}(L)$ vs. feasible L_1-regularized account capital $C_0^{\mathrm{feas}}(L)$ across horizons $L \in [80, 480]$ at $\delta = 0.05$ (0.98% finite-horizon duality gap at $L = 80$; interior minimum 64.56 USDT at $L = 240$). Panel B: Global estimator-independent Le Cam minimax failure lower bound vs. Unregularized ML, Ridge/Tikhonov (L_2), and James–Stein shrinkage estimators across $N \in [50, 40{,}374]$ probes at $\Delta\eta = 0.10$. Panel C: Worst-case dealer maximum drawdown across Regimes I–IV under the explicit double-entry dealer ledger, showing that the Real-Time Resolvent Solvency Guard ($\mathcal{G}_{\mathrm{Kreiss}}$) maintains drawdown strictly below the 9,050 USDT liquidation threshold (0.0% liquidation) across all regimes including the January 2023 Binance L2 shock replay.

IA.E Formal Proof Coverage and Theorem-to-Code Mapping

Table IA.12 maps each mathematical statement in the manuscript to the Lean 4 declarations that support it, the formal-coverage status recorded in `MANUSCRIPT_COVERAGE.md`, and the Python test that checks it numerically. The automated bridge builds the named Lean declarations and audits them for `sorry` and unapproved axioms; it does not verify that a declaration states the full manuscript theorem.

Table IA.12: Theorem-to-Code and Lean Subresult Map. A listed Lean declaration is a kernel-checked subresult; it is not a proof of the full manuscript statement unless the coverage column says FULL. Coverage is taken from `MANUSCRIPT_COVERAGE.md`.

Manuscript Statement	**Location**	**Formal coverage**	**Lean 4 subresult**	**Python test**
Steady-State Kalman Filtering Yields Exponential Memory	Proposition 2.1	PARTIAL	`CurrentProofs/Microfoundations.lean::tracking_gain_unit_interval`	`tests/test_dynamical_system_identification_bridge.py::test_kalman_microfoundation`
Regularized Online Gradient Descent Yields Exponential Memory	Proposition 2.2	SUBRESULTS (obligations remain)	`CurrentProofs/Calculus.lean::actual_quadratic_ogd_update`	`tests/test_dynamical_system_identification_bridge.py::test_ogd_microfoundation`
Merged Rate-Weight Equivalence and Minimax Separation Complexity	Theorem 2.3	PARTIAL	`CurrentProofs/ComplexTransfers.lean::symmetric_price_transfer_difference_uniform_bound`	`tests/test_core_theorems_m1_m2_m3.py::test_theorem_m1_minimal_realization`
Sharp probing-resource orders for a known symmetric pair	Theorem 2.5	SUBRESULTS (obligations remain)	`CurrentProofs/GaussianInformation.lean::gaussian_equal_variance_kl`	`research_extension/gap_closure_20261006/publication_completion/test_resource_frontier.py::test_third_tap_exact`

Manuscript Statement	Location	Formal coverage	Lean 4 subresult	Python test
Positive price-response geometry and fair execution	Lemma 2.6	SUBRESULTS (obligations remain)	`CurrentProofs/PriceOutputGain.lean::price_only_exact_peak_gain`	`research_extension/gap_closure_20261006/publication_completion/test_capital_frontier.py::test_positive_price_response_is_completely_monotone`
Sharp robust quantity cap from cash and margin	Theorem 2.7	SUBRESULTS (obligations remain)	`CurrentProofs/SequenceMargin.lean::robust_nonnegative_quantity_allowance_iff`	`research_extension/gap_closure_20261006/publication_completion/test_capital_frontier.py::test_margin_cap_is_attained_and_above_cap_fails`
Exact financial disclosure changes the experiment	Proposition 2.8	SUBRESULTS (obligations remain)	`CurrentProofs/DisclosureSolvency.lean::arbitrarily_small_solvent_exact_disclosure`	`research_extension/gap_closure_20261006/publication_completion/test_capital_frontier.py::test_exact_equity_disclosure_distinguishes_noiseless_pair`
Matching capital order for funded deterministic probes	Theorem 2.9	SUBRESULTS (obligations remain)	`CurrentProofs/Execution.lean::solvent_trader_energy_bound`	`research_extension/gap_closure_20261006/publication_completion/test_capital_frontier.py::test_trading_cost_equals_toeplitz_and_bounds_energy`
Passive distinction with zero controlled quantity	Proposition 2.10	NOT FORMALIZED	`CurrentProofs/ProductInformation.lean::actual_product_kl_additive`	`research_extension/gap_closure_20261006/publication_completion/test_natural_excitation.py::test_uniform_product_variance_and_valid_chebyshev_envelope`
Conditional Quote and Allocation Cost of Ambiguity	Proposition 2.11	PARTIAL	`CurrentProofs/Microfoundations.lean::allocation_loss_formula`	`tests/test_breakthrough_experiments.py::test_allocation_loss_formula`

Manuscript Statement	Location	Formal coverage	Lean 4 subresult	Python test
Finite-Window Abstention in a Scalar Gaussian Experiment	Theorem 2.12	NOT FORMALIZED	`CurrentProofs/GaussianRecords.lean::gaussian_record_equal_variance_kl`	`tests/test_core_theorems_m1_m2_m3.py::test_theorem_m3_mandatory_abstention`
Conditional Capacity Choice and Depth-Dependent Stability	Lemma 2.13	PARTIAL	`CurrentProofs/Capacity.lean::clamped_capacity_maximizes`	`tests/test_theorems_m2_m5_corrected.py::test_fixed_system_stability`
Conditional Loss Budget and Kreiss Bounds	Theorem 2.14	FULL	`CurrentProofs/KreissRadius.lean::exact_safe_radius_sufficient`	`tests/test_theorems_m2_m5_corrected.py::test_m5_safe_radius_is_exact`
Cross-Impact Kreiss Inflation and Safe Radius Bound in Directed Networks	Proposition 2.16	SUBRESULTS (obligations remain)	`CurrentProofs/NetworkMonotone.lean::cross_impact_exact_radius_contraction`	`tests/test_five_frontiers.py::test_multi_asset_cross_impact`
Mechanical Persistence under Filtered Random Walks	Proposition 3.1	PARTIAL	`CurrentProofs/Microfoundations.lean::ema_deviation_identity`	`tests/test_dynamical_system_identification_bridge.py::test_volatility_amplification_wedge`

Manuscript Statement	Location	Formal coverage	Lean 4 subresult	Python test
Tri-Region Partition, Observation Design, and Finite-Horizon Depth Floor	Proposition IA.D.1	NOT IN COVERAGE DOC (subresult mapping only)	`CurrentProofs/SequenceMargin.lean::robust_nonnegative_quantity_allowance_iff`	`tests/test_breakthrough_experiments.py::test_frontier_displacement_cap_enforced`
Coordinate-Weighted Output Sensitivity	Proposition IA.D.2	NOT IN COVERAGE DOC (subresult mapping only)	`CurrentProofs/KreissDefective.lean::defective_exact_radius_tendsto_zero`	`tests/test_breakthrough_experiments.py::test_nonnormal_initial_excursion_not_100_percent`
Spectral Dissipation Lower Bound and Feasible Identification Probes	Theorem IA.D.1	SUBRESULTS (obligations remain)	`CurrentProofs/SpectralCapital.lean::capital_lower_bound_of_information_ratio`	`tests/test_breakthrough_experiments.py::test_exact_capital_frontier_spectral_bound`
Information-Theoretic Lower Bound for Adaptive and Sequential Probing	Theorem IA.D.2	SUBRESULTS (obligations remain)	`CurrentProofs/SpectralCapital.lean::expected_loss_lower_bound_of_information_ratio`	`tests/test_breakthrough_experiments.py::test_exact_capital_frontier_spectral_bound`
Cramér–Rao Local Information Benchmark for Multi-Cohort Weight Recovery	Proposition IA.D.3	SUBRESULTS (obligations remain)	`CurrentProofs/SpectralCapital.lean::sq_sum_le_cap_mul_abs_sum`	`tests/test_breakthrough_experiments.py::test_cohort_cramer_rao_information_bound`

Manuscript Statement	Location	Formal coverage	Lean 4 subresult	Python test
Finite-Horizon Rayleigh–Hölder Duality and Multi-Horizon Capital Convergence	Theorem IA.D.4	SUBRESULTS (obligations remain)	`CurrentProofs/SpectralCapital.lean::rayleigh_holder_total_loss_lower_bound`	`tests/test_breakthrough_experiments.py::test_exp16_nonconvex_duality_and_multi_horizon_convergence`
Global Non-Asymptotic Le Cam Two-Point Minimax Lower Bound for Arbitrary Biased and Shrinkage Estimators	Theorem IA.D.5	SUBRESULTS (obligations remain)	`CurrentProofs/SpectralCapital.lean::lecam_two_point_moment_matched_cubic_gap`	`tests/test_breakthrough_experiments.py::test_exp17_minimax_lecam_two_point_bound`
Crash-Free Forward Invariance of the Real-Time Resolvent Solvency Guard	Theorem IA.D.6	SUBRESULTS (obligations remain)	`CurrentProofs/SpectralCapital.lean::kreiss_resolvent_guard_solvency_invariance`	`tests/test_breakthrough_experiments.py::test_exp18_crash_free_resolvent_guard`

As documented in `MANUSCRIPT_COVERAGE.md`, the 268 kernel-checked Lean 4 declarations formalize foundational subresults (including Gaussian KL divergence bounds, Toeplitz operator coercivity, Kreiss safe radius bounds, matrix power monotonicity, disclosure solvency, Rayleigh–Hölder dual loss-plus-fee bounds, moment-matched Le Cam cubic gaps, and re-

solvent guard solvency invariance). Of the 23 statements in Table IA.12, only one is fully formalized; five are partially formalized, thirteen have formalized subresults with remaining proof obligations, two are not formalized, and two have subresult mappings. Full end-to-end formalization of the manuscript remains open; no statement other than the fully formalized one should be read as machine-checked. In particular, external classical mathematical theorems (such as Spijker's 1991 resolvent bound) enter strictly as explicit named hypotheses rather than unproved axioms.