Machine Markets
Research Status and Remaining Work

Working record for author and professor review

30 October 2026

Purpose. Preserve the verified returned results, their economic interpretation and the work required before stronger claims. This record accompanies the existing papers; it does not replace them or certify their publication quality. Read the completed records C1–C6 and open tasks R1–R8 when deciding what to do next.

Contents

Returned advanced experiment: evidence and limits **2**

Research work register: 30 October 2026 **5**

Returned advanced experiment: evidence and limits

Research question and financial institution. Can information about an aggregate market response improve an exposure decision after information, execution and carry costs, while the hidden trading cohorts remain difficult to recover? The new experiment concerns a price-taking, unlevered inventory-warehousing desk. It buys a fixed inventory, holds it for $H = 10, 60, 240$ one-minute *model* periods and closes. A prequoted service premium arrives after settlement. It cannot change market feedback, quotes or other traders. This is a separate application from a dealer or venue that changes admission and thereby changes the market itself. The objective is expected service profit under a stress reserve and cash constraint, not logarithmic portfolio growth or predicted crypto returns.

Completed execution record. The author returned the Deep Research archive with base seed 20261101. All 611 manifest entries, all 288 checkpoint-to-statistics links and all 15 released source files were verified. The 288 generated economies use 24 seed clusters, 4,096 measurement repetitions per economy and N, $N = 8, 32, 128$, three information-price settings, seven measurement rules and a same-error oracle. The 20,736 policy/cost rows share readings and are not independent market experiments. All 27 implemented tests passed. Numeric policy fields reproduce the local same-seed run to about 1.25×10^{-14}. This is cross-environment reproducibility of the same implementation, not independent replication. The returned runtime used Python 3.13.18 and the pinned numerical packages; the research runner took 52.71 seconds, excluding notebook setup.

What Lean checked. Lean 4.19.0 checked seven concrete rational inequality certificates, with no reported axioms. They include a conservative positive-value numerical witness greater than 0.14 hypothetical model dollars for the original primary contract. They do not formalize the Gaussian analytic tail step, market dynamics, general separation theorem, translation from observations to economics or novelty. The earlier positive illustration is retained as a pointwise example. It must not be read as a positive average-value theorem over all economies or prices.

Economic value depends on the information bill. For the standard Gaussian rule within the independent-Gaussian economy design:

Information price	Mean advantage	95% interval	Positive / total
Free	+0.13047	[0.08458, 0.18420]	99 / 144
Baseline	−0.02953	[−0.07022, 0.01310]	59 / 144
Five times baseline	−0.17212	[−0.21856, −0.12715]	26 / 144

The comparator is the analytically optimal no-data rule with the same permitted unsafe-action probability, $\gamma = 0.025$. Values are hypothetical dollars per inventory-service decision. Baseline cost is $N(0.0005 + 0.000025H)$; its dependence on horizon and N is explicit, not an exchange calibration. The intervals use 5,000 percentile bootstrap resamples of 24 seed-cluster averages. This price-specific analysis was added after the return and is exploratory, not preregistered confirmation or uncertainty about actual market profit. The baseline-price interval crosses zero: this design does not establish a positive average benefit after costs. The high-price interval is negative under the stated simulator sampling scheme.

Protection and profitability are different outcomes. For purchased iid-Gaussian rules in iid-Gaussian economies, the mean simulated unsafe-action frequency was 0.834%, below the 2.5% allowance. Across 6,535 declared-valid purchased-rule checks, Holm correction produced zero rejections. Individual finite-sample frequencies can exceed the allowance; non-rejection is not a proof. An unsafe action means its reserve can breach the budget under a permitted worst-case shock path. It is not an observed liquidation, default or historical crash. The Gaussian and declared-variance rules coincide exactly under iid readings. Their advantage over a full-response ellipsoid is a task-specific comparison of established methods, not a new estimator.

Failures retained and interpreted.

Purchased Gaussian rule: observation/future setting	Mean unsafe action
Independent Gaussian readings, stationary future	0.834%
Correlated readings, iid variance used	14.695%
Downward-biased readings	31.169%
Future feedback changes	61.033%

These are unweighted descriptive averages of dependent configuration rows, including the zero-price setting. They are not pooled market failure probabilities. Under correlated Gaussian readings, a declared-variance rule reduces the mean to 0.409%, and a projected Student-t rule to 1.153%. Protection can improve while the information purchase remains uneconomic. The Student-t guarantee requires independent Gaussian reset totals; good observed performance on t_3 noise does not extend that guarantee. Median of means needs independent resets and an externally valid variance bound. At insufficient N it uses a paid prior-bound fallback. Neither robust interval repairs mean bias or a change in future dynamics. The full design retains 7,489 negative-advantage rows and 1,743 unaffordable purchase rows. Pooling all methods and scenarios into one success percentage would be misleading.

Information acquisition is an open economic decision. The fixed-N protocol buys information when affordable; it does not optimize the purchase. In 32 of 288 generated economies the known service premium is no greater than known linear execution/carry cost. Positive inventory then has no positive expected service surplus before any information is obtained. In the first indexed case, s000_c000 at $N = 32$ and baseline price, no-data and response-informed oracle profit both equal approximately 0.75902, because response information cannot relax the relevant position limit. Paying 0.024 for information leaves approximately 0.73500. A useful measurement need not change the financial action. A purchase rule should use only information available before purchase, include the option to abstain, and avoid selecting N or a policy after looking at hidden gains or test outcomes. That rule is still to be proved and validated; it has not been added to this returned run.

Observation boundary and contribution status. There are no historical coins, exchange tapes or crash replays in this run. The quote-data stage is **NOT_RUN**; the unit-test quote fixture is synthetic. Controlled resets, known plant/noise inputs, a feedback cap and stationarity are maintained assumptions. The finite-horizon positive-response certificate needs no adaptation-rate floor when every relevant tap is observed under that reset contract. This does not remove the

hidden-tail problem for infinite-horizon or ordinary public-tape certification. New observations do not independently verify the earlier passive, capital, AI or equilibrium claims.

The supported interpretation is a conditional financial separation: aggregate response information can support a specified exposure decision without latent cohort recovery, and positive net value requires additional economic conditions. The written general proof still needs independent review and its distinction from existing identification, data-informativity and decision theory needs specialist assessment. No publication-quality rating follows from execution.

Traceability. The immutable experiment source is repository commit `6a319b0db147`. The return audit is in commit `c0c0c010ec00`. The study directory and its relative audit directory are shown separately for readability:

```
studies/advanced_cost_separation_20261029/
validation/returned_colab_20261030/
```

The uploaded archive is:

```
MM_advanced_deep_seed20261101_v1.zip
```

Its SHA256 begins `e35c10f6b276`; the full hash, configuration, source fingerprint and exploratory price/paired audits are retained in that directory. No original return outcome has been overwritten.

Research work register: 30 October 2026

This register is a dated snapshot, not a claim that the project is complete. *Completed* means the named artifact or check exists. *Written, review open* means an argument is present but has not received independent mathematical sign-off. *Pending* means the evidence or proof is still required. Good outcomes from another experiment do not silently close a task.

Completed records to preserve

Record	Verified completion and boundary
C1	Returned archive, source fingerprint and 288 checkpoint links verified. Same-code reproducibility; no independent implementation.
C2	All 27 implementation checks passed. They do not test every theorem in the papers.
C3	Seven numerical Lean certificates passed. The full general theorem is not formalized.
C4	288-economy synthetic audit completed; every negative and unaffordable row retained. No market-data stage ran.
C5	Pointwise positive-value illustration retained; broad price-sensitive results and assumption failures recorded. No universal profitability claim.
C6	Targeted literature comparison and a written restricted separation exist. Originality and complete proof review remain open.

Core work before stronger contribution claims

R1. Establish the precise literature gap. Status: pending specialist review. Compare the exact assumptions, observation experiment, decision target and conclusion with identification theory, data informativity, task-optimal exploration, functional estimation and constrained financial decisions. Separate latent cohort recovery from learning the observable response. Do not claim that control or allocation without complete identification is new. *Completion evidence:* a source-backed comparison table identifying a specific non-overlap, or a clear statement that the result is an application or corollary; resolved reviewer comments. More simulations cannot complete R1.

R2. Independently audit the load-bearing proofs. Status: written, review open. Review the finite-cohort renewal bounds, observation-to-confidence implication, uniformity in hidden memory, fixed-query testing obstruction, all fee/cash constraints and the same-error no-data optimum. Audit the older broad certification, passive-identification and capital claims separately. Check quantifiers, constants, admissible probes, reset assumptions and counterexamples. *Completion evidence:* theorem-by-theorem assumptions and proof register, resolved objections and independent mathematical review. Record each Lean theorem's actual scope; do not substitute seven scalar facts for a general proof.

R3. Make the information purchase economically explicit. Status: pending. Define what the desk knows before buying observations and whether the acquisition objective is robust, Bayesian or another stated criterion. Include abstention and screen cases where the premium cannot cover known costs or information cannot change the constrained position. Select sample size and

price using available information, not hidden gains or test outcomes. *Completion evidence:* a precise acquisition rule, a valid decision-value or regret statement, matched-information/cost comparisons, and a frozen new validation design. Reusing this return to design a rule is exploratory.

R4. Justify the financial institution and measurement contract. Status: partial. The fixed inventory, horizon, post-settlement premium and accounting are explicit. The economic source of the premium, operating cost of reset observations, information reliability and acquisition opportunity cost need justification. An own-account investor, warehousing desk, market maker and venue supervisor have different actions and incentives. Keep their claims separate. *Completion evidence:* a coherent institutional mechanism and documented parameter/cost interpretation, with comparative statics. If logarithmic utility is introduced, rederive the objective rather than relabeling expected profit.

R5. Establish an observation bridge or restrict the empirical claim. Status: pending. Ordinary public quotes are not independent resets. Address initial state, fundamental-value proxies, endogenous order flow, changing parameters, noise covariance, sign changes and unobserved tails. A reduced-form VAR score is not automatically the structural gain in the theorem. *Completion evidence:* a justified identification/partial-identification argument under stated conditions, or a forecasting-only interpretation with certification claims removed from that application. Test the conditions rather than treating good forecasts as proof of them.

R6. Run a protected real-data evaluation. Status: NOT_RUN. Document event-level quote provenance, instruments, units, dates, fees, carry and earlier use. Freeze venues/dates and specifications before test access. Use chronological purging, volatility and OFI/depth baselines, prespecified inference units and all failure outcomes. An exposure replay must account for spread, funding, fill/latency assumptions and capital at risk. Distinguish isolated round trips from a portfolio backtest and from market-wide intervention. *Completion evidence:* licensed/reproducible data description, immutable split and code, held-out results after costs, limitations and external replay. The synthetic unit-test fixture cannot complete R6.

R7. State what happens when certification is unavailable. Status: partial diagnostics. The returned bias and future-change cases demonstrate failure. Specify a defensible response such as abstention or a prior-only restriction; determine what verified disclosure could restore the missing assumption and what it costs. Repeated use requires simultaneous or time-uniform control, not repeated use of a single-decision confidence level. *Completion evidence:* a stated contract/failure-response rule with a valid guarantee or an explicit impossibility boundary, tested under prespecified violations. Larger samples under a wrong mean do not repair the problem.

R8. Obtain independent replication and reconcile claims. Status: pending. Have another implementation reproduce the central calculation and meaningful stress cases. Reconcile abstracts, conclusions, theorem tags, old result tables, duplicate paper layouts, co-author contributions and AI-use statements. Keep the original evidence and limitations visible. Treat internal "proved" tags as statements about written arguments until independent review is recorded. *Com-*

pletion evidence: an external replication record, a claim-to-evidence map, resolved discrepancies and co-author approval for circulation.

Extensions that remain ambitions

Endogenous dealer depth and quotes, venue-wide gating, equilibrium welfare, dynamic inventories and AI-adoption externalities are separate extensions. They should enter the main claim only with a proved mechanism and appropriate evidence. The new fixed-position experiment does not establish those outcomes. A strong core contribution need not contain every extension.

Order of work and maintenance rule

Prioritize R1–R3 together: determine what is new, whether it is true, and when the information purchase changes an economically meaningful decision. R4–R6 determine the strength of the finance application. R7–R8 determine defensible use and reproducibility. Professor review can guide and challenge each item; it is not a substitute for producing the missing proof or evidence.

For each update record the date, task ID, changed claim, immutable code/data reference, result and remaining limitation. Mark an item completed only when its named evidence exists. Update this register, its Markdown/JSON source and all PDF variants together. Publication quality remains a judgment about originality, correctness and usefulness, not a score produced by the notebook.